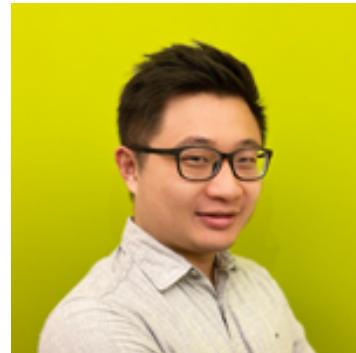


Modularizing Natural Language Processing

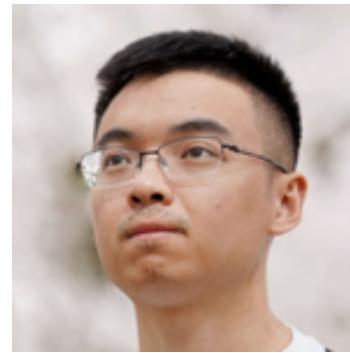
AAAI 2020 Feb 8, 2020



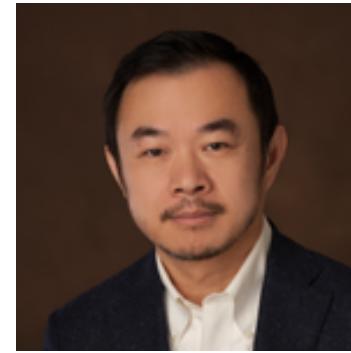
Carnegie Mellon University
School of Computer Science



Hector Liu

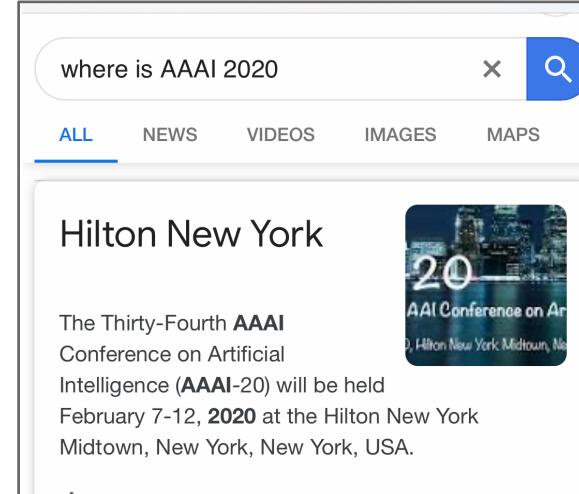
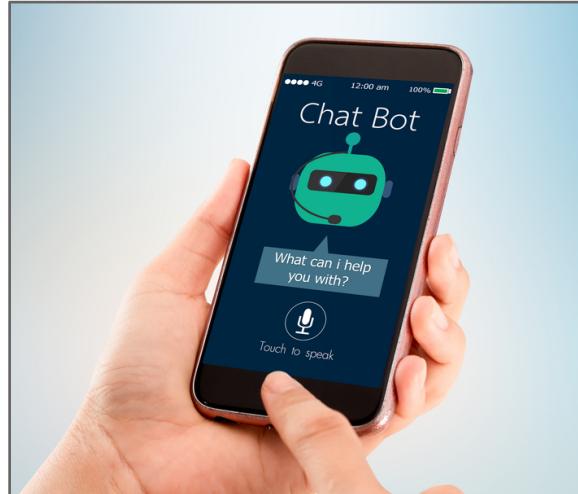
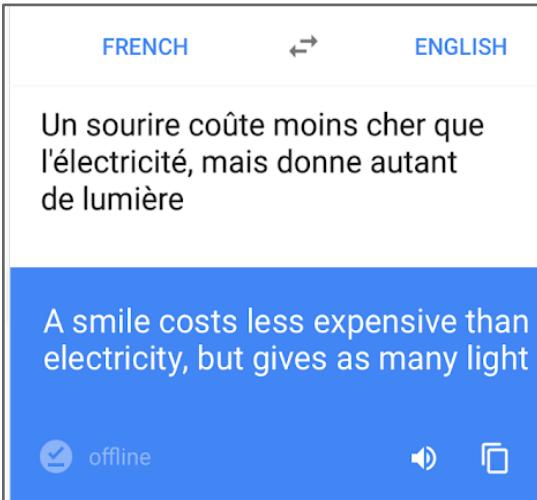


Zhitong Hu

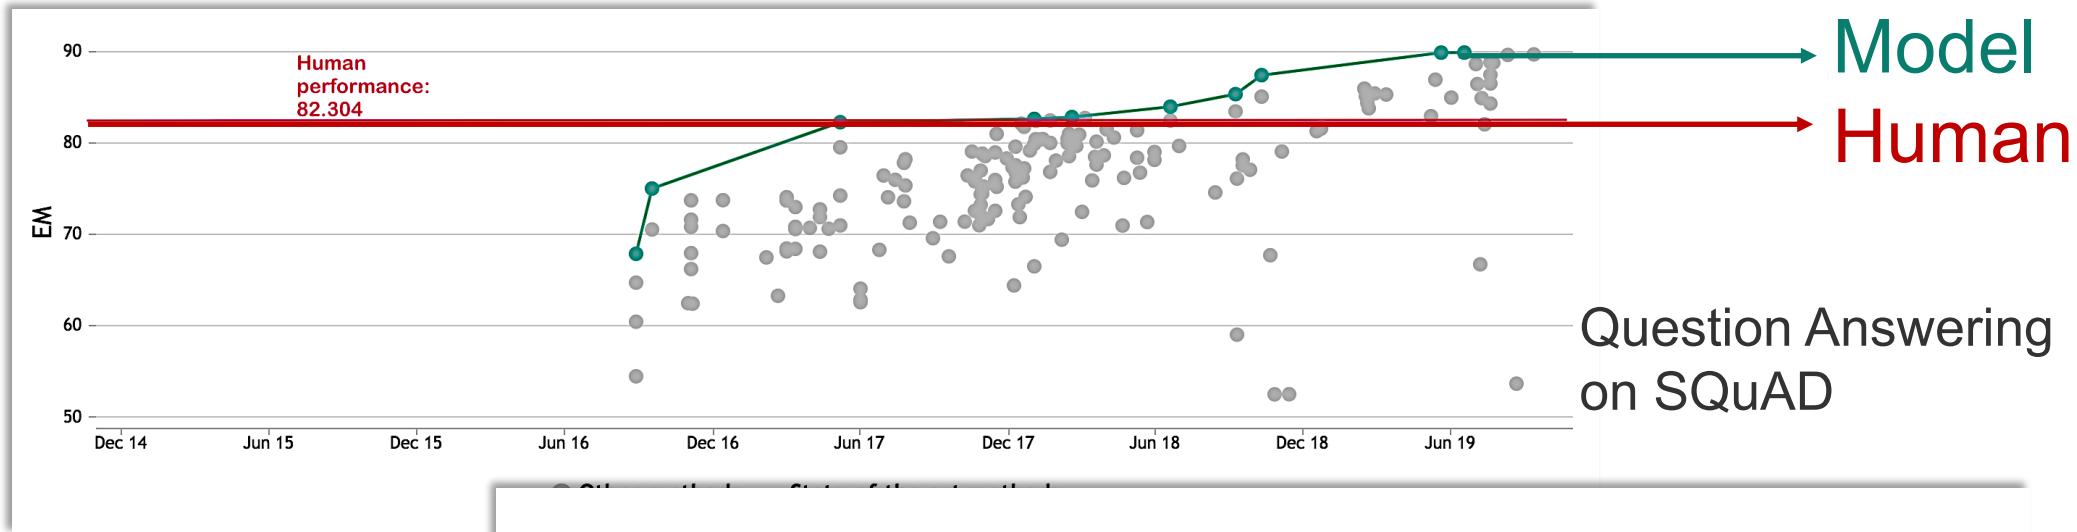


Eric Xing

NLP Applications



Inspirational Success on Benchmarks



Microsoft Achieves Human Parity on Chinese-English Machine Translation

LIKE

DISCUSS



MAR 15, 2018 • 2 MIN READ

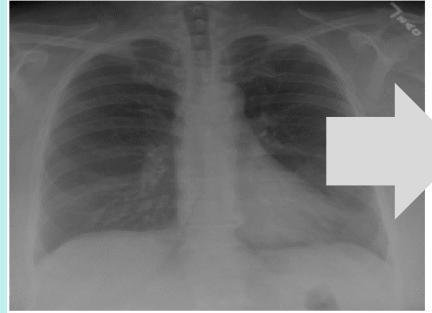
NLP in Real-world Context

The Healthcare Industry



Building a ready-to-use AI solution for this is

Extremely complex



Findings:

There are no focal areas of consolidation.
No suspicious pulmonary opacities.
Heart size within normal limits.
No pleural effusions.
There is no evidence of pneumothorax.
Degenerative changes of the thoracic spine.

Impression:

No acute cardiopulmonary abnormality.

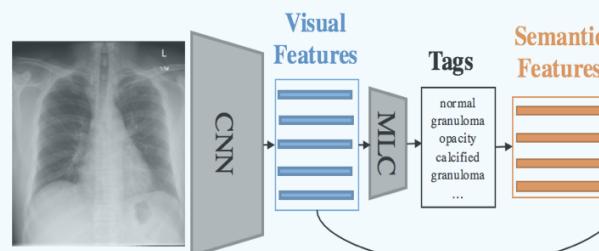
Task: Automatic Medical Report Generation

Requires inter-operation between diverse components

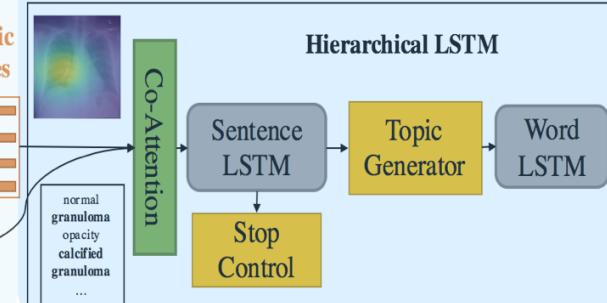
Raw Data Cleansing



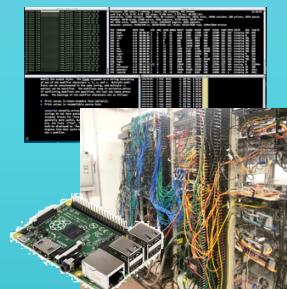
Data Enrichment



Model/Algorithm

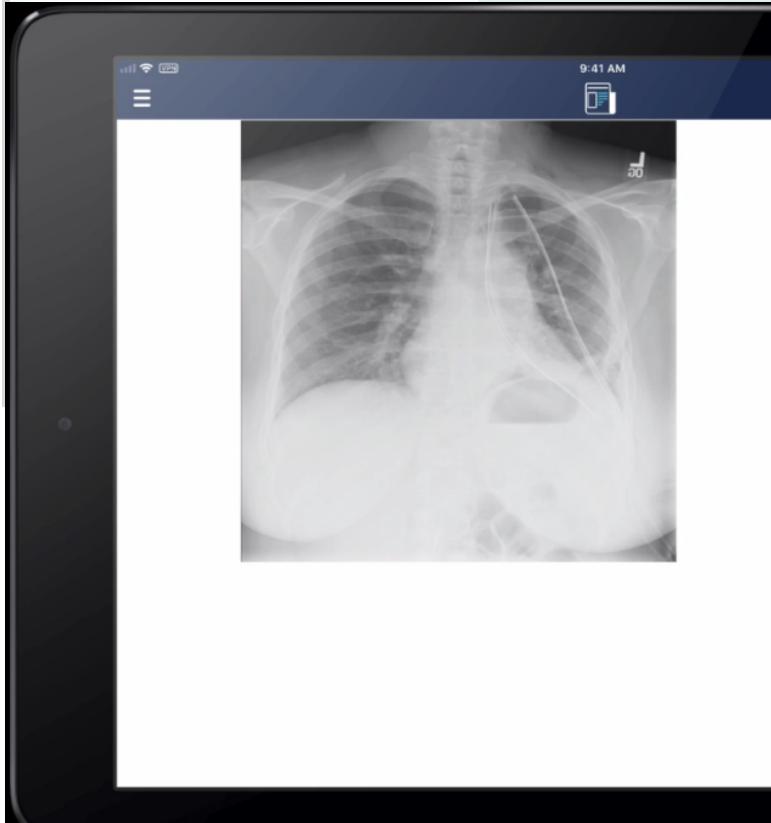


System/Infra



Building a ready-to-use AI solution for this is

User interface for Doctors



Petuum Med

Name Here Name Here

Name Here

The patient is a 62 year old male with a history of mild COPD who complains of cough, shortness of breath, fatigue, and fever progressively worsening for the past week. Today he measured a fever of 101 F. The productivity of his cough has progressively increased over the past two days. He has been using his albuterol inhaler two to three times daily but it is helping only minimally. He has a history of COPD, which he believes is only mild. He states that he is treated with antibiotics for a case of bronchitis or pneumonia almost every year by his primary care provider. He has never been hospitalized for pneumonia. He received approx 80mg IV of Lasix at that time and was 2.4L negative in 24 hours. He has never been hospitalized for pneumonia. He denies any known sick contacts recently. He denies chest pain but admits to some chest tightness and an increase in heart rate when he coughs a lot and is short of breath. He denies any recent weight changes or lower extremity pain or swelling. He denies any recent travel. He denies a history of lung disease, heart disease, or diabetes. He currently is a non-smoker but did smoke a pack a day for approximately 15 years prior to quitting five years ago.

Name Here

Discharge Medications:

1 Furosemide 20 mg PO Tablet (2 times a day).

Critical Information Extraction

Heart Failure

History Comorbidities
 Symptoms Cardiac Tests
 Lab Tests Medications

Symptoms

cough 3 , shortness of breath 1 , fatigue 2 , fever 1 , chest pain 1 , chest tightness 2 , increase in heart rate 2 , weight change 1 , lower extremity pain 2 , swelling 3

Medications

Furosemide 2 , antibiotics 3

Show context when clicked Clear

Lasix, Vol: 80 mg, Usage: IV

Furosemide, Vol: 20mg, Usage: PO (2 times a day)

Building a ready-to-use AI solution for this is

Extremely complex



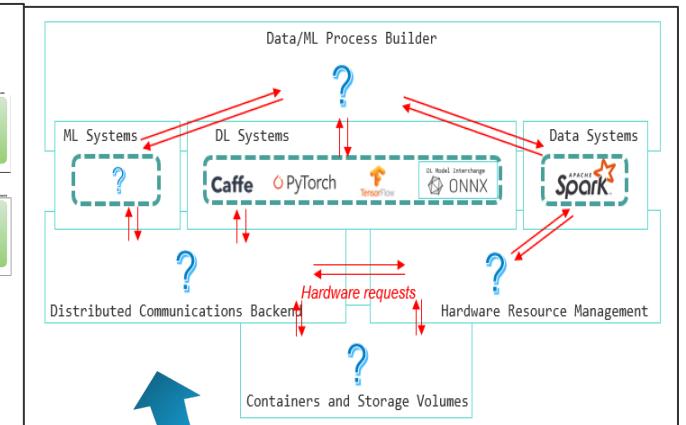
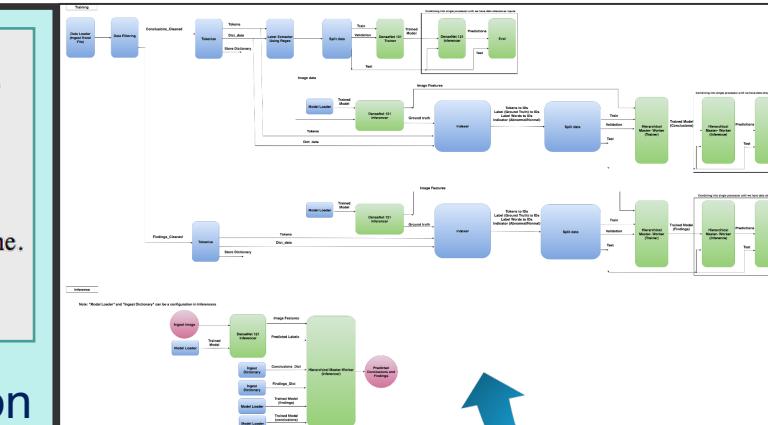
Findings:

There are no focal areas of consolidation.
No suspicious pulmonary opacities.
Heart size within normal limits.
No pleural effusions.
There is no evidence of pneumothorax.
Degenerative changes of the thoracic spine.

Impression:

No acute cardiopulmonary abnormality.

Task: Automatic Medical Report Generation

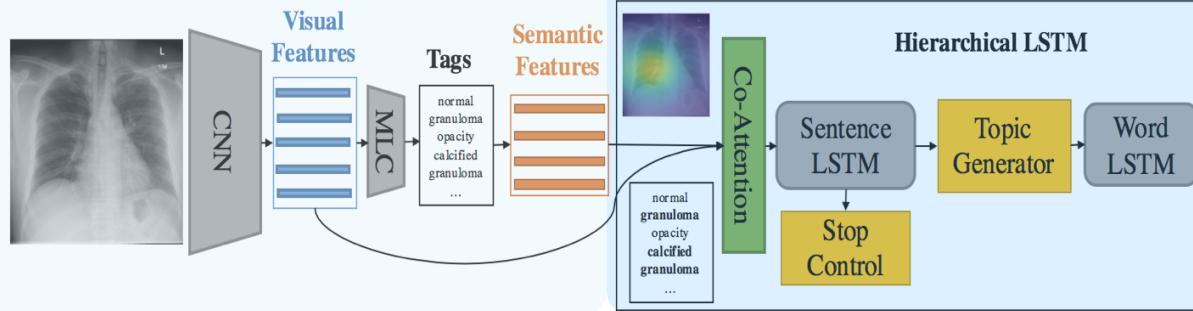


Requires inter-operation between diverse components

Raw Data Cleansing

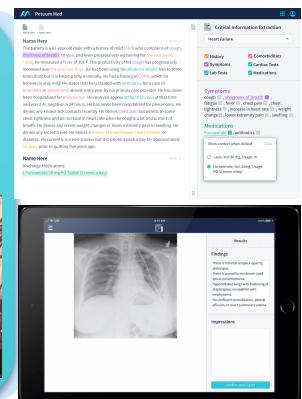


Data Enrichment



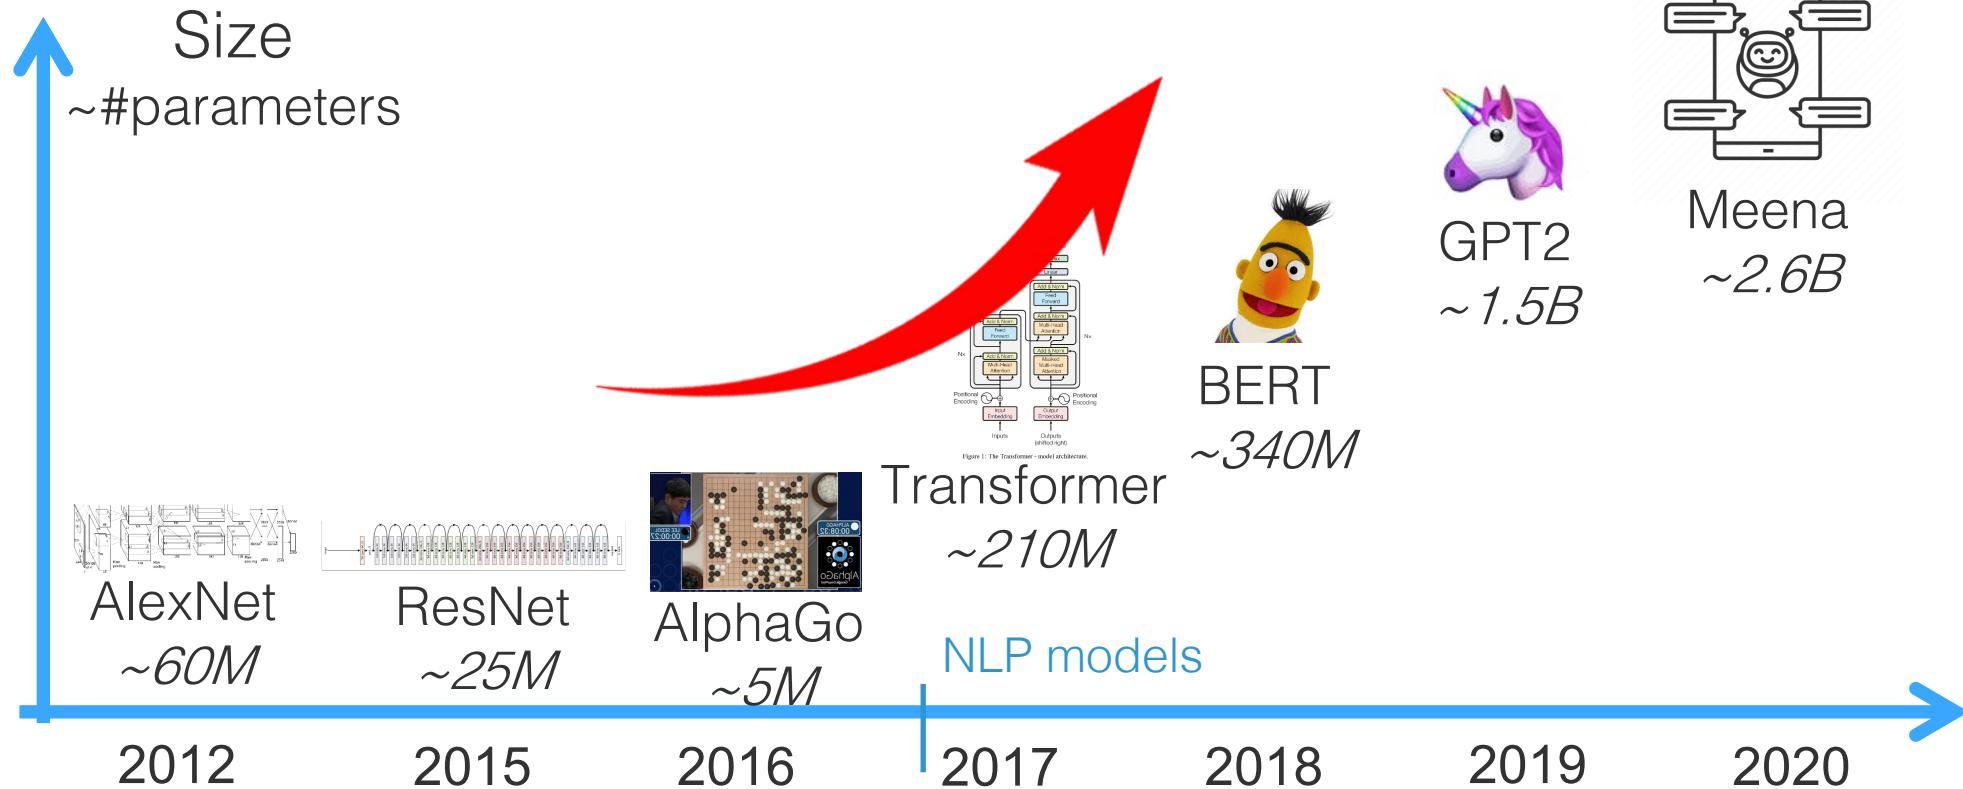
Model/Algorithm

System/Infra

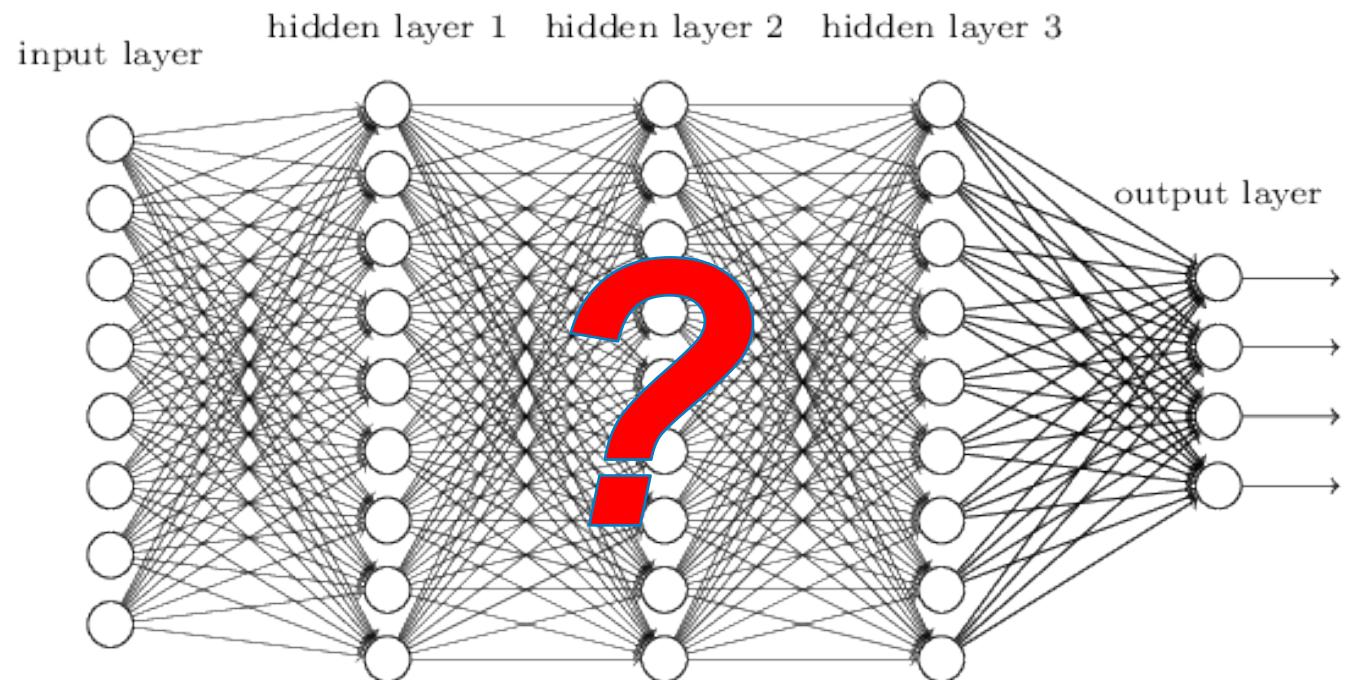
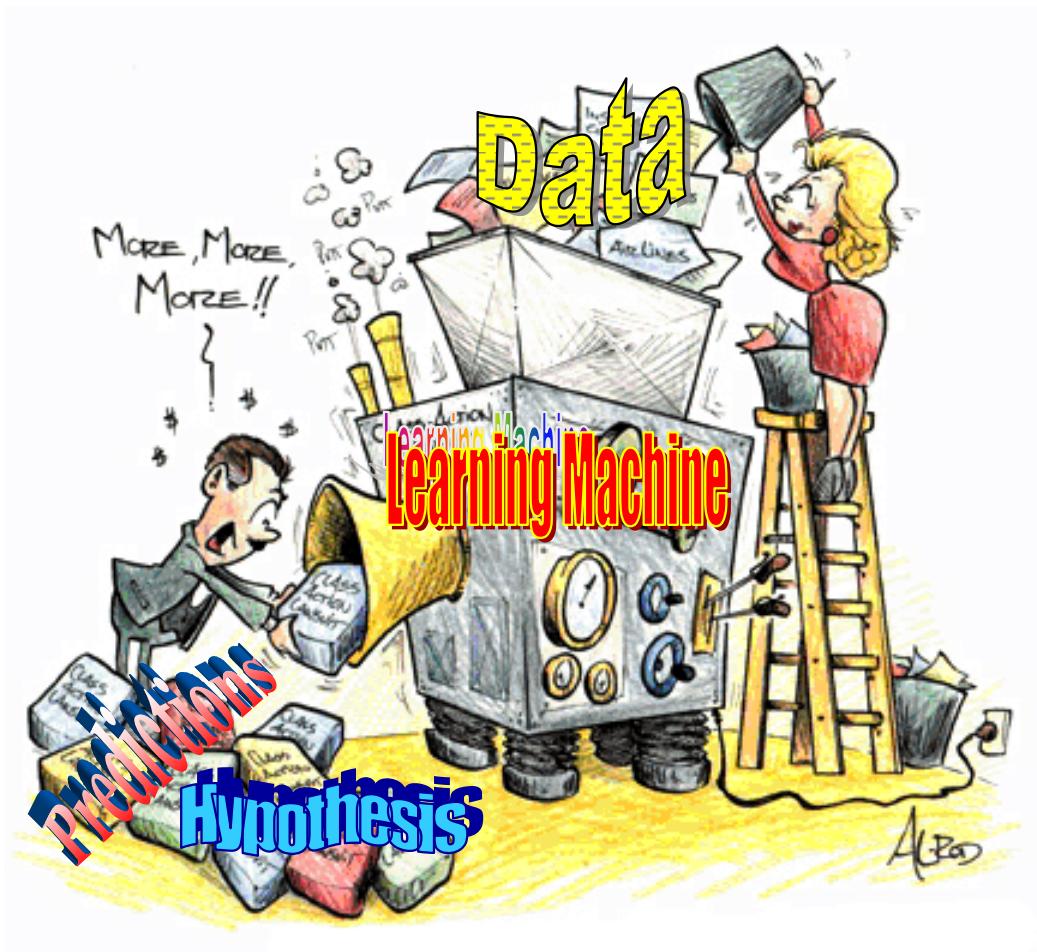


Single Models with Increasing Size and Performance

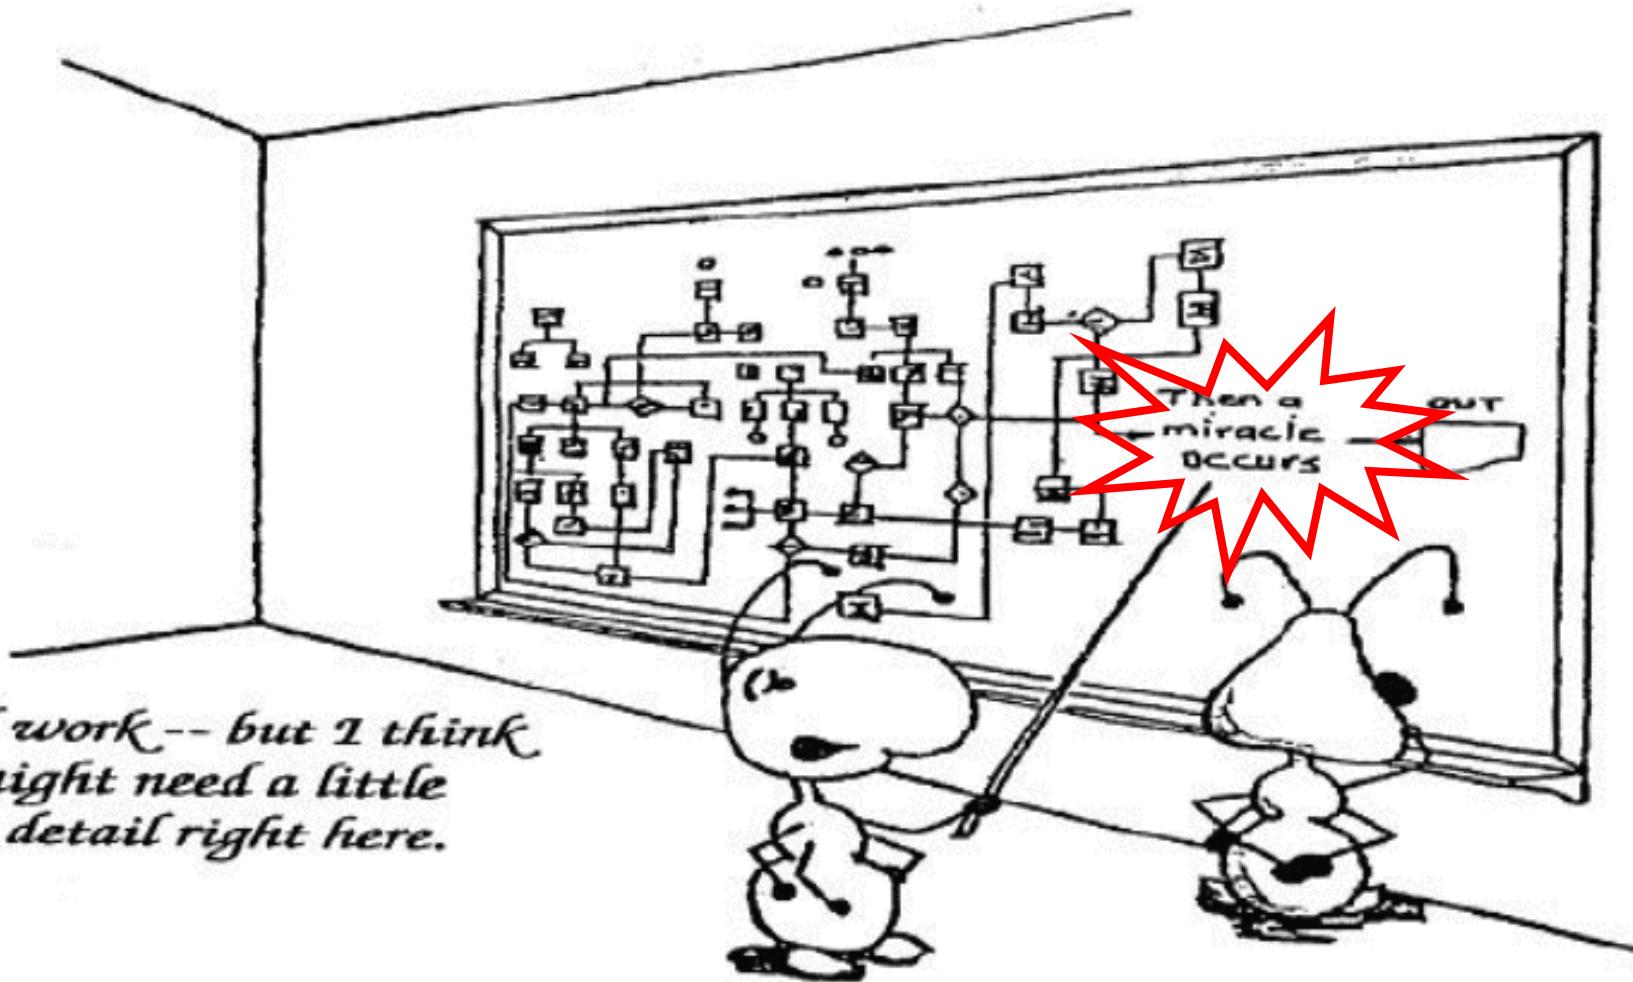
- Increasingly large black-box neural networks
- Good, even super-human performance on some tasks



Single Giant Models Enough?

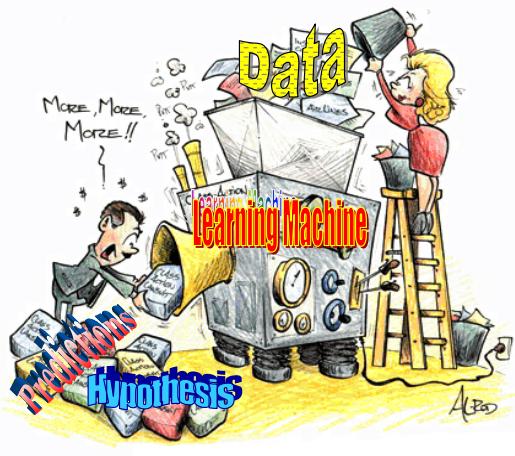
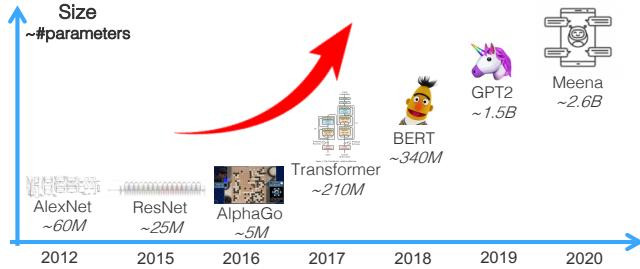


Difficulties of Single Giant Models

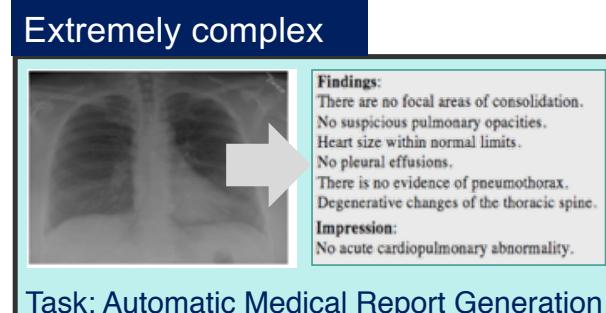


- Explainability
- Debugging
- Maintenance
- Upgrade
- Scalability
- ...

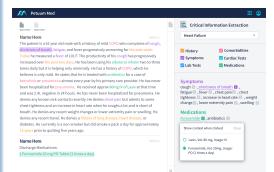
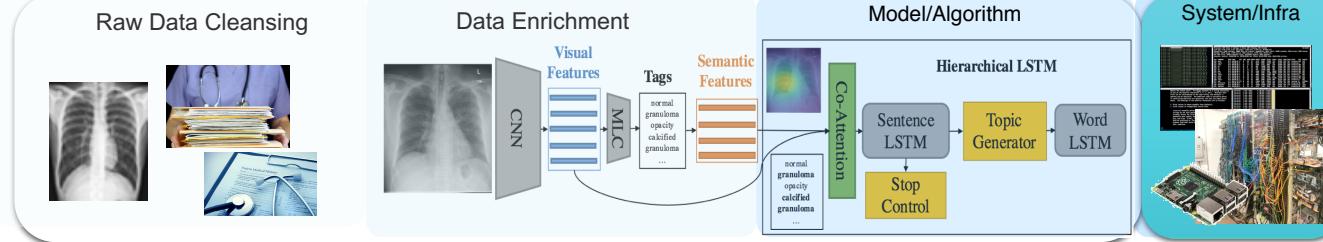
Far from Solving Real Complex Problems



V.S.



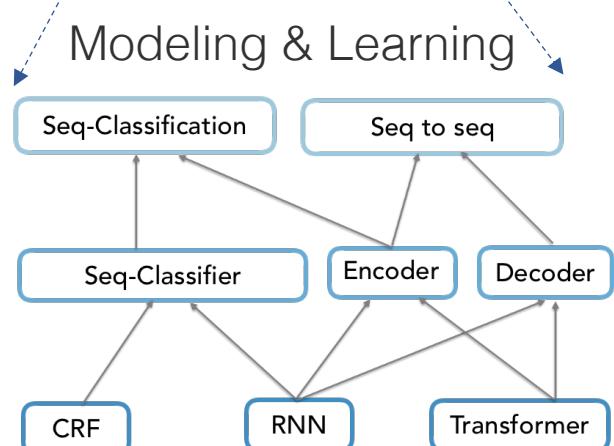
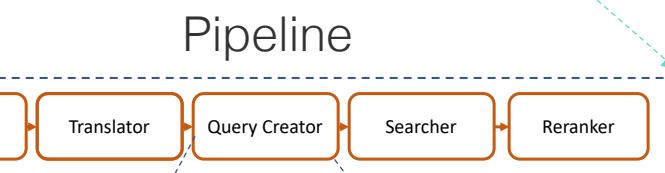
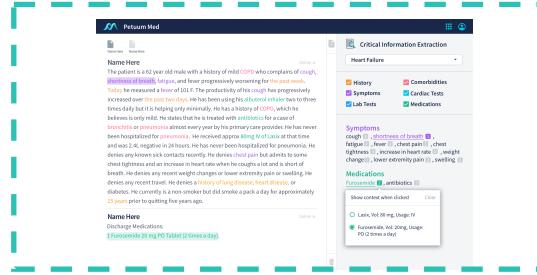
Requires inter-operation between diverse components



Modularization & Standardization



Application



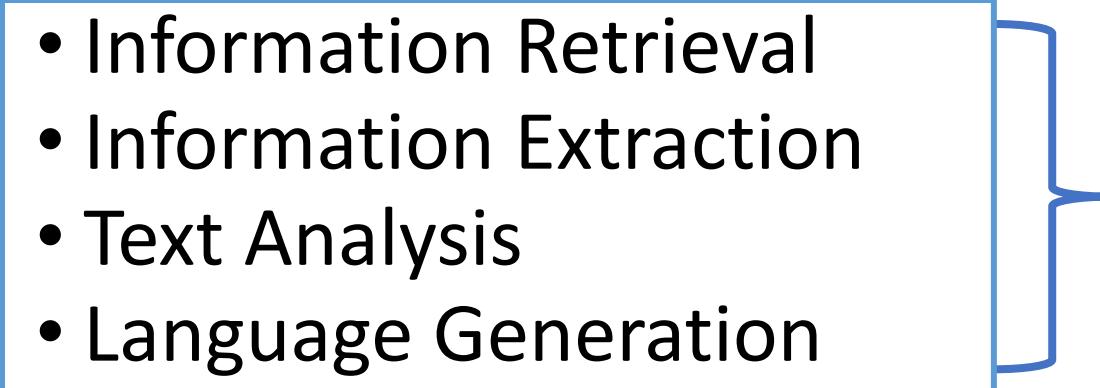
Agenda

- Natural Language Processing Overview (10mins)
- Modularizing NLP Pipeline (40mins)
 - Complexity of NLP pipeline
 - A standardized view of NLP pipeline
 - A standardized implementation of NLP pipeline
- Modularizing NLP Model & Learning (30mins)
 - Composable ML
- QA (10mins)

Agenda

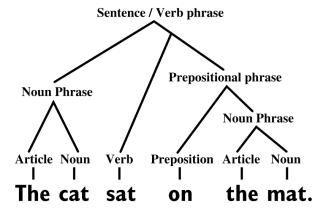
- Natural Language Processing Overview (10mins)
- Modularizing NLP Pipeline (40mins)
 - Complexity of NLP pipeline
 - A standardized view of NLP pipeline
 - A standardized implementation of NLP pipeline
- Modularizing NLP Model & Learning (30mins)
 - Composable ML
- QA (10mins)

What's in NLP?

- Informally, NLP is about interacting with or processing human languages with computers
 - Broadly, NLP contains/relates to many fields
 - Speech Processing
 - Information Retrieval
 - Information Extraction
 - Text Analysis
 - Language Generation
 - Speech Synthesize
- 
- Text Processing**

We'll talk about these today.

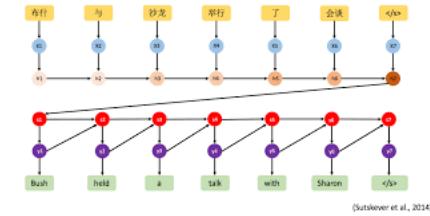
Categorization of Text Processing



Language Understanding



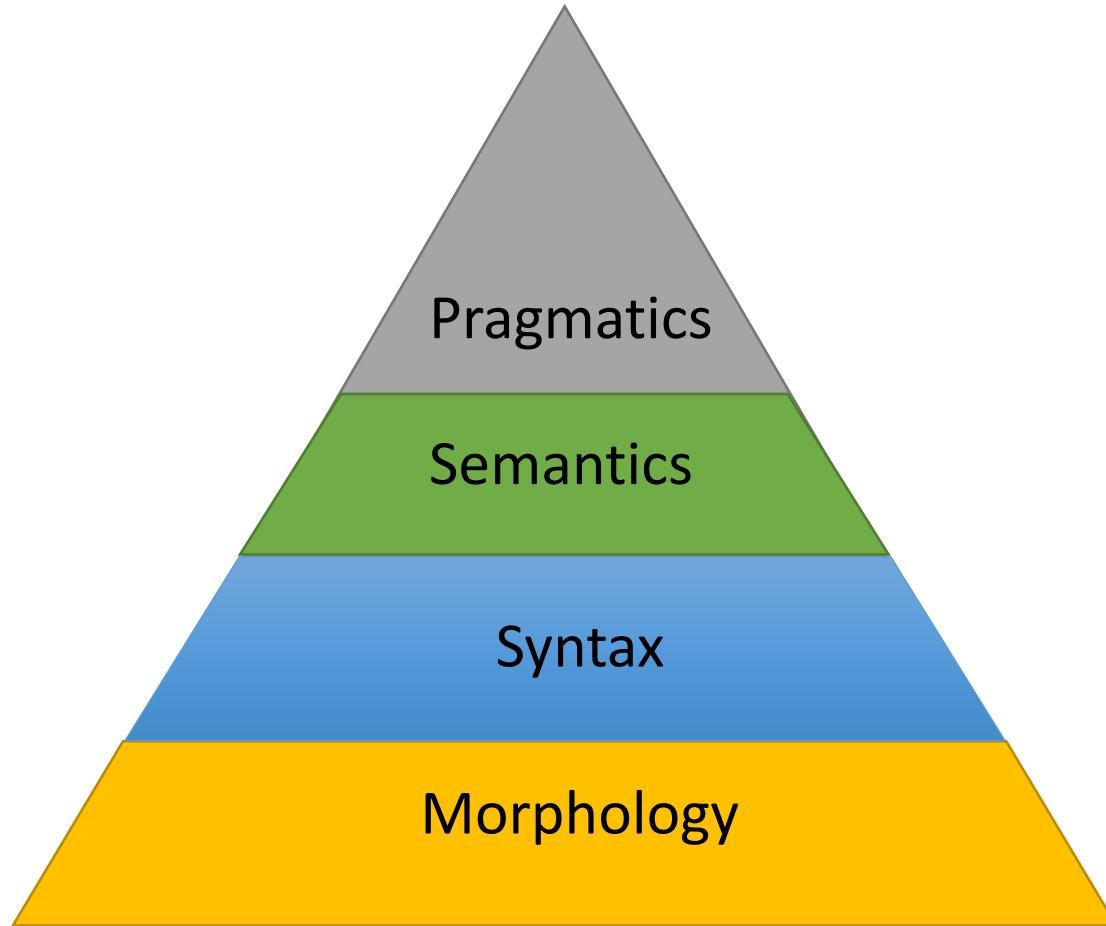
Text Retrieval



Language Generation

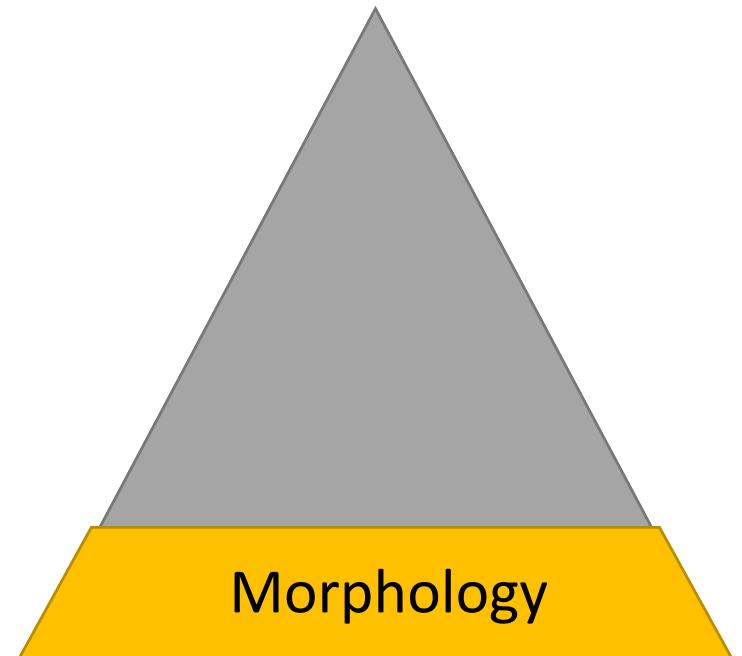
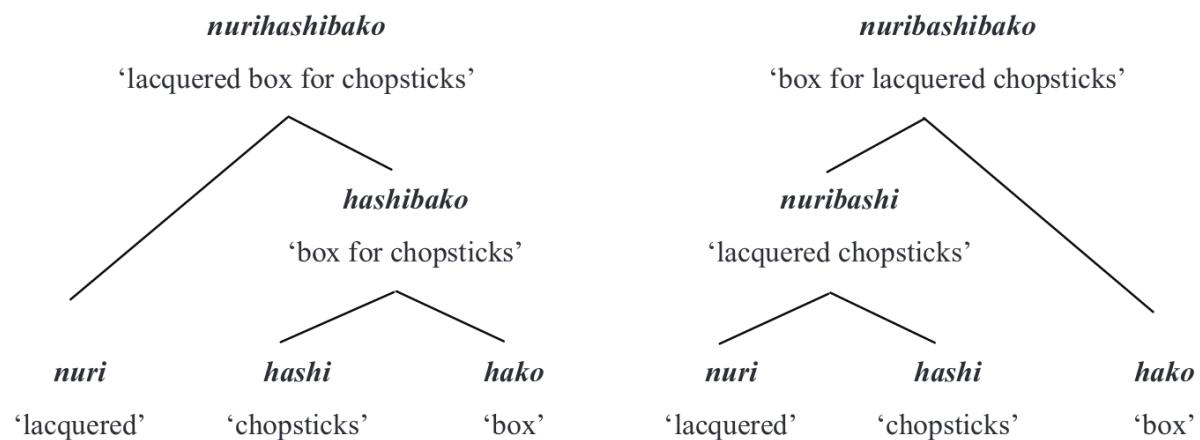
Many different levels of tasks

Language Understanding Pyramid



Morphology

Morphology Analysis



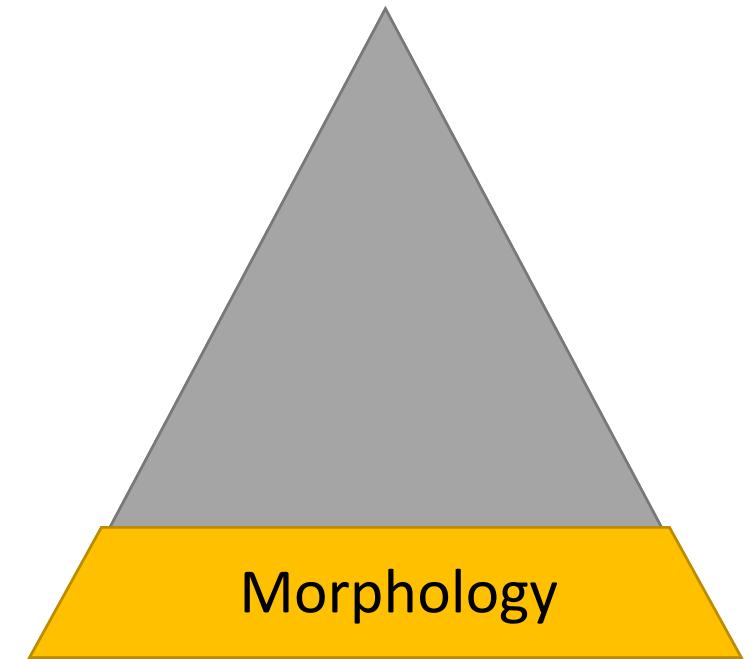
Morphology

Stemming

adjustable → adjustable
formality → formality
formaliti → formaliti
airliner → airliner

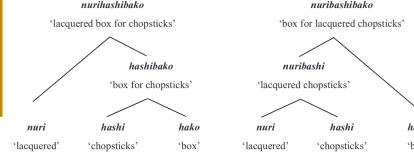
Lemmatize

was → (to) be
better → good
meeting → meeting



Language Understanding Pyramid

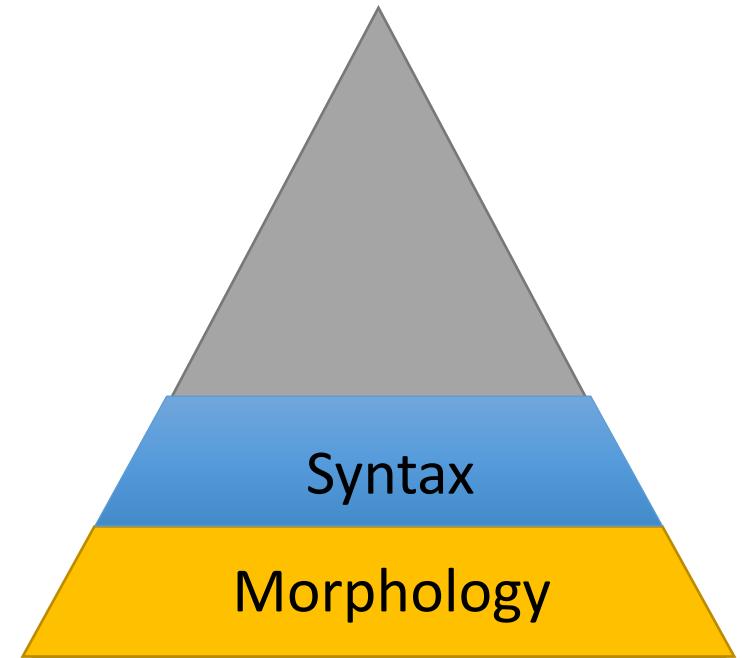
Morphology



Lemmatize

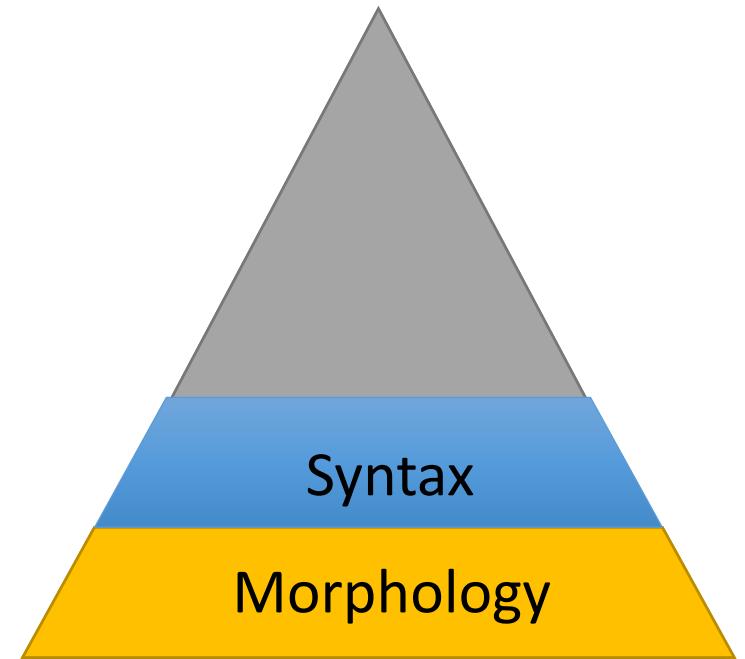
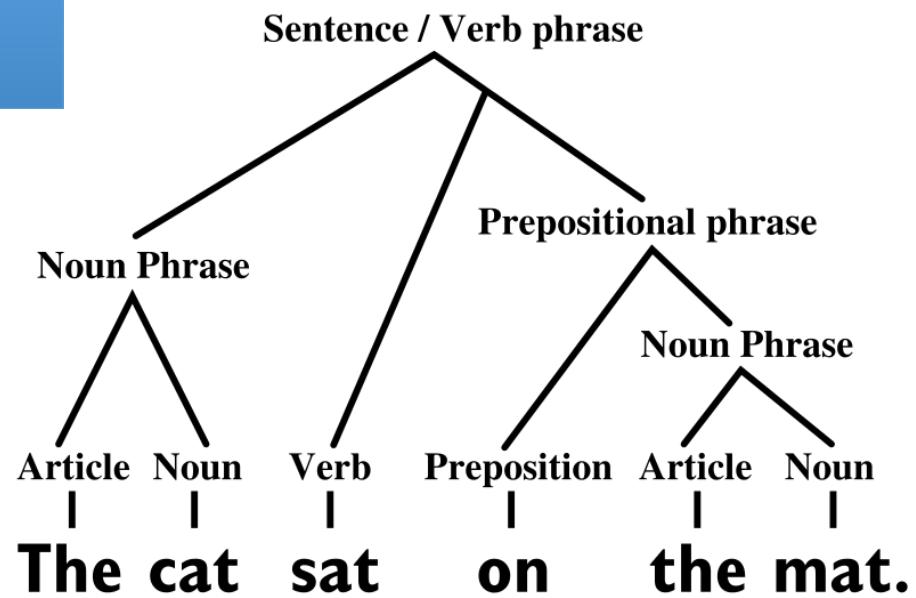
ning
→ adjust
formaliti
→ formal
· airlin △

Lemmatization
was → (to) be
better → good
meeting → meeting



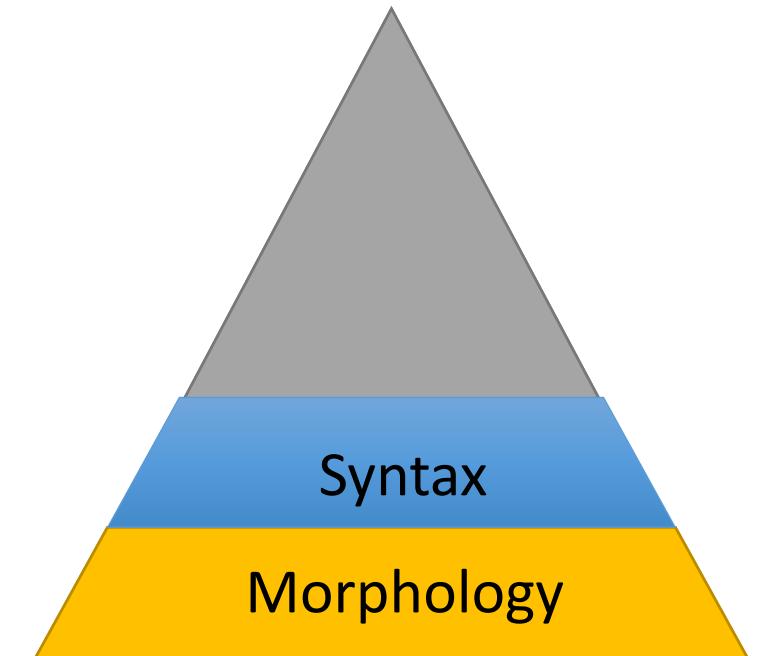
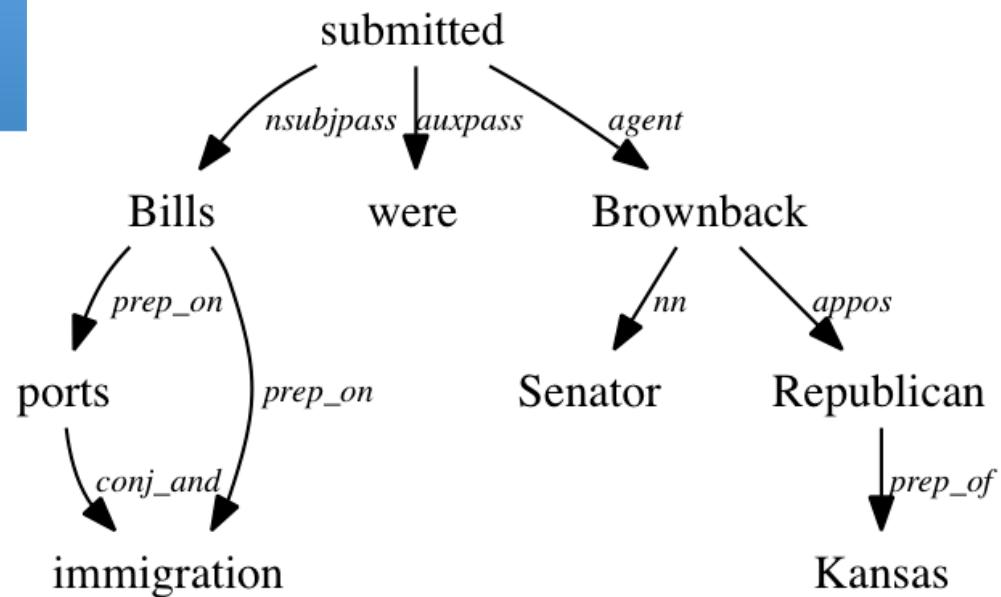
Syntax

Constituent Parse

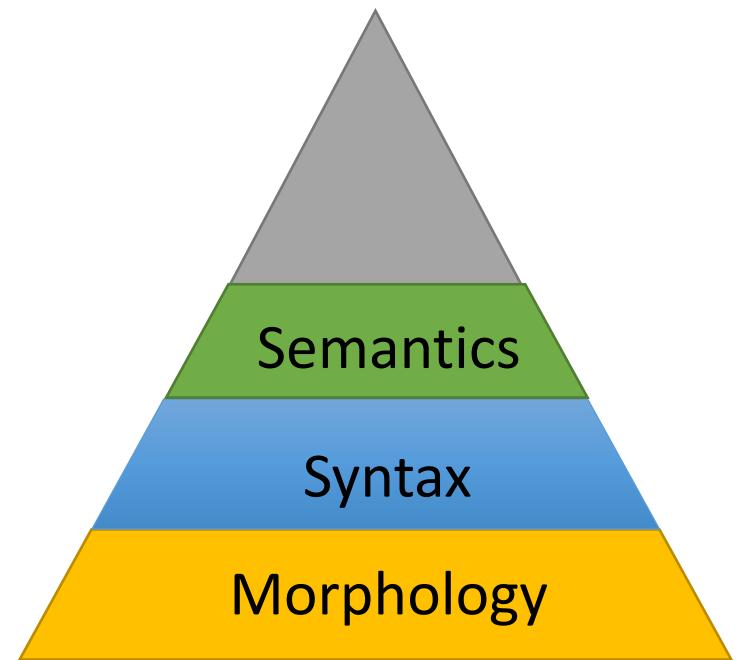


Syntax

Dependency Parse

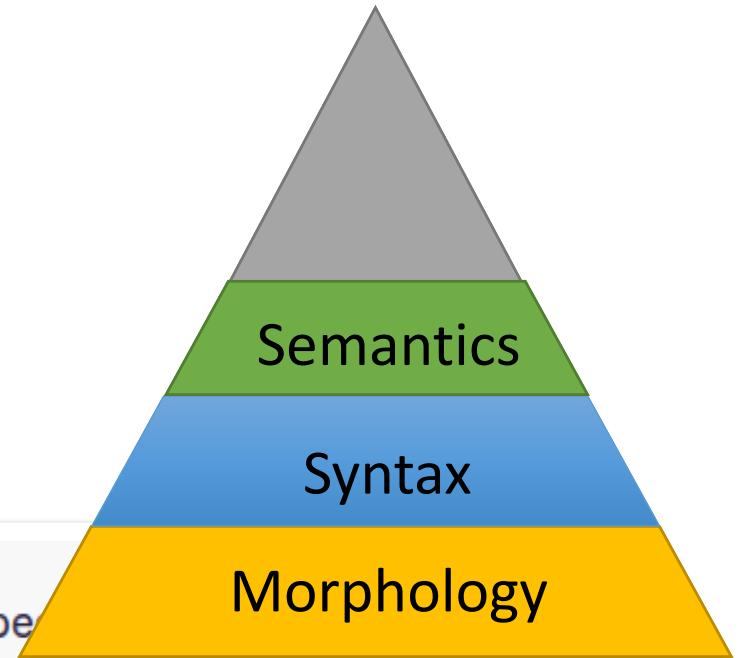


Language Understanding Pyramid



Semantic

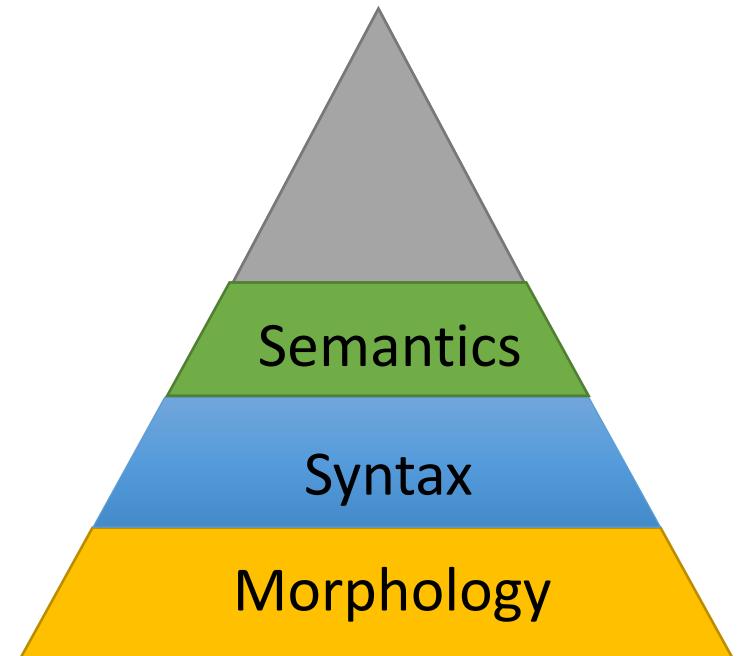
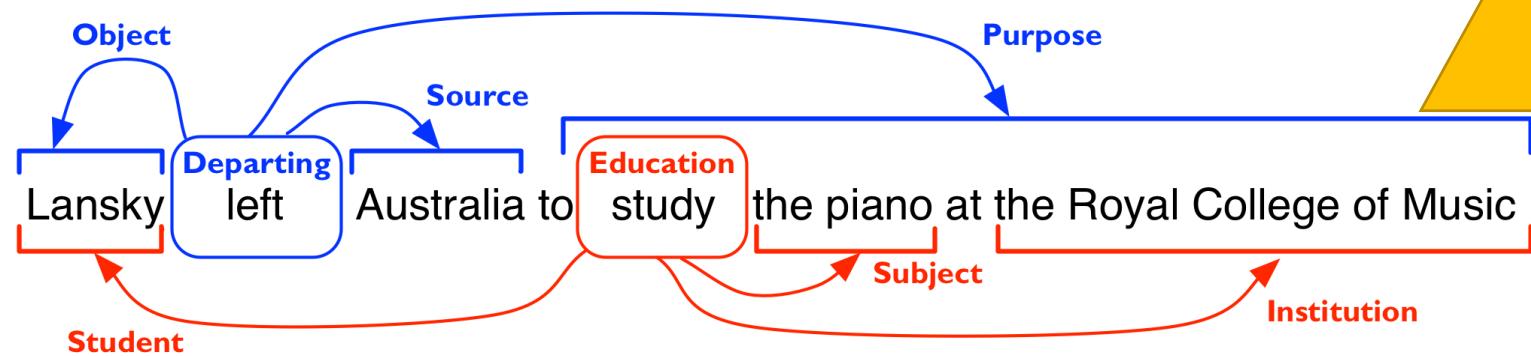
Named Entity Recognition



When Sebastian Thrun PERSON started at Google ORG in 2007 DATE, few people Morphology took him seriously. “I can tell you very senior CEOs of major American NORP car companies would shake my hand and turn away because I wasn’t worth talking to,” said Thrun PERSON, now the co-founder and CEO of online higher education startup Udacity, in an interview with Recode ORG earlier this week DATE.

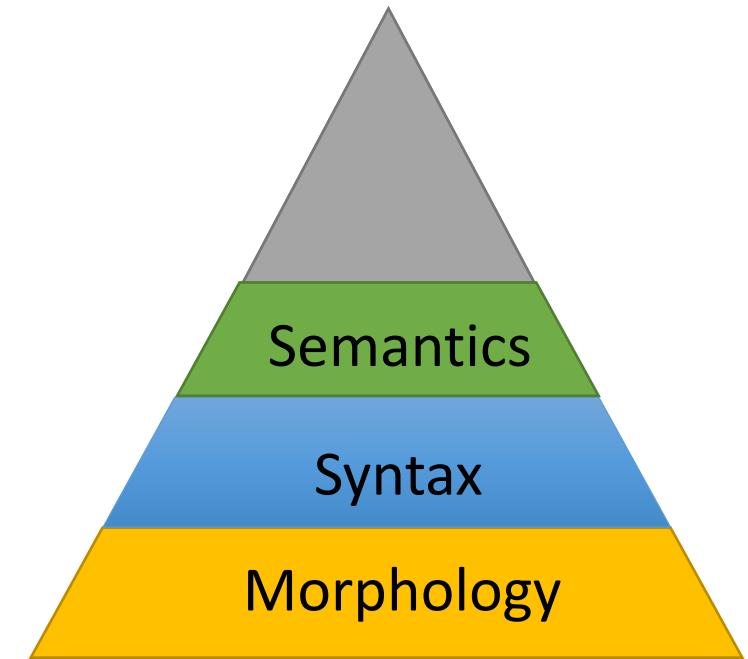
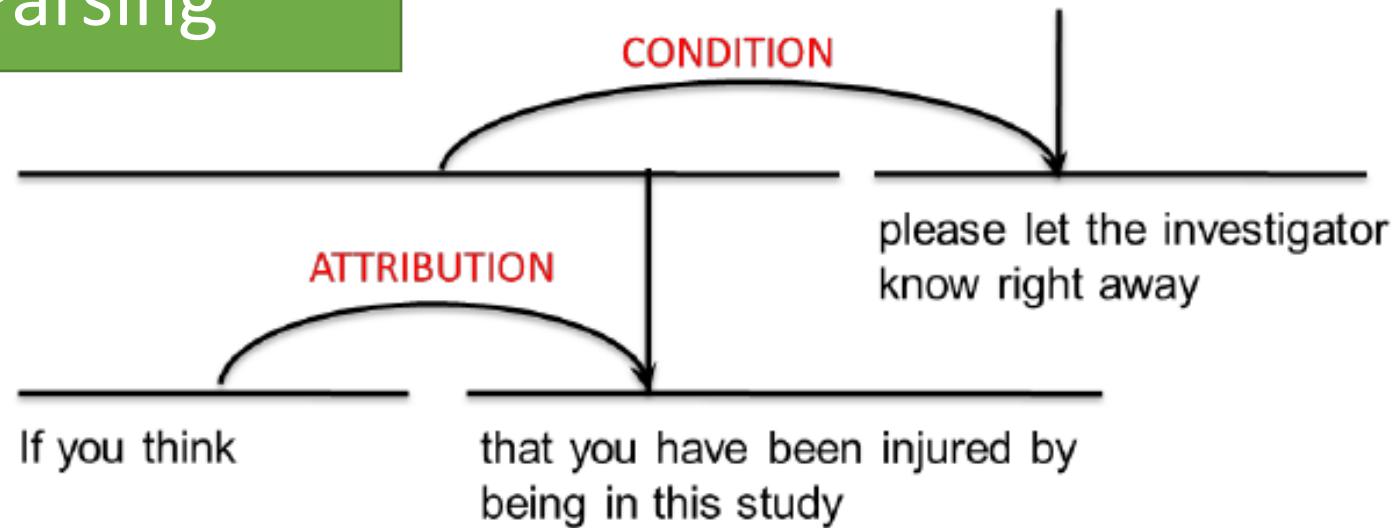
Semantic

Semantic Roles



Semantic

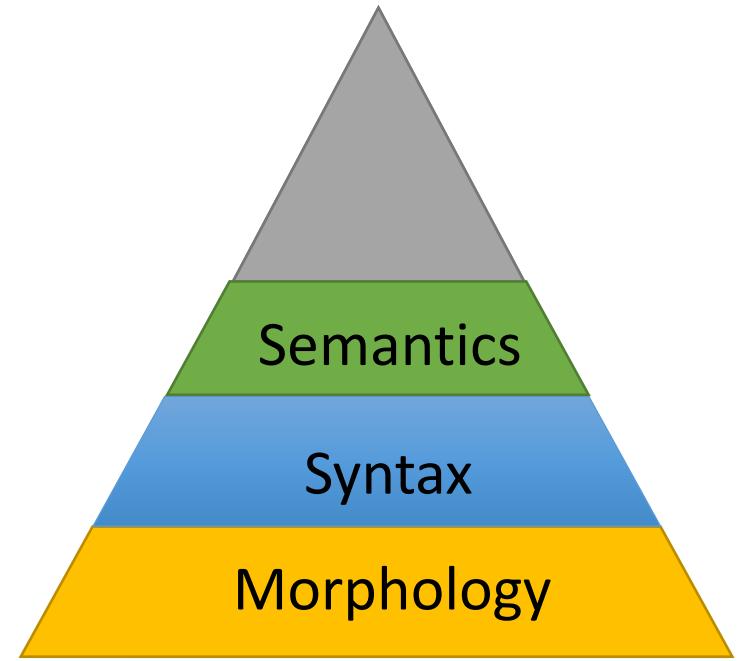
Discourse Parsing



Semantic



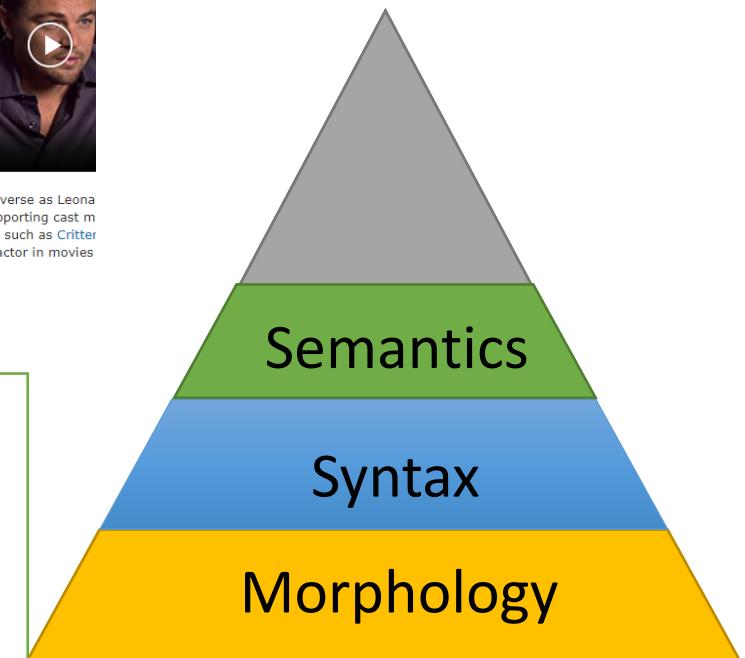
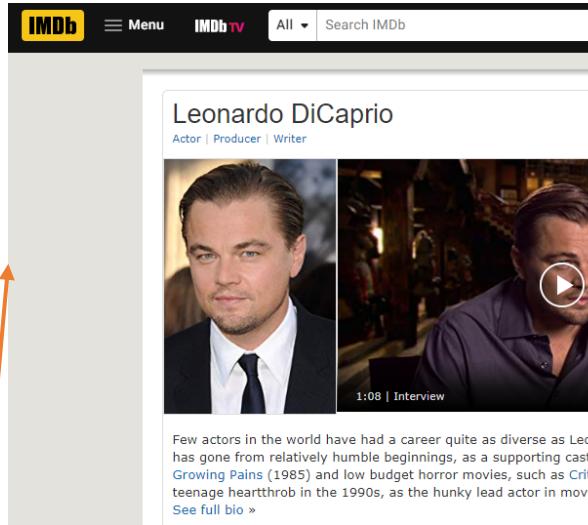
“I voted for Nader because he was most aligned with my values,” she said.



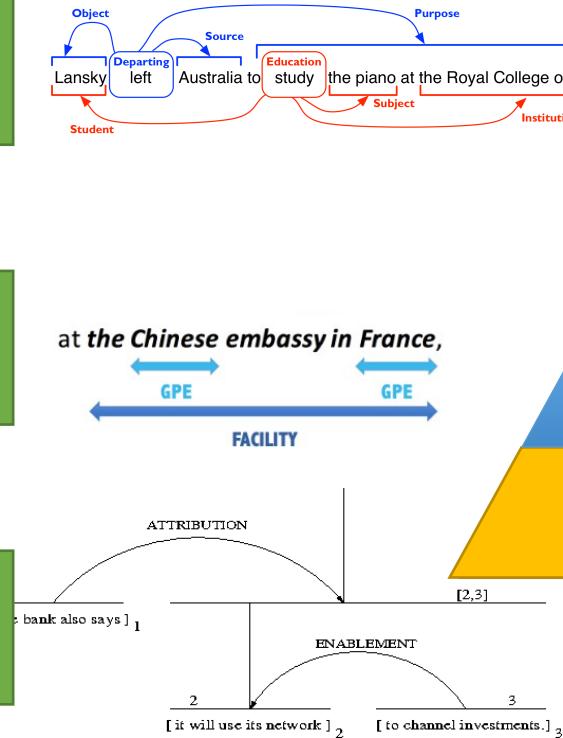
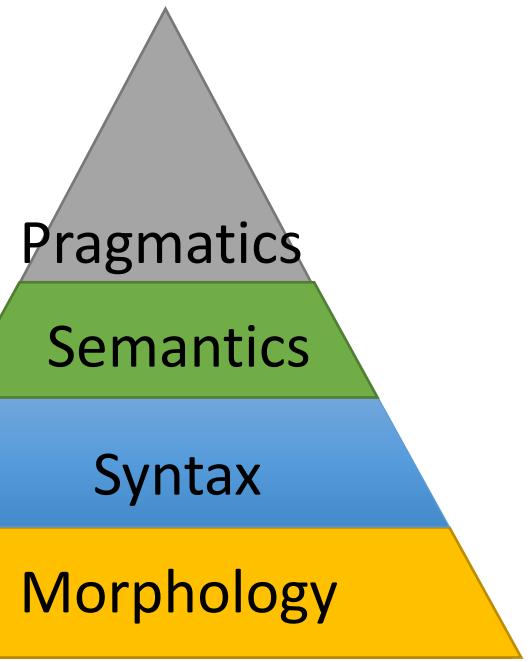
Semantic

Entity Linking

Kate Winslet and Leonardo DiCaprio
have definitely created a timeless classic.



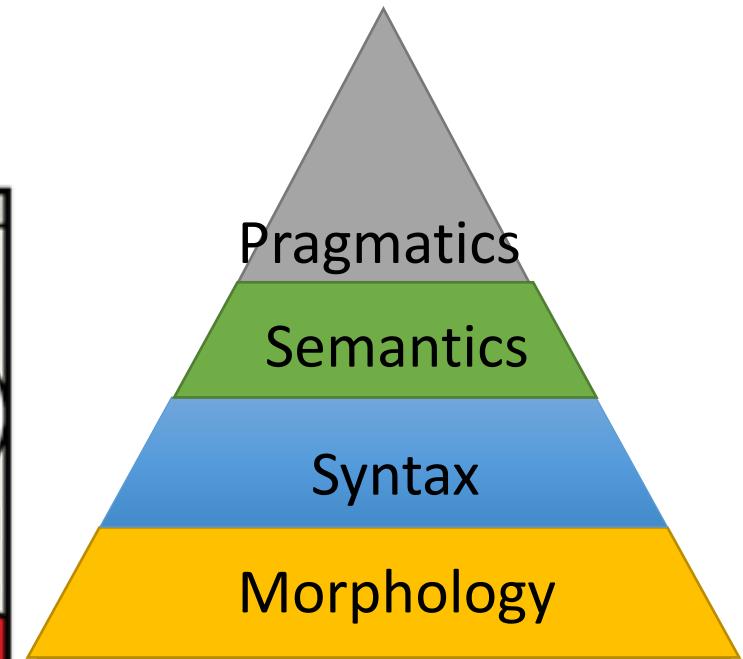
Language Understanding Pyramid



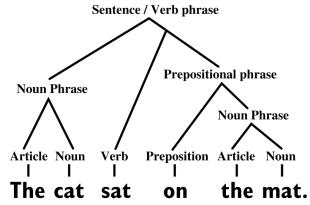
Pragmatics



WWW.PHDCOMICS.COM



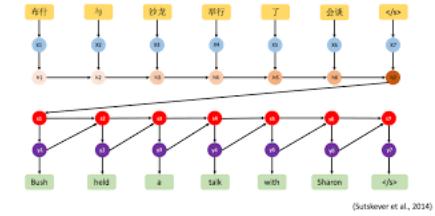
Categorization of Text Processing



Language
Understanding



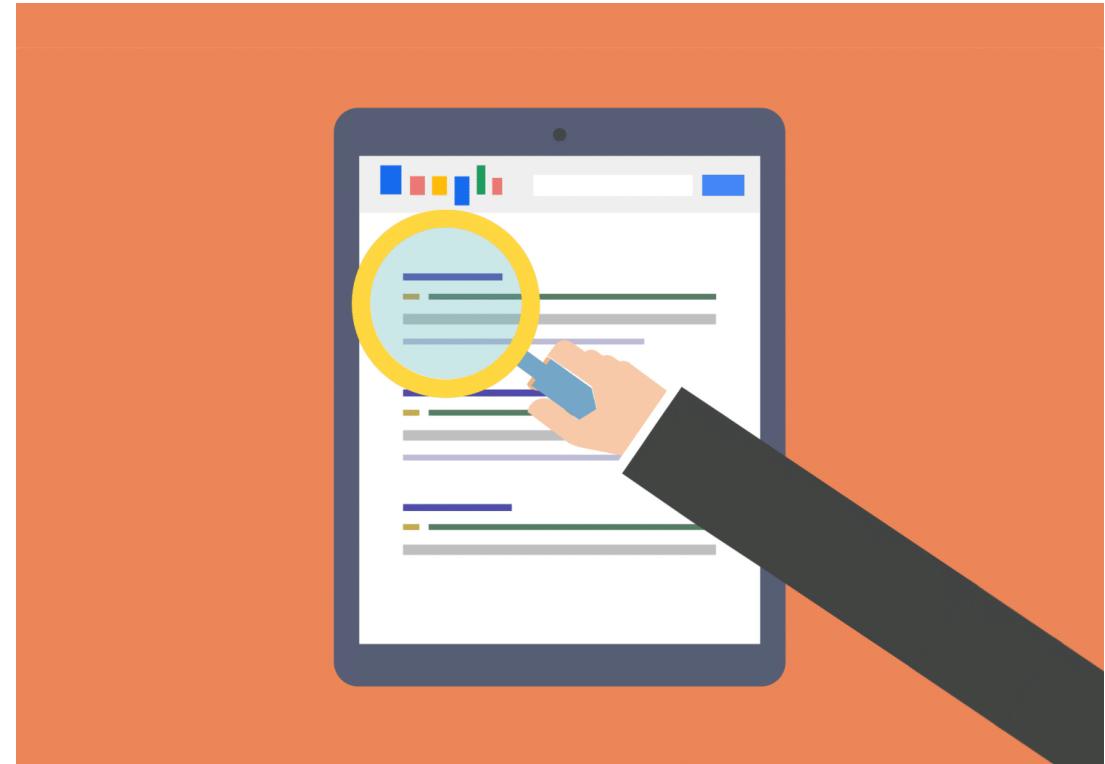
Text Retrieval



Language
Generation

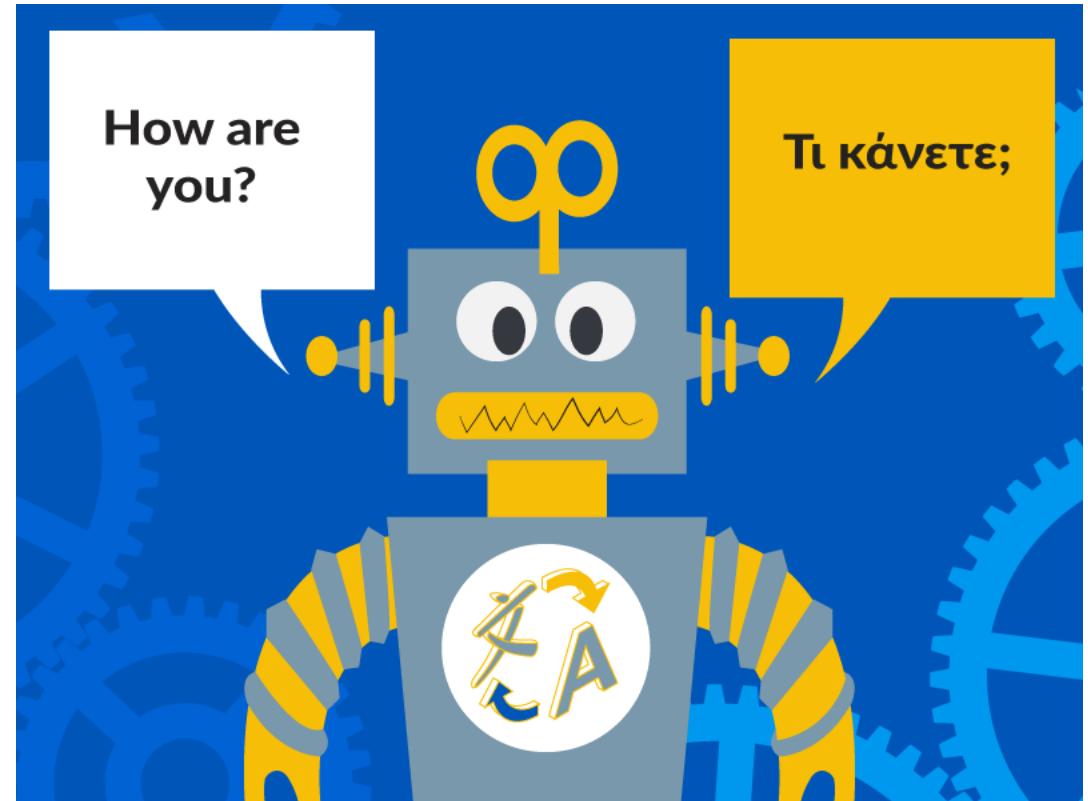
Information Retrieval

- Retrieve relevant documents based on user query
- Some IR subs-tasks:
 - The Search Step: quickly get relevant documents as a rank list based on an efficient Index.
 - The Reranking Step: fine-tune the rank list to create better ranking



Text Generation

- Machine Translation
- Summarization
- Dialogue Response Generation

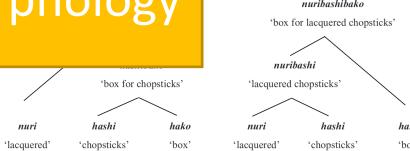


Agenda

- Natural Language Processing Overview (10mins)
- Modularizing NLP Pipeline (40mins)
 - Complexity of NLP pipeline
 - A standardized view of NLP pipeline
 - A standardized implementation of NLP pipeline
- Modularizing NLP Model & Learning (30mins)
 - Composable ML
- QA (10mins)

NLP tasks are Complex

Morphology



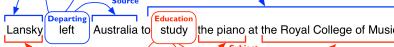
Lemmatize

formaliti → formal
airliner → airlin

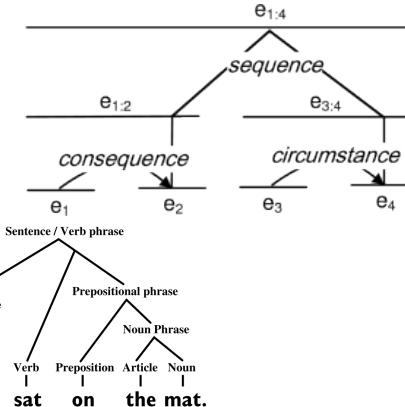
Lemmatization

was → (to) be
better → good
meeting → meeting

Semantic Parsing



Discourse Parsing

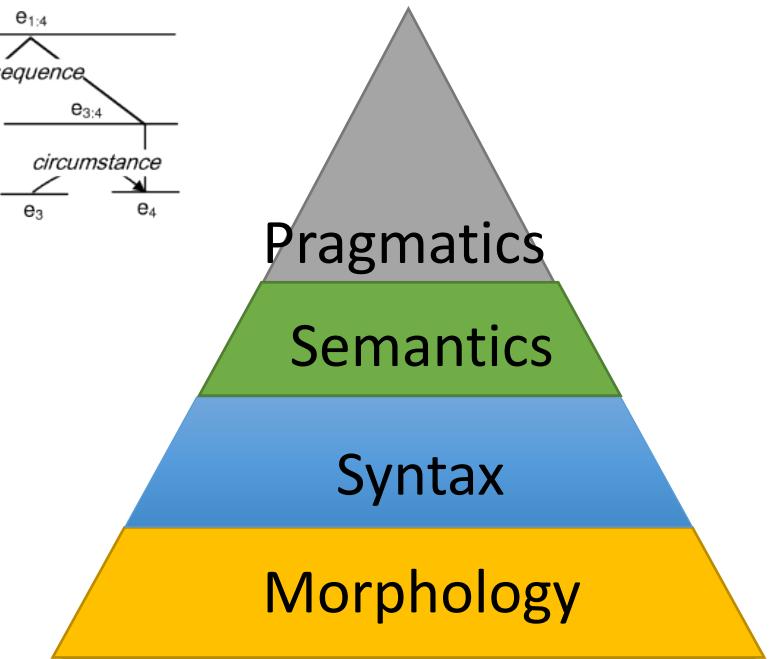


NER

at the Chinese embassy in France,
GPE ↔ GPE
FACILITY

Syntax Parse

Many Different Tasks



Retrieval

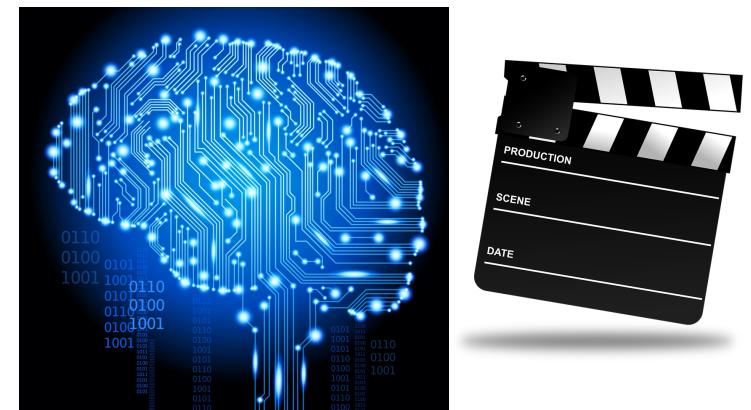
Understand

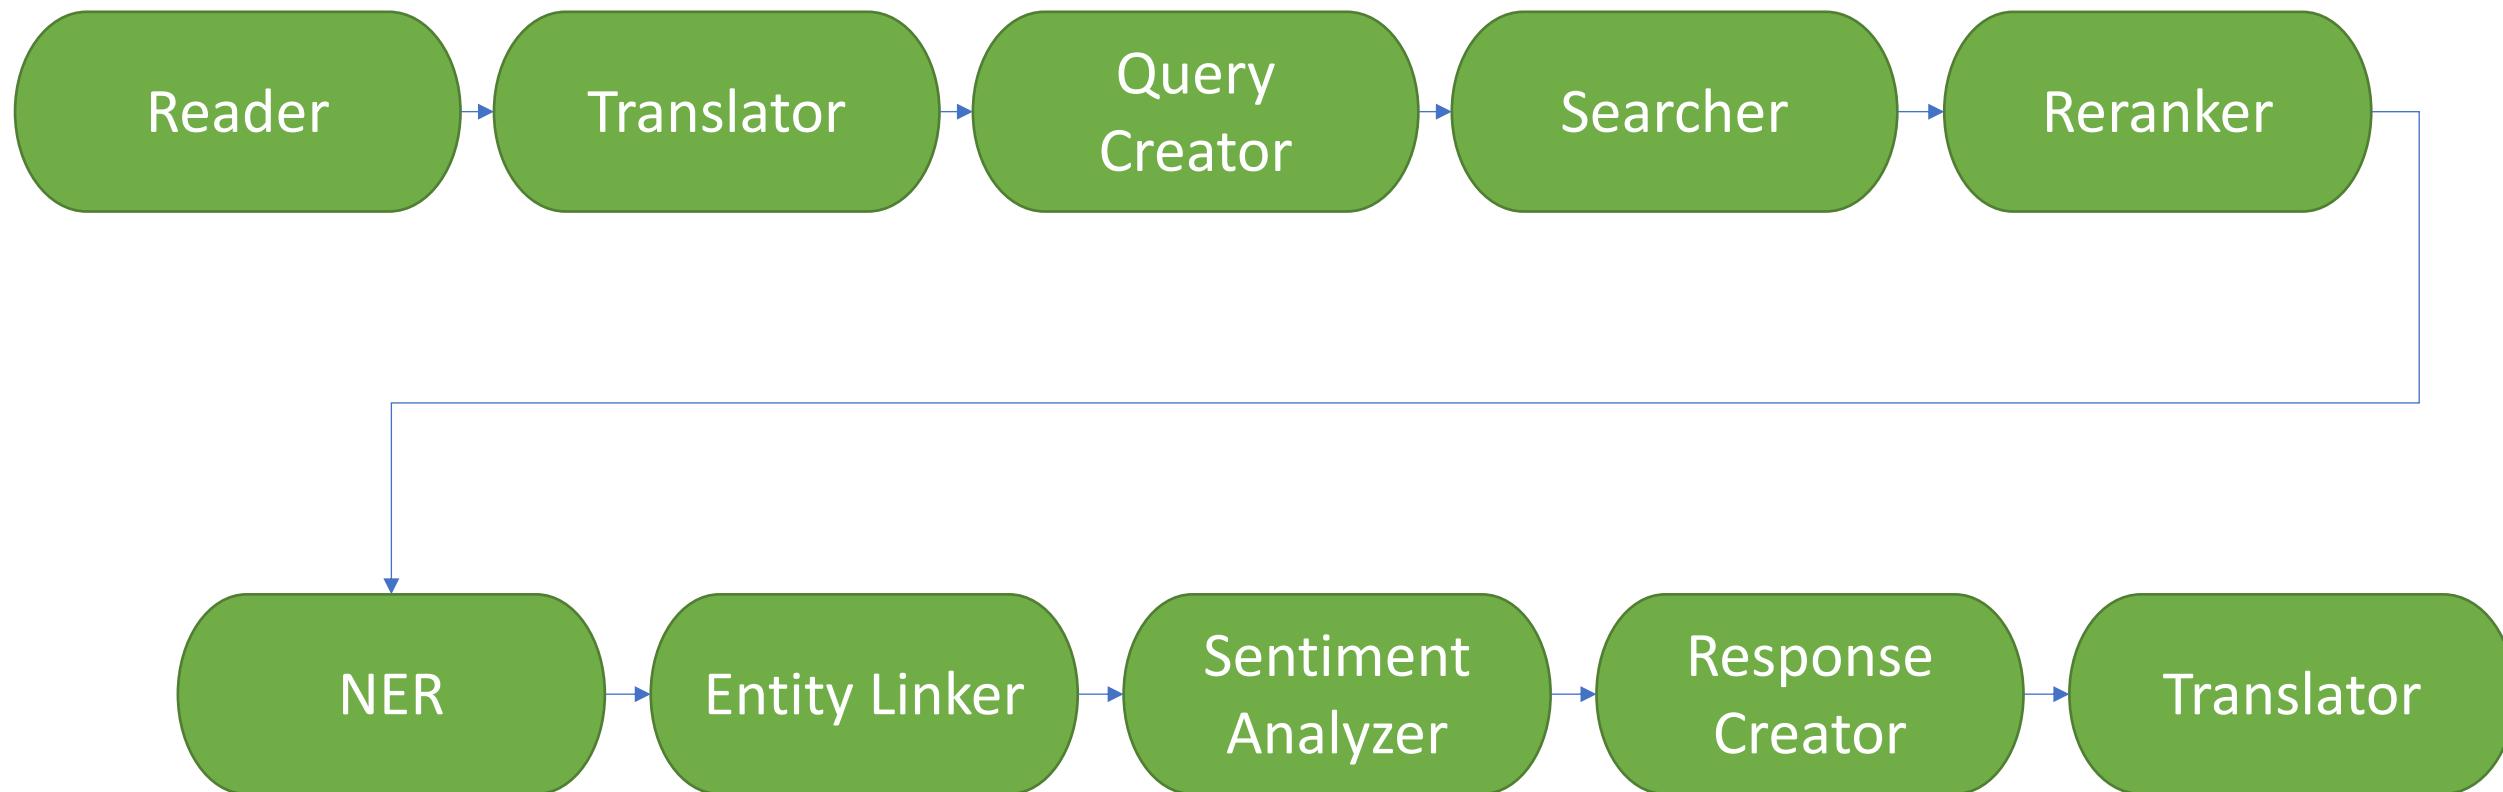
Generation

Different Levels

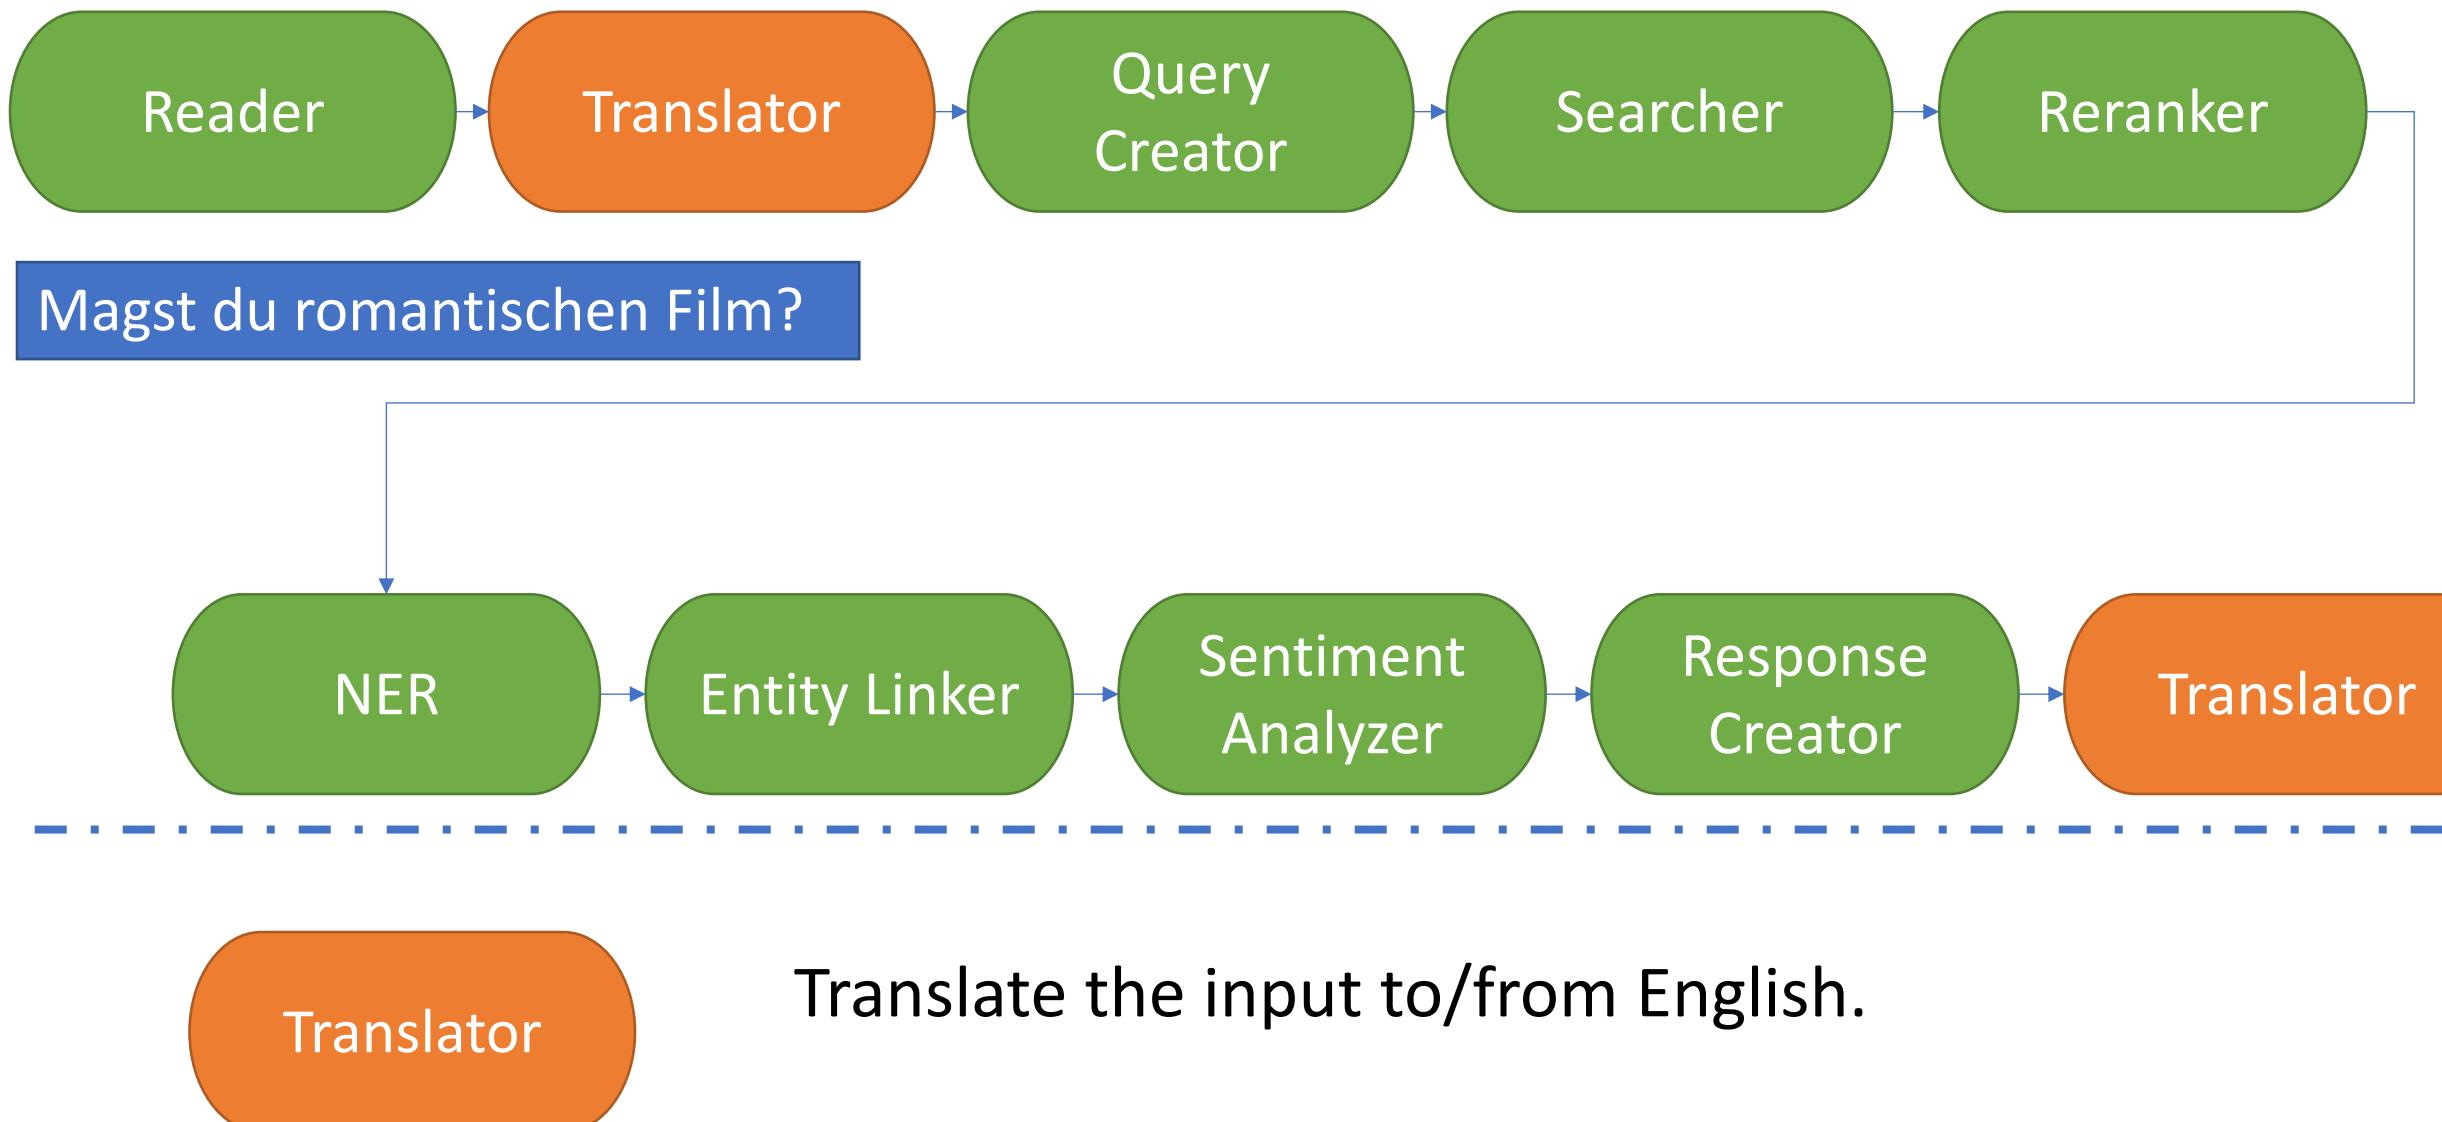
Real World Applications are Complex

- A user **speaks German** but would like to find good **romantic movies**.
- We have a **corpus of English movie reviews**.
- What should we do?



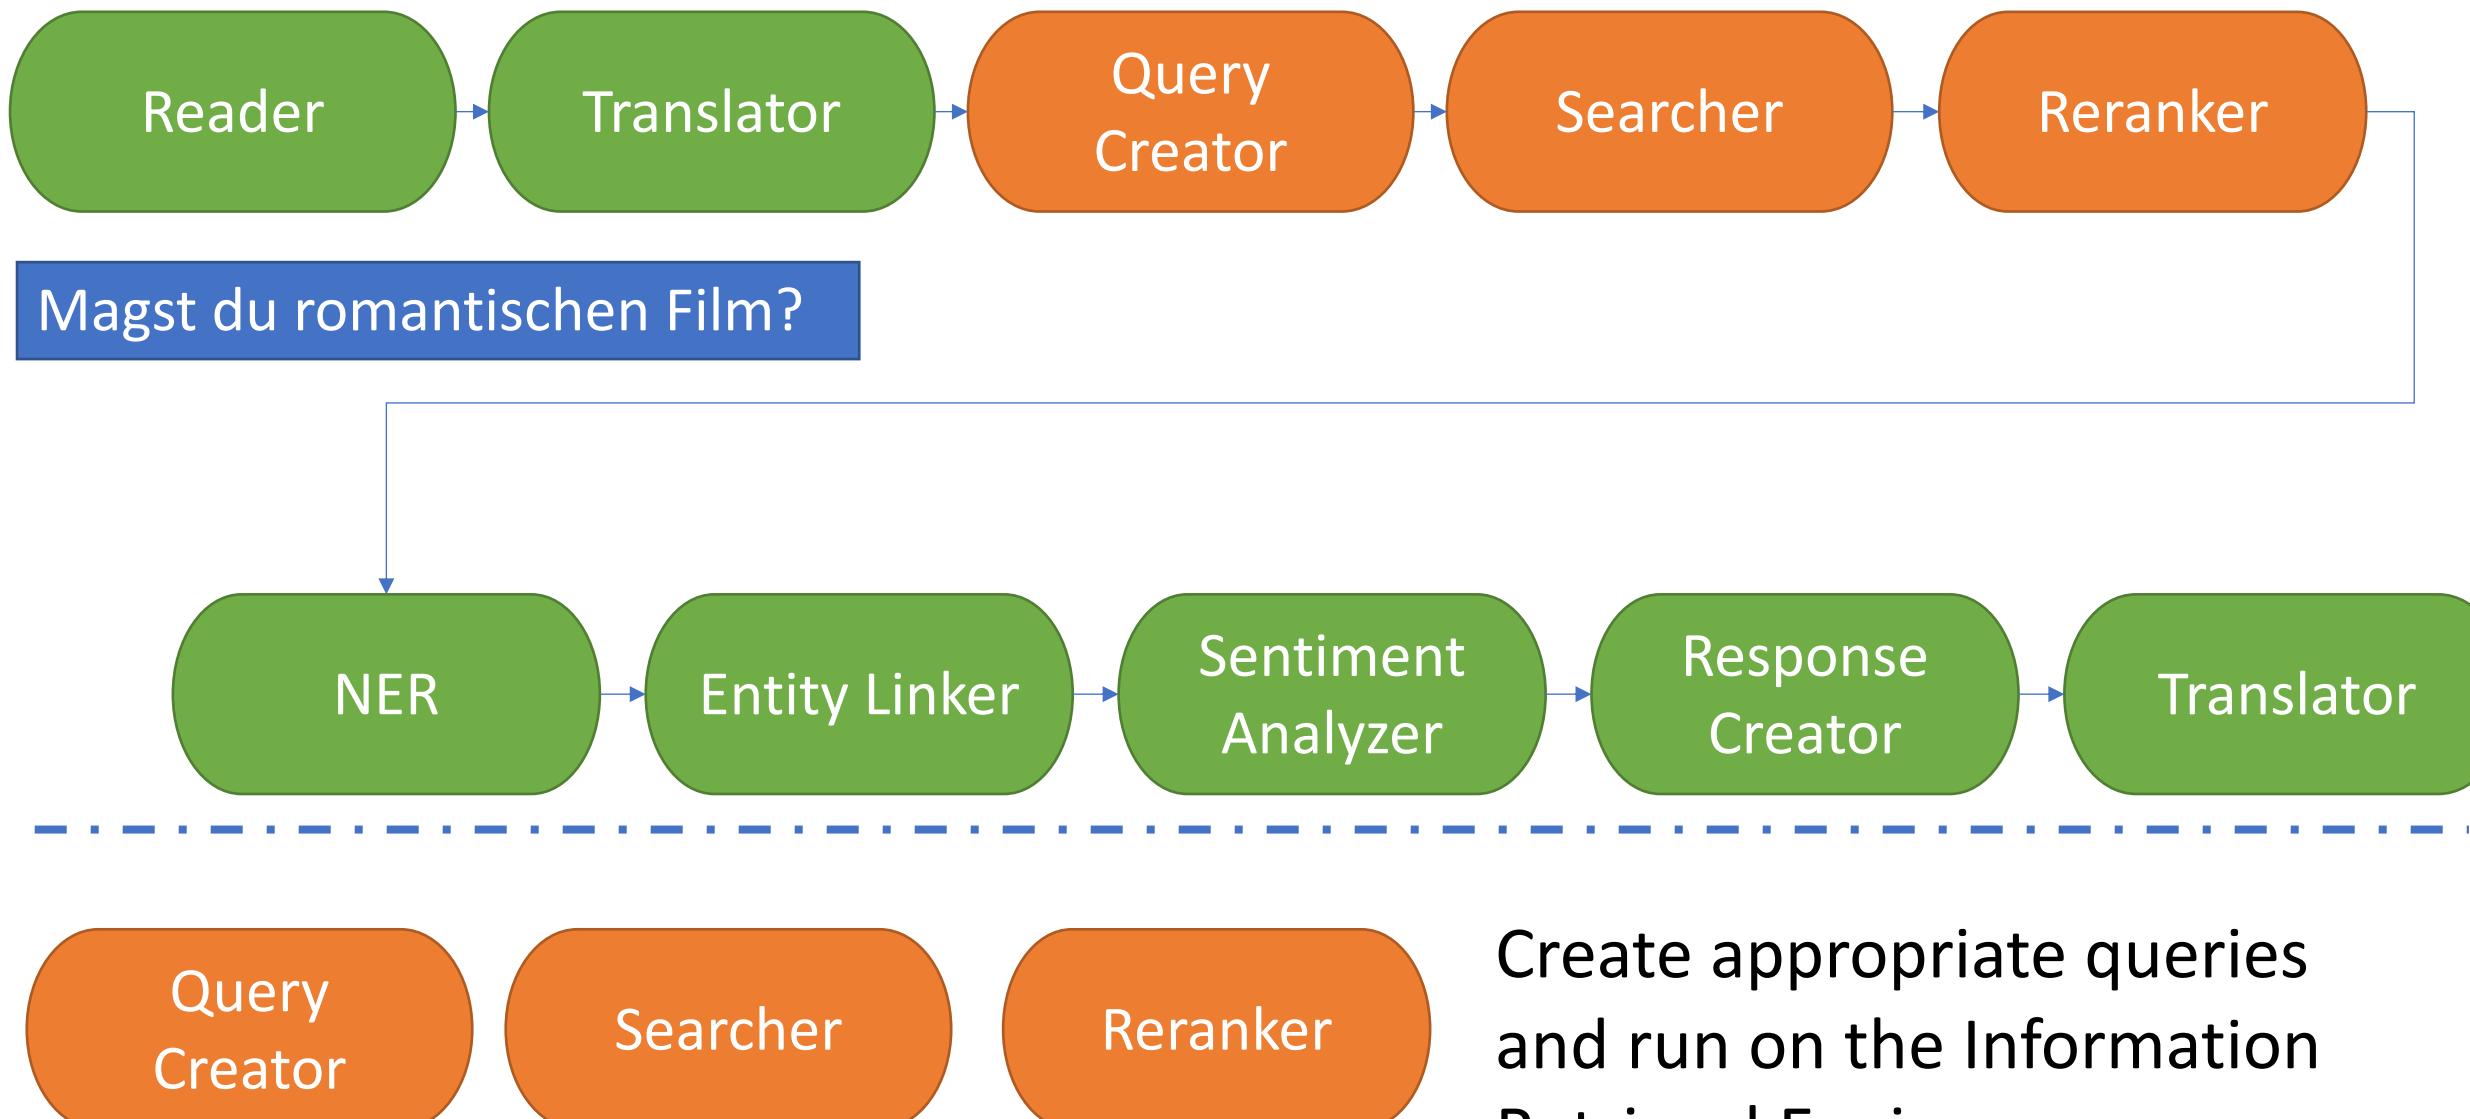


Here is one possible solution pipeline



Translator

Translate the input to/from English.





Magst du romantischen Film?



NER

Kate Winslet and Leonardo Dicaprio have definitely created a timeless classic.

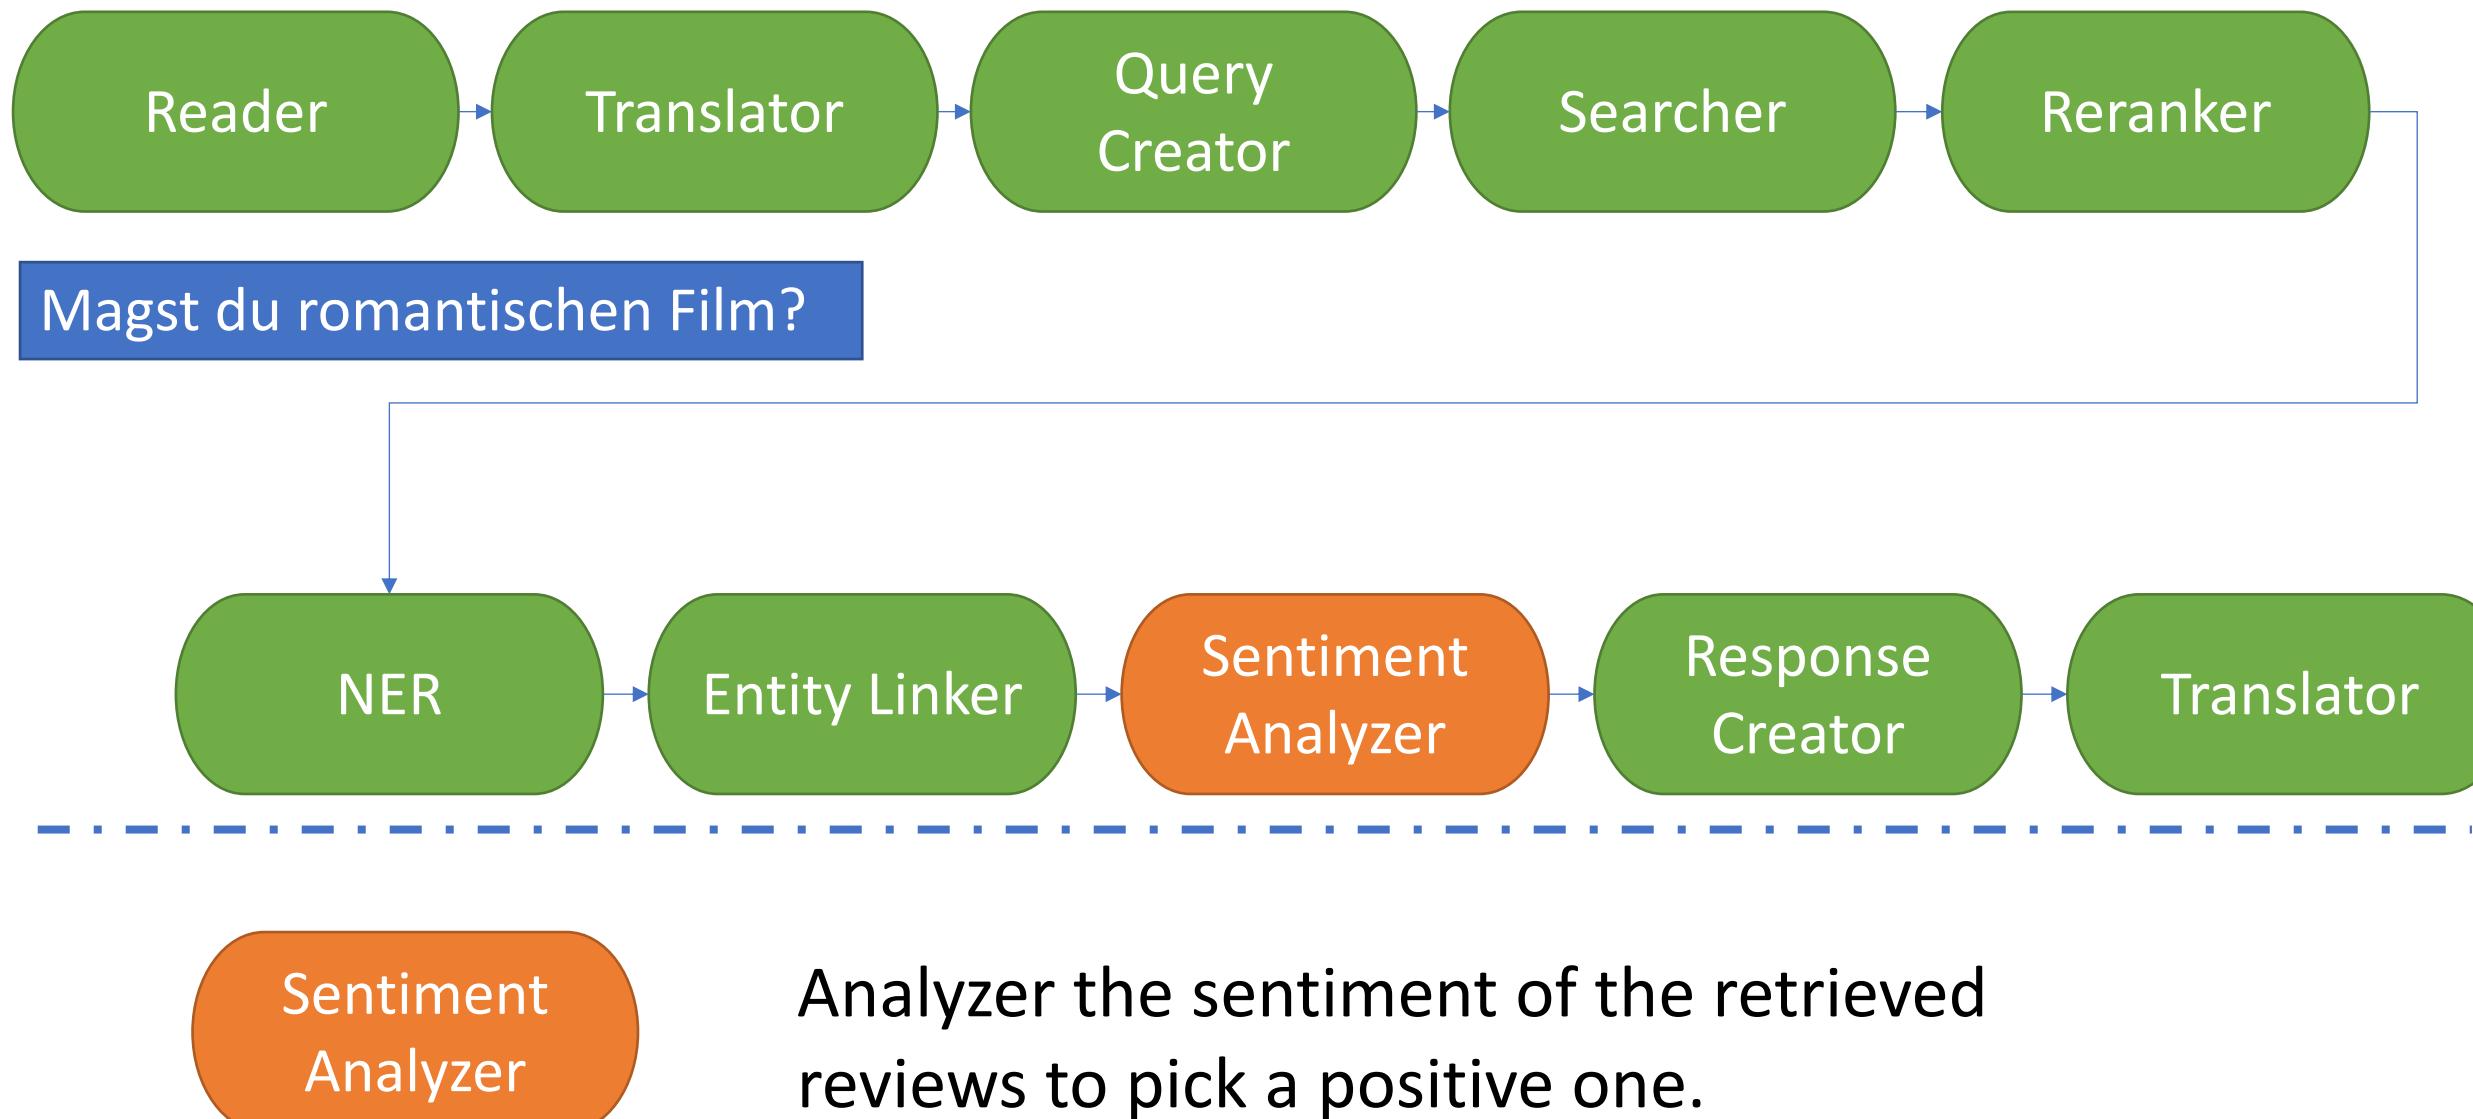


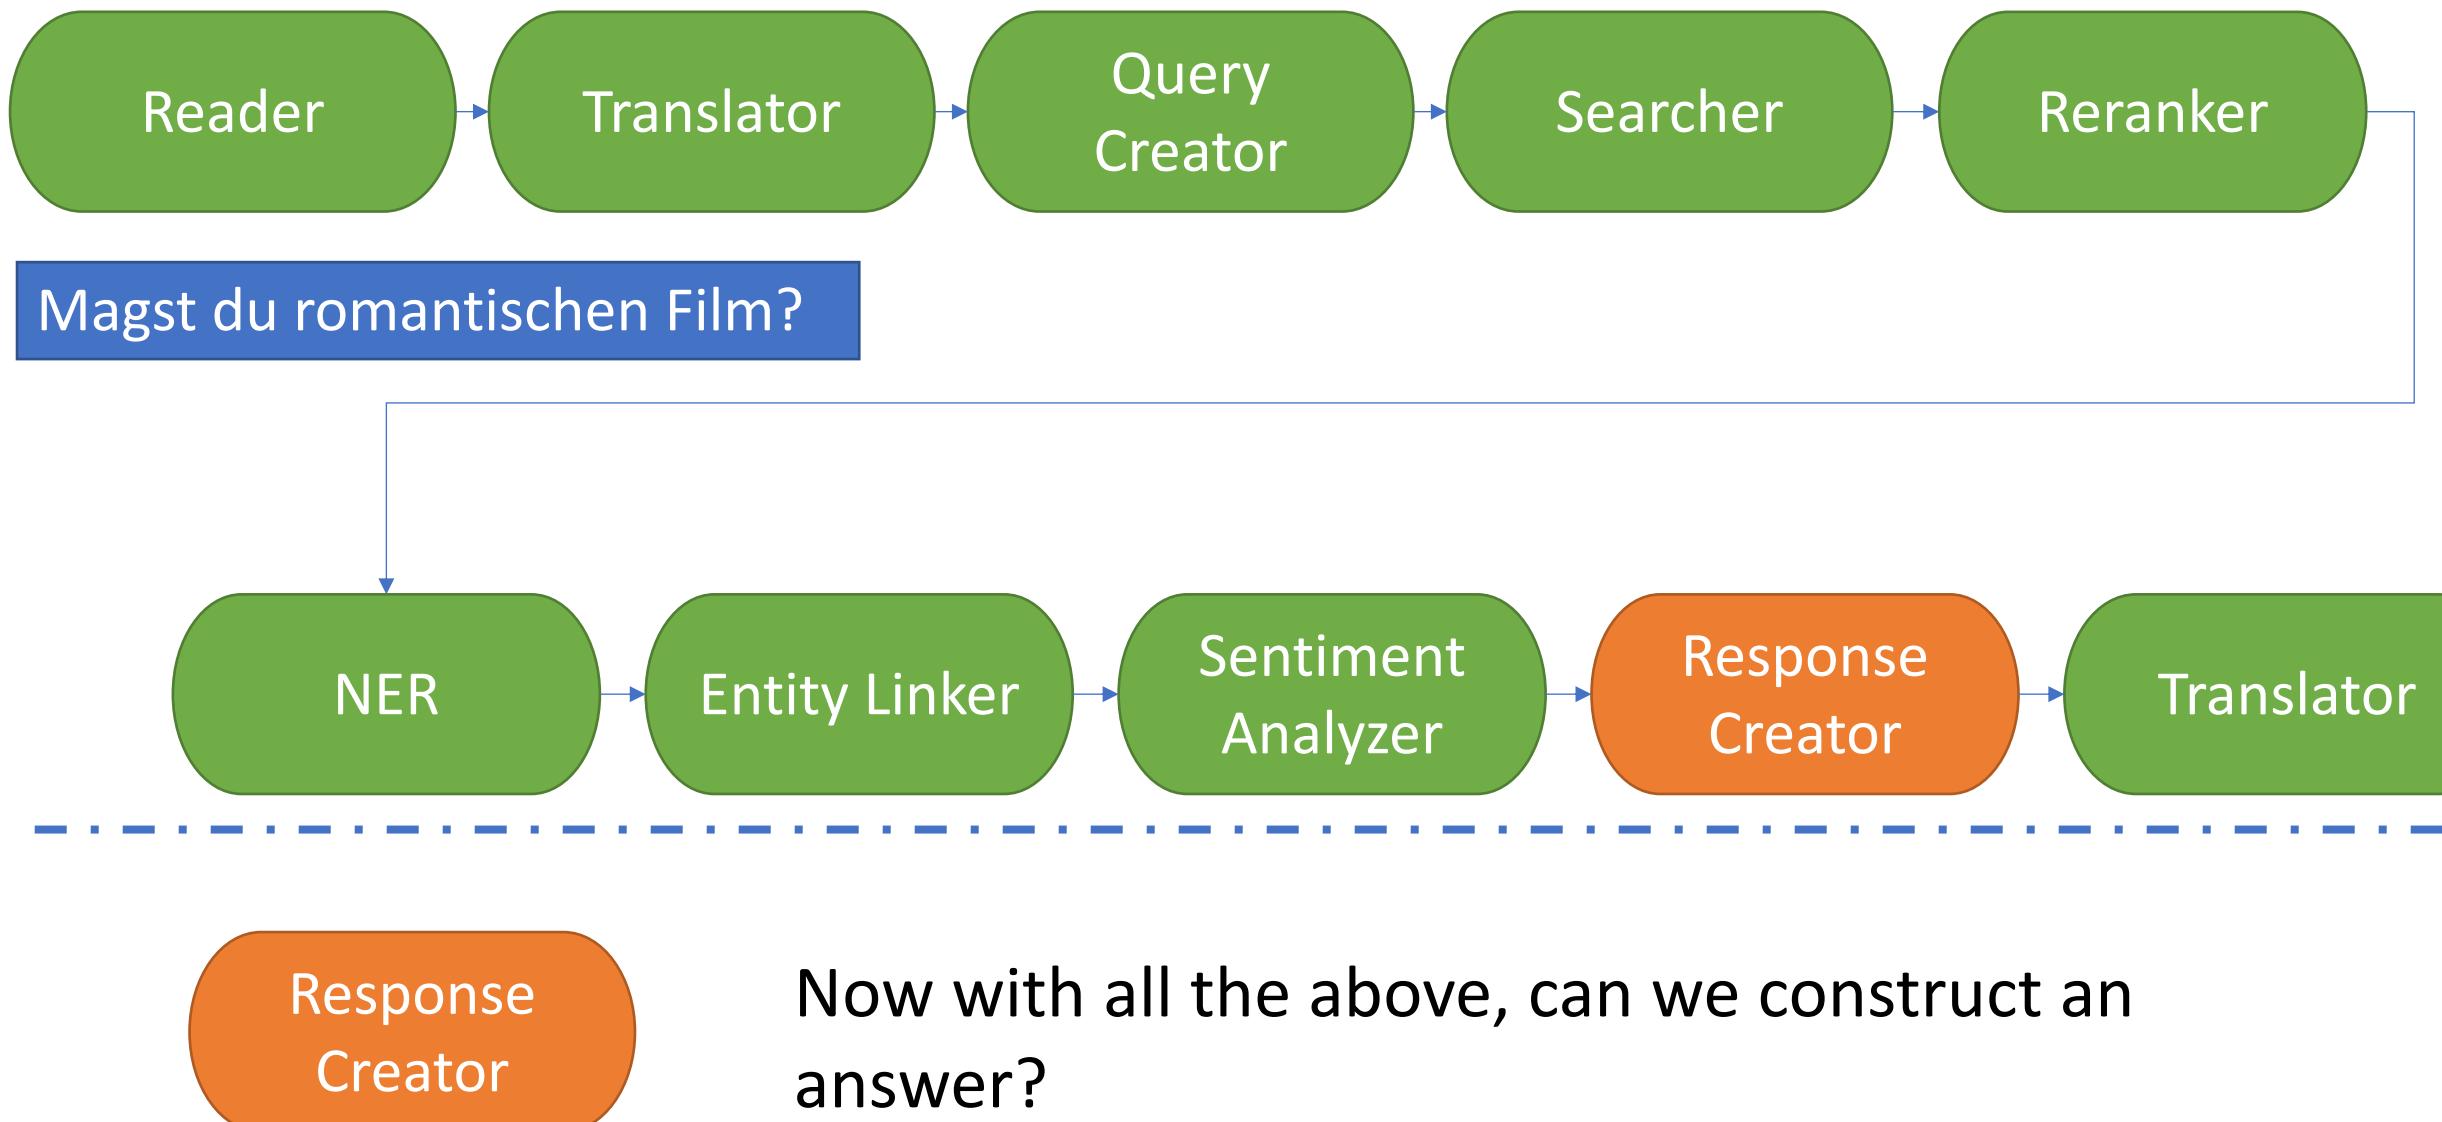
Magst du romantischen Film?



Entity Linker

Link “well-known” mentions to backend knowledge base, such as “Leonardo Dicaprio”



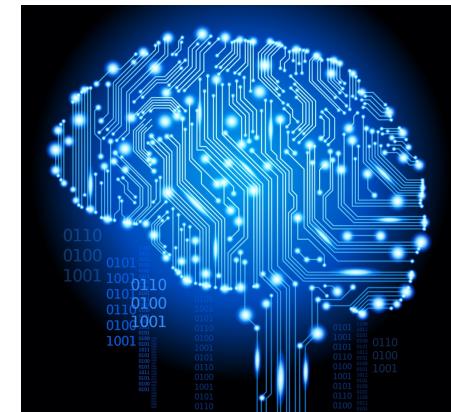


Response
Creator

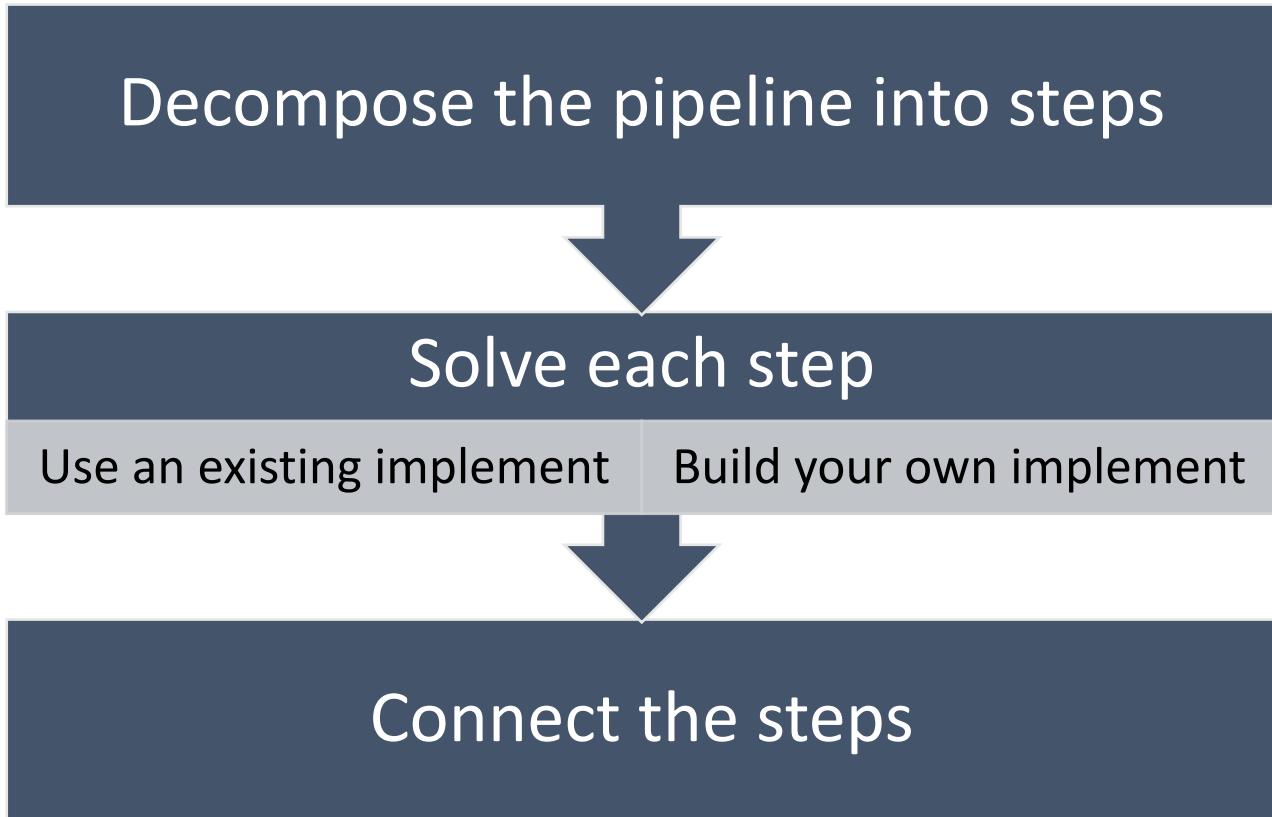
Now with all the above, can we construct an answer?

Inter-Operation Across the Pipeline

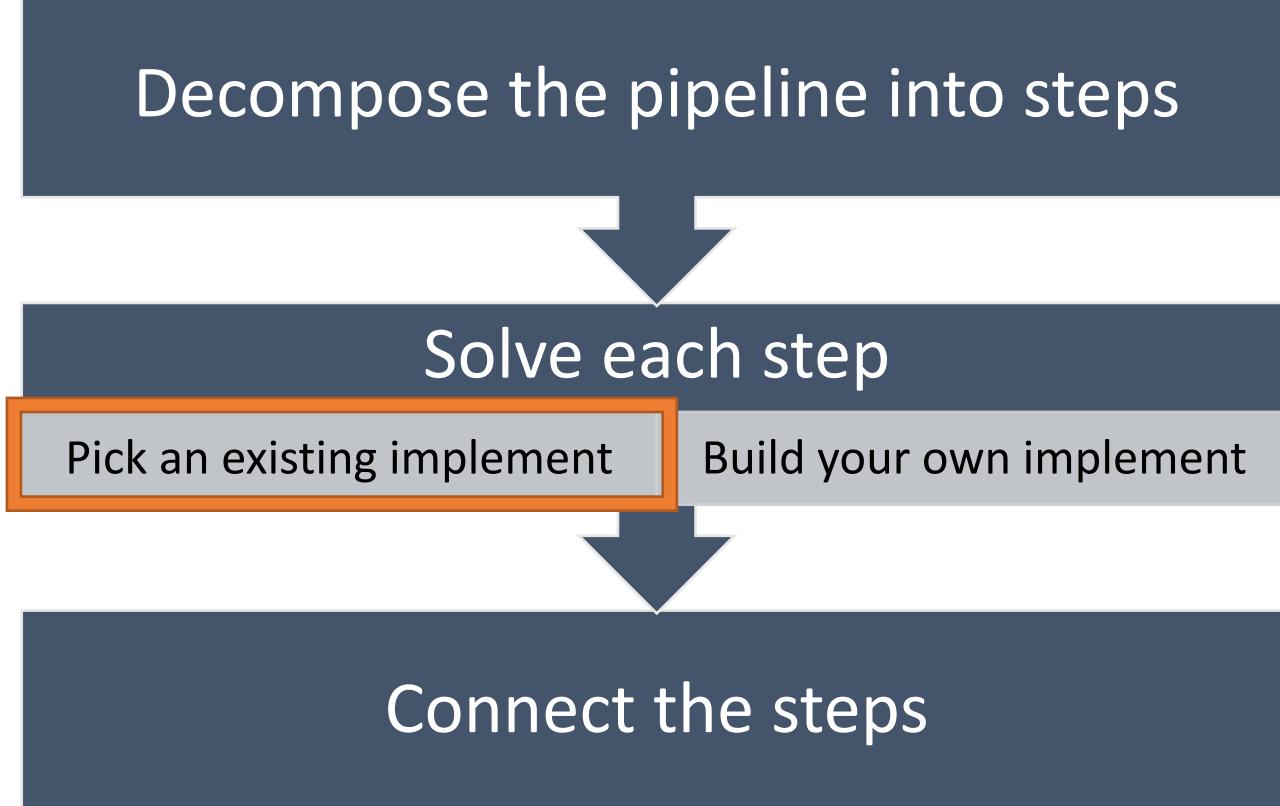
- Every Step in the Pipeline produces useful information
- Can we easily access and utilize these for the final goal?
- Let's review how we normally build such a pipeline



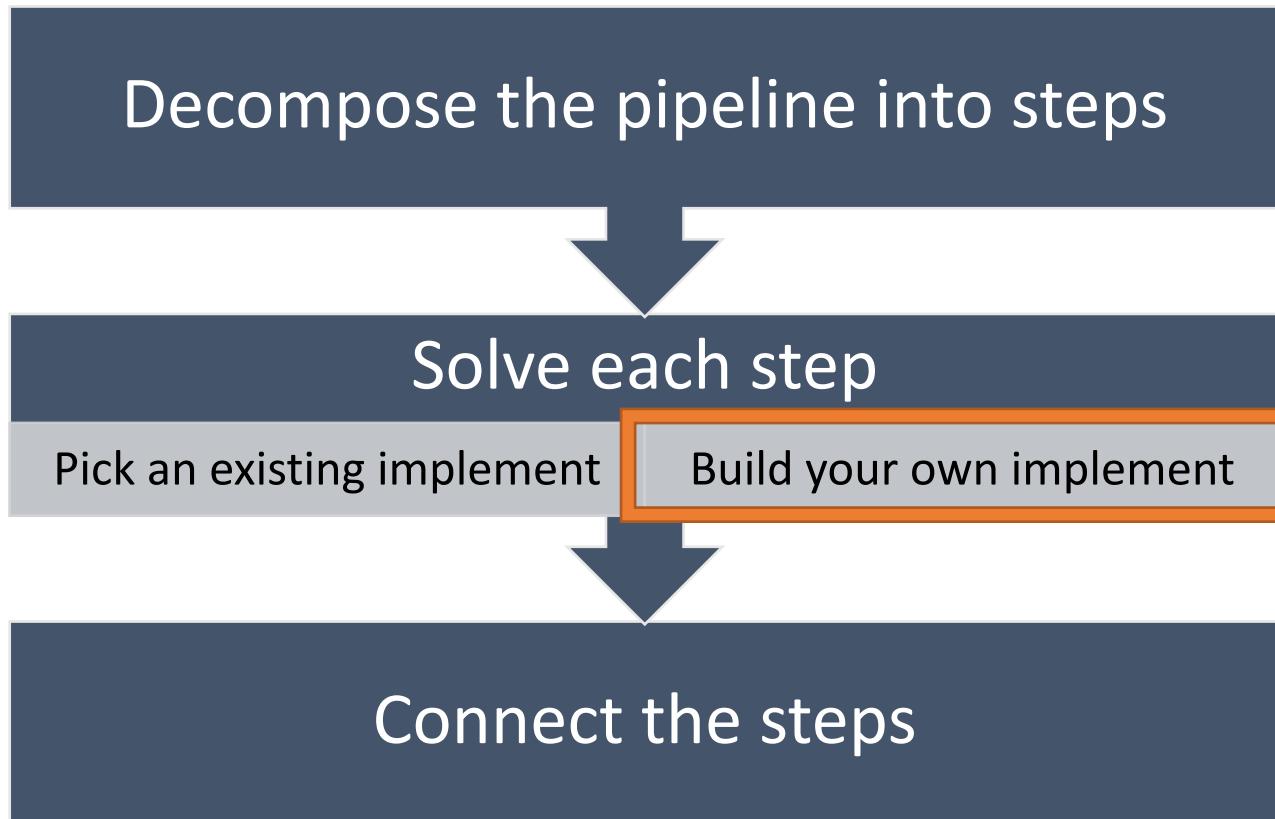
How to build complex NLP pipelines?



How to build complex NLP pipelines?



How to build complex NLP pipelines?



How to build complex NLP pipelines?

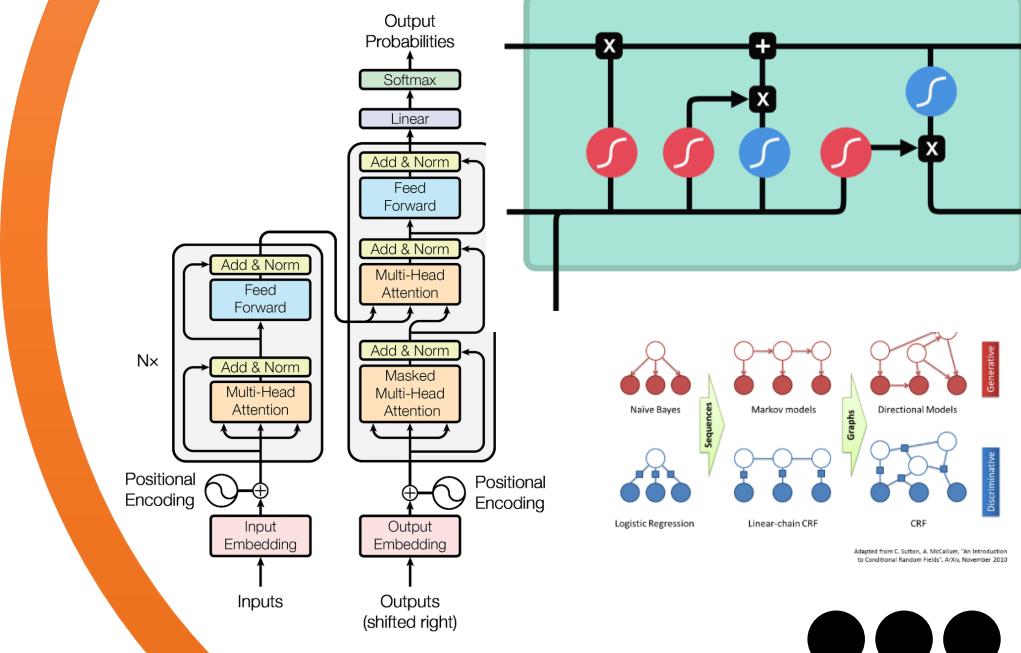
Decompose the pipeline into steps

Solve each step

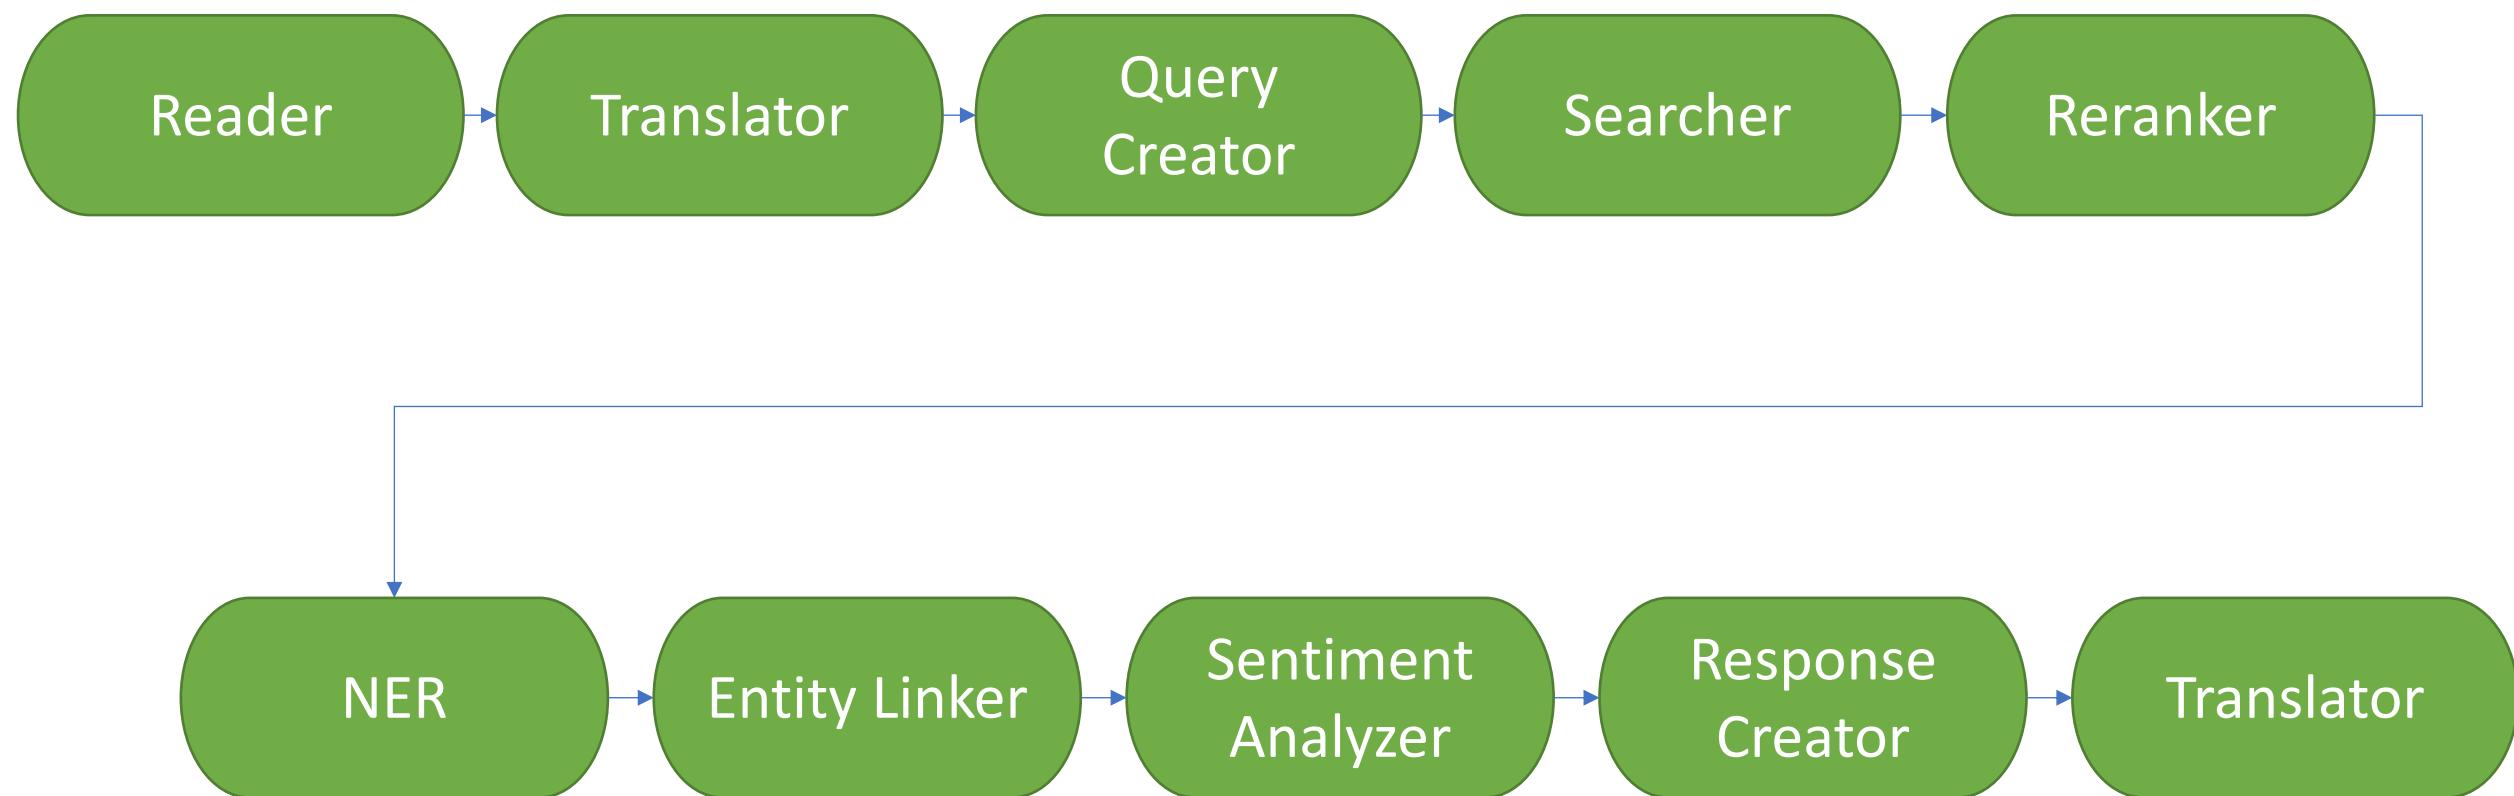
Pick an existing implement

Build your own implement

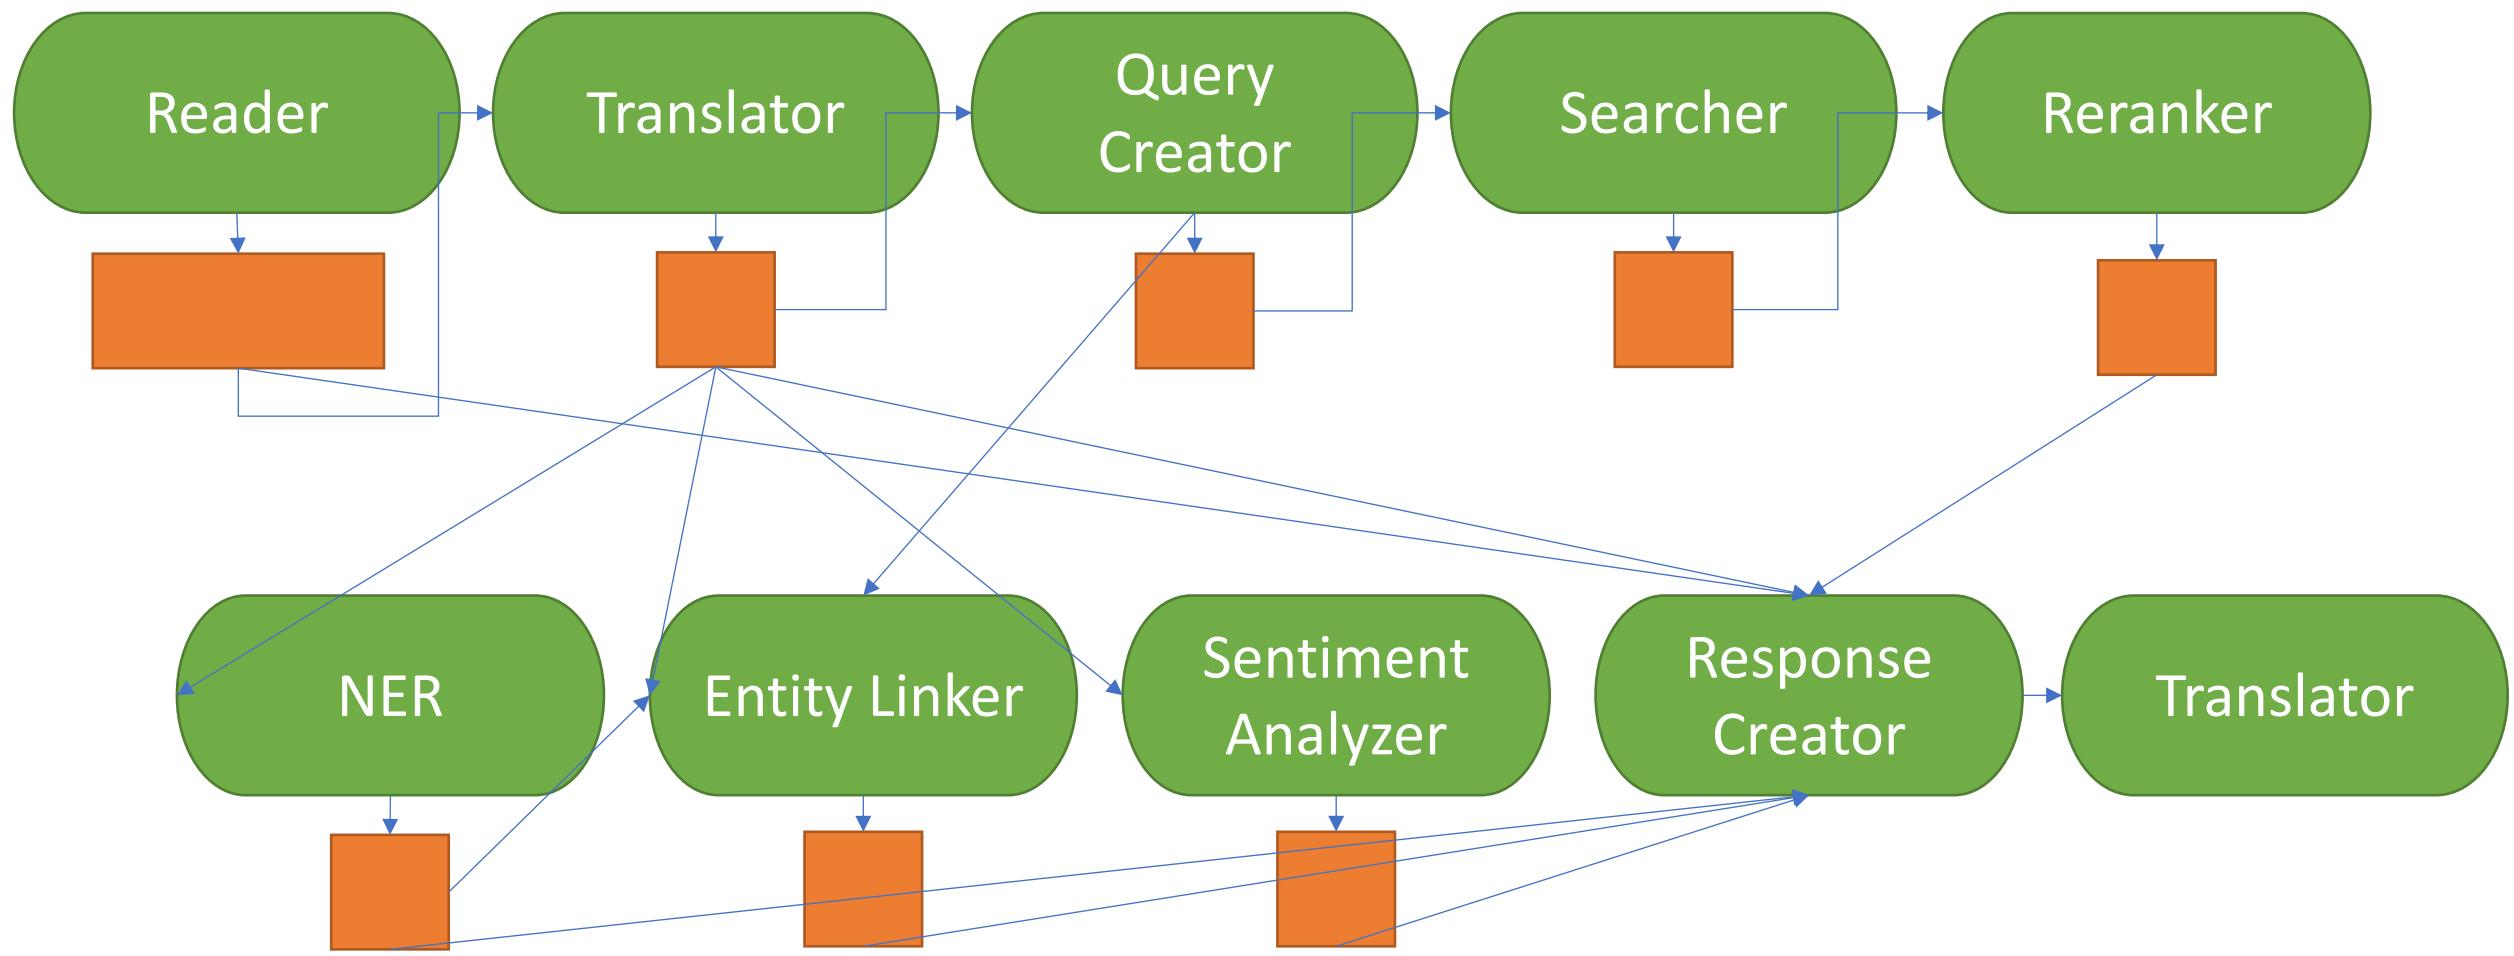
Connect the steps



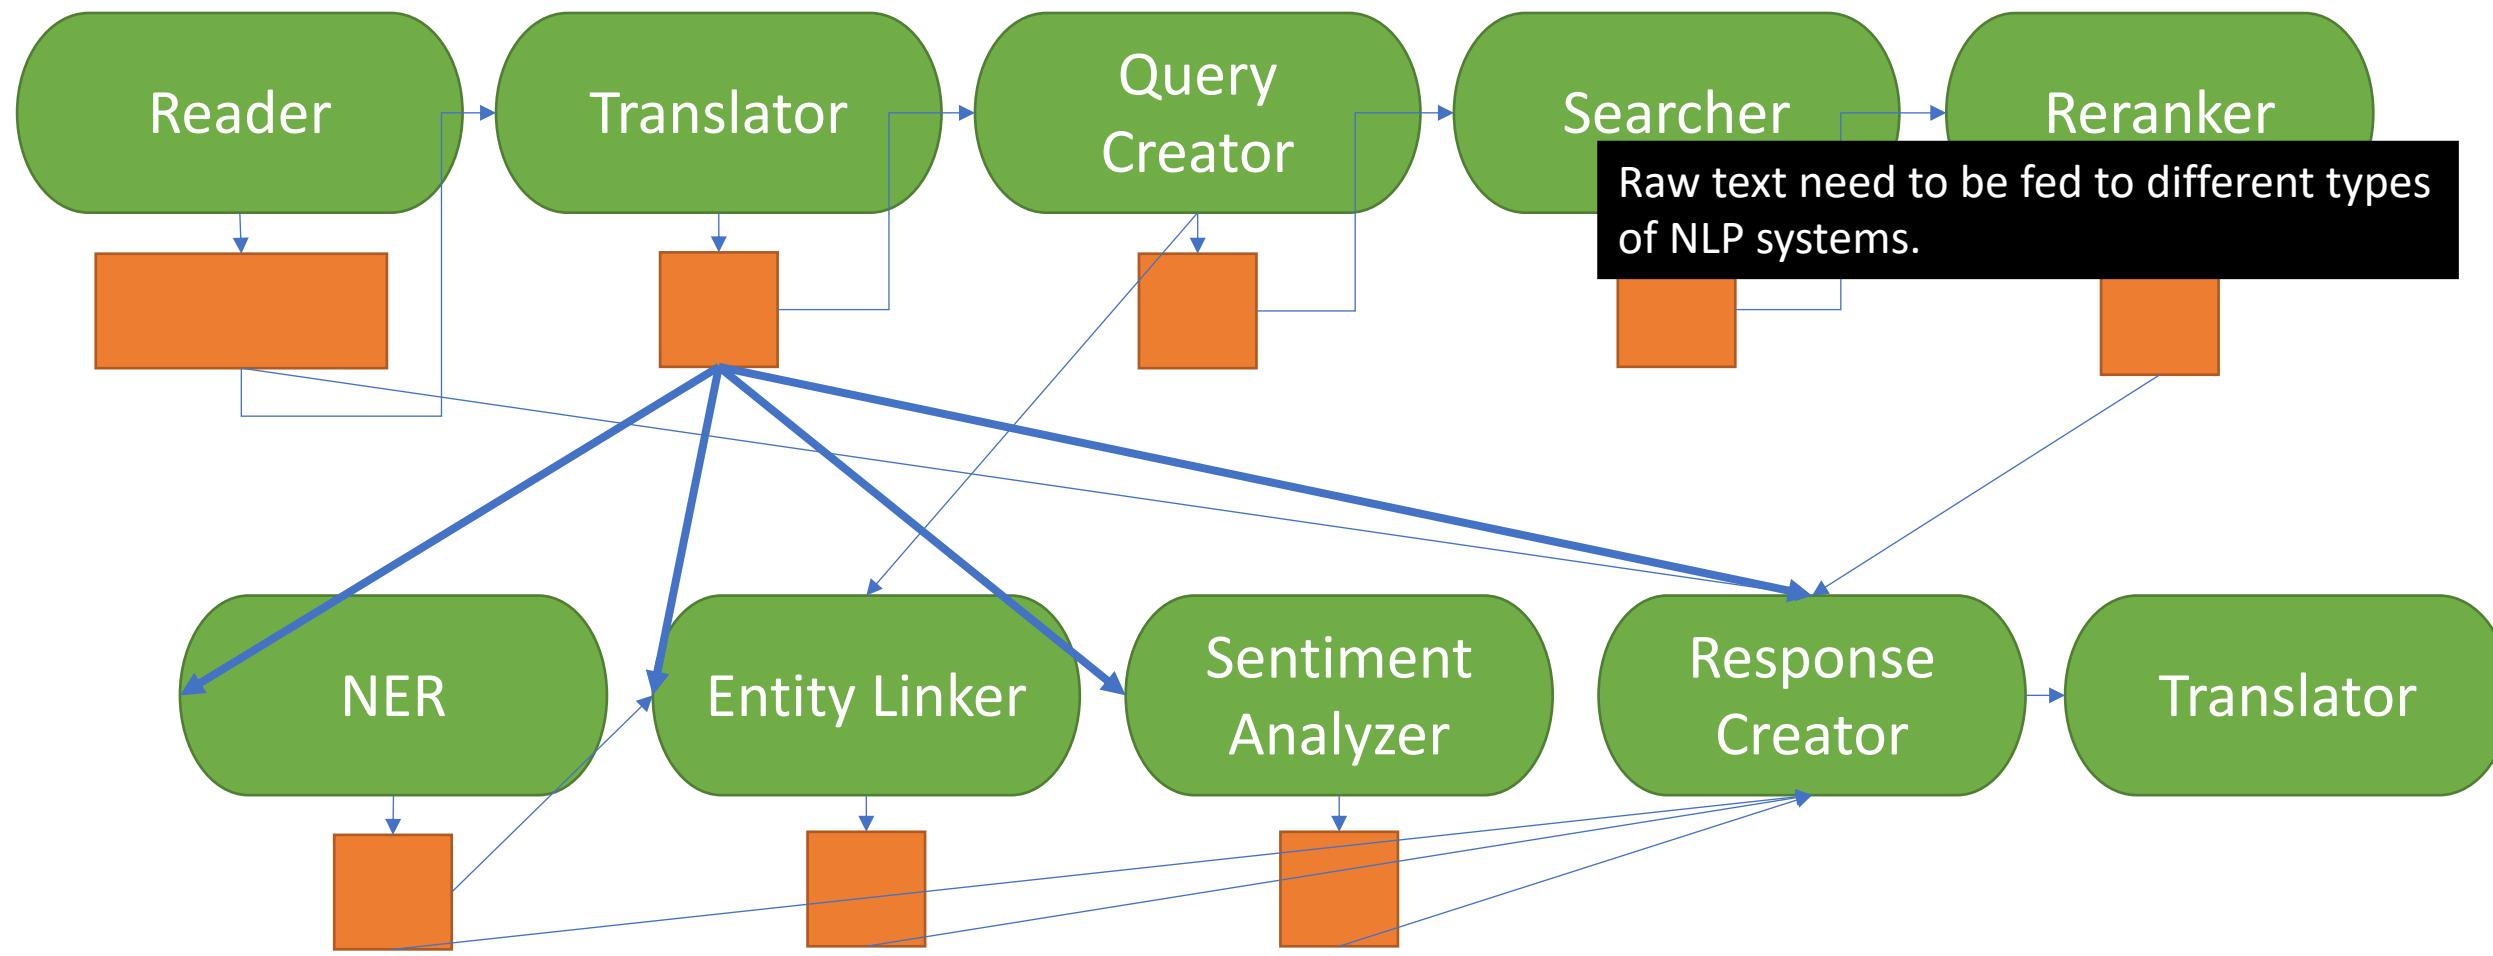
So here is the “expected” pipeline



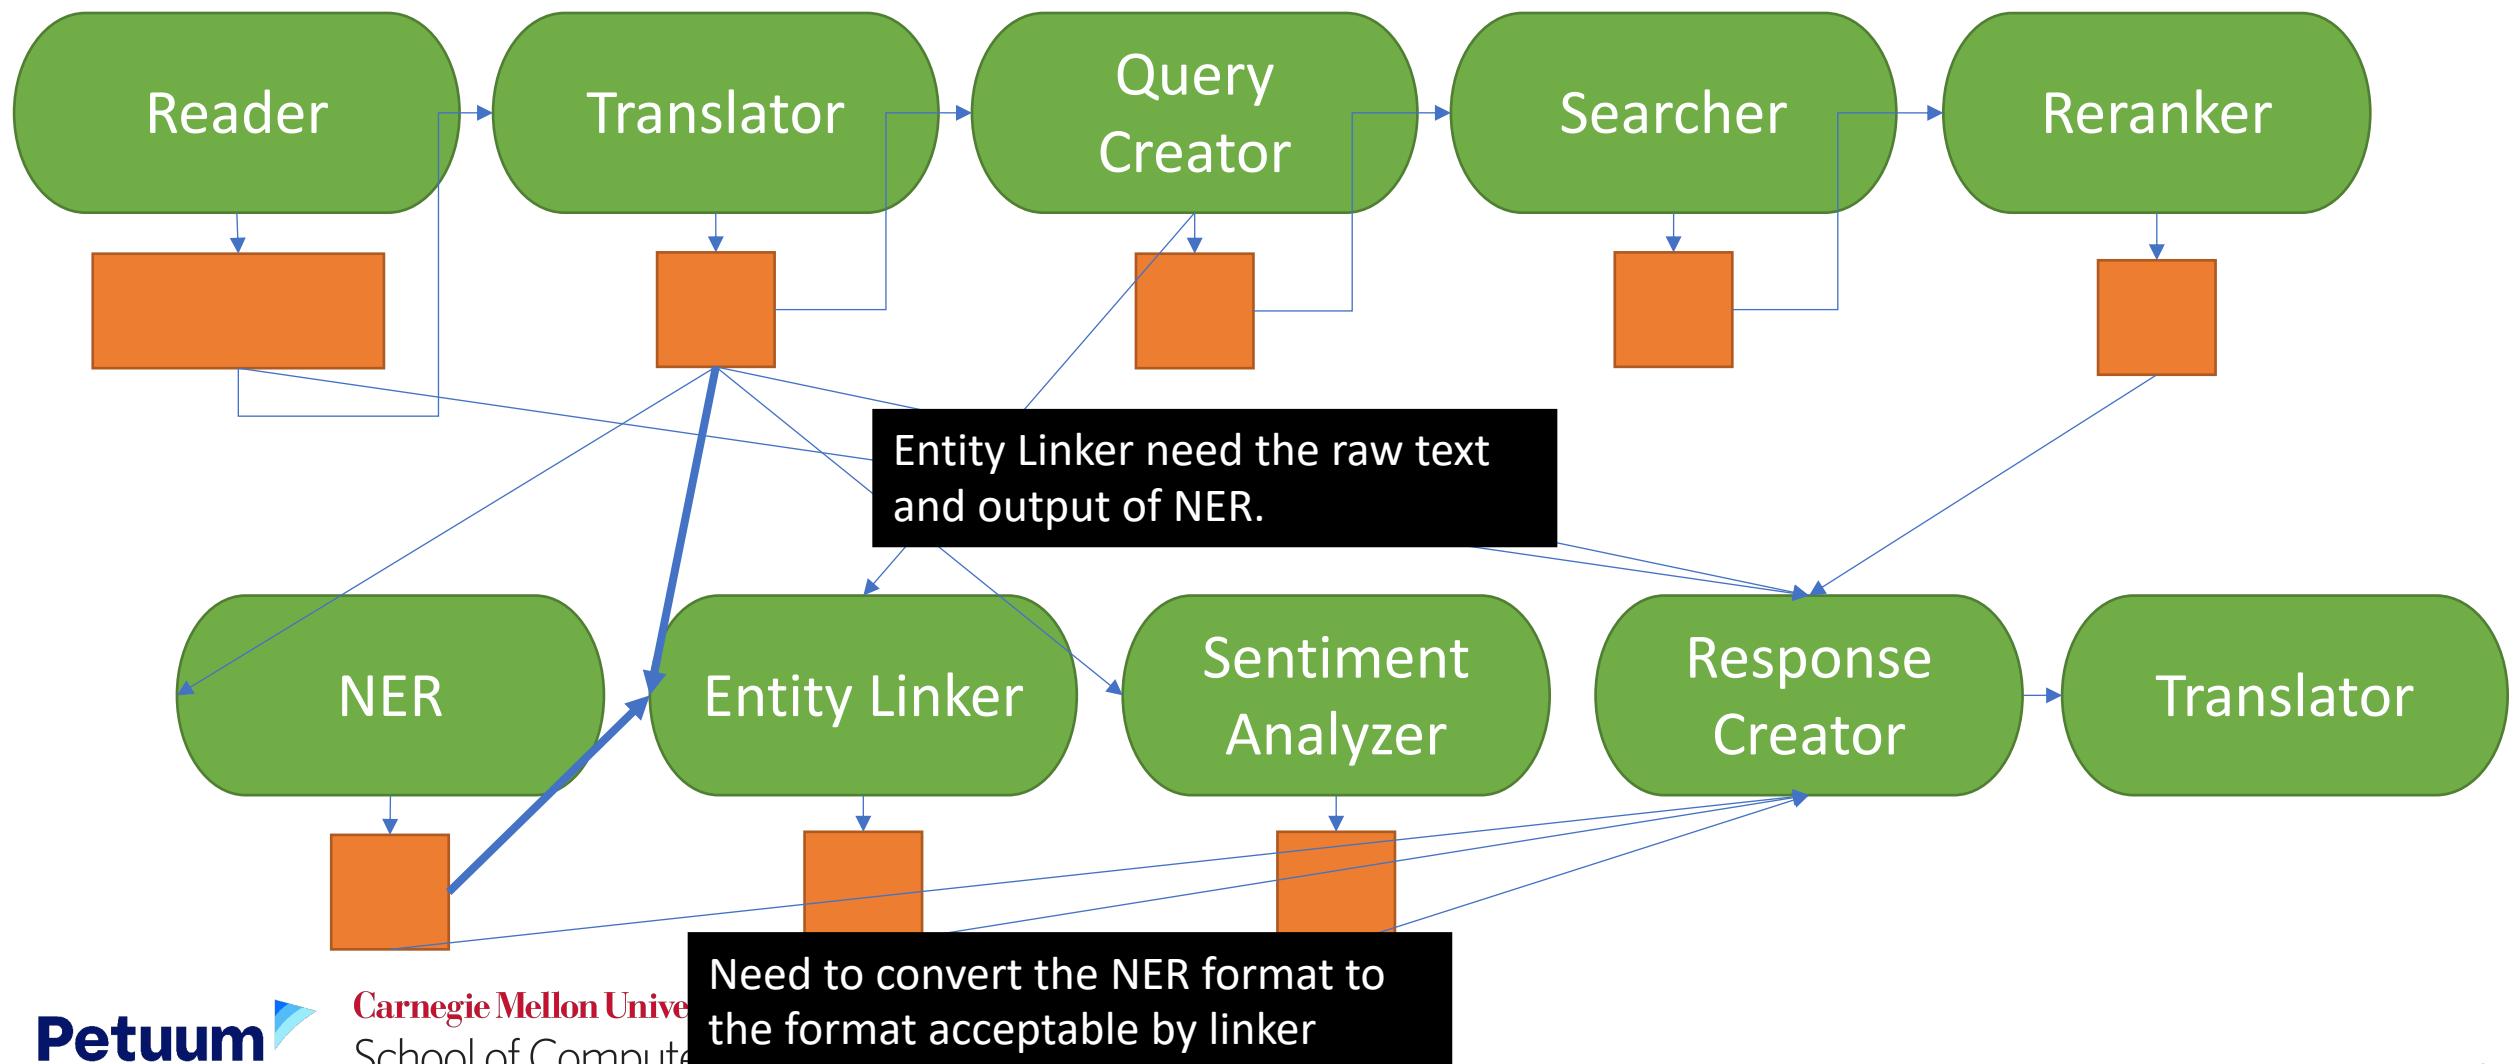
What it might look like



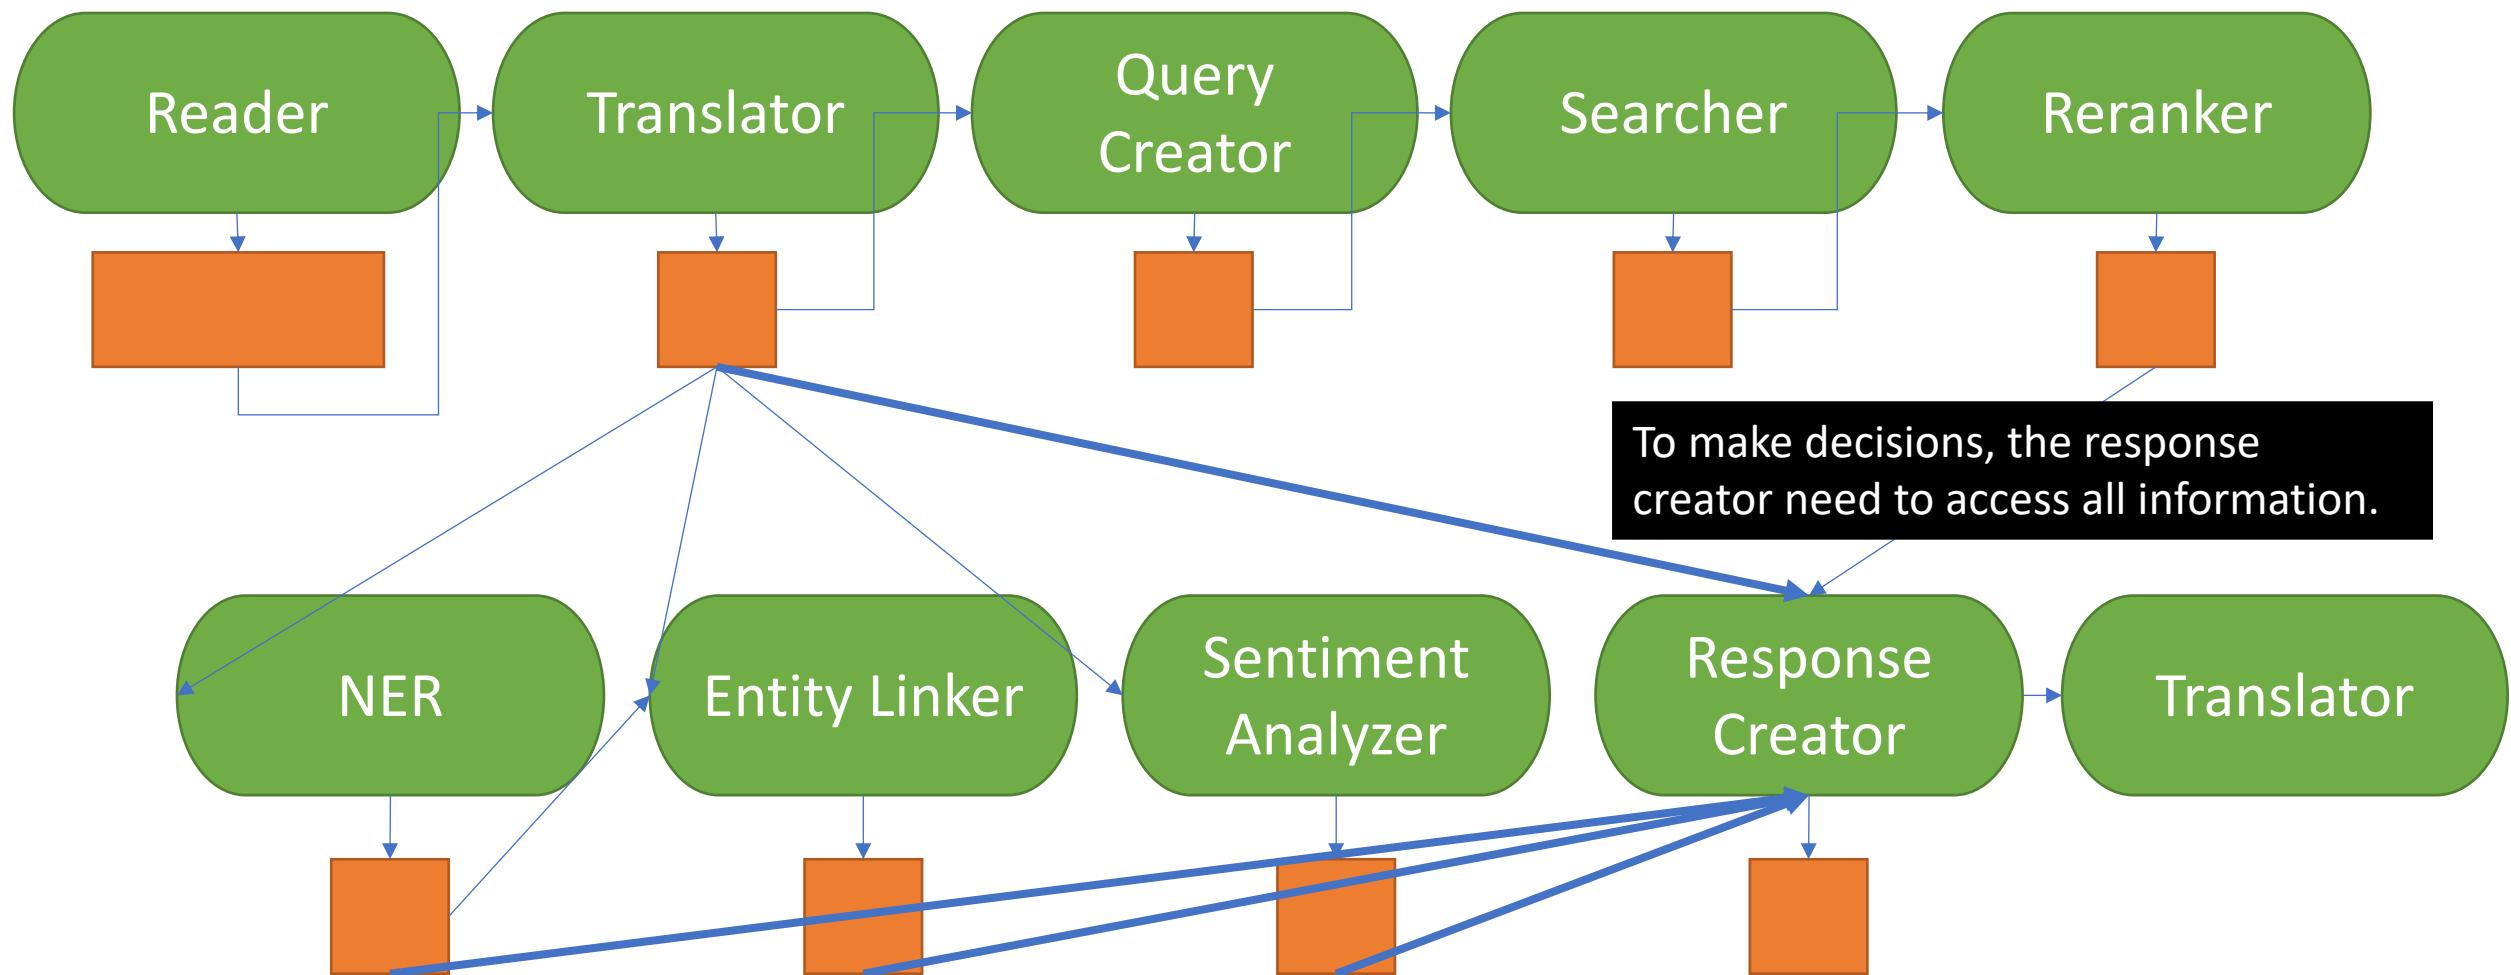
What it might look like



What it might look like



What it might actually look like



What it might look like

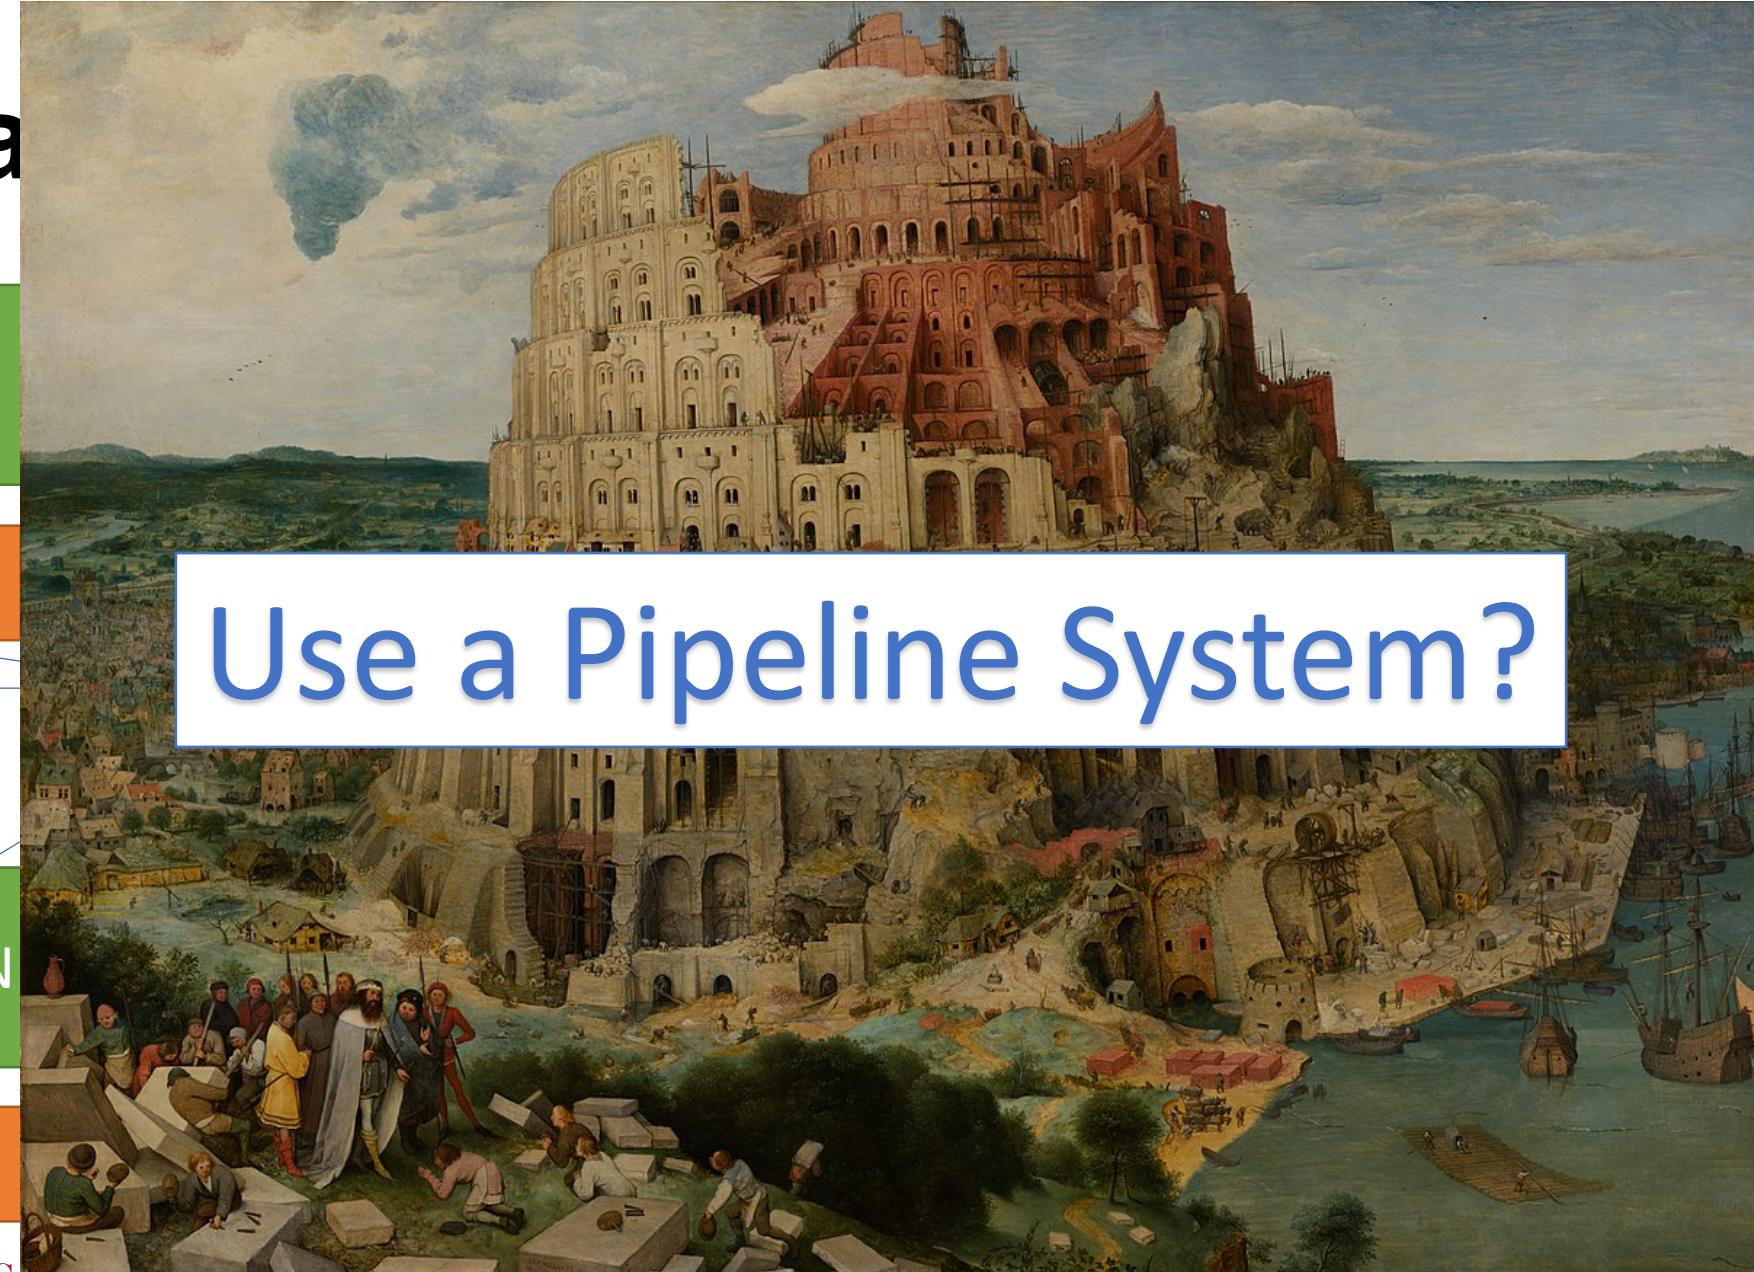


What

Reader

Anker

Use a Pipeline System?



N

Translator

Examples of Some Pipeline Systems

- Illinois Curator
 - Support many different NLP tasks
 - Especially strong in different flavors of SRL
 - Allow composing many different tasks

Curator Demo [check status?](#) [about?](#)

Enter some text in the box below. Then, check the boxes next to the types of annotation you want Curator to provide. Click "process" to send your text to be annotated with those resources. When the selected annotation services have finished, the outputs of those services will be displayed on this page.

You can replace this placeholder text with the sentences you'd like the Curator to annotate.

process

Select the annotation services you'd like to be applied to your text:

- Sentences -- Illinois sentence-level segmenter
- Tokens -- Illinois token-level segmenter
- Part-of-Speech -- Illinois Part-of-Speech tagger
- Shallow Parse chunks -- Illinois Chunker (a.k.a. Shallow Parser)
- Named Entities -- Illinois Named Entity Recognizer 2014 CoNLL (PER/LOC/ORG/MISC)
- Named Entities -- Illinois Extended Named Entity Recognizer 2014 Ontonotes (18 types)
- Quantities -- Illinois Quantity Recognizer
- Semantic Roles (verbs) -- older SRL

Examples of Some Pipeline Systems

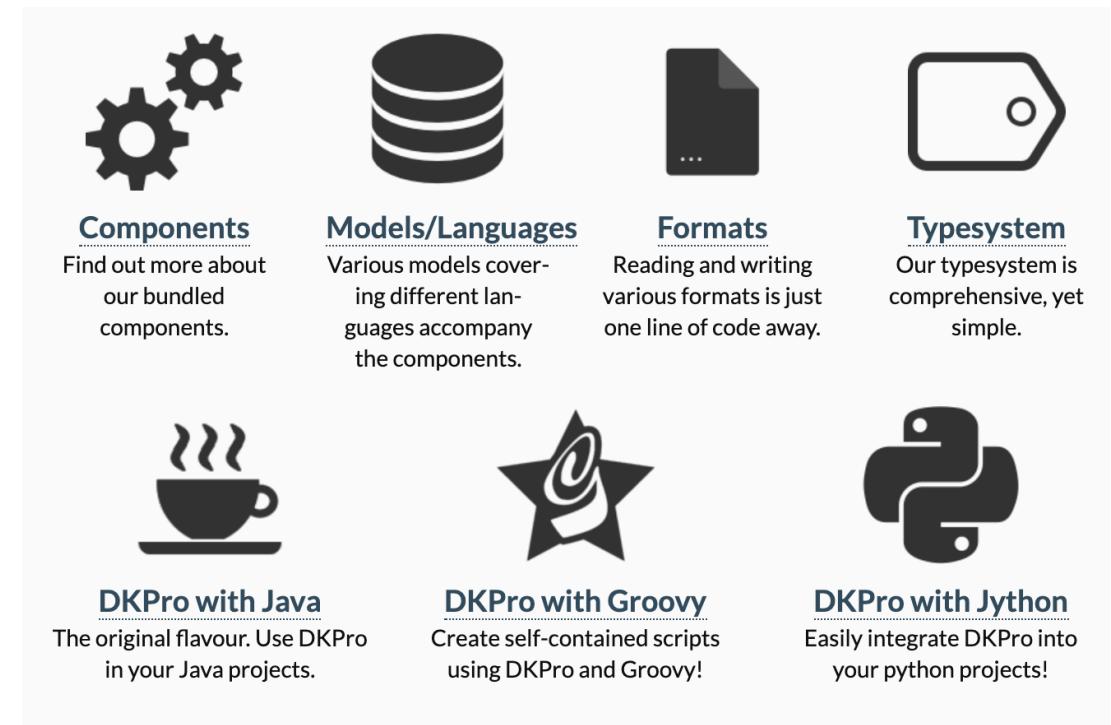
- Stanford
 - Features in a good coverage in core NLP tasks
 - Provide utilities on these tasks
 - Strong dependency between tasks (e.g. parsing depends on tokenization)

The screenshot shows the Stanford CoreNLP website with a dark red header bar containing the logo, the text "Stanford CoreNLP", and links for "Github repo" and "Quick links". Below the header, the page title "CoreNLP" and version "version 3.9.2" are displayed. A sidebar on the left has dropdown menus for "Overview", "Usage", and "Annotators", with "Annotators" currently selected. The main content area is titled "Annotators" and includes a "Table of Contents" with links to "Annotator Descriptions", "Annotator Dependencies", and "Sub-Annotators". Below this is a section titled "Annotator Descriptions" with a table. The table has columns for NAME, ANNOTATOR CLASS, GENERATED ANNOTATION, and DESCRIPTION. It lists two entries: " tokenize" (TokenizerAnnotator) which generates TokensAnnotation and describes tokenization; and " cleanxml" (CleanXmlAnnotator) which generates XmlDocumentAnnotation and describes cleaning XML.

NAME	ANNOTATOR CLASS	GENERATED ANNOTATION	DESCRIPTION
tokenize	TokenizerAnnotator	TokensAnnotation (list of tokens); CharacterOffsetBeginAnnotation, CharacterOffsetEndAnnotation, TextAnnotation (for each token)	Tokenizes the text splits the text into "words", using methods suitable for the language being processed. Sometimes tokens split up words in ways suitable for further NLP-ing, for example becomes "is" and "are". The tokenizer sets beginning and end character offsets of tokens in the input
cleanxml	CleanXmlAnnotator	XmleContextAnnotation	Remove xml tokens

Examples of Some Pipeline Systems

- DKPro Core
 - Support large number of tools
 - Tools are very loosely coupled
 - Use a universal data format based on Typesystem
 - Easy pipeline composing
 - A transparent data flow



What's Provided Now?

The current approach



Useful NLP Tools



Pre-trained Models



Pipeline Systems

Decompose the pipeline into steps



Solve each steps

Pick an existing implement

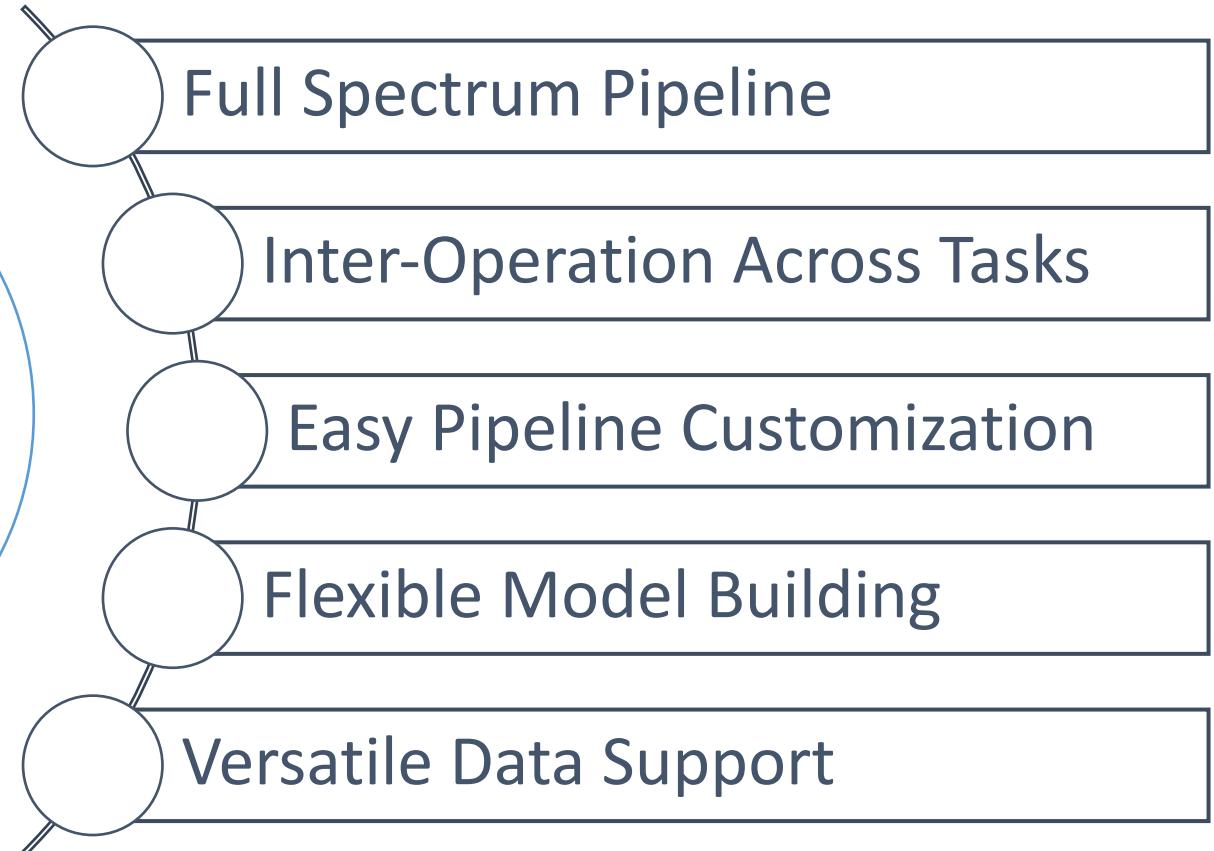
Build your own implement

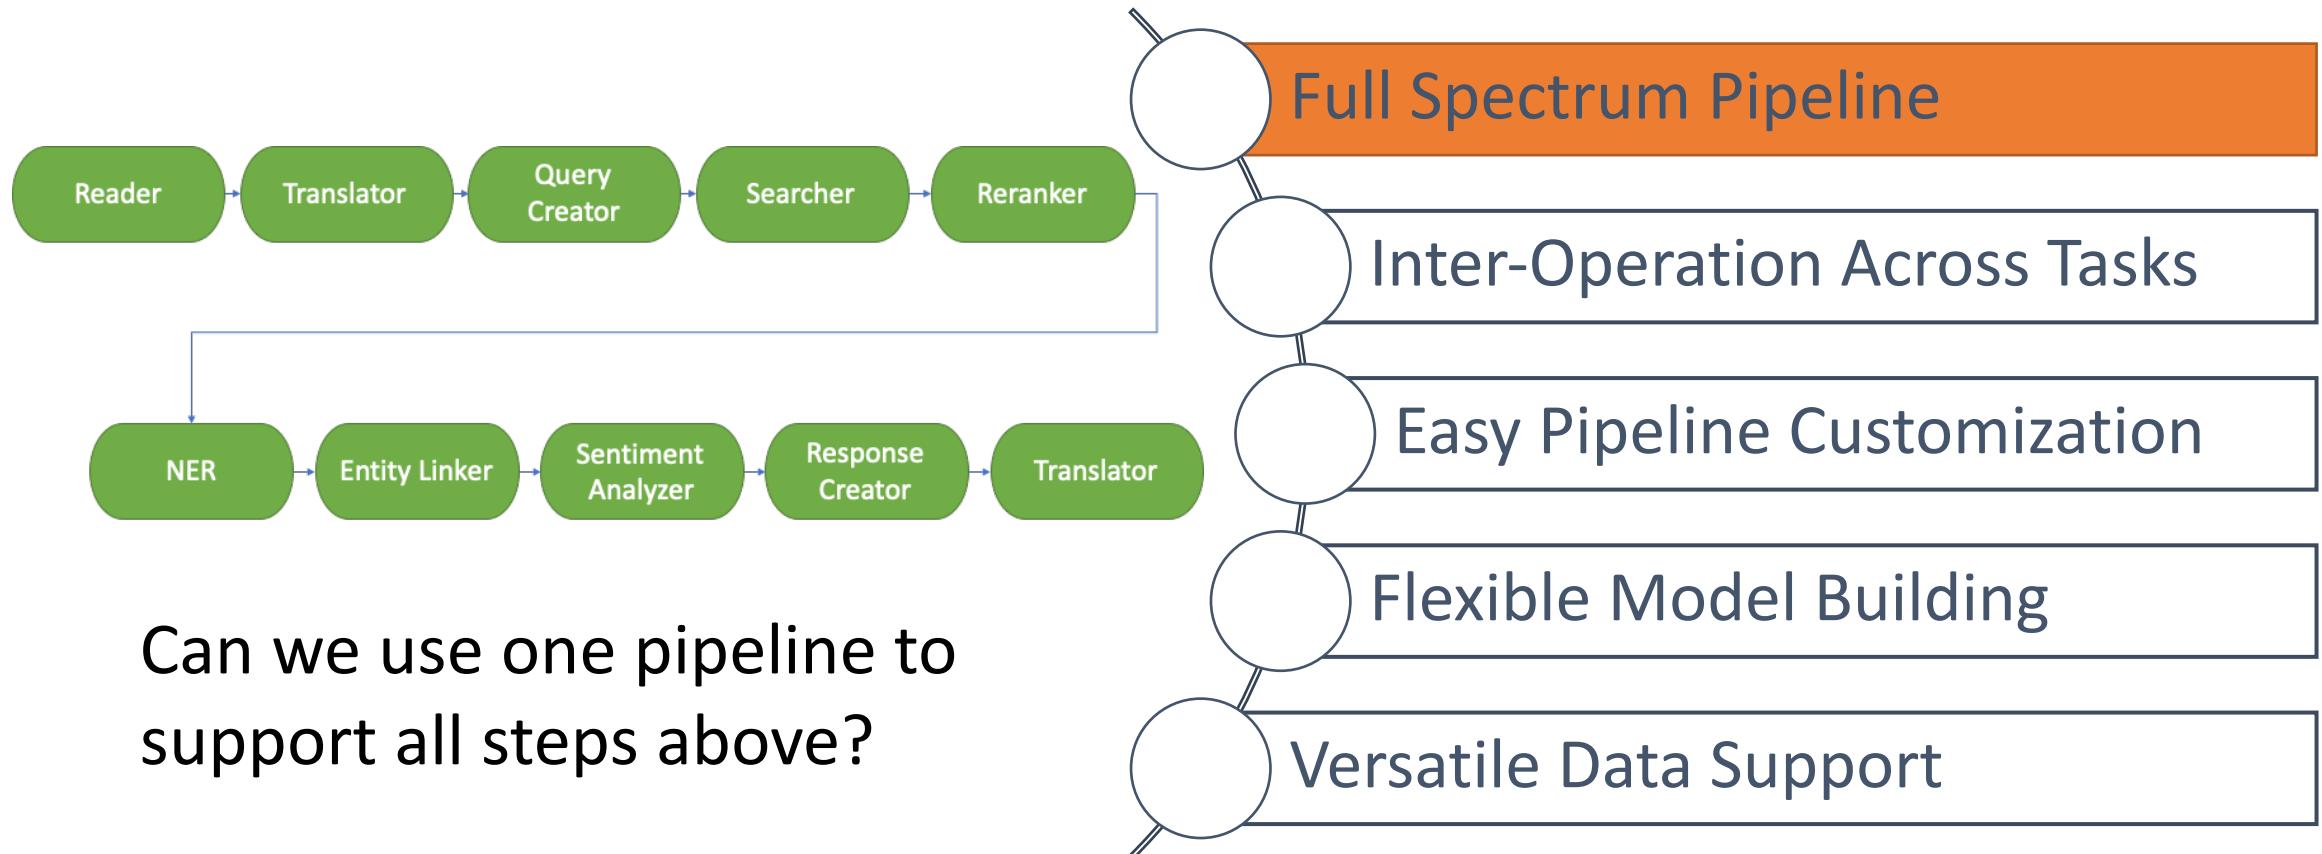


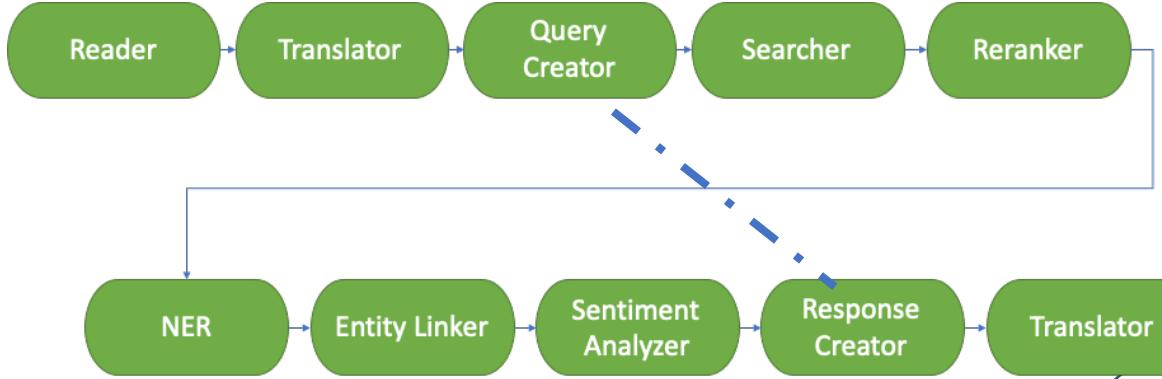
Connect the steps



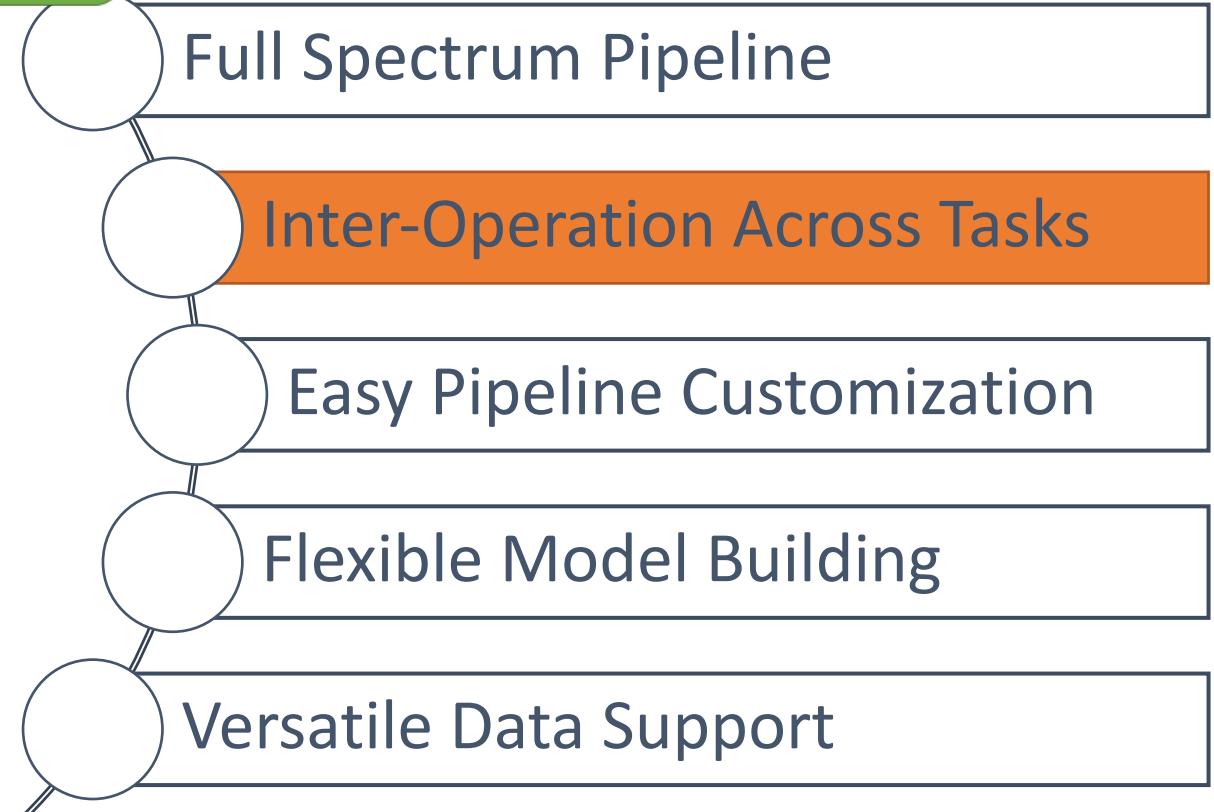
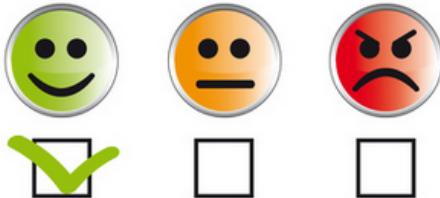
What do
we want to
achieve?







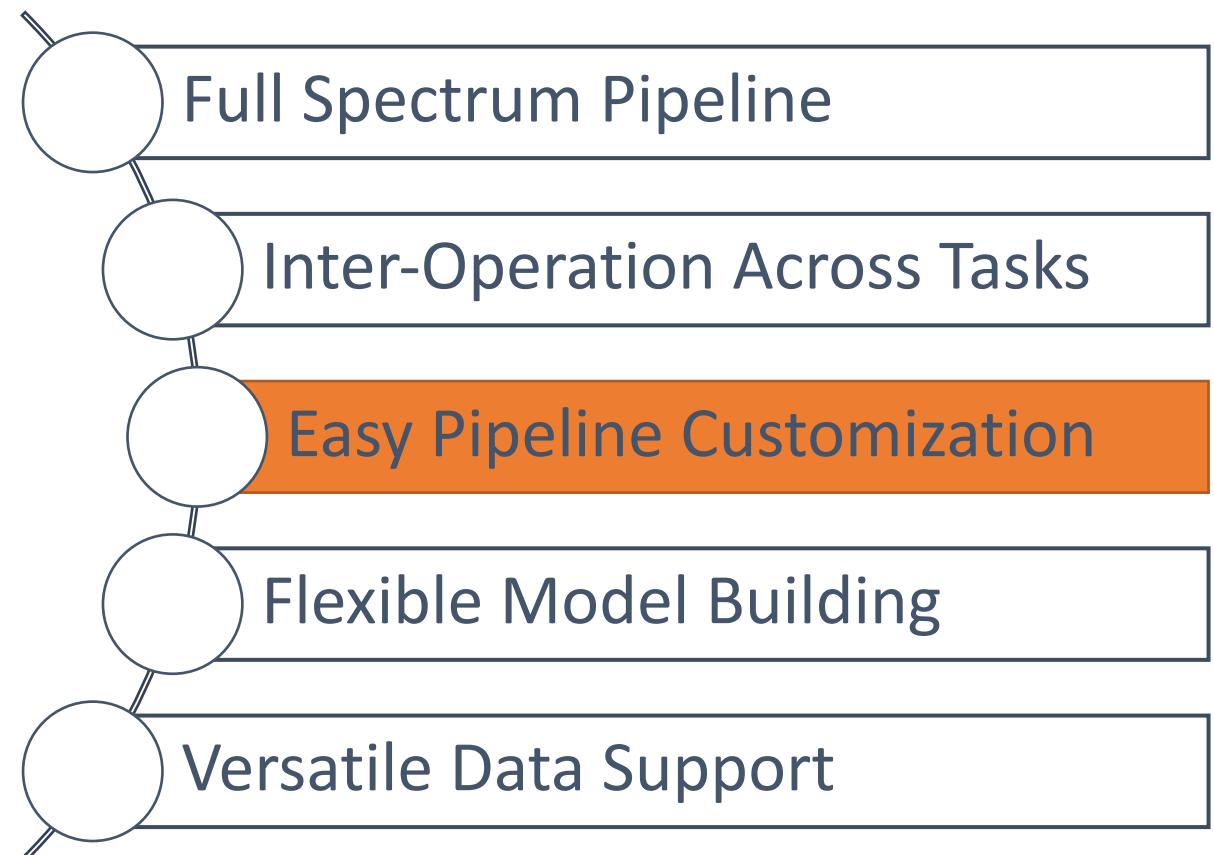
- Downstream tasks can use the results from the previous tasks
- E.g. Using sentiment analysis result in the answer selection step



Can we freely change the operations in a pipeline?

For example:

- Plug-in a tokenizer.
- Replace the translator.
- Remove lemmatization.

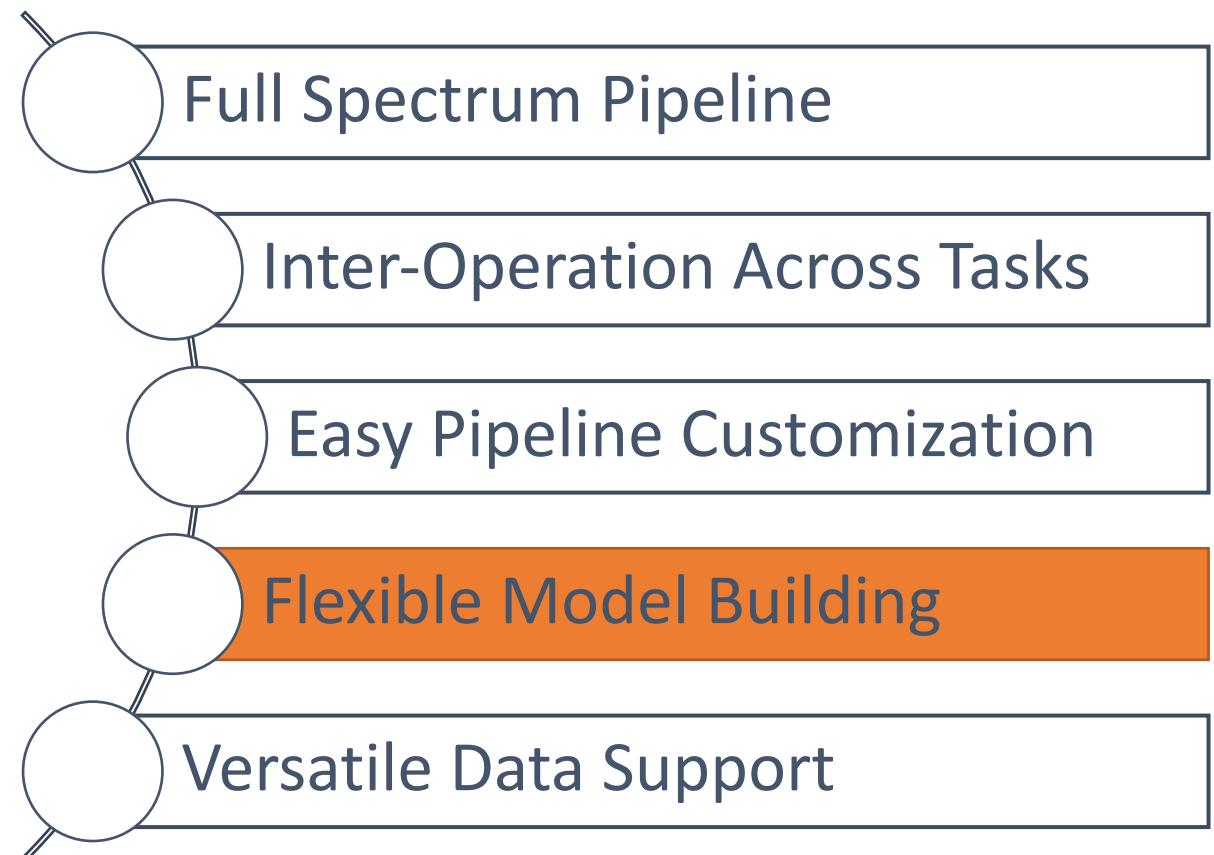


Can we directly reuse a model for
a new task quickly?

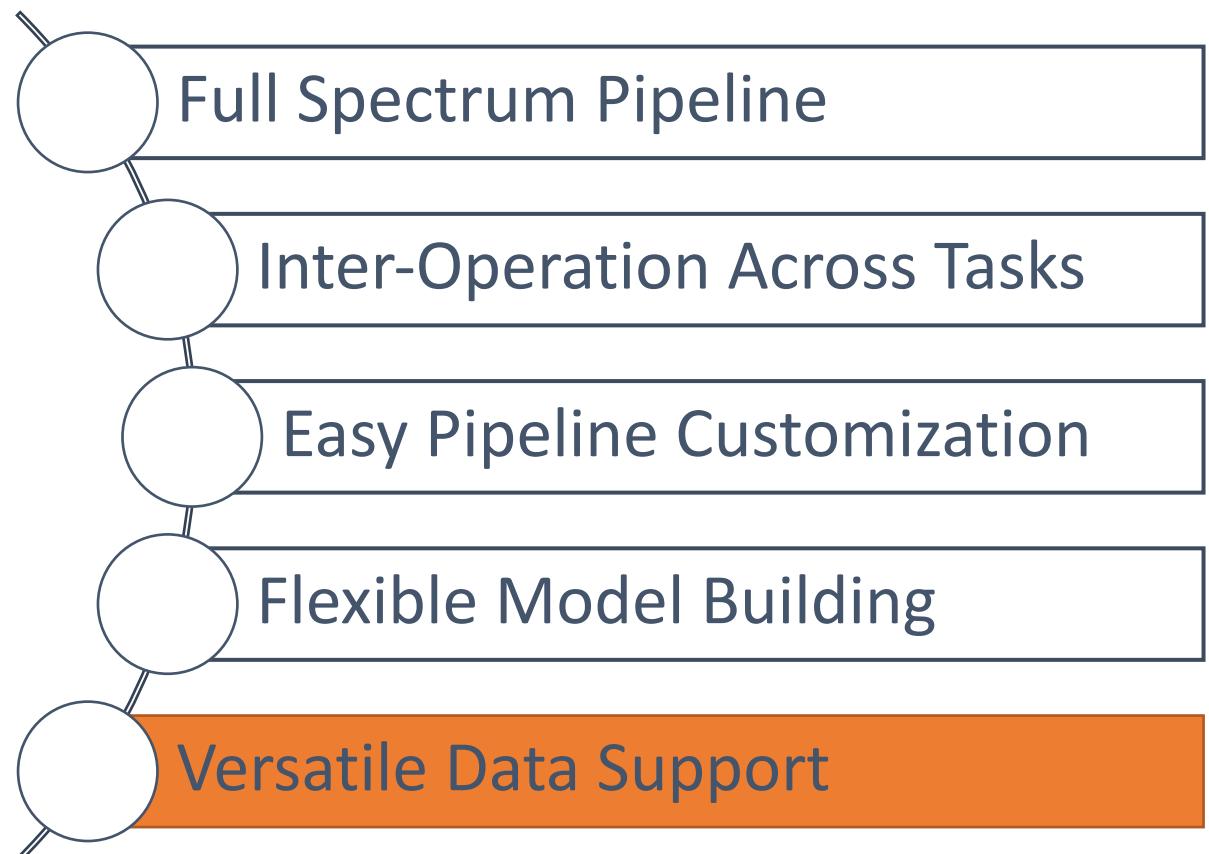
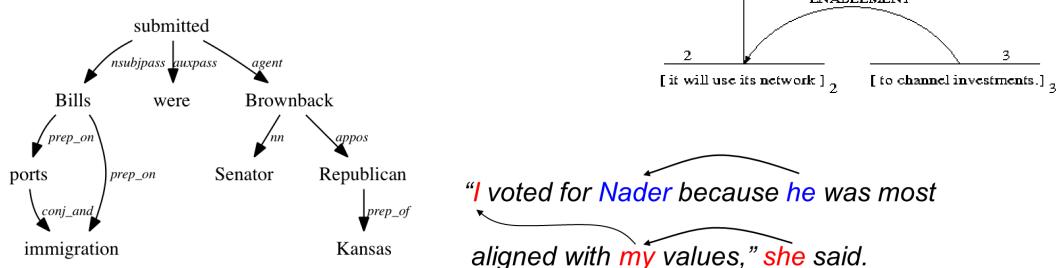
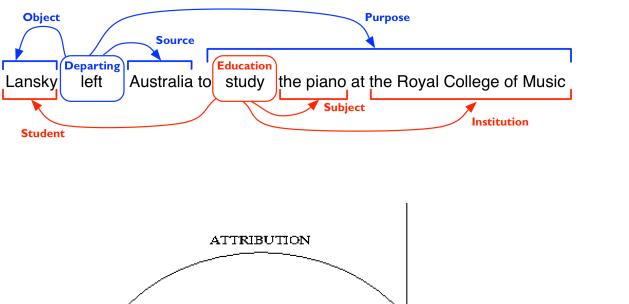
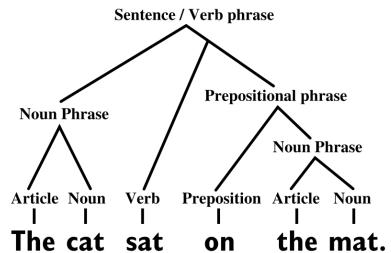


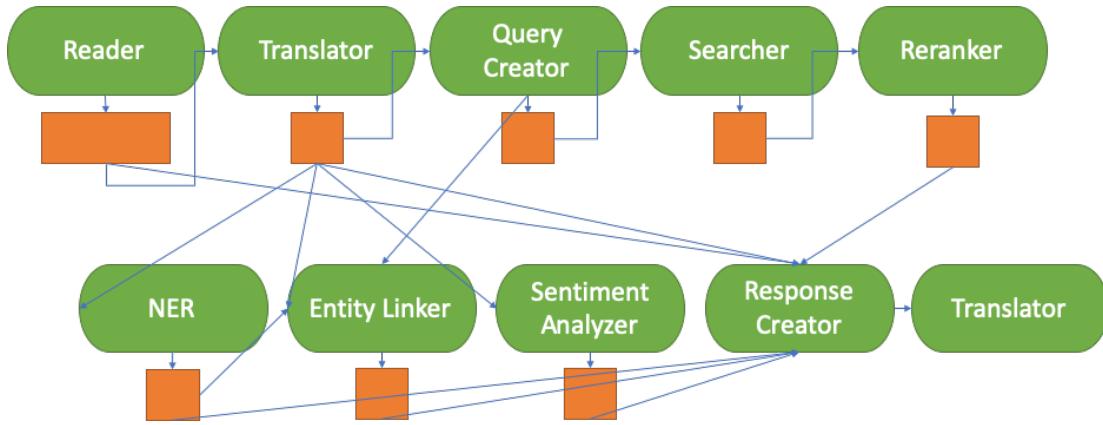
Tom and Bill work at the same company.
PER O PER O O O O O

He is going to the bookshop.
PR V V IN DT N

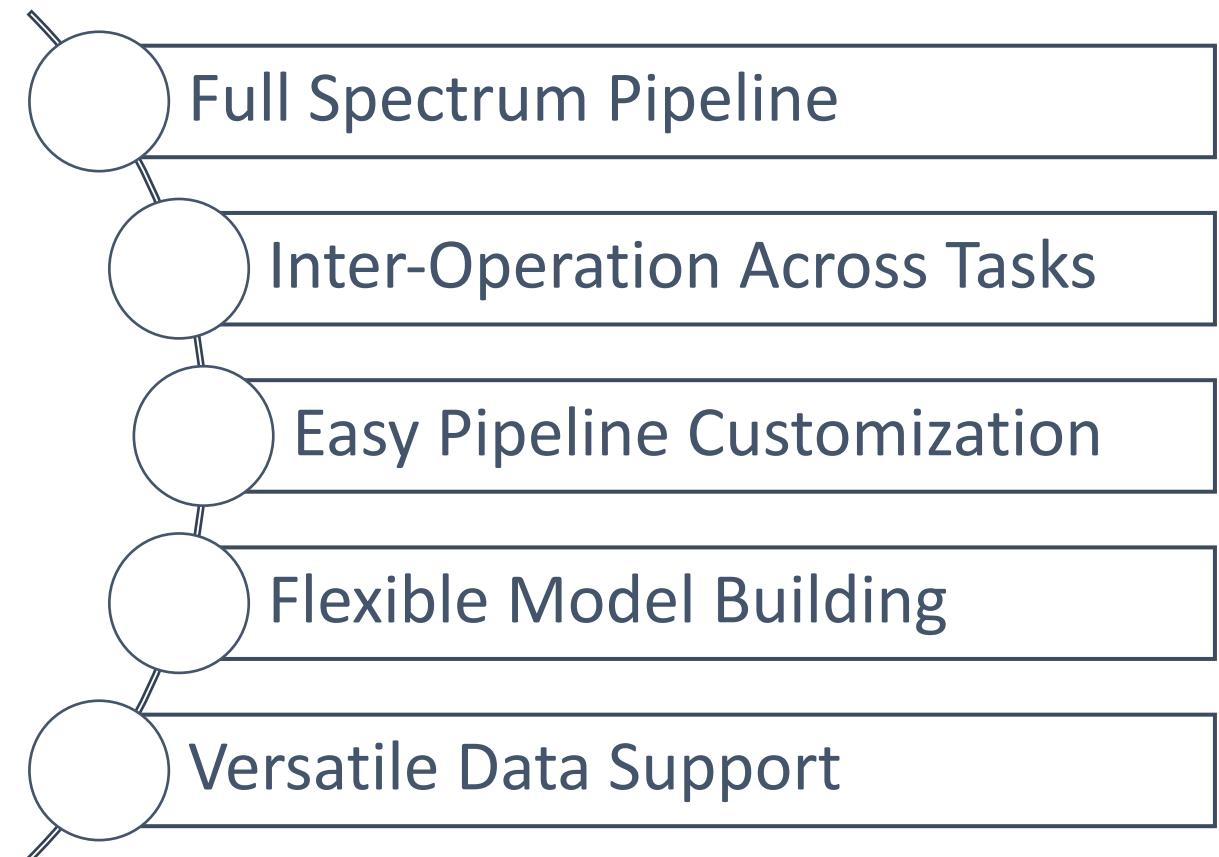
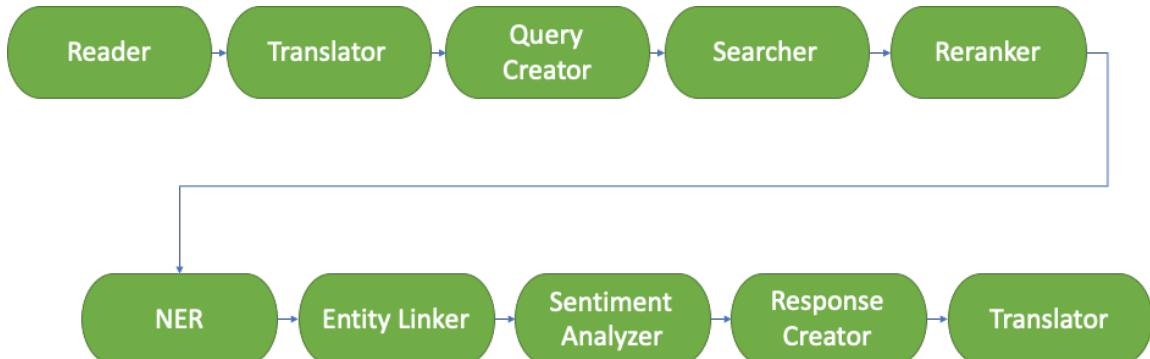


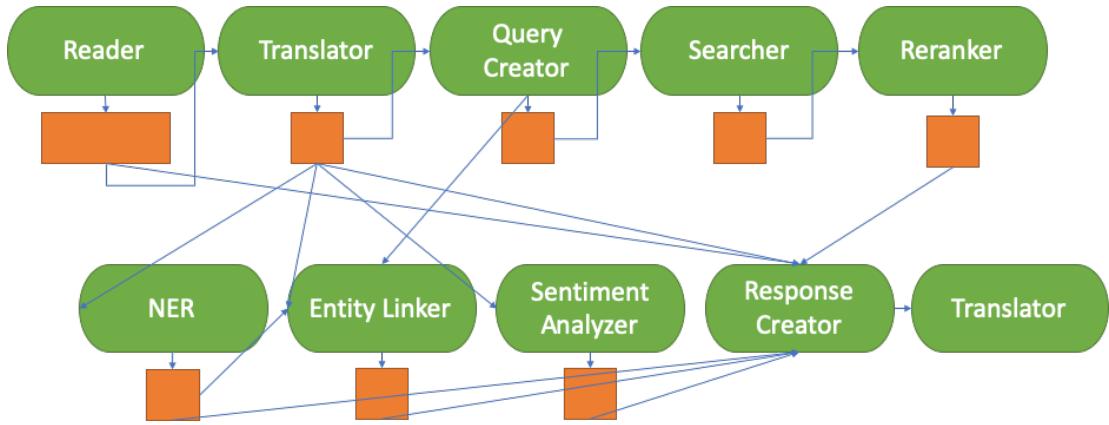
E.g. Different NLP data formats of different tasks



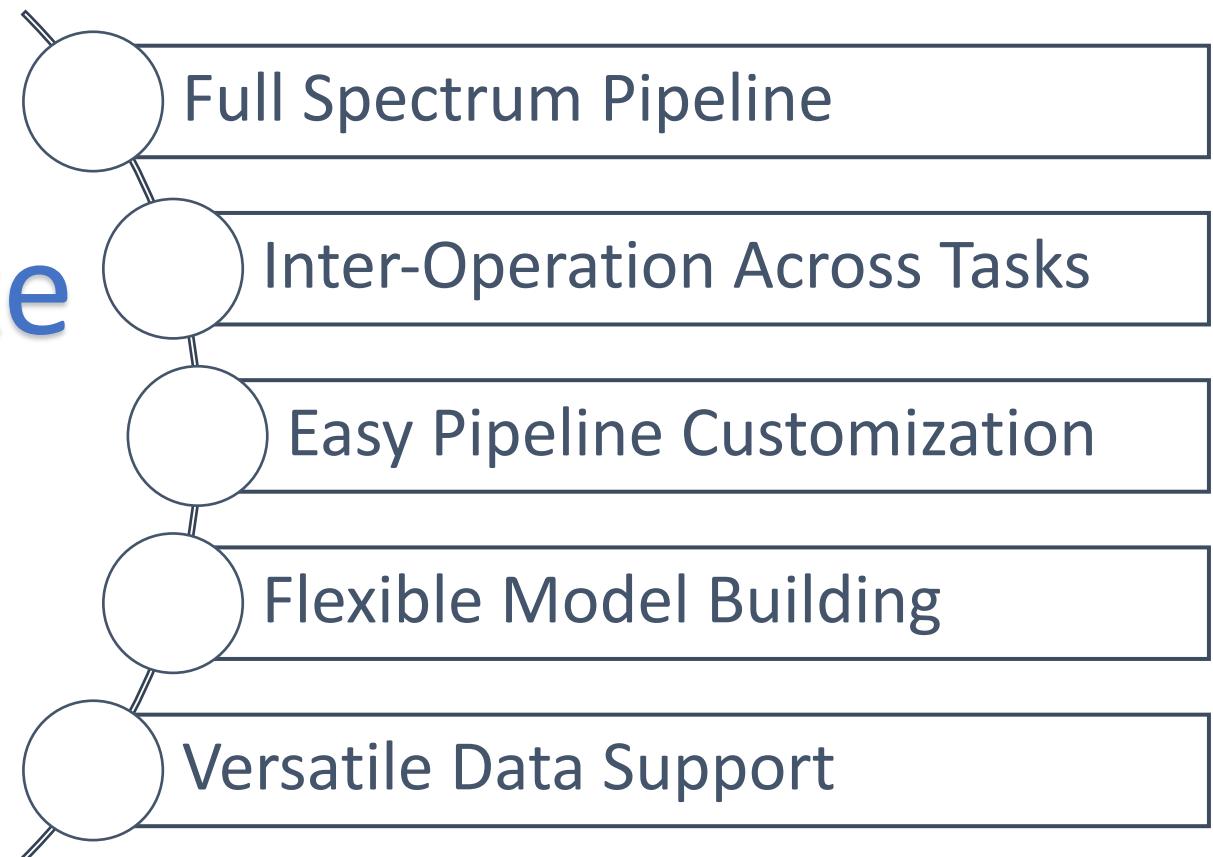
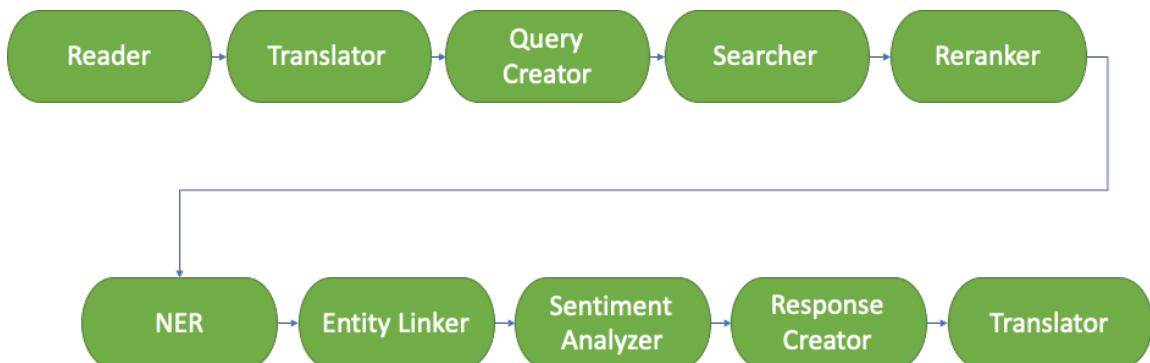


How?

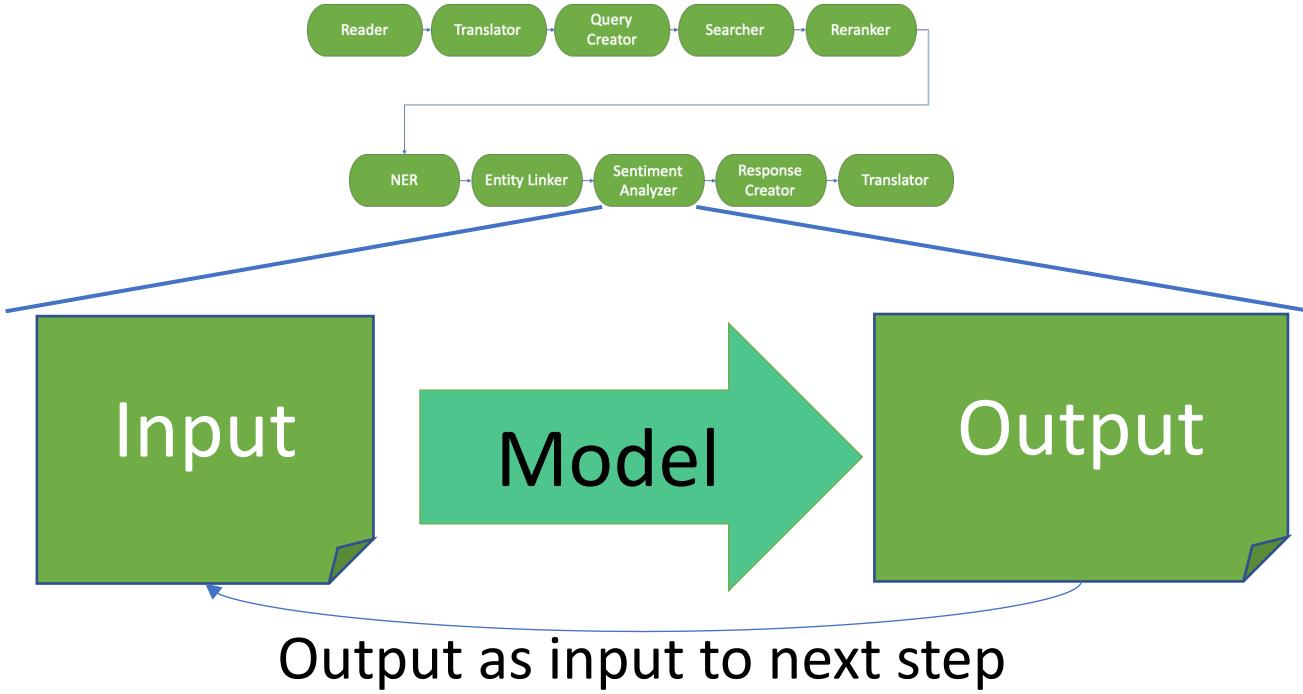




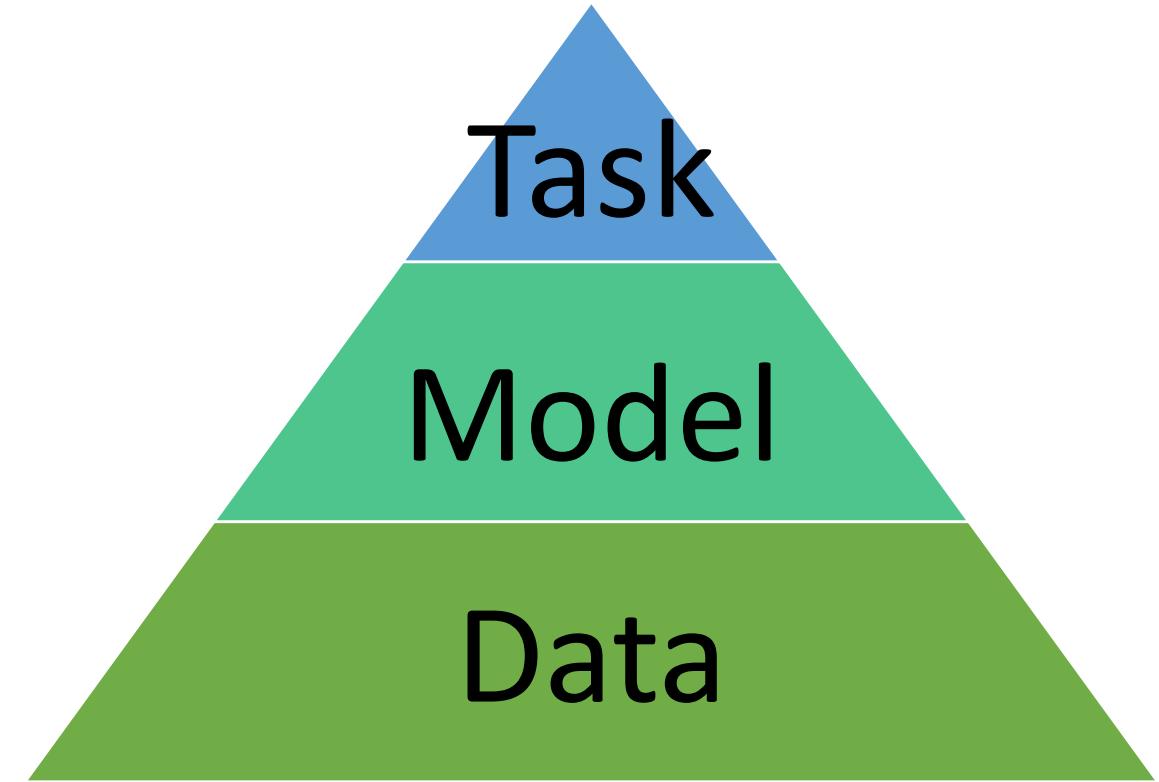
Standardize NLP



Standardize NLP Interfaces



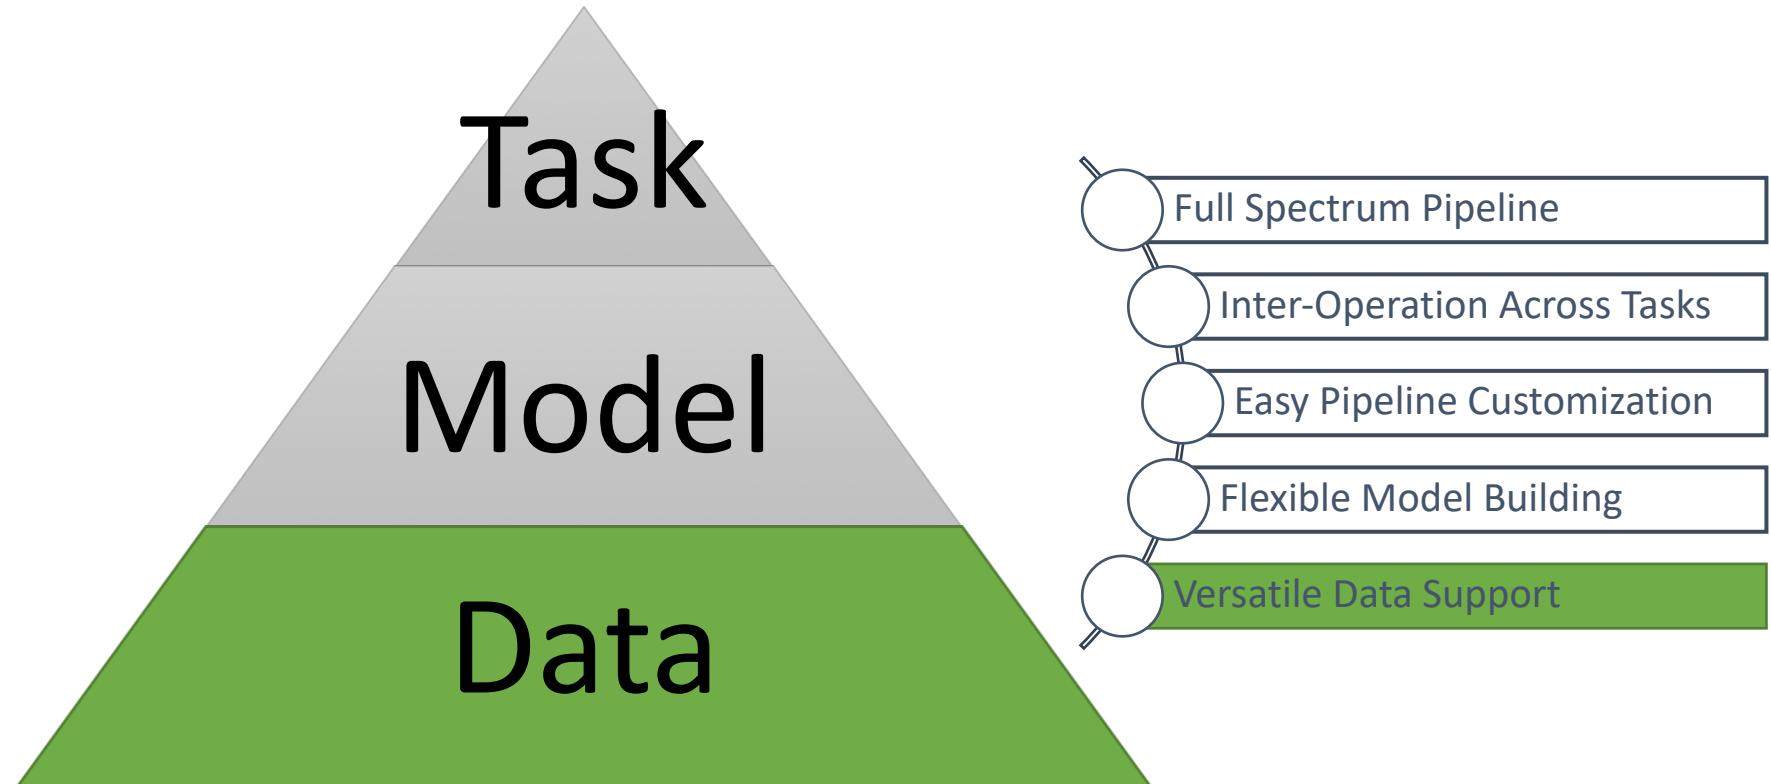
Operationalized View



Conceptual View

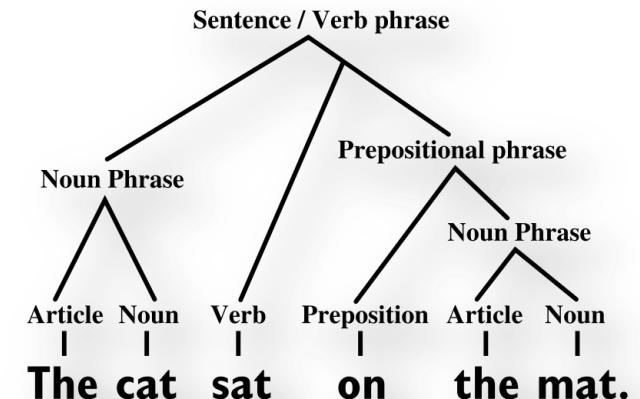
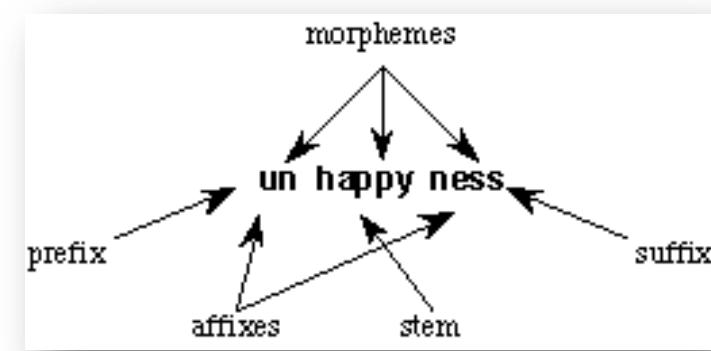
Standardize interfaces between the 3 levels

Considerations of Standardizing Data Representation

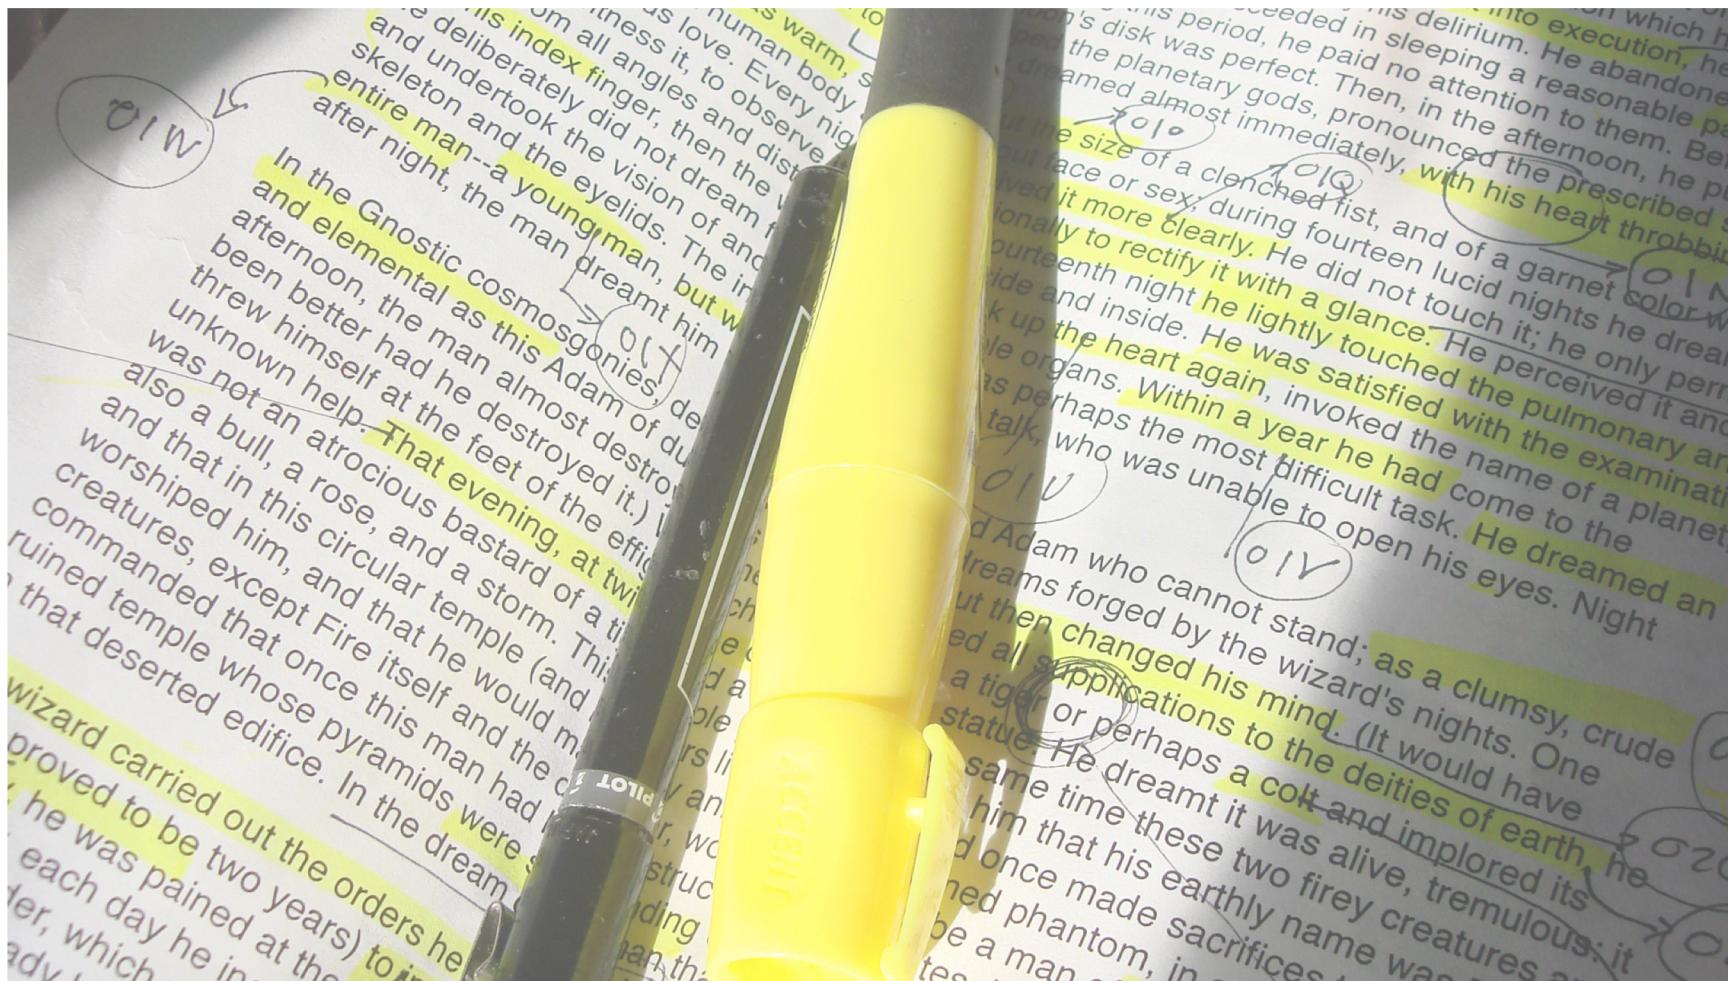


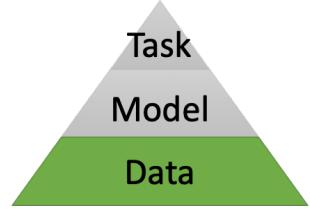
Data Granularity

- NLP data have different granularity:
 - **Character** level (Chinese Radical)
 - **Token** level (Morphology)
 - **Sentence** level (Syntax Parsing)
 - **Document** level (Entity Coreference)
 - **Corpus** level (Information Retrieval)
- Representation should handle flexible granularities



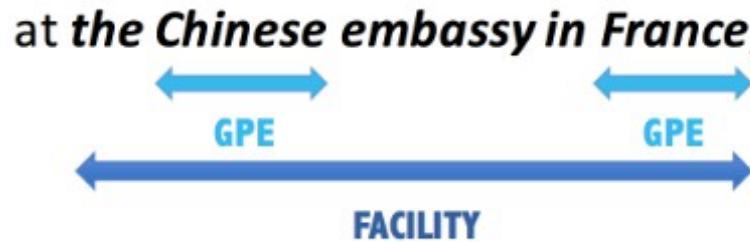
Data Structures





Data Structures

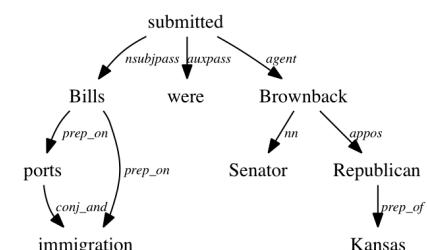
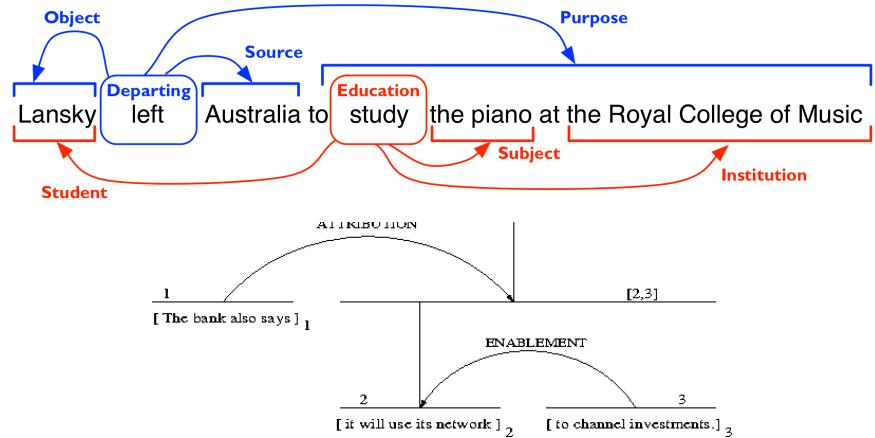
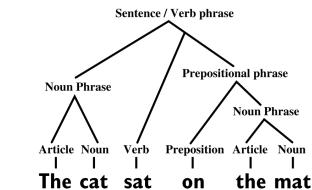
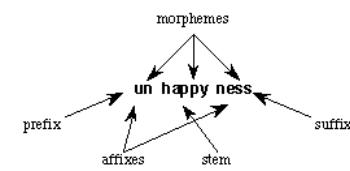
- Let's again look at the tasks, any patterns?



Stemming
adjustable → adjust
formality → formaliti
formaliti → formal
airliner → airlin

Lemmatization
was → (to) be
better → good
meeting → meeting

"I voted for Nader because he was most aligned with my values," she said.



Data Structures

- Let's again look at the tasks, any patterns?

Sequence Structure

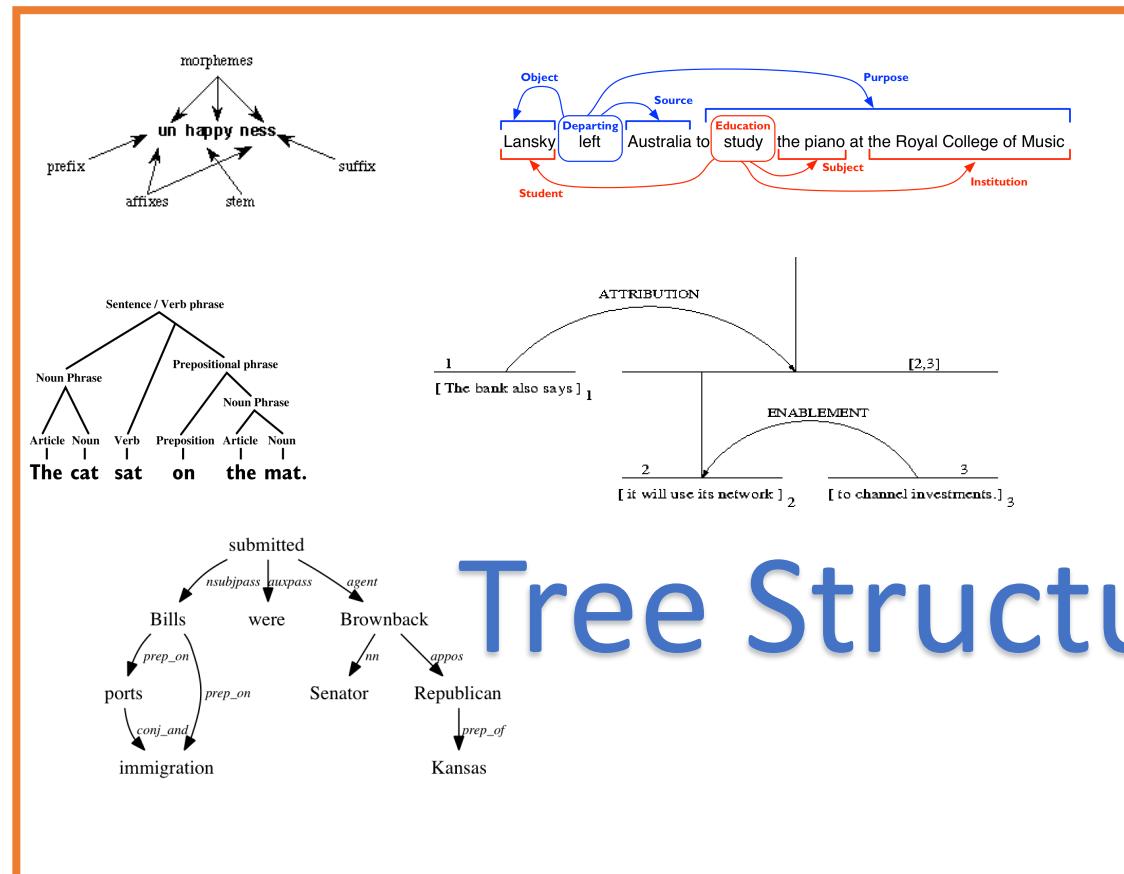
at the *Chinese embassy in France*,
 GPE GPE
 FACILITY

Stemming
 adjustable → adjust
 formality → formaliti
 formaliti → formal
 airliner → airlin

Lemmatization
 was → (to) be
 better → good
 meeting → meeting

Cluster Structure

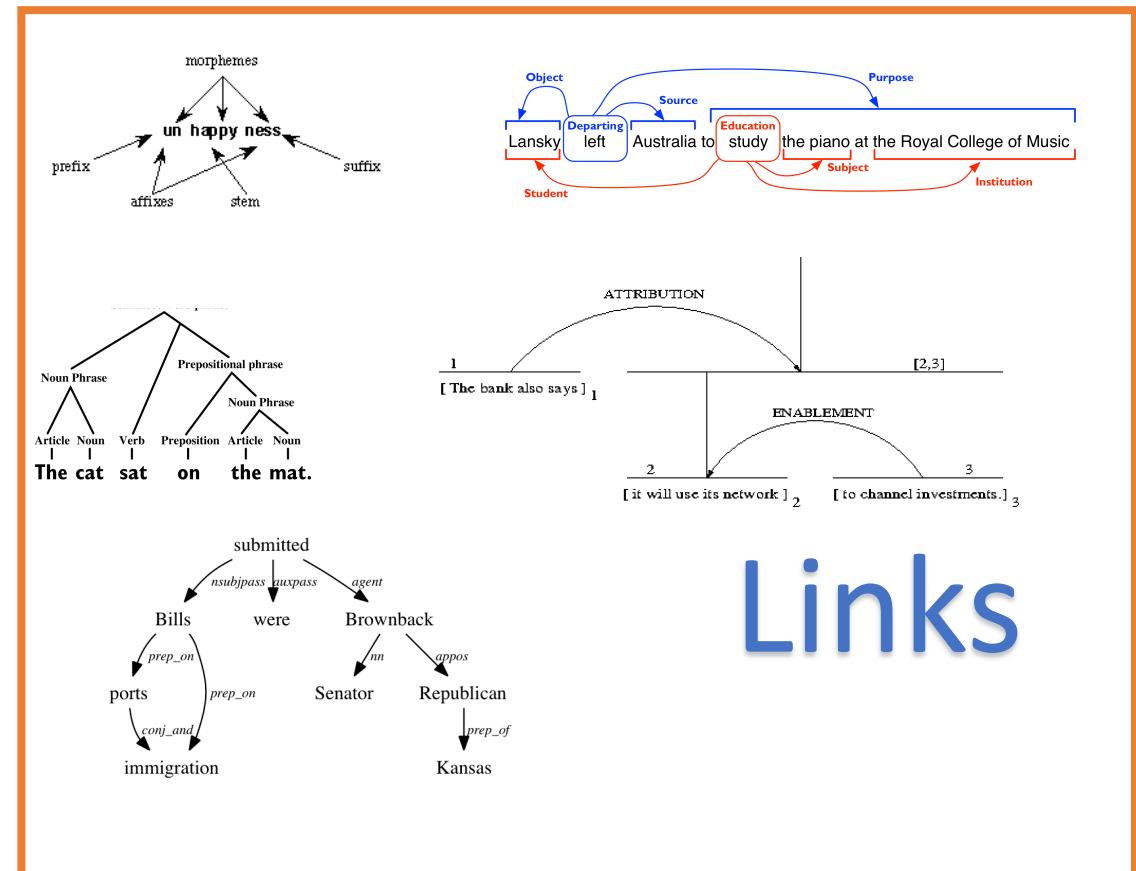
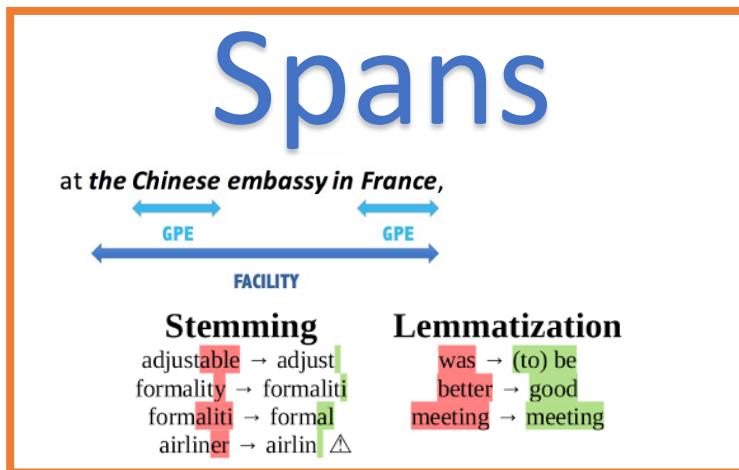
"I voted for Nader because he was most
 aligned with my values," she said.

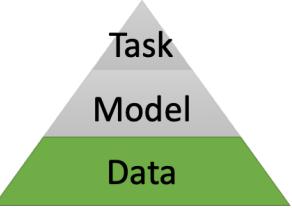


Tree Structure

Data Structures

- Let's again look at the tasks, any patterns?





Spans

additional fields, e.g., embeddings, ner tag

Span = (type, begin, end, ...)

E1 = (EntityMention, 25, 35)

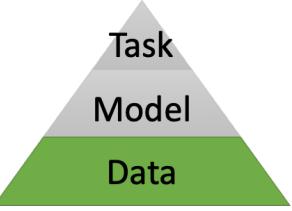
EntityMention

P2 = (Predicate, 42, 46)

Predicate

EntityMention

Last year, it was Rams quarterback Jared Goff, who failed to spot a wide-open Brandin Cooks when the NFC Champions ran in the second half a play that had sprung Cooks free in the first half.



Links

additional fields, e.g., dependency type

Link = (Type, Parent, Child, ...)

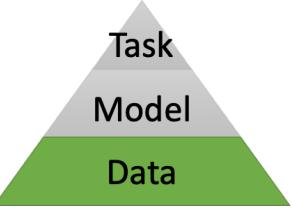
S1= (SemanticRoleLink, P1, E1)

EntityMention

Predicate

EntityMention

Last year, it was Rams quarterback Jared Goff, who failed to spot a wide-open Brandin Cooks when the NFC Champions ran in the second half a play that had sprung Cooks free in the first half.



Groups

additional fields, e.g., group type, score

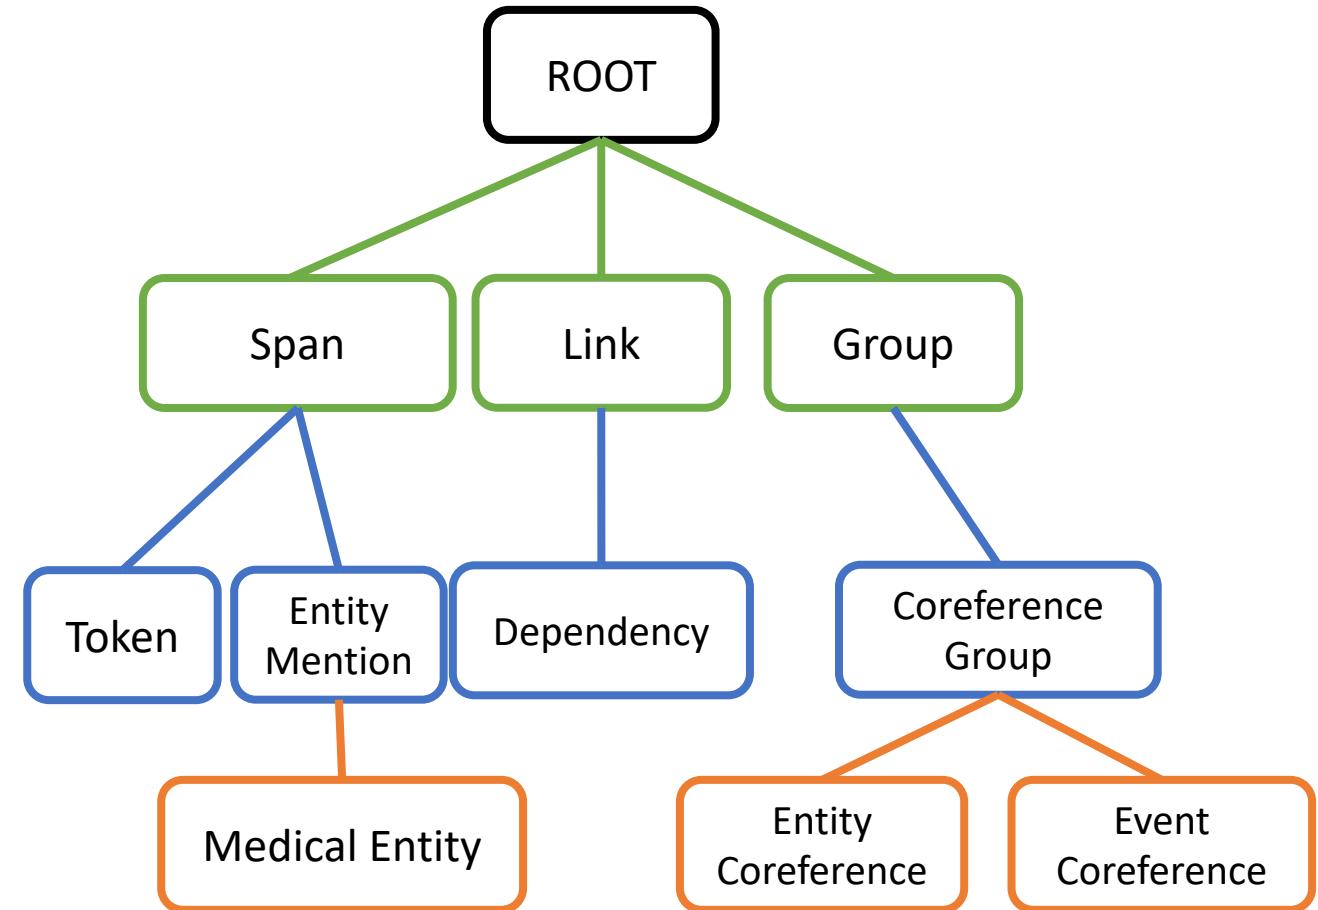
Group = (Type, Member List, ...)

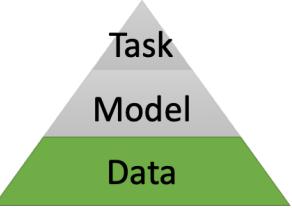
G1= (EntityCoreferenceGroup, Member=[E2,E3])

Last year, it was Rams quarterback Jared Goff, who failed to spot a wide-open Brandin Cooks when the NFC Champions ran in the second half a play that had sprung Cooks free in the first half.

Ontology of Data Structures

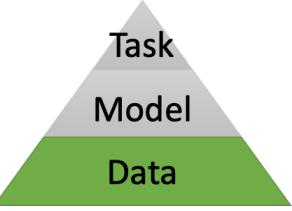
- Types of data structures can be organized as an ontology tree.
- The ontology can be customized for new domains.





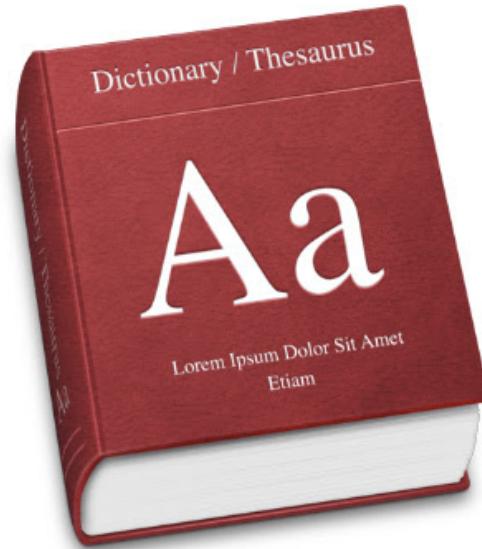
Practical Data Considerations

- Serialization (S) and Deserialization (D)
 - Loseless: $D(S(\text{data})) = \text{data}$
 - Can be passed around (e.g. networks channels)
- Readable, Interpretable
- Meta Data
 - Keep source information
 - Automatic record creation time, creator, etc.

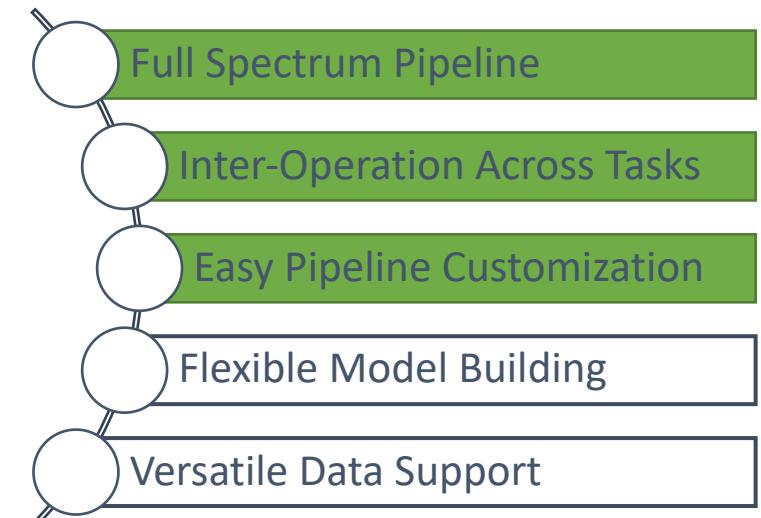
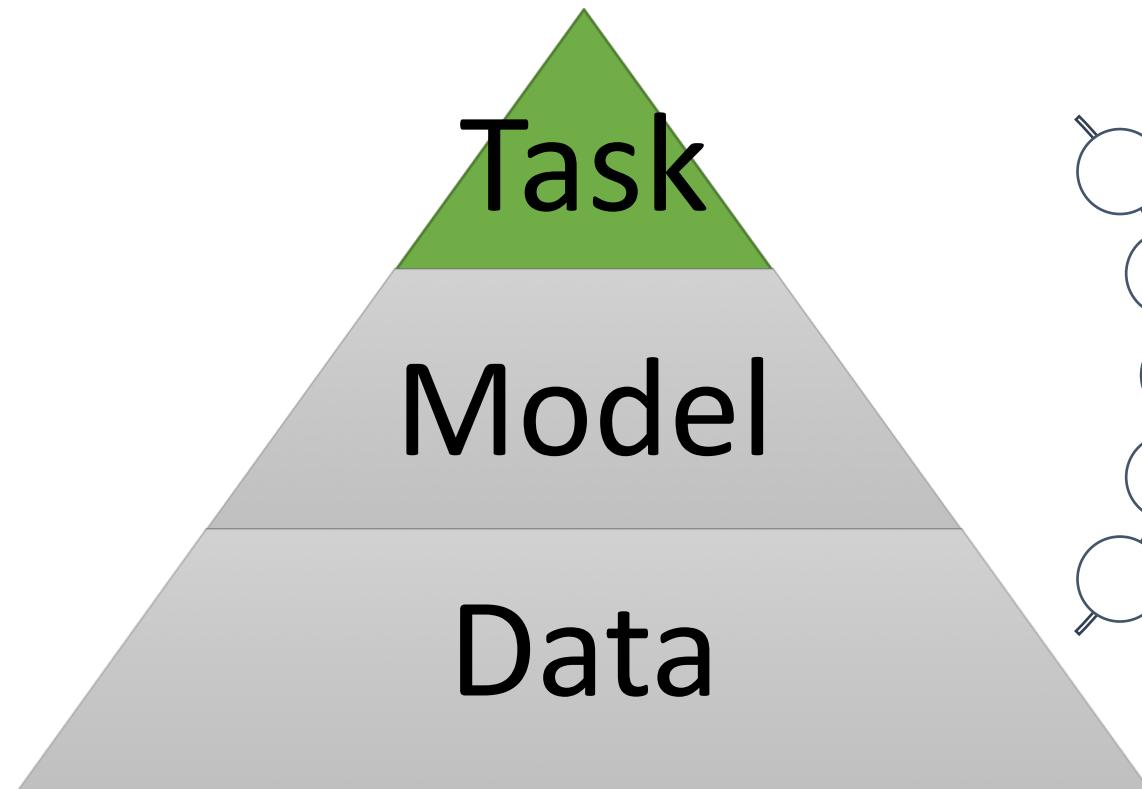


Global Data

- Vocabulary
- Embeddings (e.g. Word2Vec, Bert)

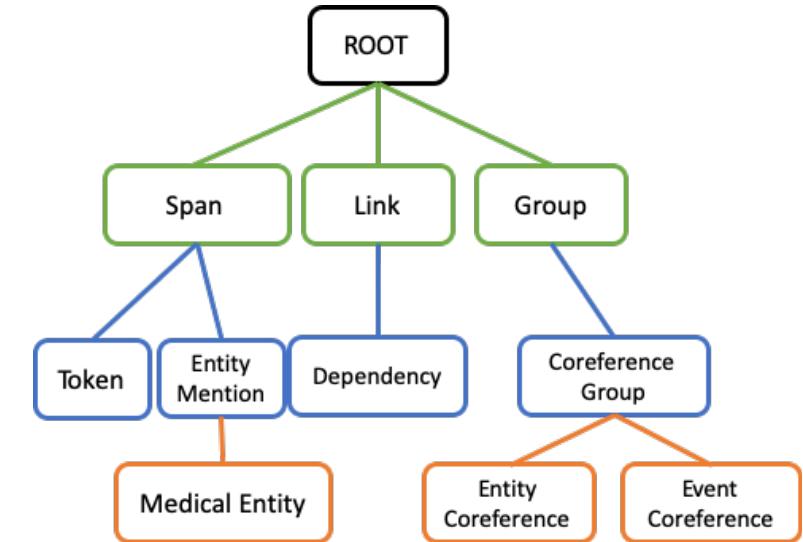


How about Tasks?



Data as Interface

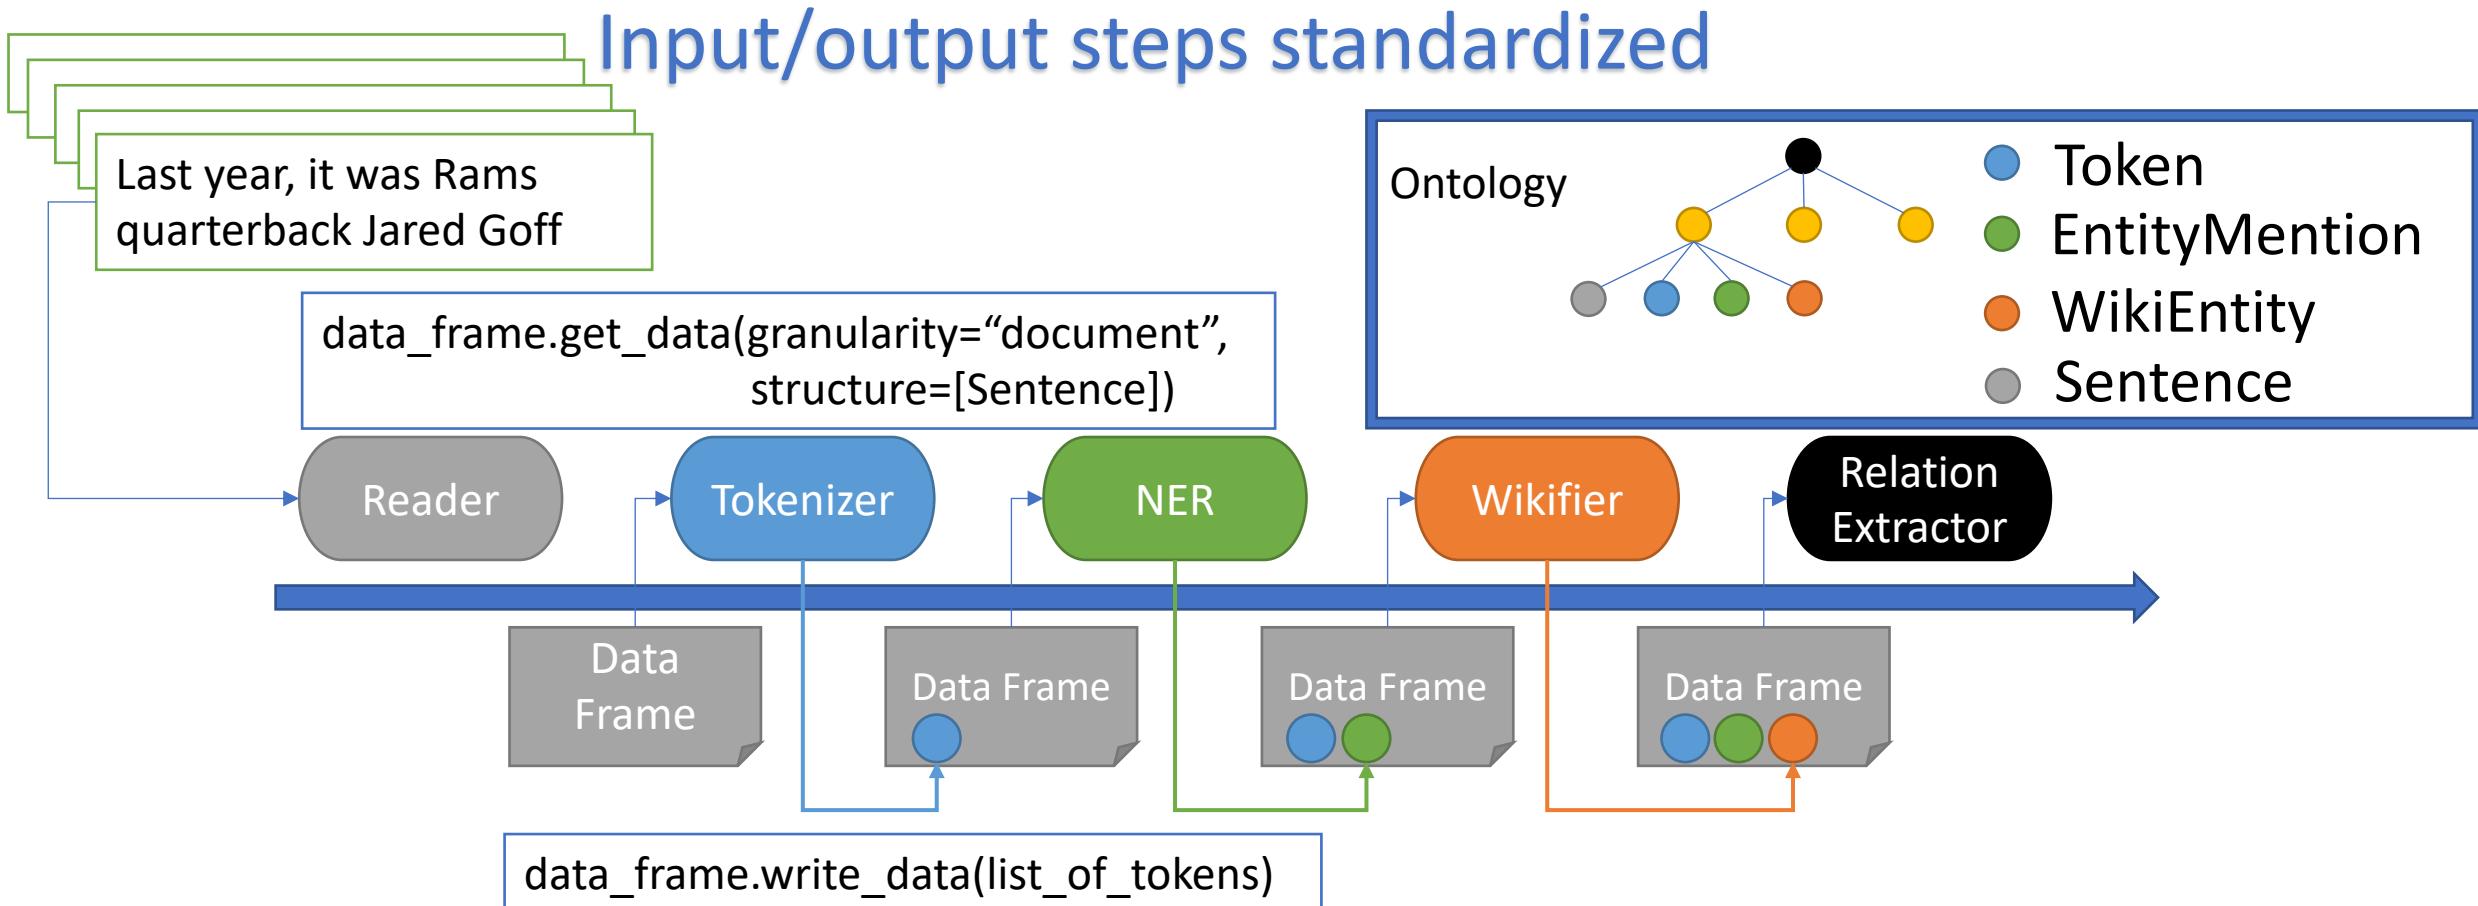
- A standard interface can be defined given the data representation.
- This helps task modularization.



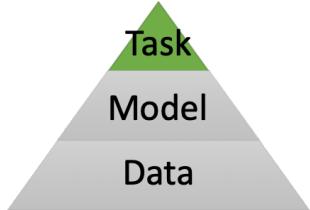
Accessing data using the granularity and structure

```
get_data(granularity="sentence",
          structure=[PredicateMention, EntityMention])
```

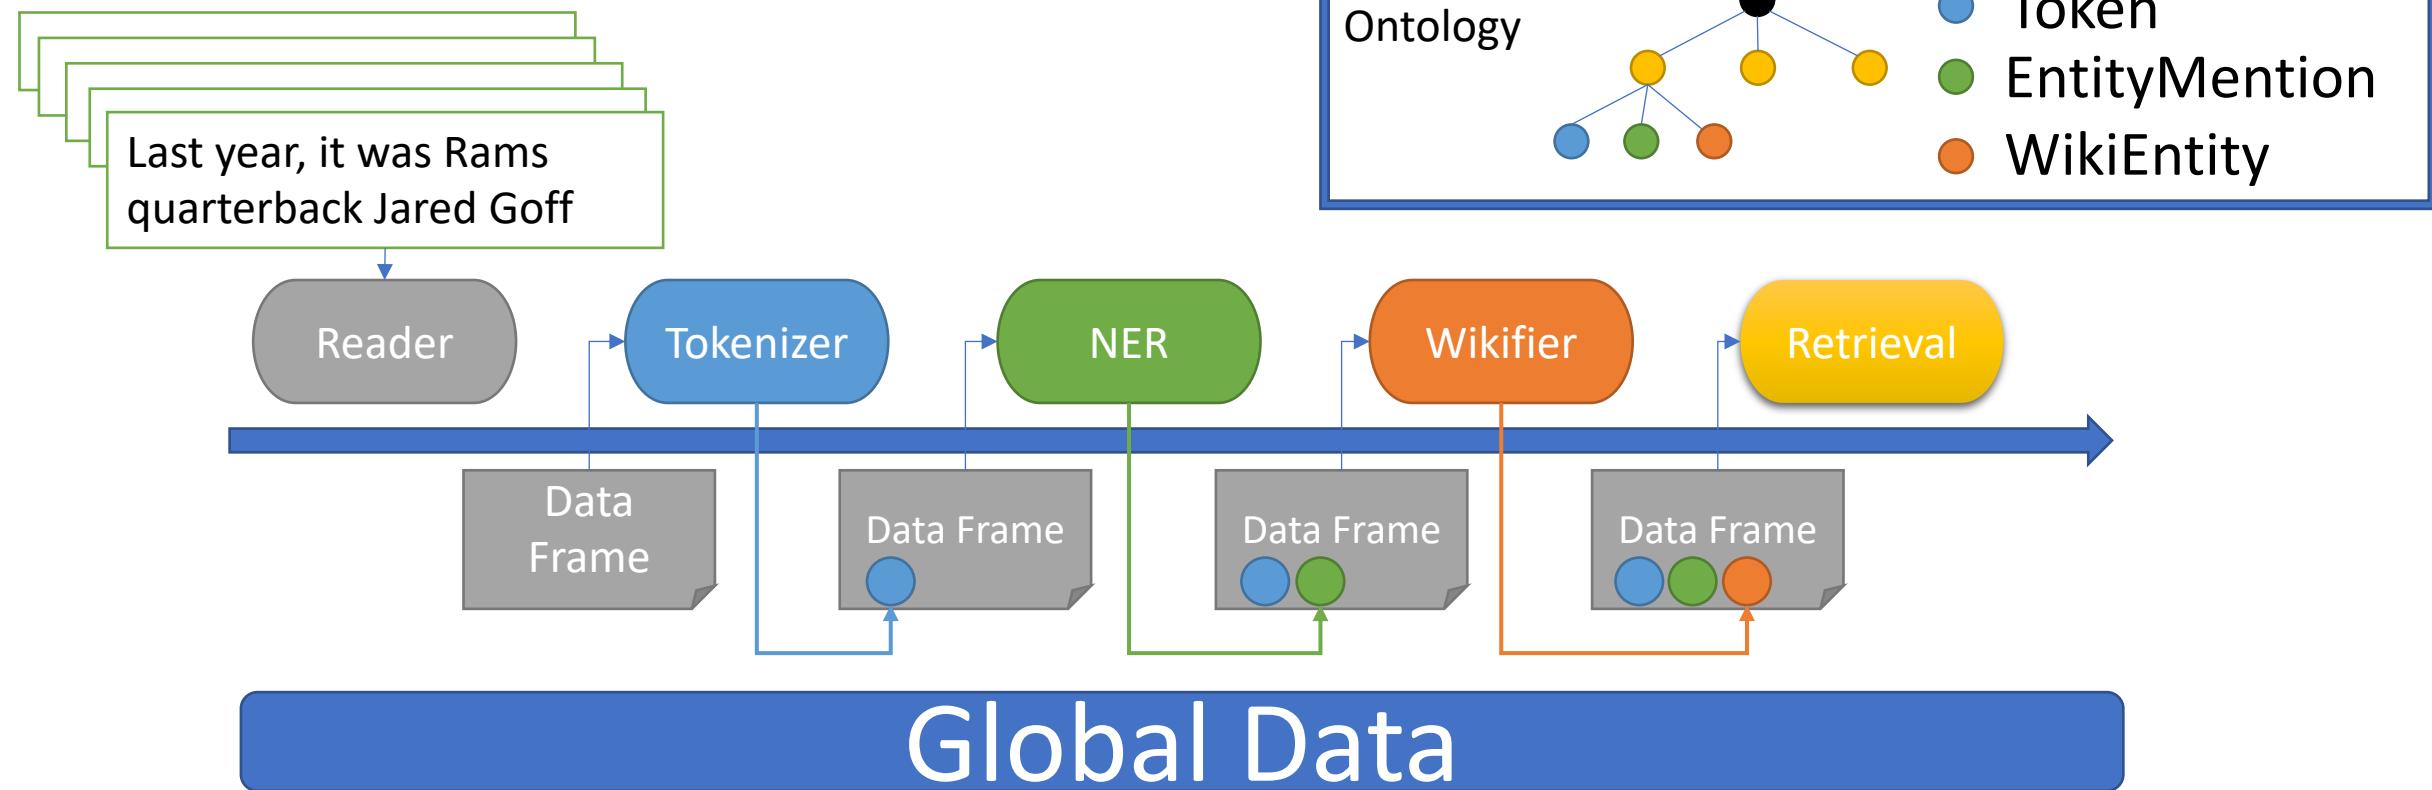
- With Standardized Data and Interfaces:
 - Tasks can already be organized in a modular way



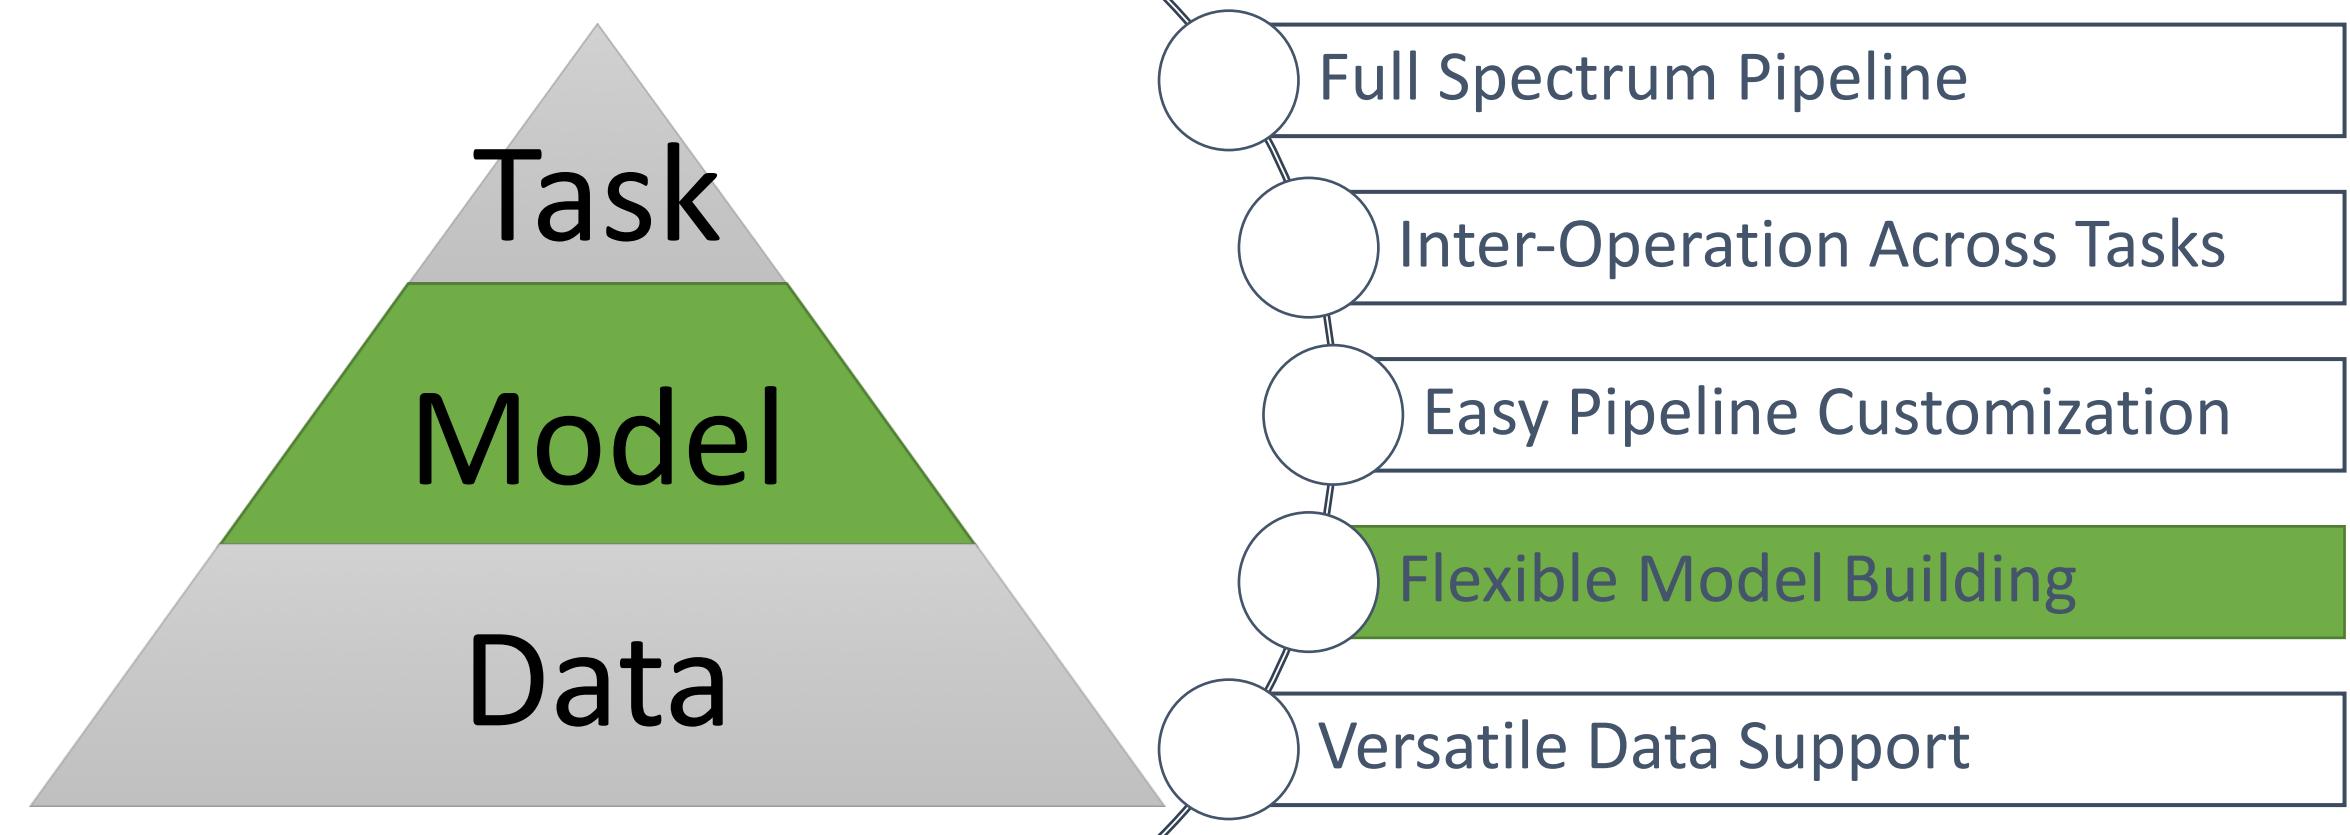
Global Data

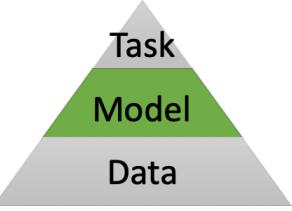


- Full Spectrum Pipeline Customization: We can insert/replace an Information Retrieval into the pipeline
- Inter Operation: Complex queries can be built using the NER and Wikifier information

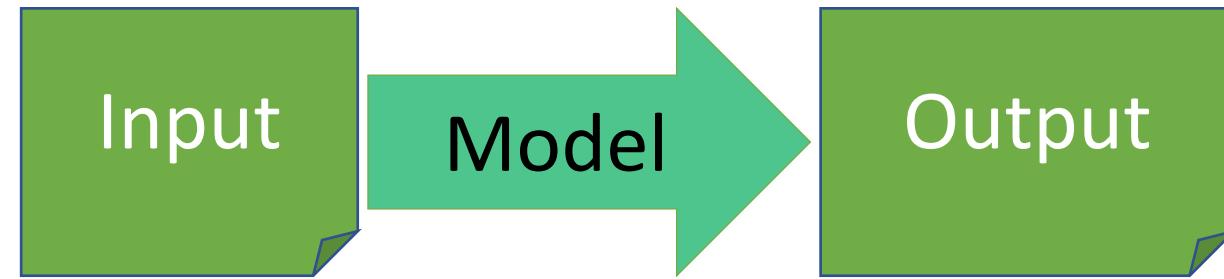


How to Standardize Model?



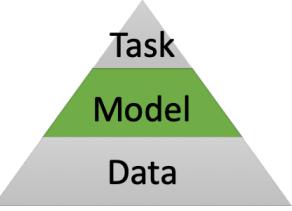


He is **going** to the book shop.



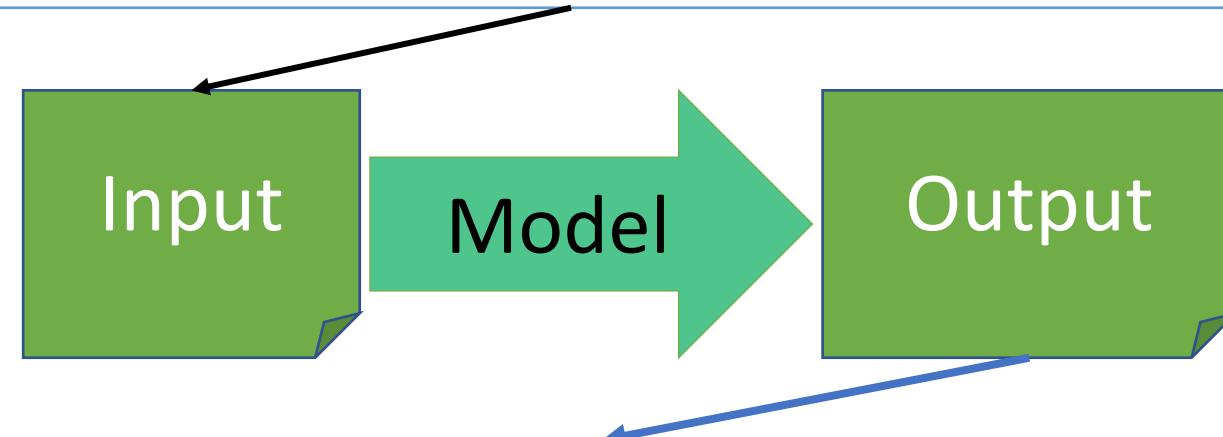
ARG0 ARGM-LOC

He is **going** to the book shop.



He is **going** to the book shop.

```
data_frame.get_data(granularity=sentence,  
                    structure=[PredicateMention, EntityMention])
```



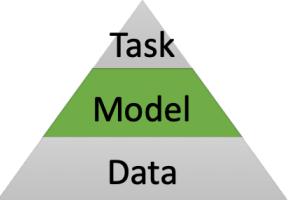
We already have a standard interface between the data and the model

```
data_frame.write_data(granularity=sentence, structure=PredicateLink)
```

ARG0 ARG-M-LOC

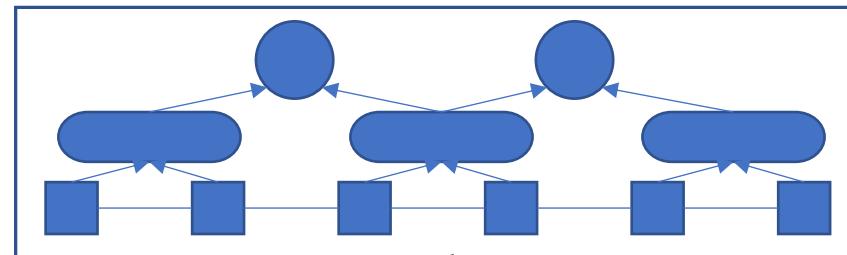


He is **going** to the book shop.



Model-Data Interface

By changing a few parameters, we can fit to different models.



Span-Link Model



LSTM

```
get_data(context=sentence,
         predicate_link,
         field=[role])
```

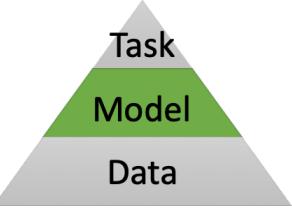
```
get_data(context=sentence,
         relation_link,
         field=[relation_type])
```

```
get_data(context=sentence,
         token, fields=[pos])
```

```
get_data(context=sentence,
         token, fields=[ner_tag])
```

ARGO
He is **going** to the **bookshop**.
PR V V IN DT N

Colleague
Tom and **Bill** work at the same company.
PER O PER O O O O O



Other Model Support

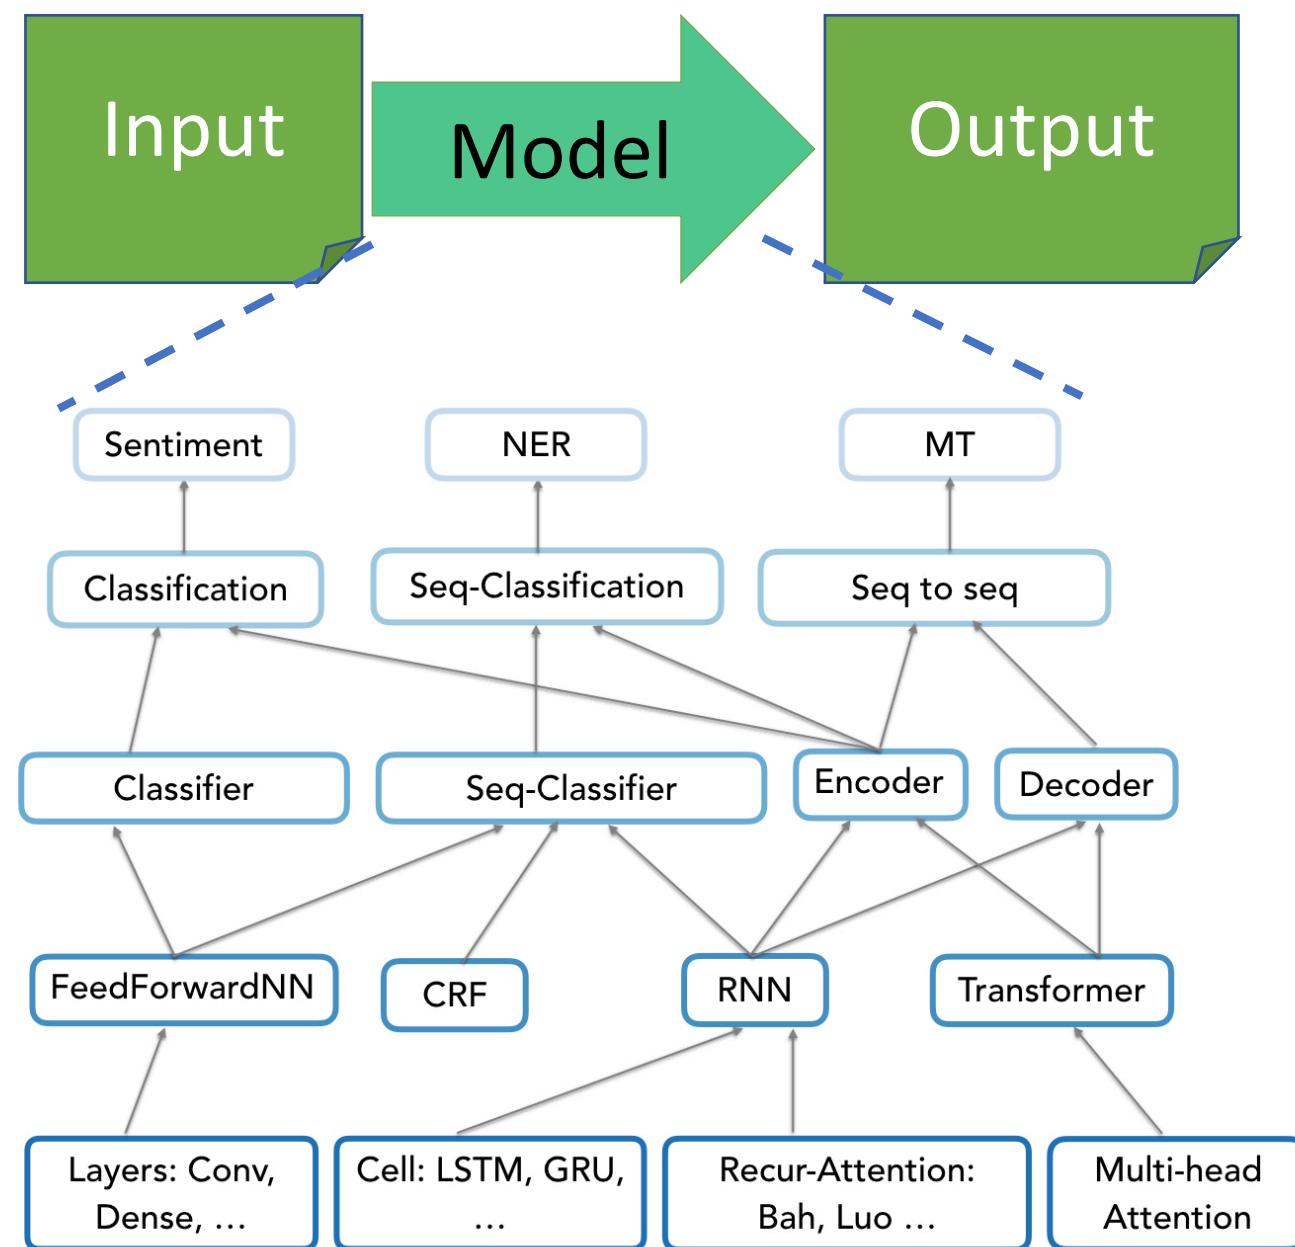
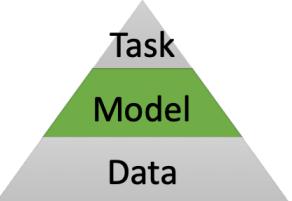
- Feature Extraction
 - E.g. Embedding support
- Training Loop Support
 - Data Iterator
 - Batcher
- Learning Library Interface

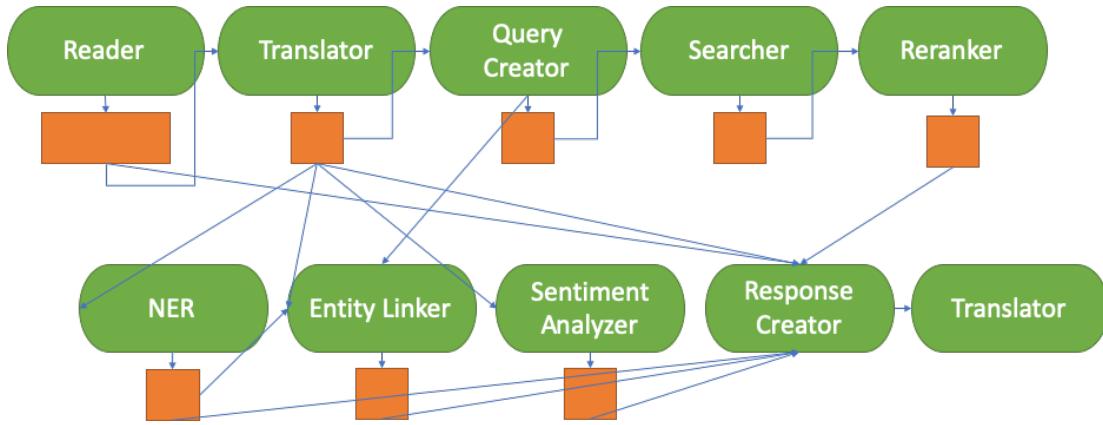
```
vector = sentence.embedding(type='BERT')
```

```
batcher = FixSizeBatcher(num_instance=8)

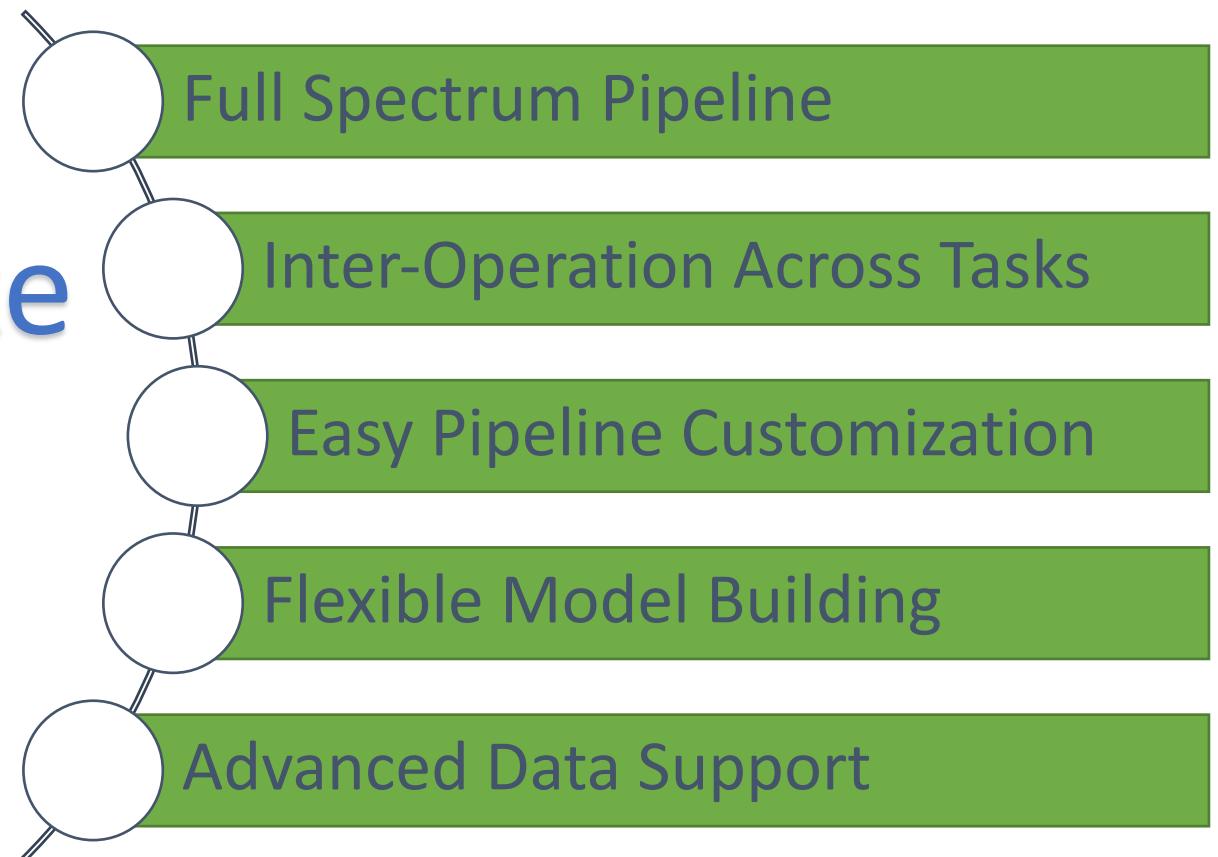
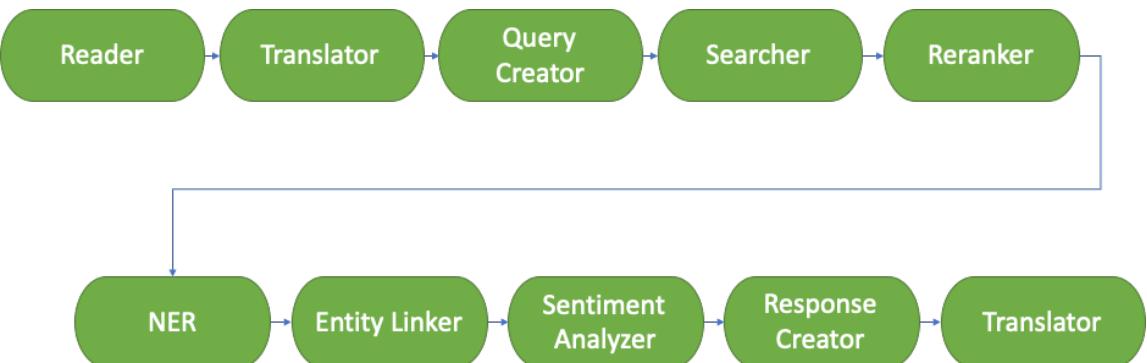
for data in get_data(context=sentence, relation_link,
                     field=[relation_type]):
    batcher.add(data)

    if batcher.ready():
        batcher.yield()
```



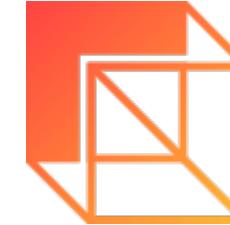


Standardize NLP



Case Study: Modularization with *Forte*

Forte: a flexible and Powerful NLP Pipeline FOR TExt.



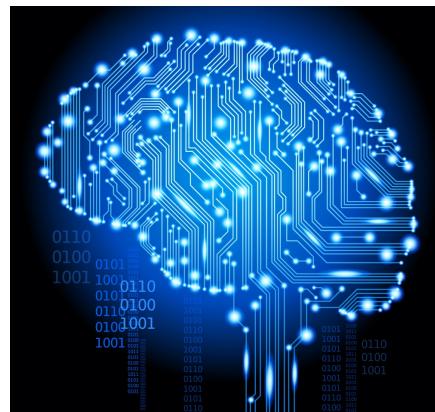
Forte

<https://github.com/asymly/forte>

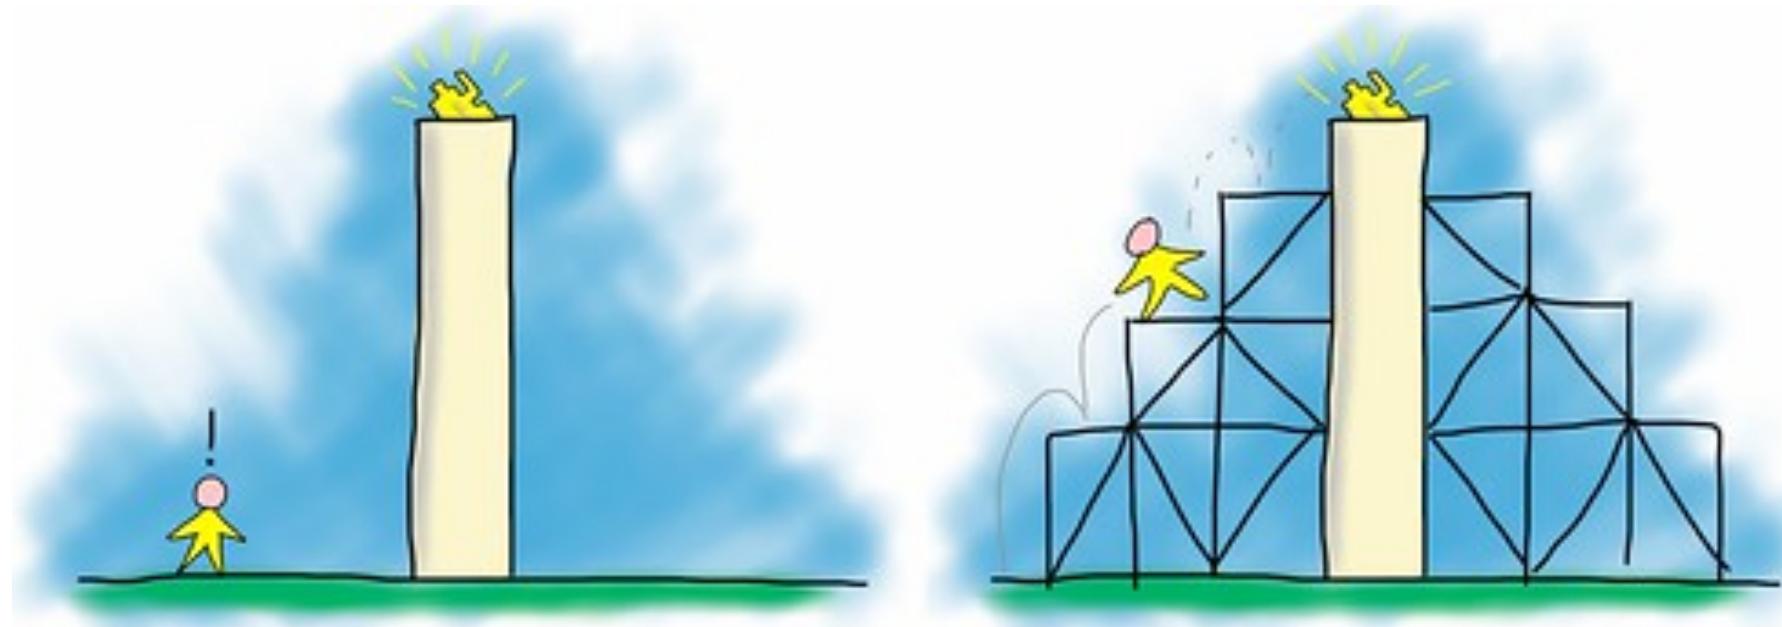
- We built **Forte**, a system that validates these ideas:
 - Universal Data Flow
 - Pipeline Construction
 - Full Spectrum: Information Retrieval, Text Analysis, Generation
 - Abstract Data, Model, Task Interfaces
- Other features:
 - Batching
 - Bookkeeping
 - ...

Revisit the Problem

- A user **speaks German** but would like to find good **romantic movies**.
- We have a **corpus of English movie reviews**.
- What can we do?



Forte: Scaffold for the Goal



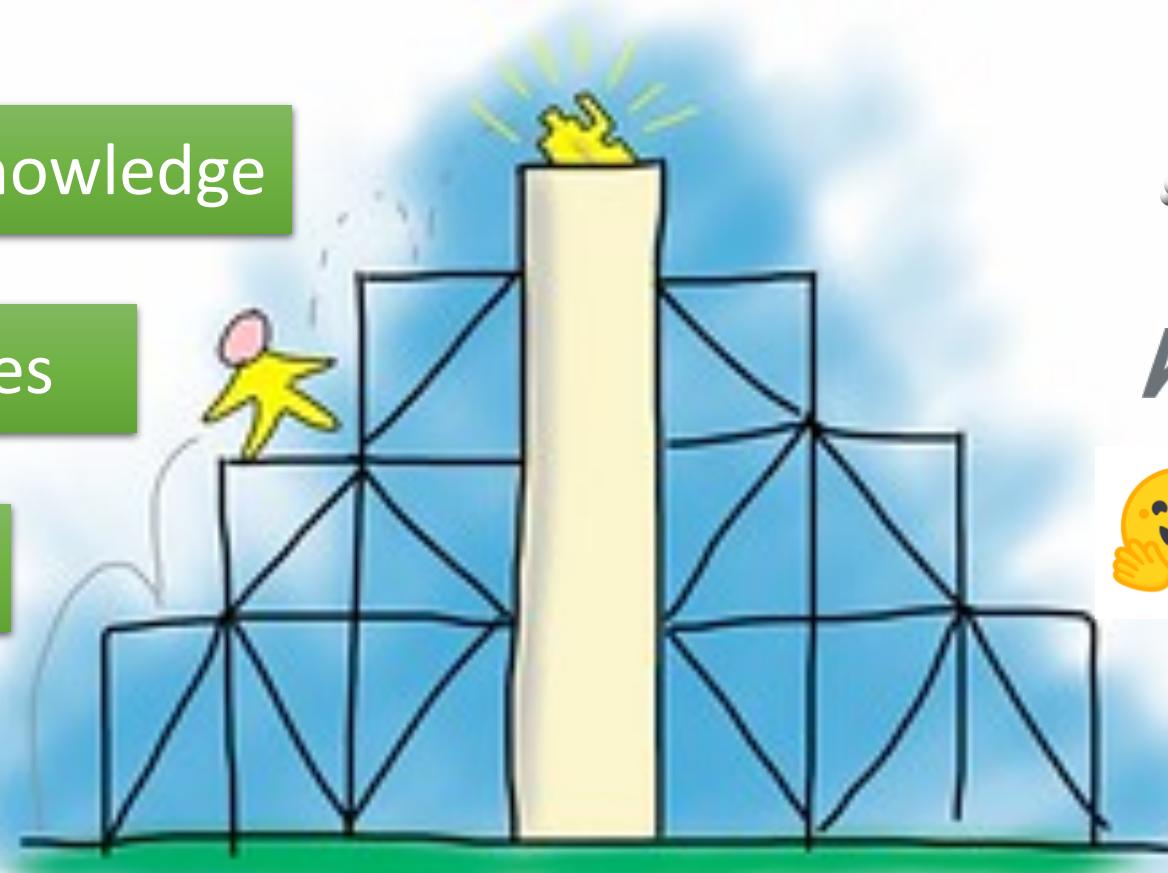
An Open, Flexible Scaffold for NLP

Domain Knowledge

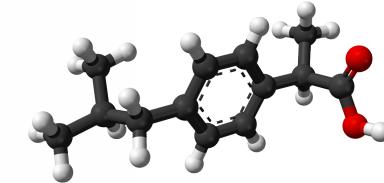
NLP Techniques

Pre-train Models

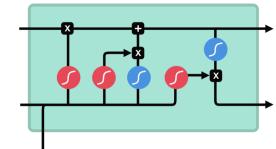
ML Library



Data Representation



AllenNLP



Texar

PYTORCH

	Forte	spaCy	CoreNLP	DKPro	AllenNLP	Curator
Universal Data Flow	😊	😐	😊	😐	😡	😊
Extendable Ontology	😊	😡	😡	😊	😡	😡
Lossless Serialization	😊	😊	😡	😊	😡	😊
Pipeline Construction	😊	😊	😊	😊	😡	😊
Easy Processor Replacement	😊	😐	😐	😊	😡	😊
Inter Operation	😊	😐	😊	😡	😡	😊
Retrieval	😊	😡	😊	😡	😡	😡
Generation	😊	😡	😡	😡	😊	😡
Standard Model Interface	😊	😡	😡	😡	😐	😊
Deep Learning Integration	😊	😊	😐	😡	😊	😡
Integrated ML Support	😊	😐	😐	😡	😊	😊



Well Supported!



Some Support



No Support



Standardization with Forte



Petuum Med

Name Here Name Here

Name Here

The patient is a 62 year old male with a history of mild COPD who complains of cough, shortness of breath, fatigue, and fever progressively worsening for the past week. Today he measured a fever of 101 F. The productivity of his cough has progressively increased over the past two days. He has been using his albuterol inhaler two to three times daily but it is helping only minimally. He has a history of COPD, which he believes is only mild. He states that he is treated with antibiotics for a case of bronchitis or pneumonia almost every year by his primary care provider. He has never been hospitalized for pneumonia. He received approx 80mg IV of Lasix at that time and was 2.4L negative in 24 hours. He has never been hospitalized for pneumonia. He denies any known sick contacts recently. He denies chest pain but admits to some chest tightness and an increase in heart rate when he coughs a lot and is short of breath. He denies any recent weight changes or lower extremity pain or swelling. He denies any recent travel. He denies a history of lung disease, heart disease, or diabetes. He currently is a non-smoker but did smoke a pack a day for approximately 15 years prior to quitting five years ago.

Name Here

Discharge Medications:

1 Furosemide 20 mg PO Tablet (2 times a day).

Critical Information Extraction

Heart Failure

History Comorbidities
 Symptoms Cardiac Tests
 Lab Tests Medications

Symptoms

cough 3 , shortness of breath 1 , fatigue 2 , fever 1 , chest pain 1 , chest tightness 4 , increase in heart rate 2 , weight change 1 , lower extremity pain 2 , swelling 3

Medications

Furosemide 2 , antibiotics 3

Show context when clicked Clear

Lasix, Vol: 80 mg, Usage: IV

Furosemide, Vol: 20mg, Usage: PO (2 times a day)

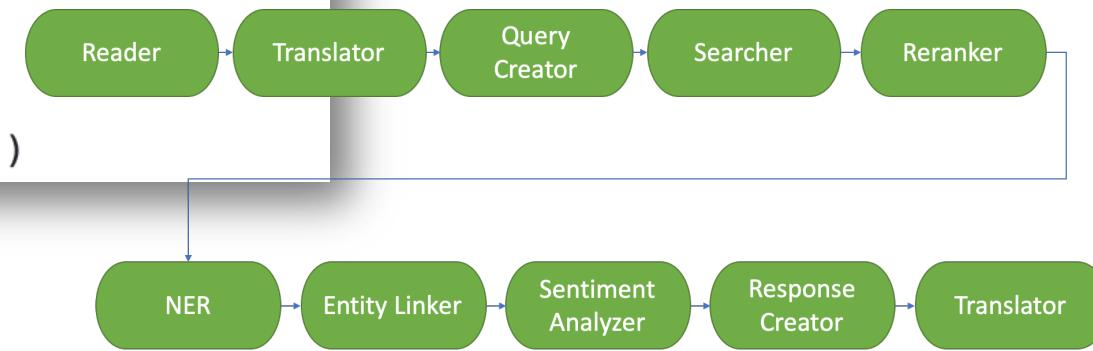
Chaining with Forte

```

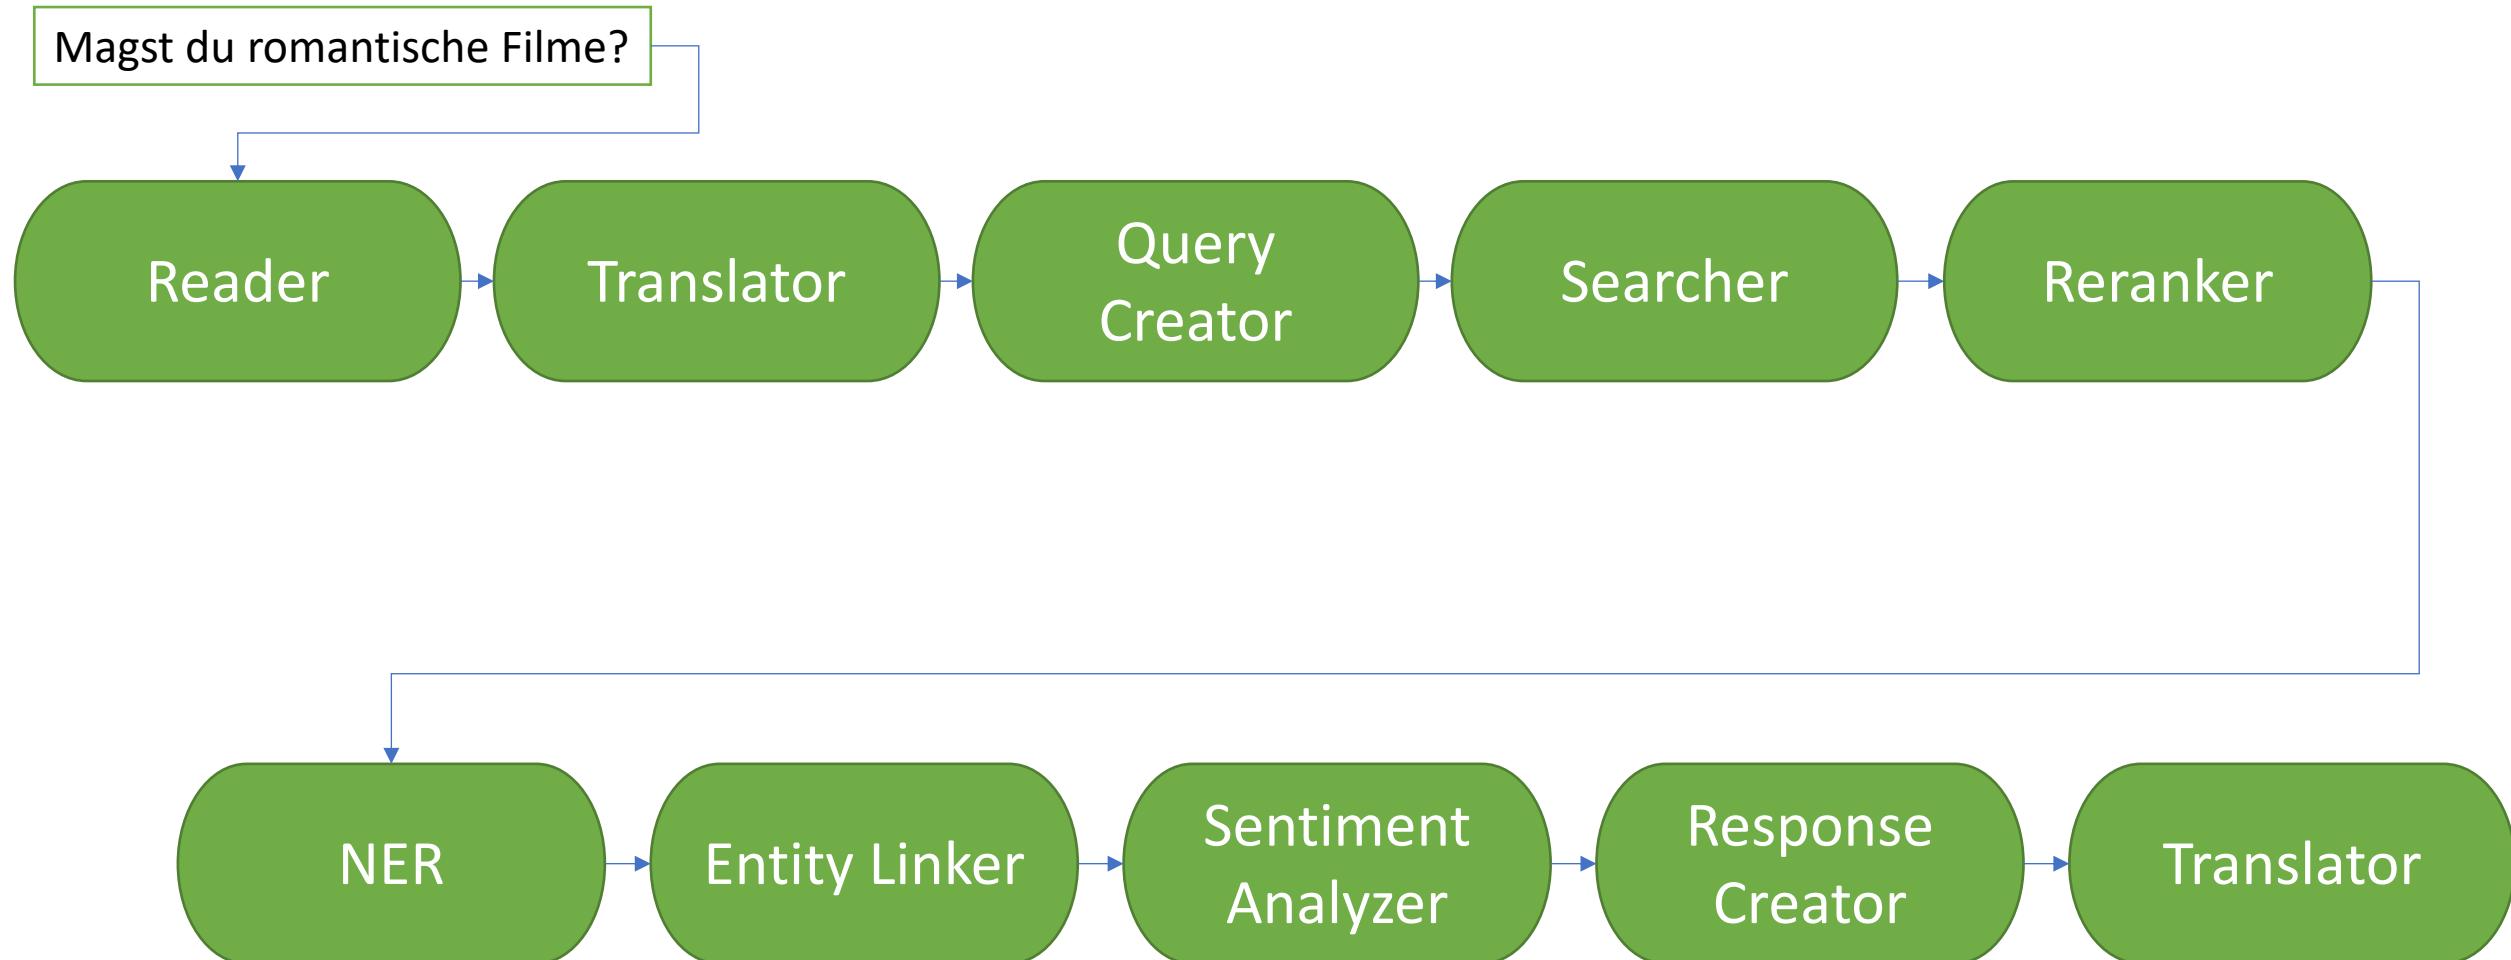
query_pipeline = Pipeline(resource=resource)
query_pipeline.set_reader(
    reader=MultiPackTerminalReader(), config=config.reader)

query_pipeline.add_processor(
    processor=MicrosoftBingTranslator(), config=config.translator)
query_pipeline.add_processor(
    processor=BertBasedQueryCreator(), config=config.query_creator)
query_pipeline.add_processor(
    processor=SearchProcessor(), config=config.indexer)
query_pipeline.add_processor(
    processor=CoNLLNERPredictor(), config=config.ner,
    selector=NameMatchSelector(
        select_name=config.indexer.response_pack_name[0]))

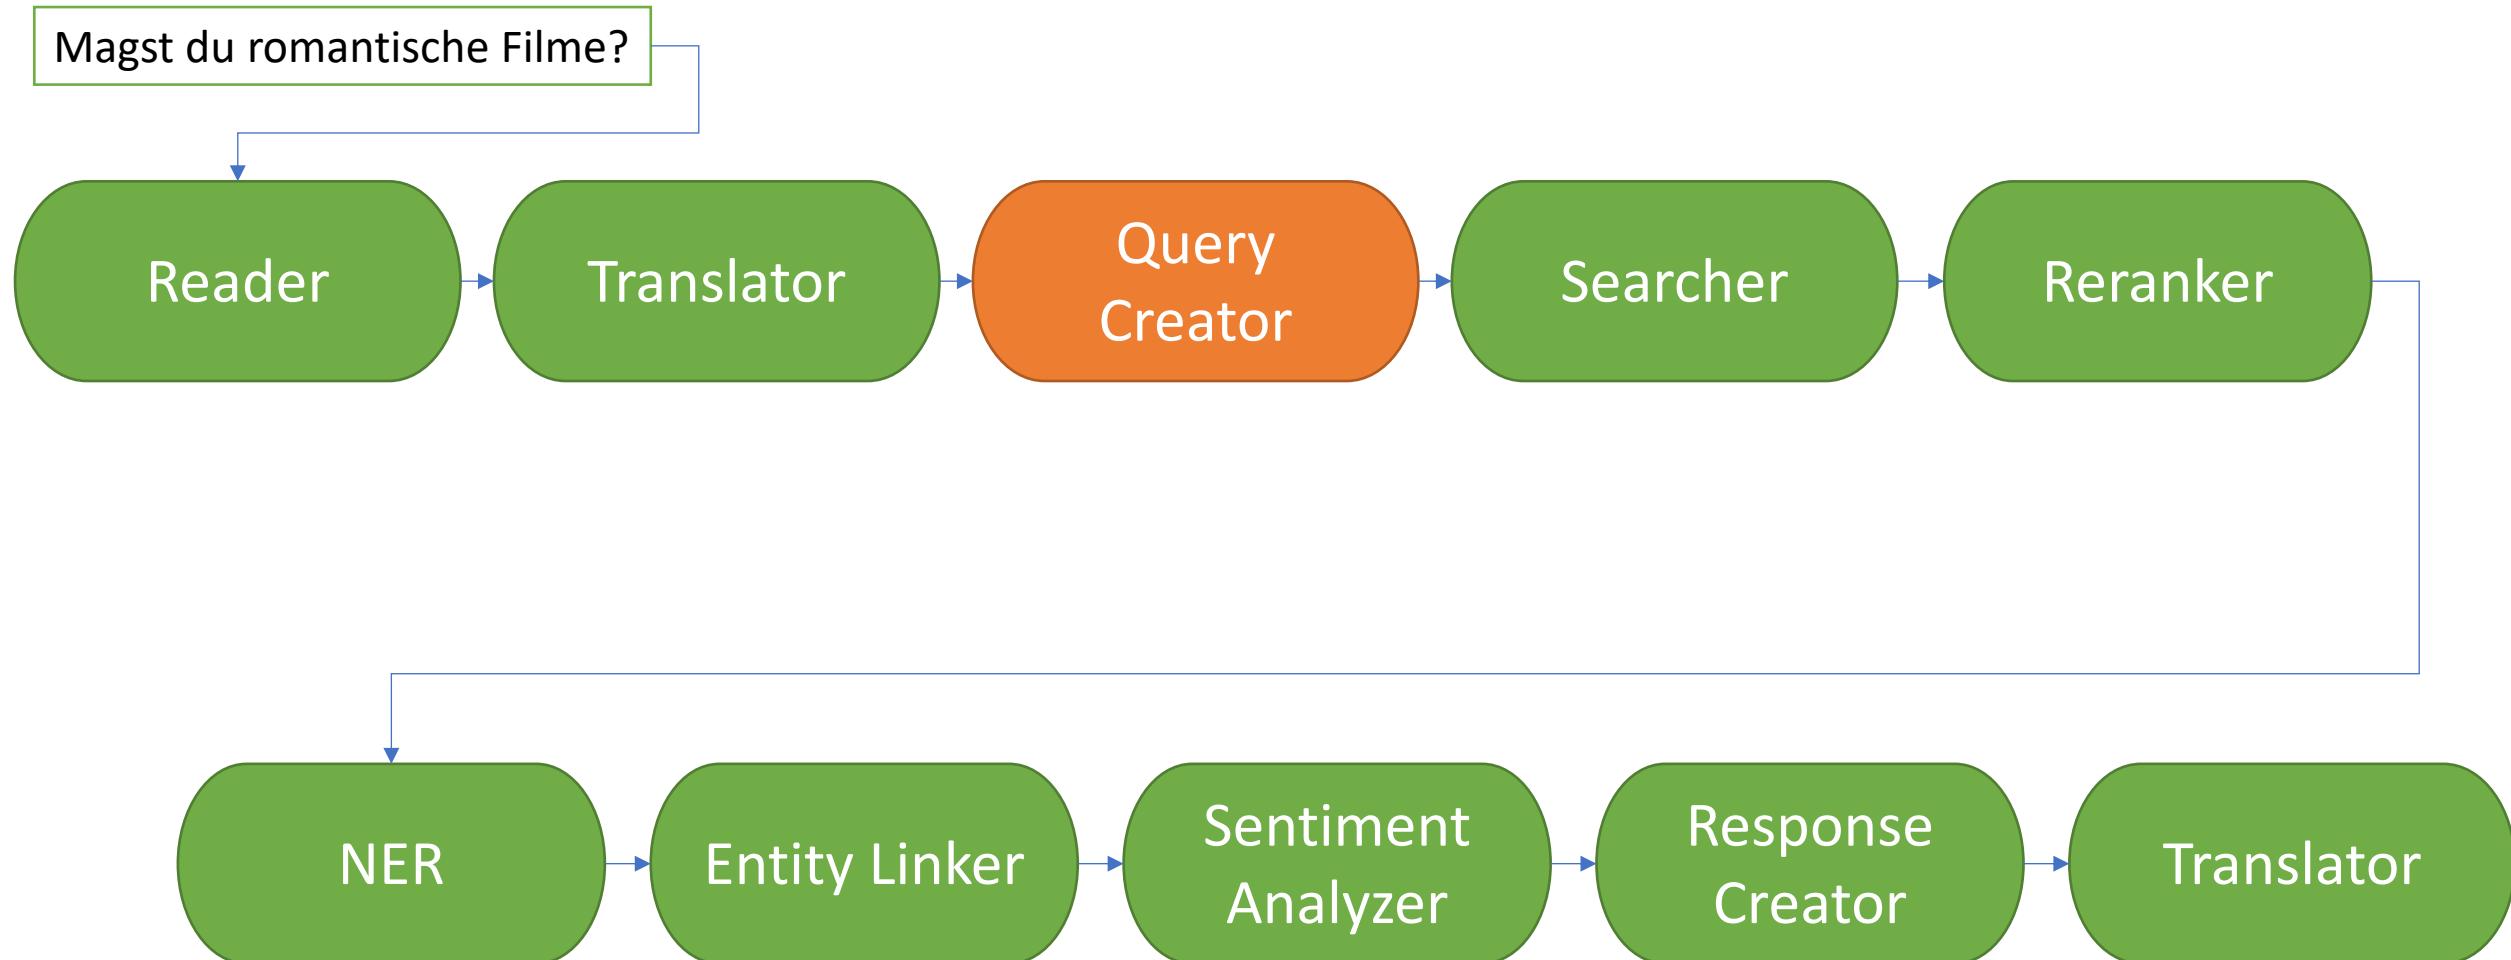
```



Implement Processors



Implement Processors



Query Creator

Here we create a simple text-based query.

Build Simple Query

```
class ElasticSearchQueryCreator(QueryProcessor[MultiPack]):  
  
    def _build_query(self, text: str) -> Dict[str, Any]:  
        """Constructs Elasticsearch query that will be consumed by  
        Elasticsearch processor.  
  
        Args:  
            text: str  
                A string which will be looked up for in the corpus under field  
                name `field`. `field` can be passed in a `config` during  
                :meth:`ElasticSearchQueryCreator::initialize`. If `config` does  
                not contain the key `field`, we will set it to "content"  
            ...  
            size = self.config.size or 1000  
            field = self.config.field or "content"  
            return {"query": {"match": {field: text}}, "size": size}
```

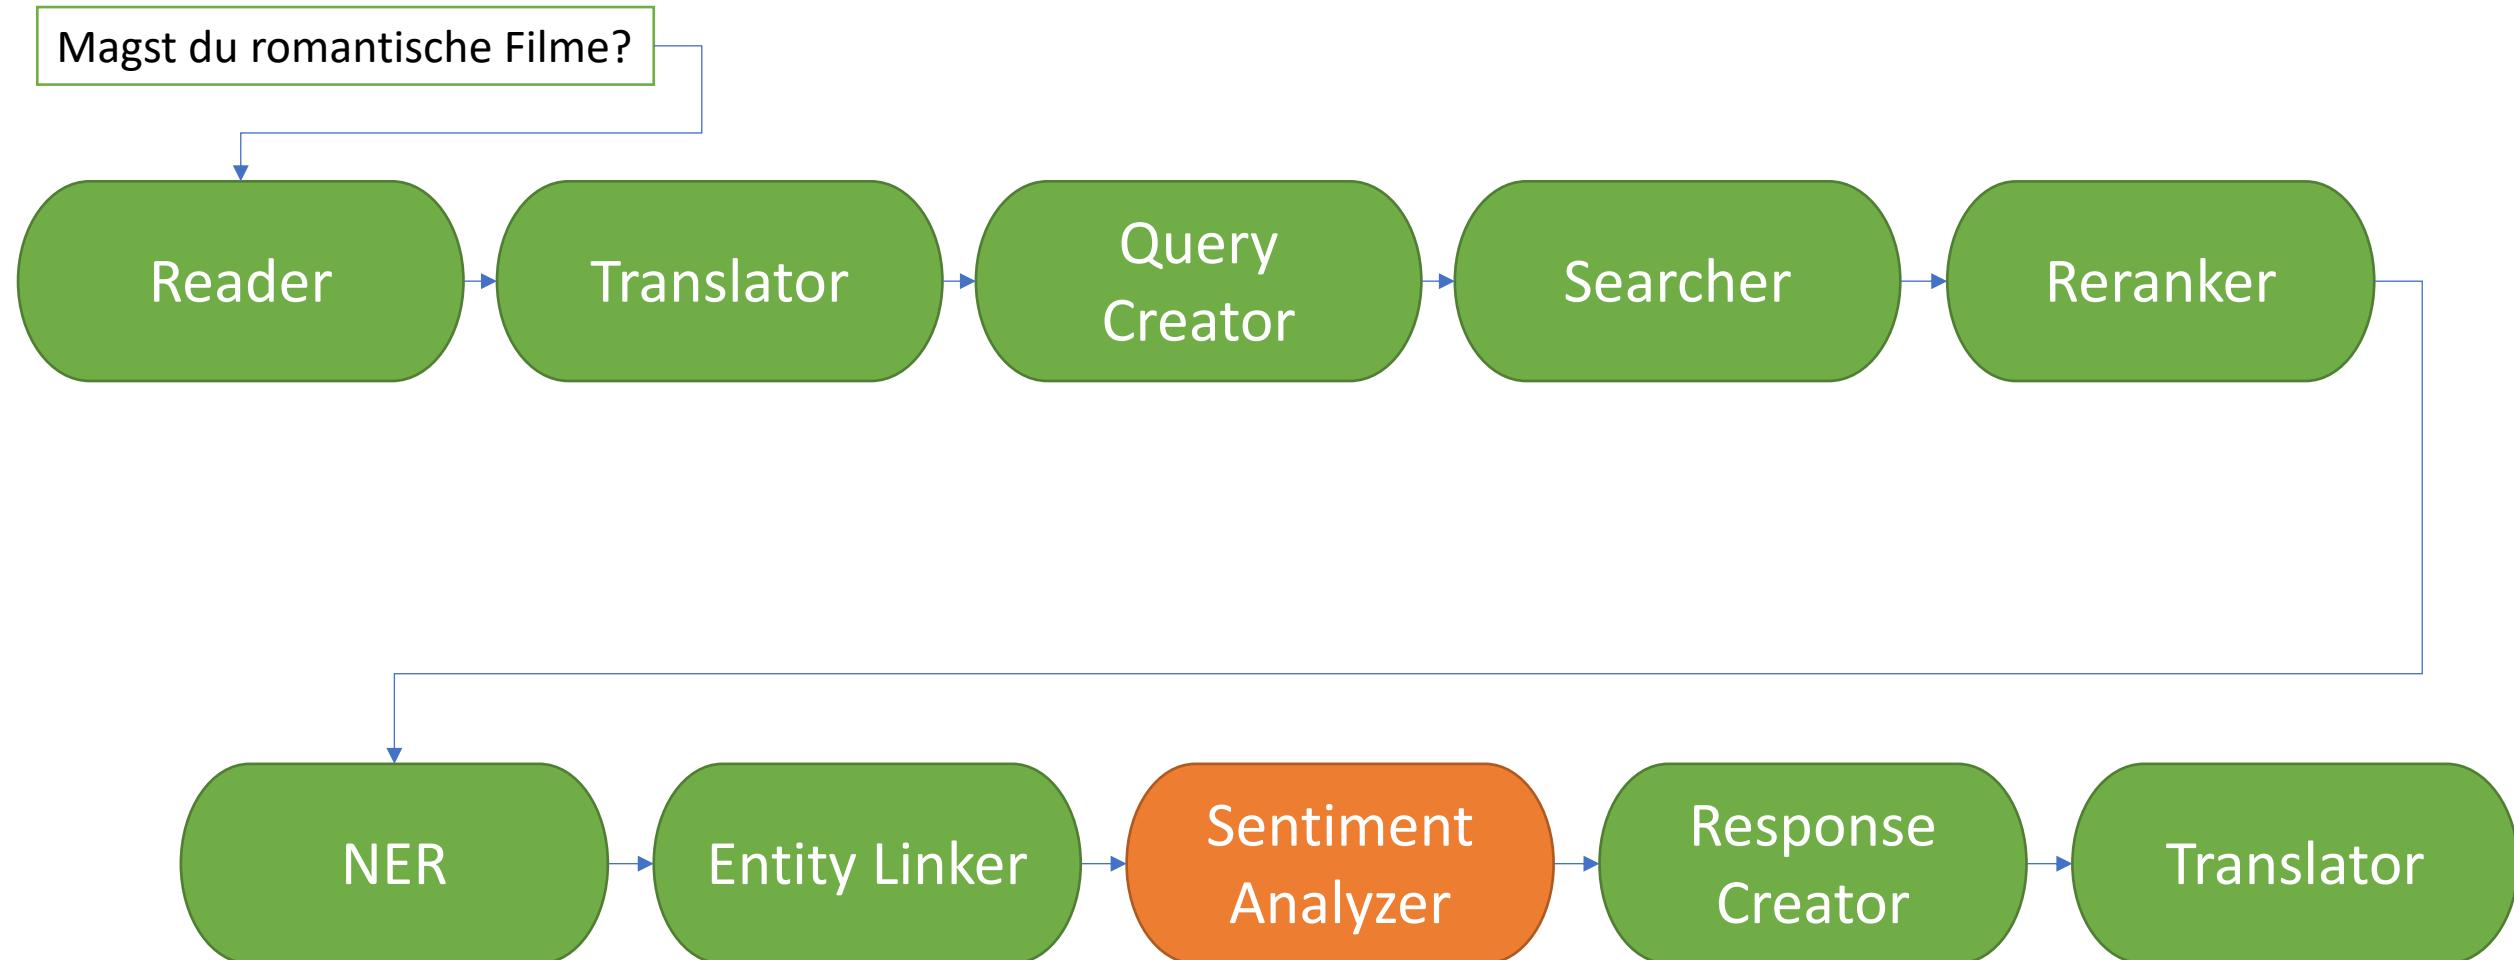
Query Creator

Alternatively, we can create an embedding based query.

Build query by using the embedding

```
class BertBasedQueryCreator(QueryProcessor[MultiPack]):  
    """This processor searches relevant documents for a query"""  
  
    @torch.no_grad()  
    def get_embeddings(self, inputs, sequence_length, segment_ids):  
        output, _ = self.encoder(inputs=inputs,  
                               sequence_length=sequence_length,  
                               segment_ids=segment_ids)  
        cls_token = output[:, 0, :]  
  
        return cls_token  
  
    def _build_query(self, text: str) -> np.ndarray:  
        input_ids, segment_ids, input_mask = \  
            self.tokenizer.encode_text(  
                text, max_seq_length=self.config.max_seq_length)  
        input_ids = torch.LongTensor(input_ids).unsqueeze(0).to(self.device)  
        segment_ids = torch.LongTensor(segment_ids).unsqueeze(0).to(self.device)  
        input_mask = torch.LongTensor(input_mask).unsqueeze(0).to(self.device)  
        sequence_length = (1 - (input_mask == 0)).sum(dim=1)  
        query_vector = self.get_embeddings(inputs=input_ids,  
                                         sequence_length=sequence_length,  
                                         segment_ids=segment_ids)  
        query_vector = torch.mean(query_vector, dim=0, keepdim=True)  
        query_vector = query_vector.cpu().numpy()  
        return query_vector
```

Implement Processors



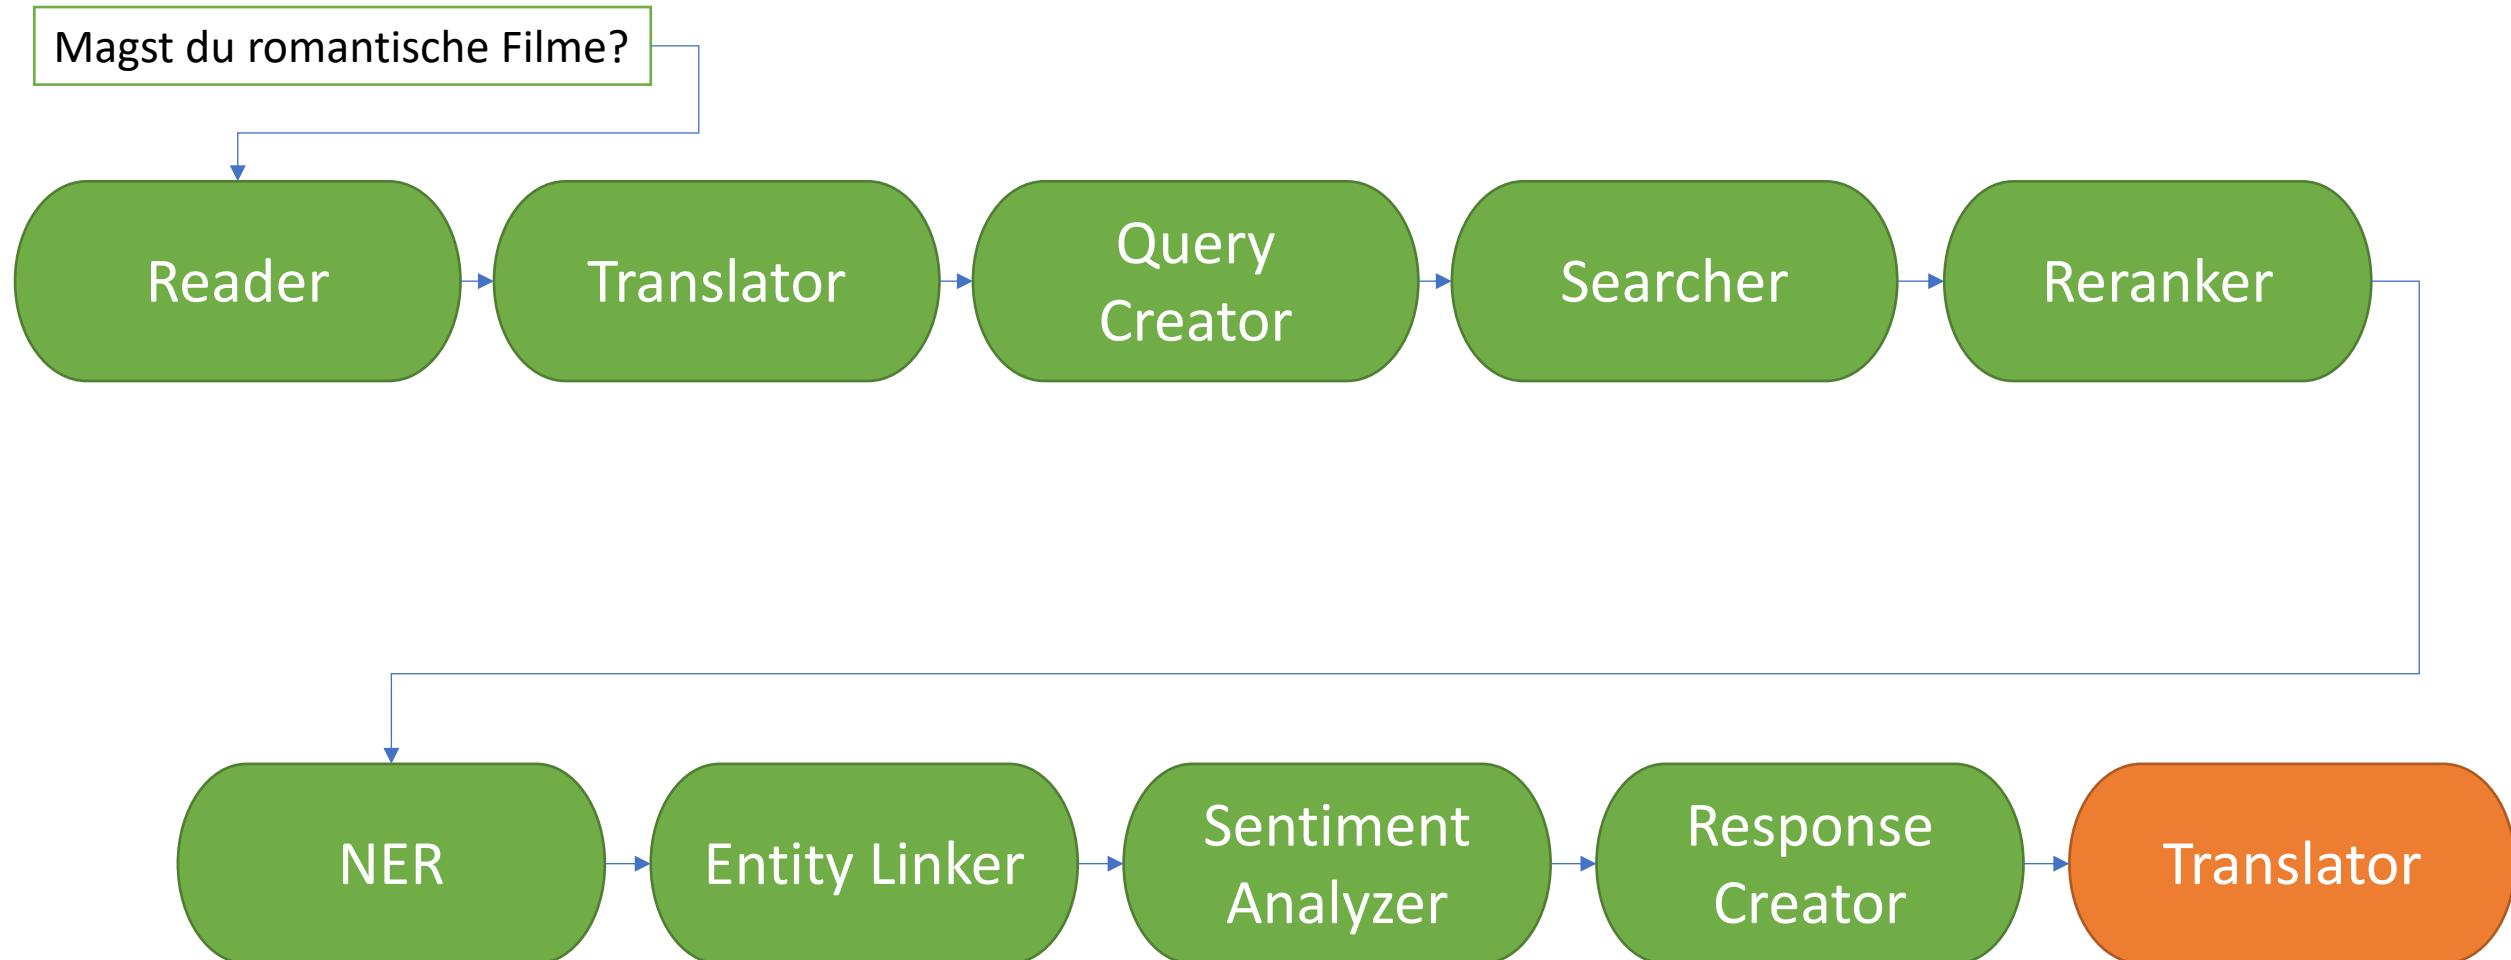
Adding Sentiment Analysis

- We can easily wrap an external sentiment analysis system

Assign the scores to each sentence.

```
class VaderSentimentProcessor(PackProcessor):  
    def __init__(self):  
        super().__init__()  
        self.sentence_component = None  
        self.analyzer = SentimentIntensityAnalyzer()  
  
    def initialize(self, resource: Resources, configs: HParams):  
        self.sentence_component = configs.get('sentence_component')  
  
    def _process(self, input_pack: DataPack):  
        sentence: Sentence  
        for sentence in input_pack.get(entry_type=Sentence,  
                                       component=self.sentence_component):  
            scores = self.analyzer.polarity_scores(sentence.text)  
            sentence.sentiment = scores
```

Implement Processors



Adding a Translator

Get translation output
with Bing API

Add translated results to
data

```

def _process(self, input_pack: MultiPack):
    query = input_pack.get_pack(self.in_pack_name).text
    params = '?' + urlencode(
        {'api-version': '3.0',
         'from': self.src_language,
         'to': [self.target_language]}, doseq=True)
    microsoft_constructed_url = self.microsoft_translate_url + params

    response = requests.post(
        microsoft_constructed_url, headers=self.microsoft_headers,
        json=[{"text": query}])

    if response.status_code != 200:
        raise RuntimeError(response.json()['error']['message'])

    text = response.json()[0]["translations"][0]["text"]
    pack = DataPack()

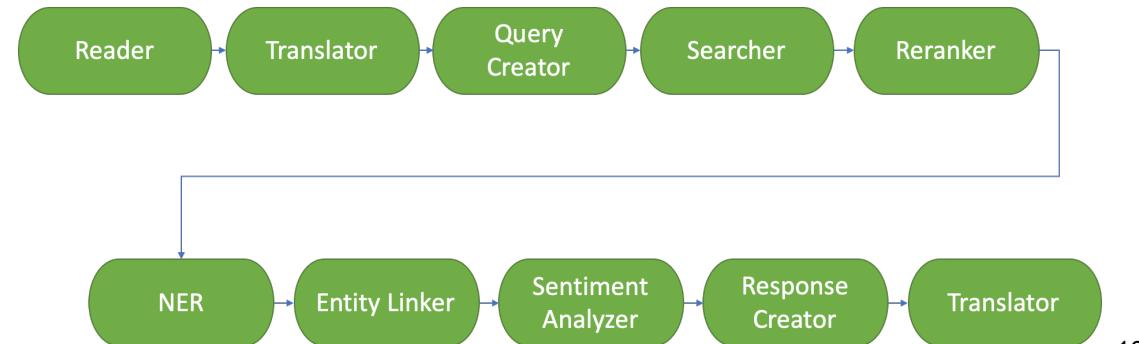
    document = Document(pack, 0, len(text))
    utterance = Utterance(pack, 0, len(text))
    pack.add_entry(document)
    pack.add_entry(utterance)

    pack.set_text(text=text)
    input_pack.update_pack({self.out_pack_name: pack})

```

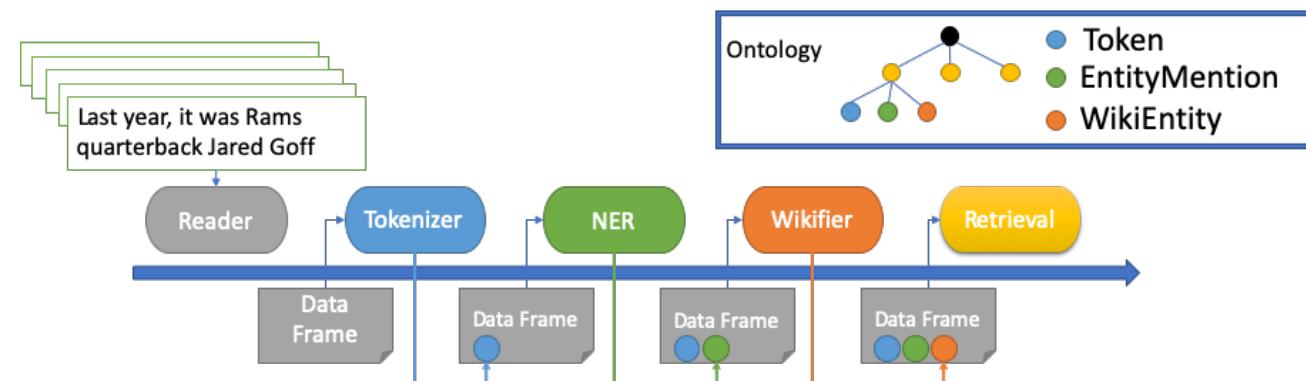
Inter-operation

- User: Kennst du ein paar gute romantische Filme?
- Response: *Yes, Titanic. Kate Winslet and Leonardo Dicaprio have definitely created a timeless classic.*
 - Response selected by inter-operation
 - It contains the actors (NER + Entity Linking)
 - It contains the sentiment (Sentiment Analyzer)
- Informed decisions can be made with a well-designed pipeline



Some Take-home Messages

- Use standard and shared data representation
- NLP concepts can be categorized
- Enrich data, don't delete data
- Keep consistent interfaces between models and tasks
- Understand your domain

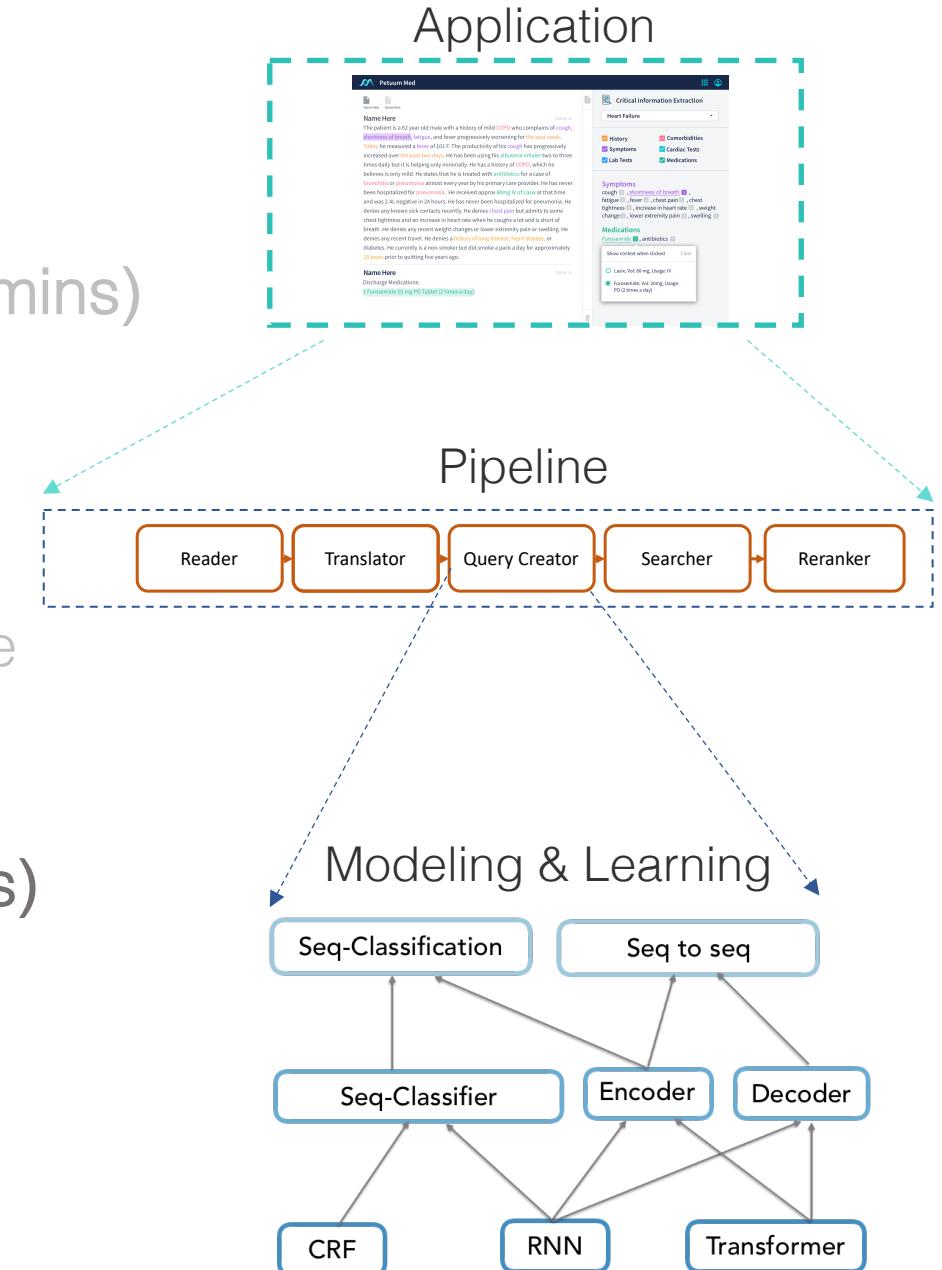


Agenda

- Natural Language Processing Overview (10mins)
- Modularizing NLP Pipeline (35mins)
 - Complexity of NLP pipeline
 - A standardized view of NLP pipeline
 - A standardized implementation of NLP pipeline
- Short break & QA (5mins)
- Modularizing NLP Model & Learning (30mins)
 - Composable ML
- QA (10mins)

Agenda

- Natural Language Processing Overview (10mins)
- Modularizing NLP Pipeline (35mins)
 - Complexity of NLP pipeline
 - A standardized view of NLP pipeline
 - A standardized implementation of NLP pipeline
- Short break & QA (5mins)
- Modularizing NLP Model & Learning (30mins)
 - Composable ML
- QA (10mins)

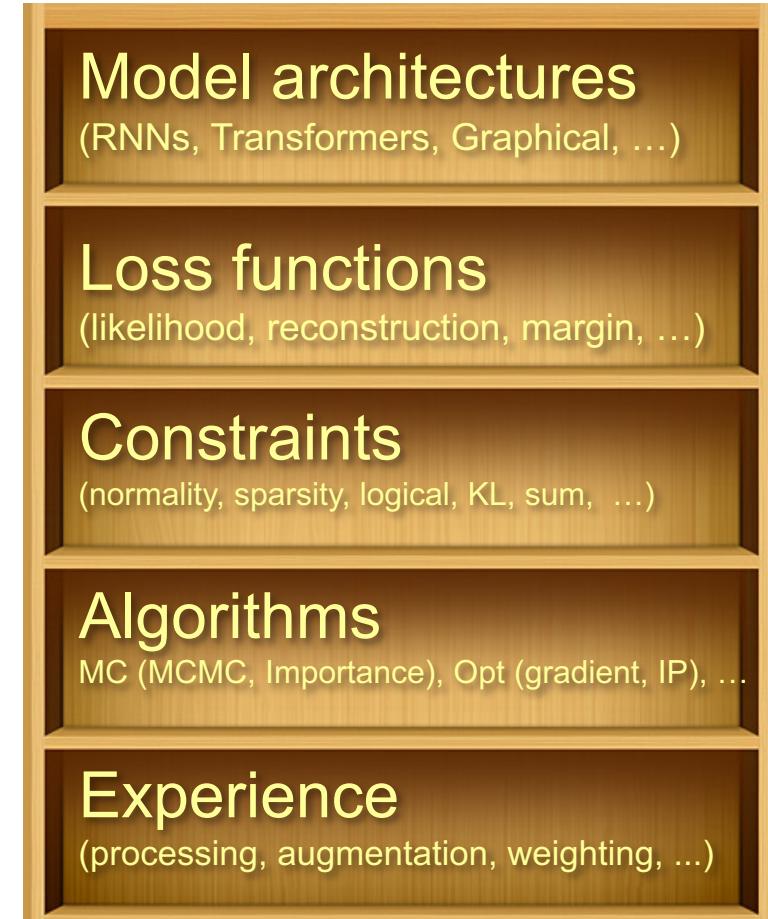
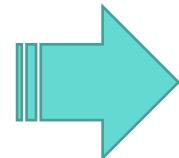
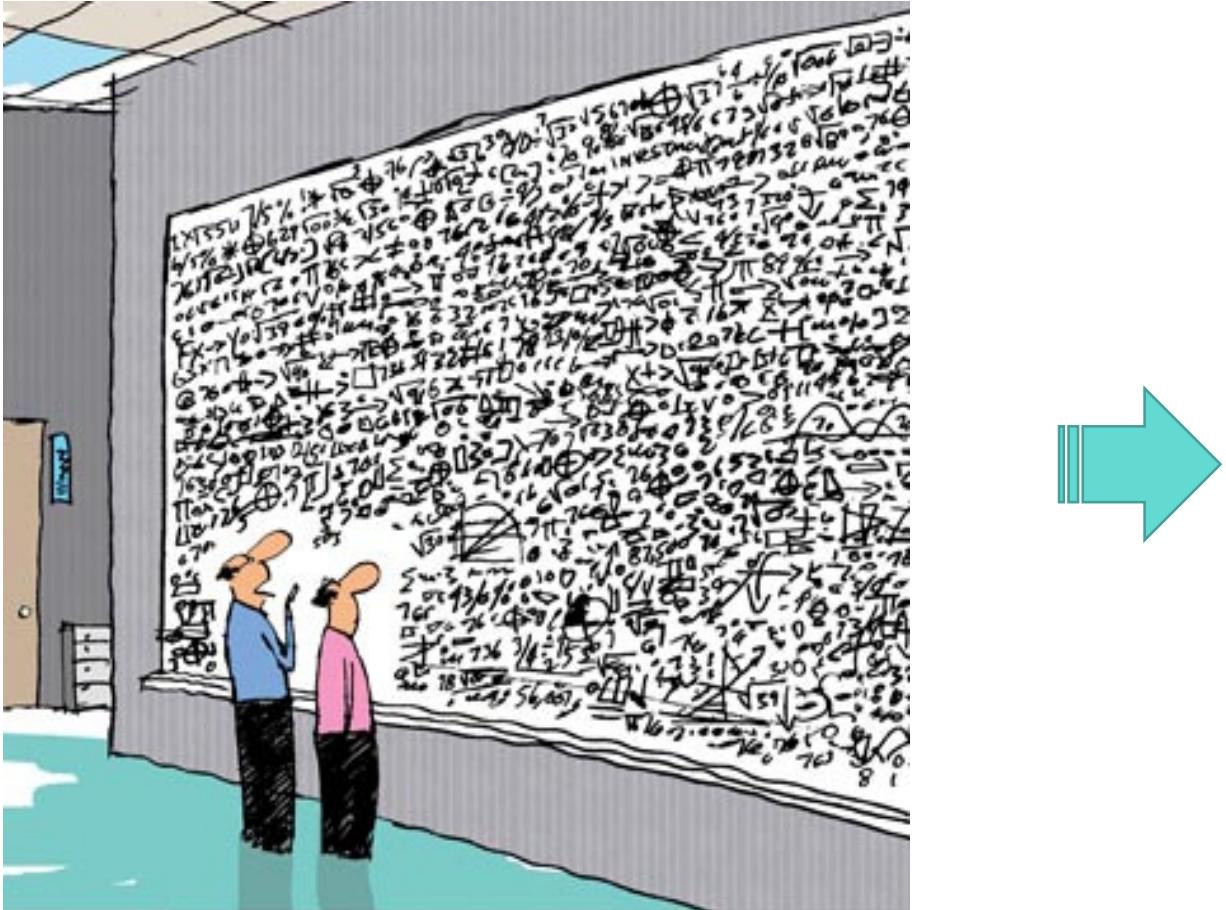


Composable ML

- Build ML models more easily, via pick-and-choose
- Stop writing same one-off code again and again
 - More reliable and easier to debug
 - Easier to onboard new developers



Decomposing Machine Learning



Decomposing Machine Learning

Machine Learning:

Computational methods that enable machines to learn concepts and improve performance from experience

Decomposing Machine Learning

Machine Learning:

Computational methods that enable machines to learn concepts and improve performance from experience

$$\min_{\theta} \mathcal{L}(\theta, \mathcal{E}) + \Omega(\theta)$$

Decomposing Machine Learning

Machine Learning:

Computational methods that enable machines to learn concepts and improve performance from experience

$$\min_{\theta} \mathcal{L}(\theta, \mathcal{E}) + \Omega(\theta)$$

$$\mathbf{y} \sim p_{\theta}(\mathbf{y}|\mathbf{x})$$



model architecture/
inference procedure

Decomposing Machine Learning

Machine Learning:

Computational methods that enable machines to learn concepts and improve performance from experience

$$\min_{\theta} \mathcal{L}(\theta, \mathcal{E}) + \Omega(\theta)$$

(x^*, y^*)

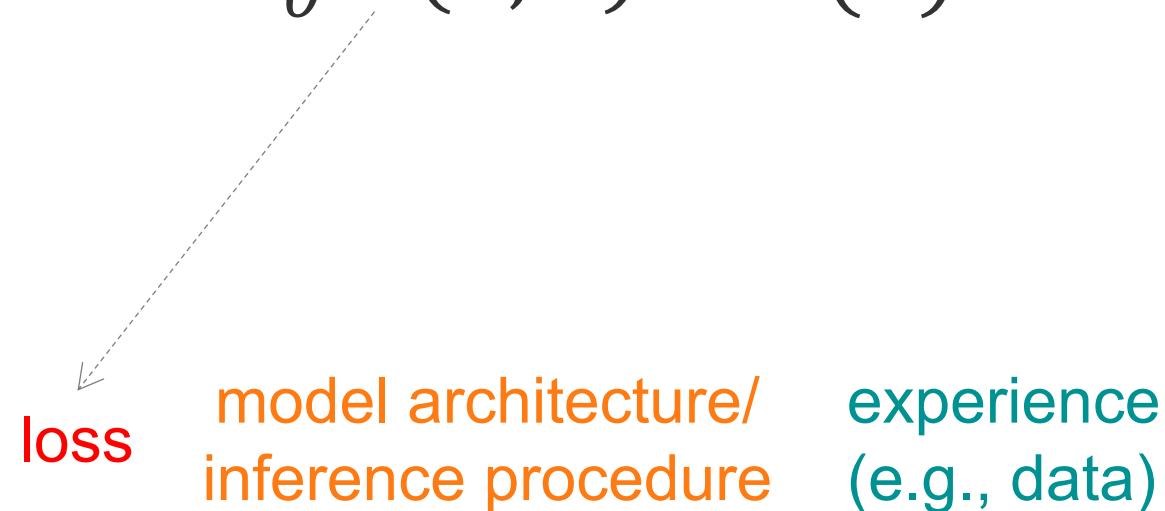
model architecture/
inference procedure experience
(e.g., data)

Decomposing Machine Learning

Machine Learning:

Computational methods that enable machines to learn concepts and improve performance from experience

$$\min_{\theta} \mathcal{L}(\theta, \mathcal{E}) + \Omega(\theta)$$

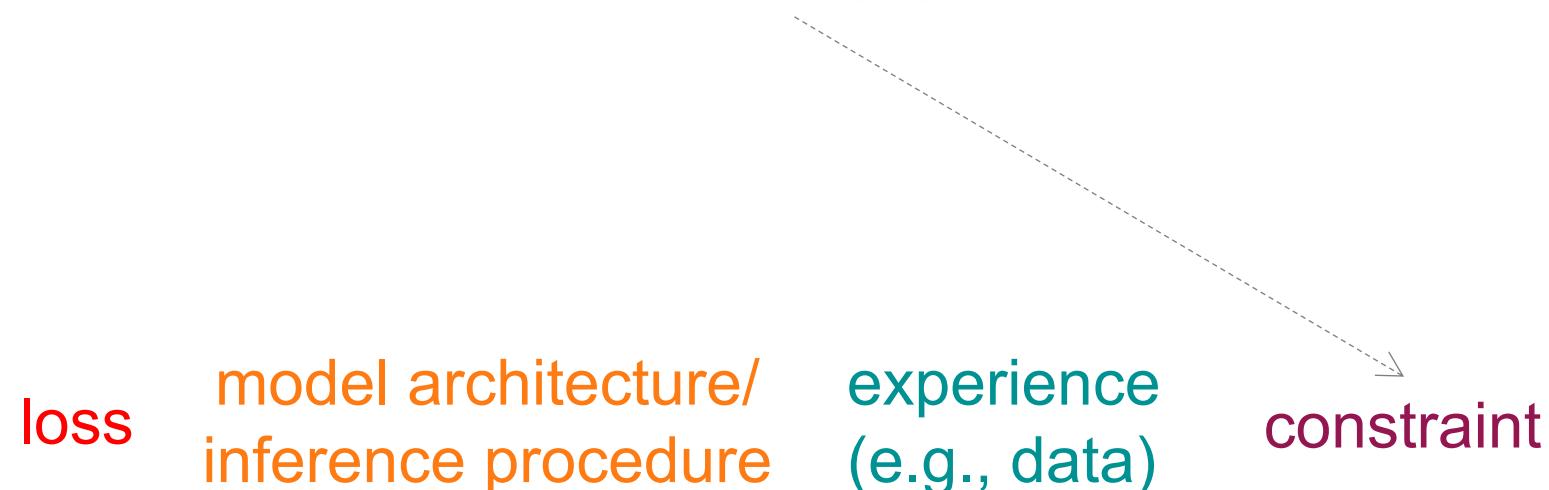


Decomposing Machine Learning

Machine Learning:

Computational methods that enable machines to learn concepts and improve performance from experience

$$\min_{\theta} \mathcal{L}(\theta, \mathcal{E}) + \Omega(\theta)$$



Decomposing Machine Learning

Machine Learning:

Computational methods that enable machines to learn concepts and improve performance from experience

$$\min_{\theta} \mathcal{L}(\theta, \mathcal{E}) + \Omega(\theta)$$

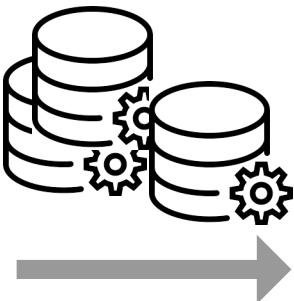
learning procedure loss model architecture/
inference procedure experience
(e.g., data) constraint

$$\min_{\theta} \mathcal{L}(\theta, \mathcal{E}) + \Omega(\theta)$$

Running Example: Machine Translation



raw data



cleaning
tokenizing
vocabulary
truncation
...

source.dat

I like this movie.
Lovely and poignant
Insanely hilarious!
...

target.dat

Ich mag diesen film.
Schön und ergreifend
Wahnsinnig witzig!
...

clean data

*evaluation
post-processing*



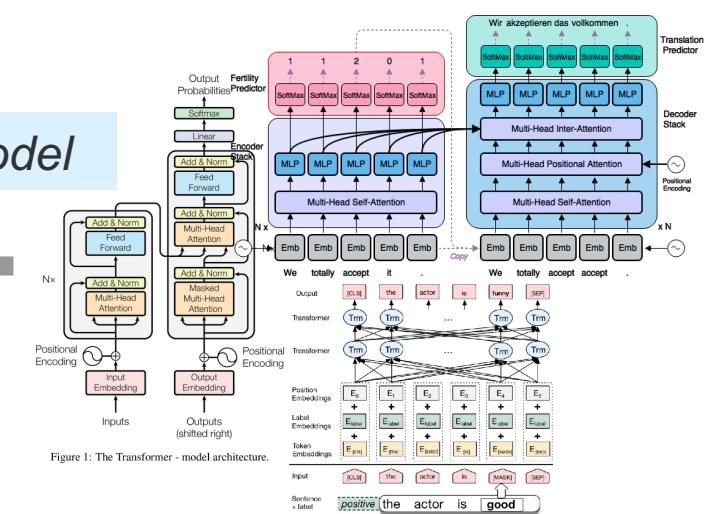
training

**Maximum likelihood
training**

**Reinforcement
Adversarial learning
Finetuning**



model



ML Components



Constraint

Loss

Learning

Inference

Architecture

ML Components



Constraint

Loss

Learning

Inference

Architecture

Architecture (1): Language Model



- Calculates the probability of a sentence:

- Sentence:

Example:

$$\mathbf{y} = (y_1, y_2, \dots, y_T)$$

(I, like, this, ...)



Architecture (1): Language Model

- Calculates the probability of a sentence:

- Sentence:

Example:

$$\mathbf{y} = (y_1, y_2, \dots, y_T)$$

(I, like, this, ...)

$$p_{\theta}(\mathbf{y}) = \prod_{t=1}^T p_{\theta}(y_t | \mathbf{y}_{1:t-1})$$

... $p_{\theta}(\text{like} | I) p_{\theta}(\text{this} | I, \text{like}) \dots$



Architecture (1): Language Model

- Calculates the probability of a sentence:
 - Sentence:

$$\mathbf{y} = (y_1, y_2, \dots, y_T)$$

Example:

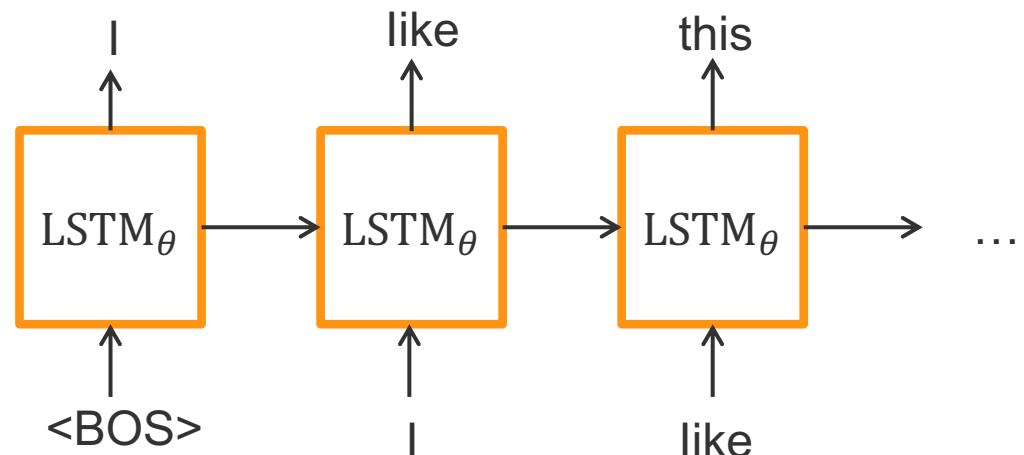
(I, like, this, ...)

$$p_{\theta}(\mathbf{y}) = \prod_{t=1}^T p_{\theta}(y_t | \mathbf{y}_{1:t-1})$$

$$\dots p_{\theta}(\text{like} | I) p_{\theta}(\text{this} | I, \text{like}) \dots$$

Architecture (1.1)

LSTM RNN



Architecture (1): Language Model



- Calculates the probability of a sentence:
 - Sentence:

$$\mathbf{y} = (y_1, y_2, \dots, y_T)$$

Example:

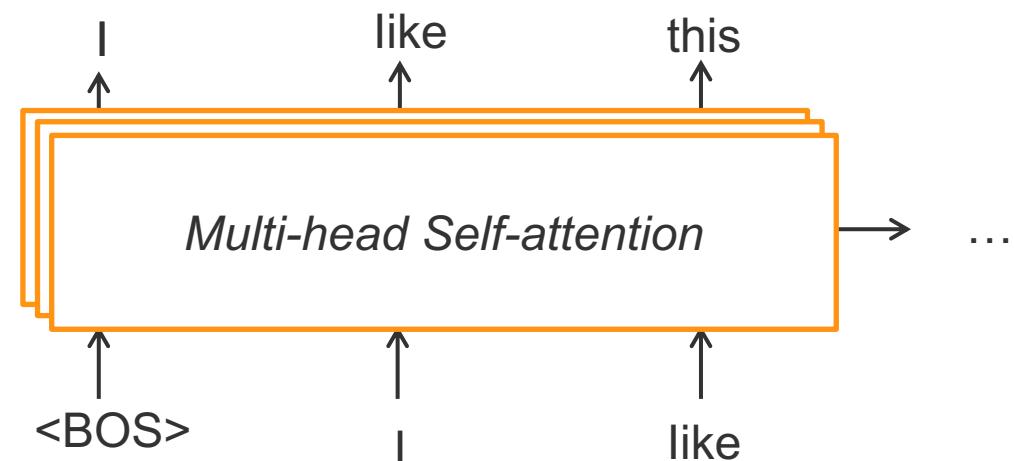
(I, like, this, ...)

$$p_{\theta}(\mathbf{y}) = \prod_{t=1}^T p_{\theta}(y_t | \mathbf{y}_{1:t-1})$$

$$\cdots p_{\theta}(\text{like} | I) p_{\theta}(\text{this} | I, \text{like}) \cdots$$

Architecture (1.2)

Transformer



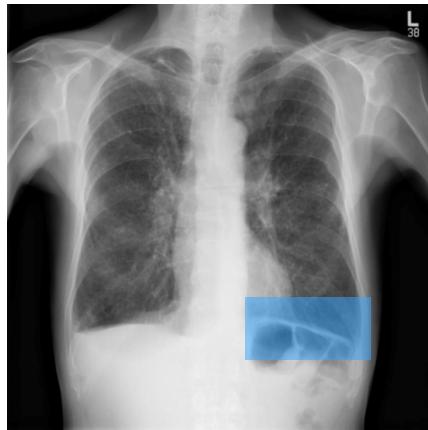
Architecture (2): Conditional Language Model



- Conditions on additional task-dependent context x
 - Machine translation: source sentence

I like this movie. → Ich mag diesen film.

- Medical image report generation: medical image



... There is chronic pleural-parenchymal scarring within the lung bases. No lobar consolidation is seen. ...

Architecture (2): Conditional Language Model



- Conditions on additional task-dependent context \mathbf{x}

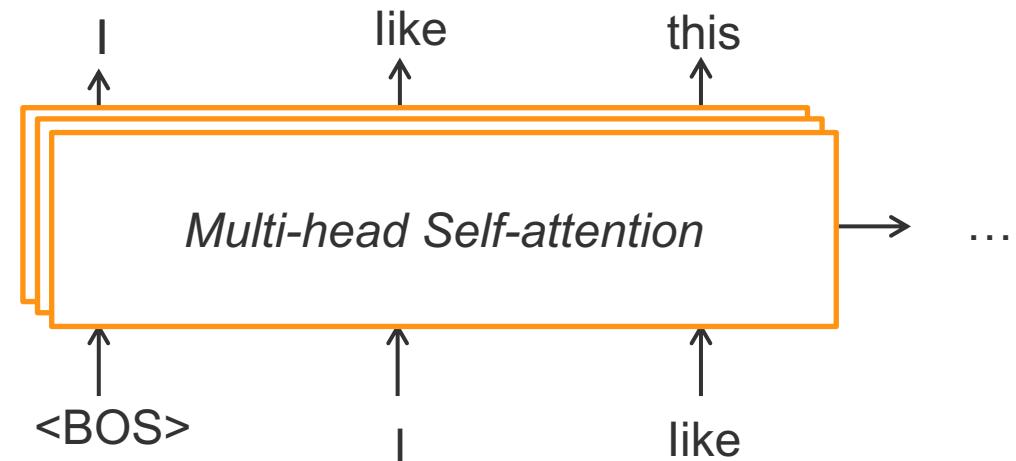
$$p_{\theta}(y | \mathbf{x}) = \prod_{t=1}^T p_{\theta}(y_t | y_{1:t-1}, \mathbf{x})$$

Architecture (2): Conditional Language Model



- Conditions on additional task-dependent context \mathbf{x}

$$p_{\theta}(y | \mathbf{x}) = \prod_{t=1}^T p_{\theta}(y_t | y_{1:t-1}, \mathbf{x})$$



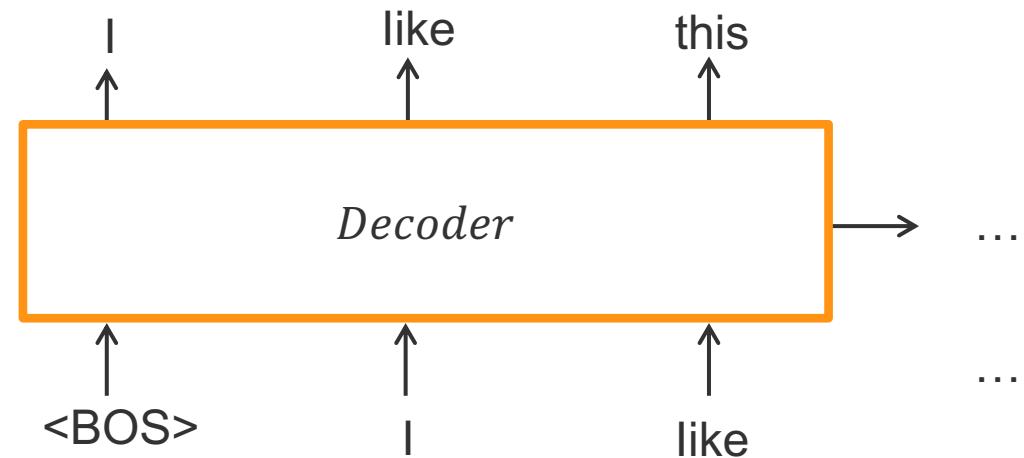
Architecture (2): Conditional Language Model



- Conditions on additional task-dependent context \mathbf{x}

$$p_{\theta}(y | \mathbf{x}) = \prod_{t=1}^T p_{\theta}(y_t | y_{1:t-1}, \mathbf{x})$$

- Language model as a **decoder**



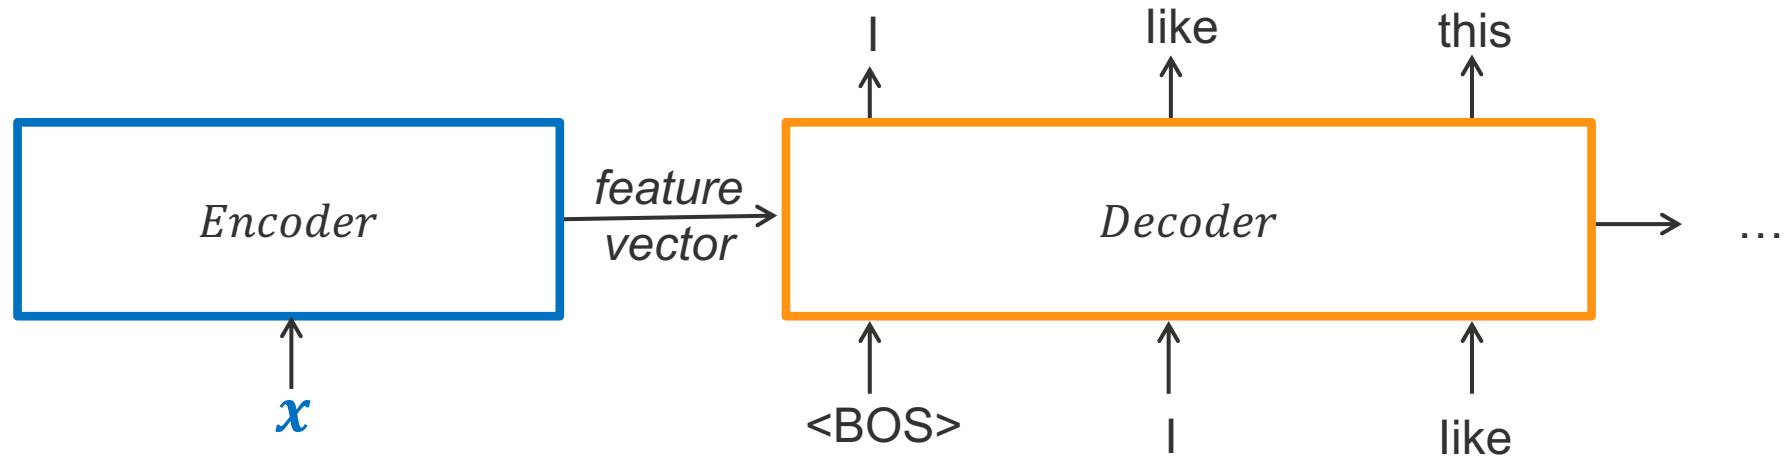
Architecture (2): Conditional Language Model



- Conditions on additional task-dependent context \mathbf{x}

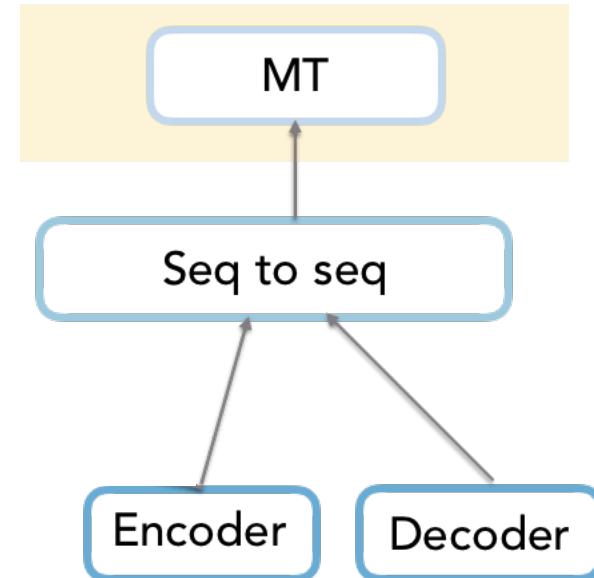
$$p_{\theta}(y | \mathbf{x}) = \prod_{t=1}^T p_{\theta}(y_t | y_{1:t-1}, \mathbf{x})$$

- Language model as a **decoder**
- Encodes context with an **encoder**



Architecture Graph

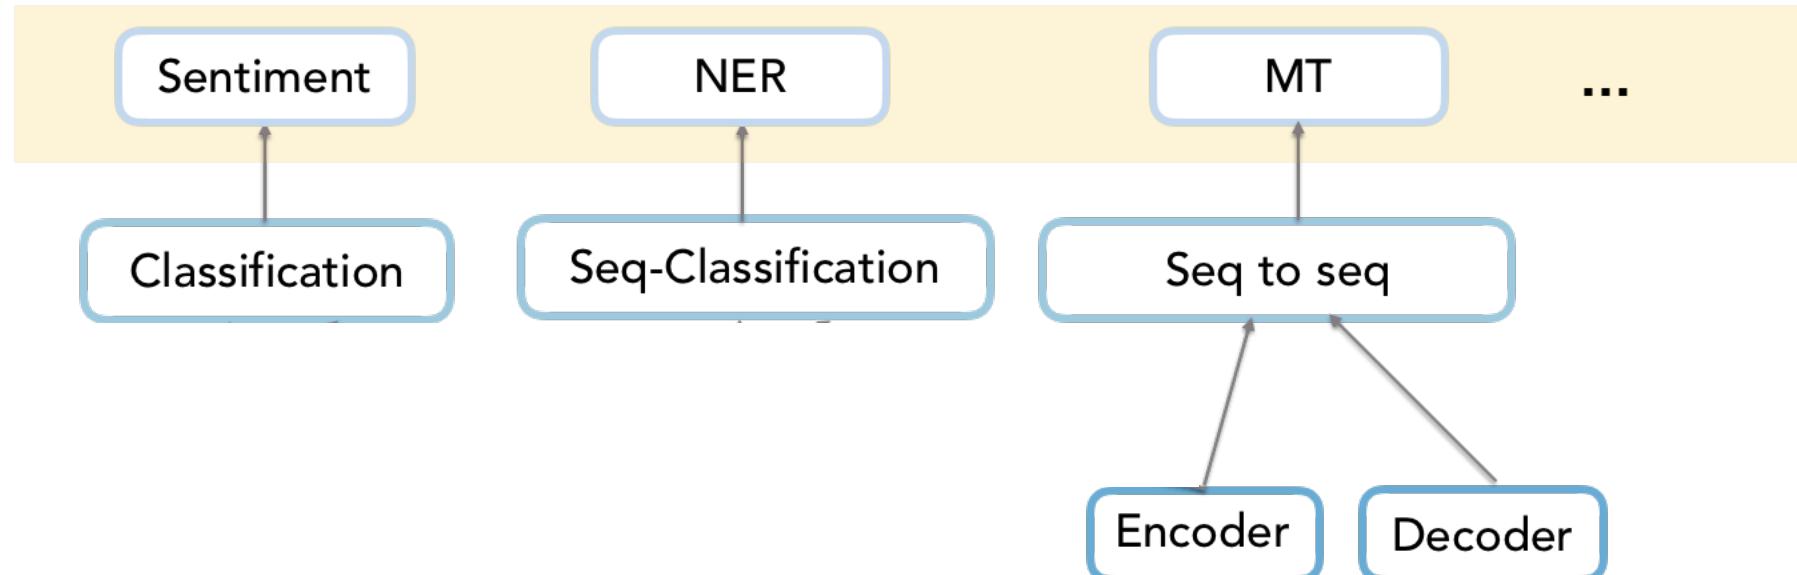
Task:



Architecture Graph



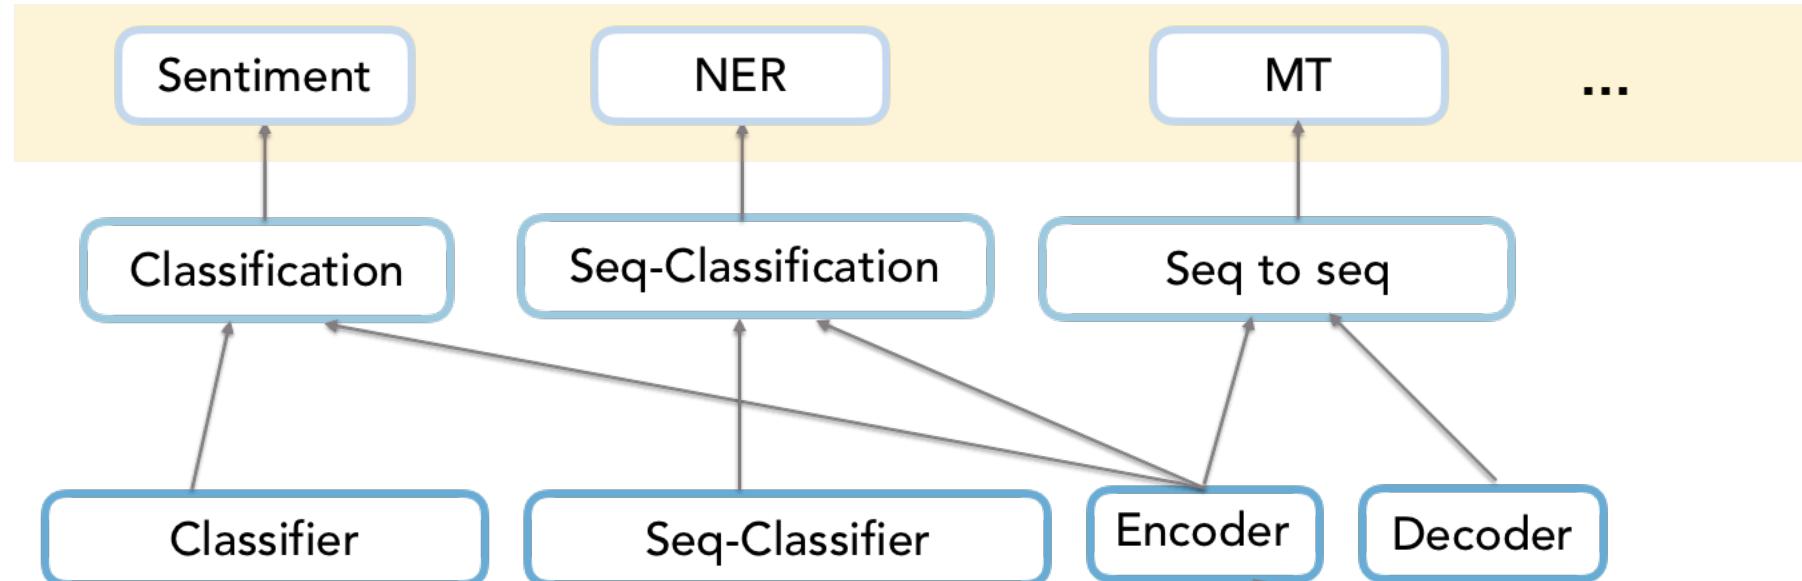
Task:



Architecture Graph



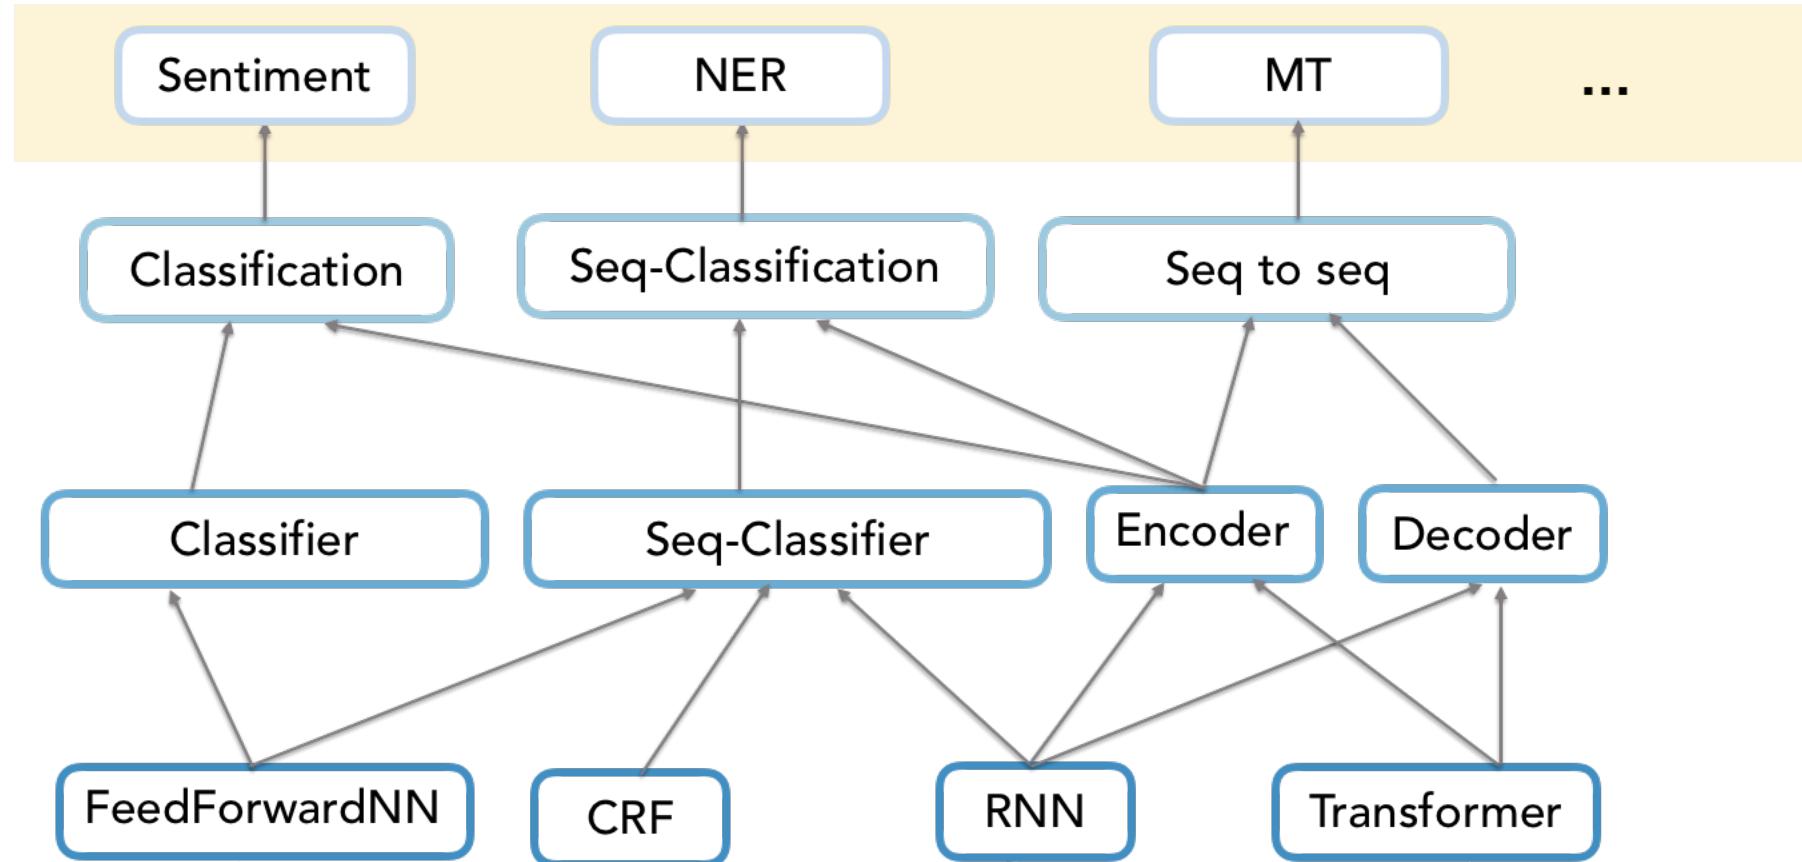
Task:



Architecture Graph



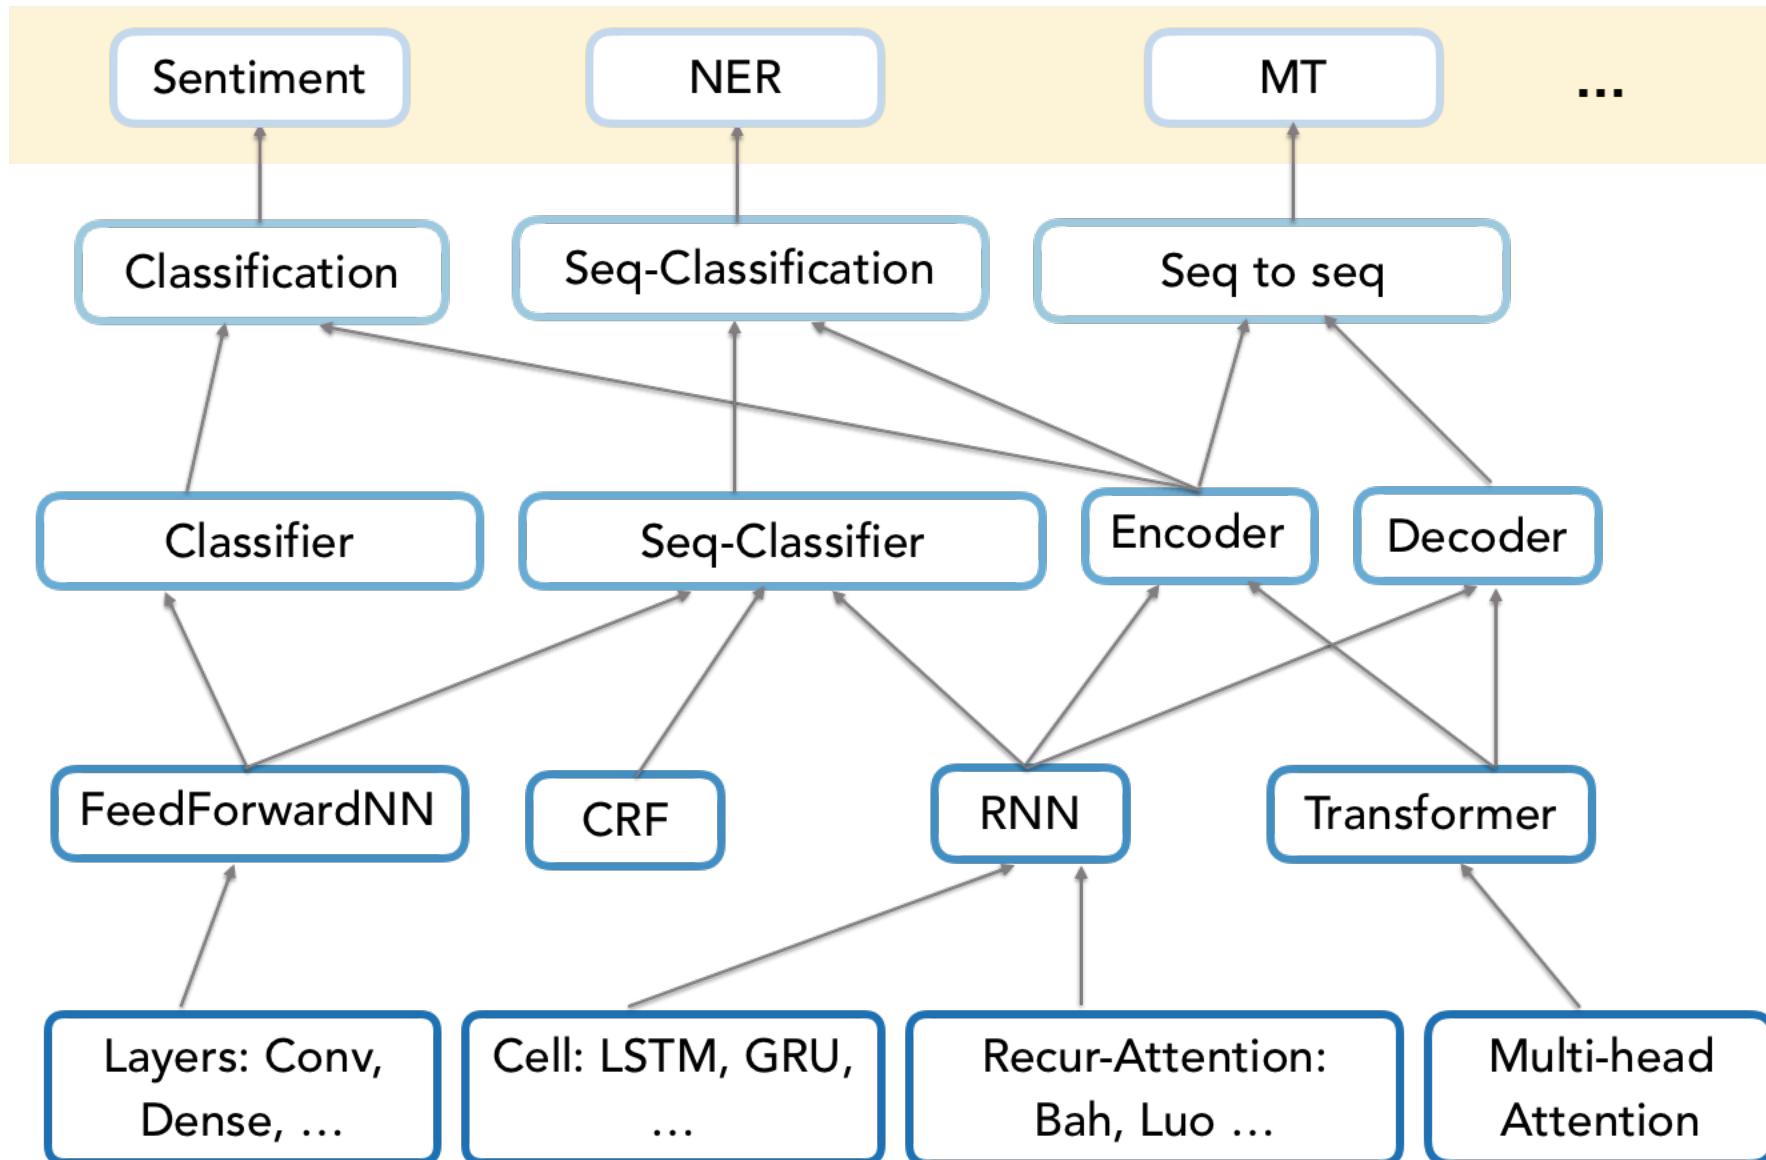
Task:



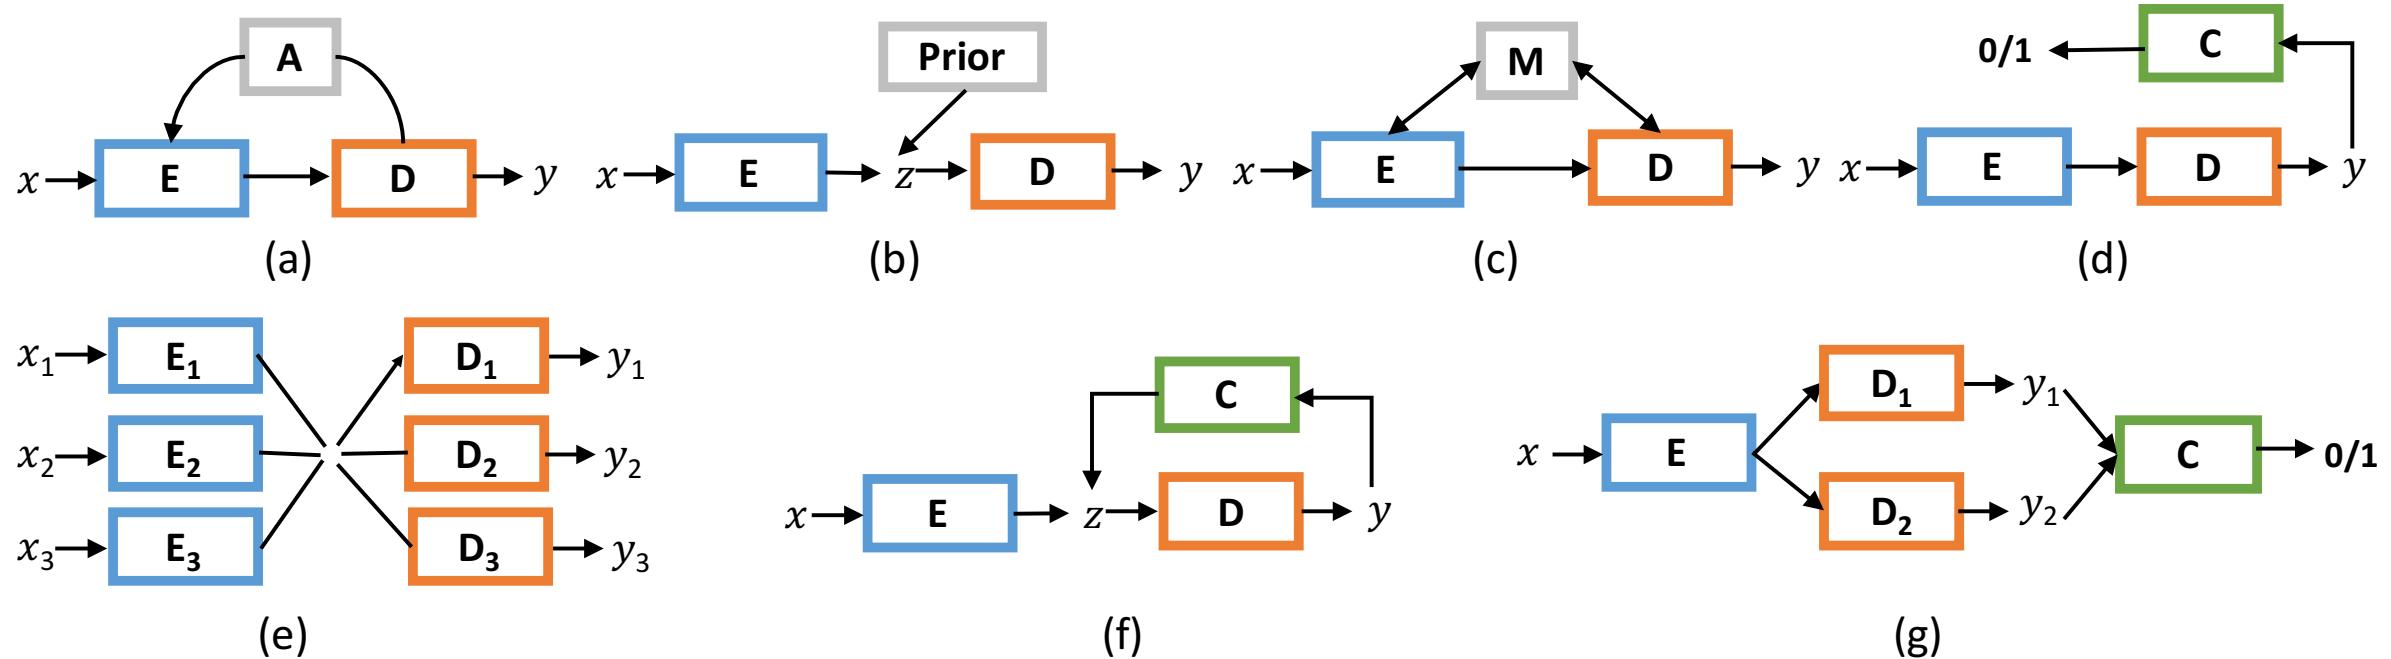
Architecture Graph



Task:



Complex Composite Architectures



E refers to encoder, D to decoder, C to Classifier, A to attention, Prior to prior distribution, and M to memory

ML Components



Constraint

Loss

Learning

Inference

Architecture

decoder

LSTM RNN

Attention RNN

Transformer

...

encoder

classifier

...

ML Components



Constraint

Loss

Learning

Inference

Architecture

decoder

LSTM RNN

Attention RNN

Transformer

...

encoder

classifier

...

Learning, Inference & Loss: Many Variations



Supervision forms:	Data examples	Reward	Auxiliary model
active learning		imitation learning	
weak/distant supervision		intrinsic reward	adversarial domain adaptation
reward-augmented MLE		inverse RL	GANs
data re-weighting		actor-critic	
data augmentation		RL as inference	knowledge distillation
maximum likelihood estimation		softmax policy gradient	prediction minimization
		policy gradient	energy-based GANs

Learning, Inference & Loss (1): Maximum Likelihood Estimation



Learning

Given data examples $\mathcal{D} = \{(\mathbf{x}^*, \mathbf{y}^*)\}$

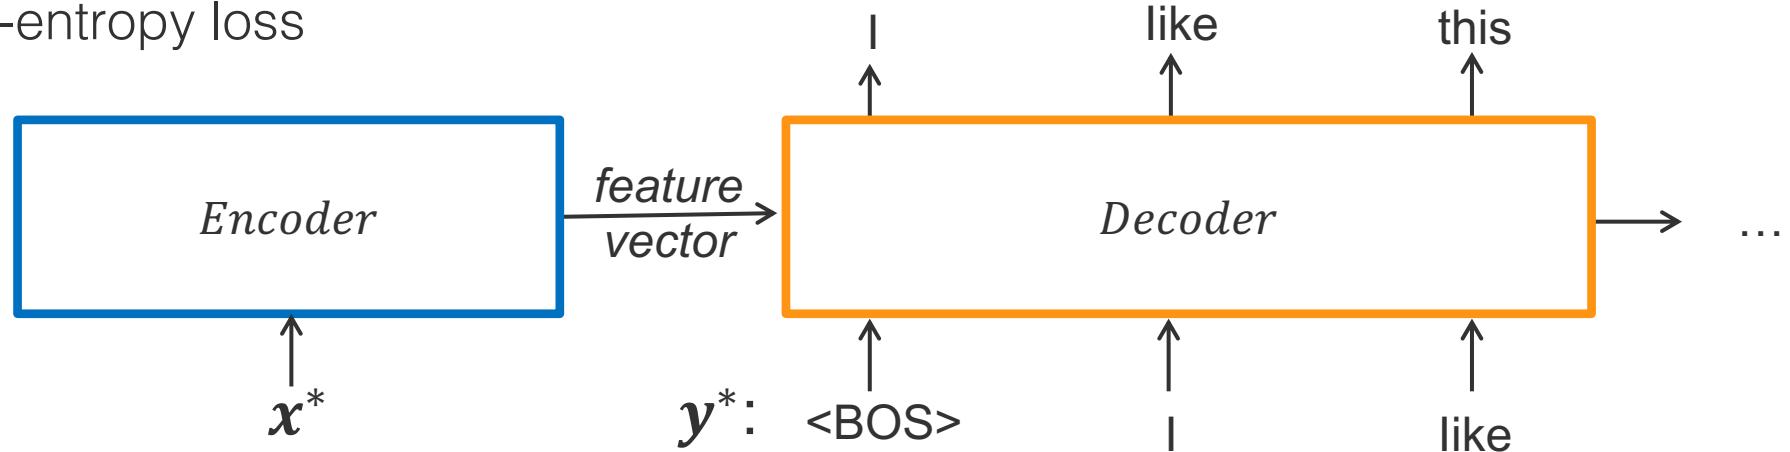
$$\max_{\theta} \mathbb{E}_{(\mathbf{x}^*, \mathbf{y}^*) \sim \mathcal{D}} [\log p_{\theta}(\mathbf{y}^* | \mathbf{x}^*)]$$

\downarrow

$$\prod_{t=1}^T p_{\theta}(y_t^* | y_{1:t-1}^*, \mathbf{x}^*)$$

Loss

Cross-entropy loss



Inference

Teacher-forcing decoding:

For every step t , feeds in the previous ground-truth tokens $y_{1:t-1}^*$ to decode next step

Learning, Inference & Loss (2): Policy Gradient



Learning

Optimizes expected task reward $R(\hat{\mathbf{y}}, \mathbf{y}^*)$

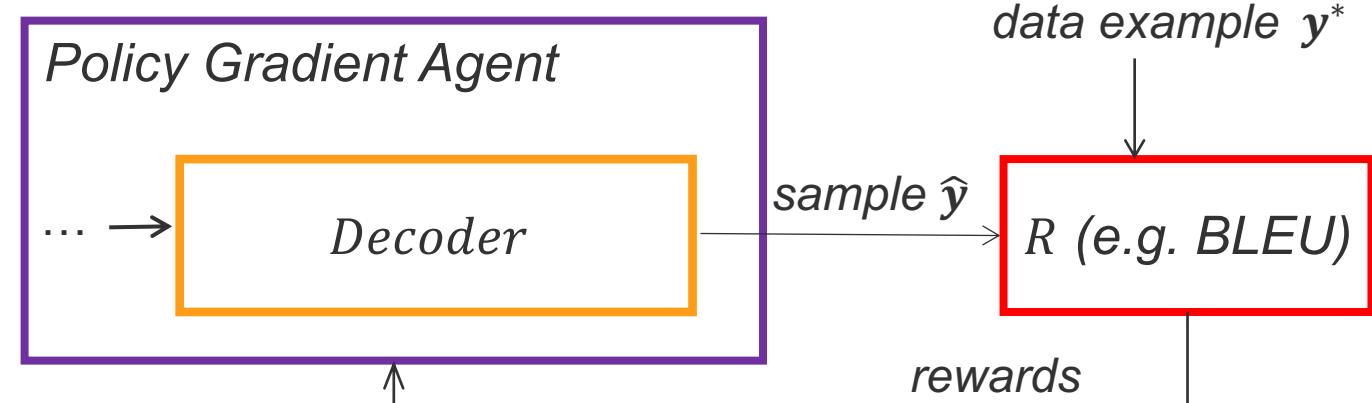
$$\max_{\theta} \mathbb{E}_{\hat{\mathbf{y}} \sim p_{\theta}(\mathbf{y} | \mathbf{x})} [R(\hat{\mathbf{y}}, \mathbf{y}^*)]$$

Loss

- Policy gradient loss
- Policy gradient loss w/ baseline
- ...

Inference

- Greedy decoding
- Sampling decoding
- Beam search decoding
- Top- k / Top- p decoding
- ...



Learning, Inference & Loss (3): Adversarial Learning

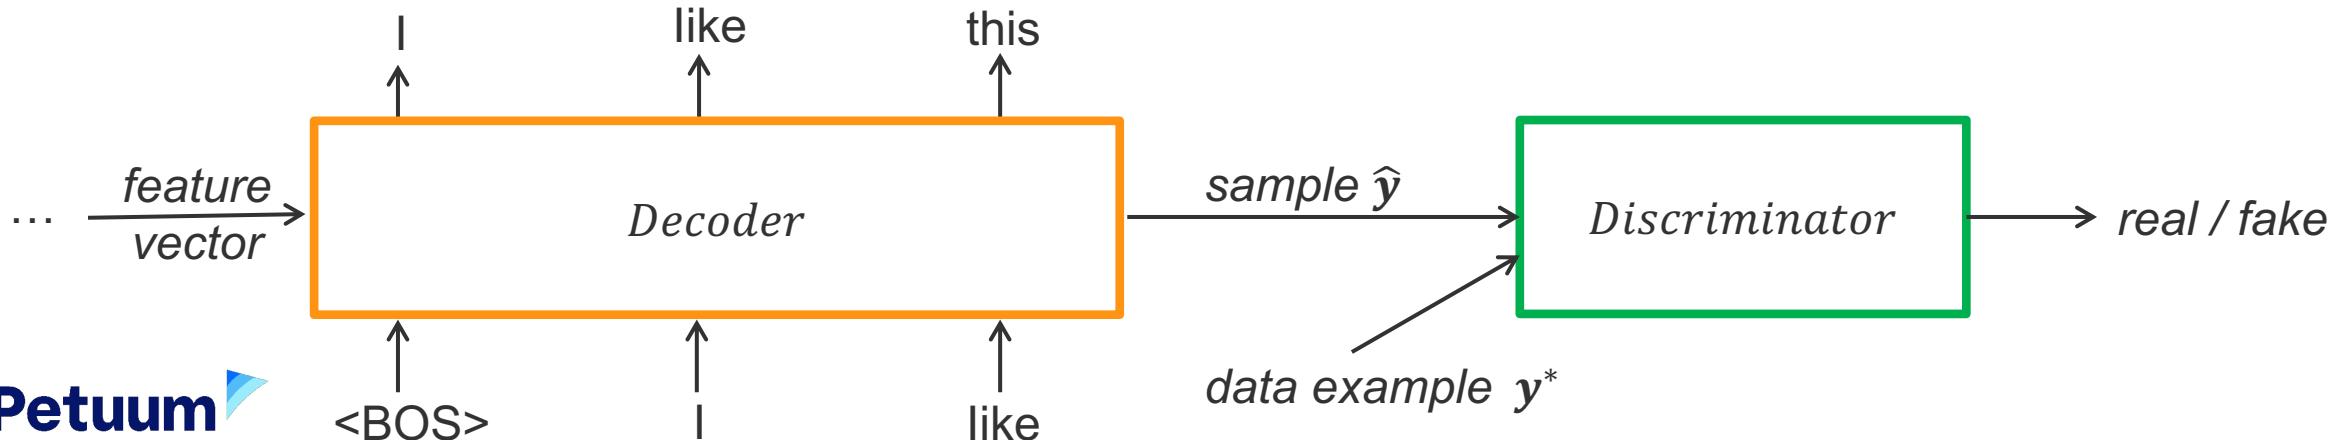


Learning

- A **discriminator** is trained to distinguish b/w *real* data examples and *fake* generated samples
- The model is trained to fool the discriminator

Loss

- Binary adversarial loss
- Feature-matching adversarial loss



Inference

Gumbel-softmax decoding:
Uses a differentiable approximation of sample $\hat{\mathbf{y}}$ for gradient backpropagation

$$\frac{\partial \mathcal{L}(\hat{\mathbf{y}})}{\partial \theta} = \frac{\partial \mathcal{L}(\hat{\mathbf{y}})}{\partial \hat{\mathbf{y}}} \frac{\partial \hat{\mathbf{y}}}{\partial \theta}$$

ML Components



Constraint

Loss

Learning

Inference

Architecture

Cross-entropy

MLE

Teacher-forcing

decoder

Binary Adv loss

Adversarial

Gumbel-softmax

LSTM RNN

Matching Adv loss

Reinforcement

Sample

Attention RNN

PG loss

...

Greedy

Transformer

PG loss + baseline

...

Beam-search

...

Top-k sample

encoder

classifier

...

ML Components



Constraint

Loss

Learning

Inference

Architecture

Cross-entropy

MLE

Teacher-forcing

decoder

Binary Adv loss

Adversarial

Gumbel-softmax

LSTM RNN

Matching Adv loss

Reinforcement

Sample

Attention RNN

PG loss

...

Greedy

Transformer

PG loss + baseline

...

Beam-search

...

Top-k sample

encoder

classifier

...

Constraint (1): Conventional Constraints



Many choices for get different statistical properties:

- Normality, Sparsity, KL, sum, ...

Constraint (2): Structured Knowledge



Structured knowledge as constraints

Sentiment classification:

- ``Food was good, **but** the service was very disappointing.''

Logic rule:

- Sentence x with structure A -**but**- $B \Rightarrow$ sentiment of B dominates

Constraint function: $f(x = \text{sentence}, y = \text{sentiment}) = \text{truth value}$

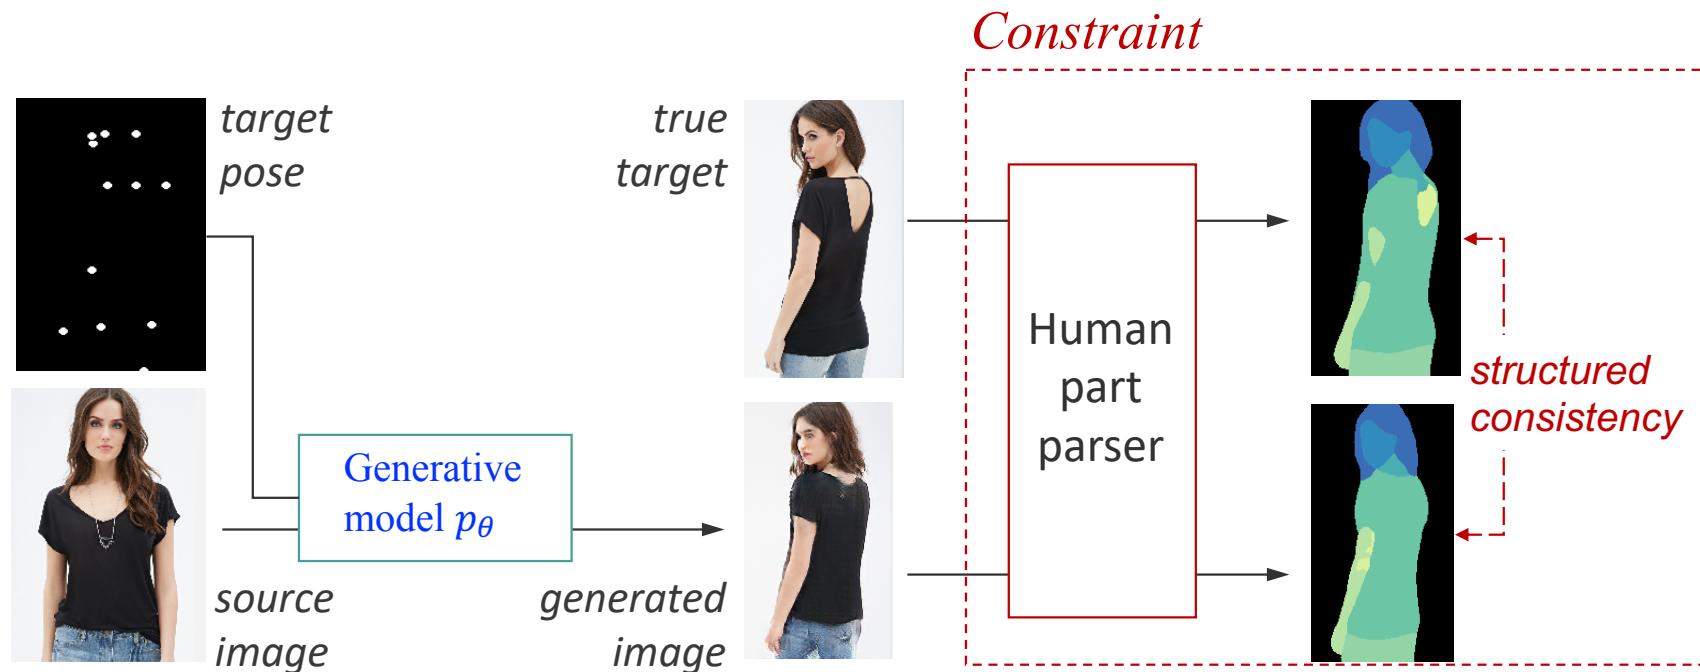
Constraint (2): Structured Knowledge



Structured knowledge as constraints

Human Image Generation

Constraint function: $f(\mathbf{y} = \text{generated}, \mathbf{o} = \text{ground truth}) = \text{match score}$





Constraint (2): Structured Knowledge

- Constraint function: $f(\mathbf{x}, \mathbf{y}) \in \mathbb{R}$
 - Higher f value, better (\mathbf{x}, \mathbf{y}) in the light of the knowledge
- Model: $p_{\theta}(\mathbf{y}|\mathbf{x})$
- Posterior Regularization [Hu et al., 2018, 2016; Zhu et al., 2014; Ganchev et al. 2010]

$$\min_{\theta, q} \mathcal{L}(\theta, q) = \mathcal{L}(\theta) + \text{KL}(q(\mathbf{y}|\mathbf{x}) \parallel p_{\theta}(\mathbf{y}|\mathbf{x})) + \xi$$

s.t. $\mathbb{E}_{q(\mathbf{y}|\mathbf{x})}[f(\mathbf{x}, \mathbf{y})] \geq 1 - \xi$

Regular loss
(e.g., cross-entropy loss)
Constraint

- Related: constraint-driven learning [Chang et al., 2007], generalized expectation [Mann & MacCallum, 2007], learning from measurements [Liang et al., 2009]



Constraint (2): Structured Knowledge

$$\mathcal{L}(\theta, q) = \mathcal{L}(\theta) + \text{KL}(q(y|x) || p_\theta(y|x)) + \xi$$

$$s.t. \quad \mathbb{E}_{q(y|x)}[f(x, y)] \geq 1 - \xi$$

- $\min_q \mathcal{L}(\theta, q) \rightarrow q^*(y|x) \propto p_\theta(y|x) \exp \left\{ f(x, y) \right\}$ Combines the **model** and the **knowledge**
— teacher model
- $\min_{\theta} \mathcal{L}(\theta, q^*) \rightarrow \max_{\theta} \mathbb{E}_{q^*(y|x)} \left[\log p_\theta(y|x) \right]$ The **model** imitates the teacher model predictions
— student model

ML Components



Constraint

Loss

Learning

Inference

Architecture

L1 / L2

Cross-entropy

MLE

Teacher-forcing

decoder

Logical

Binary Adv loss

Policy Gradient

Gumbel-softmax

LSTM RNN

Structured

Matching Adv loss

Adversarial

Sample

Attention RNN

...

PG loss

Greedy

Transformer

PG loss + baseline

Beam-search

...

...

Top-k sample

encoder

...

classifier

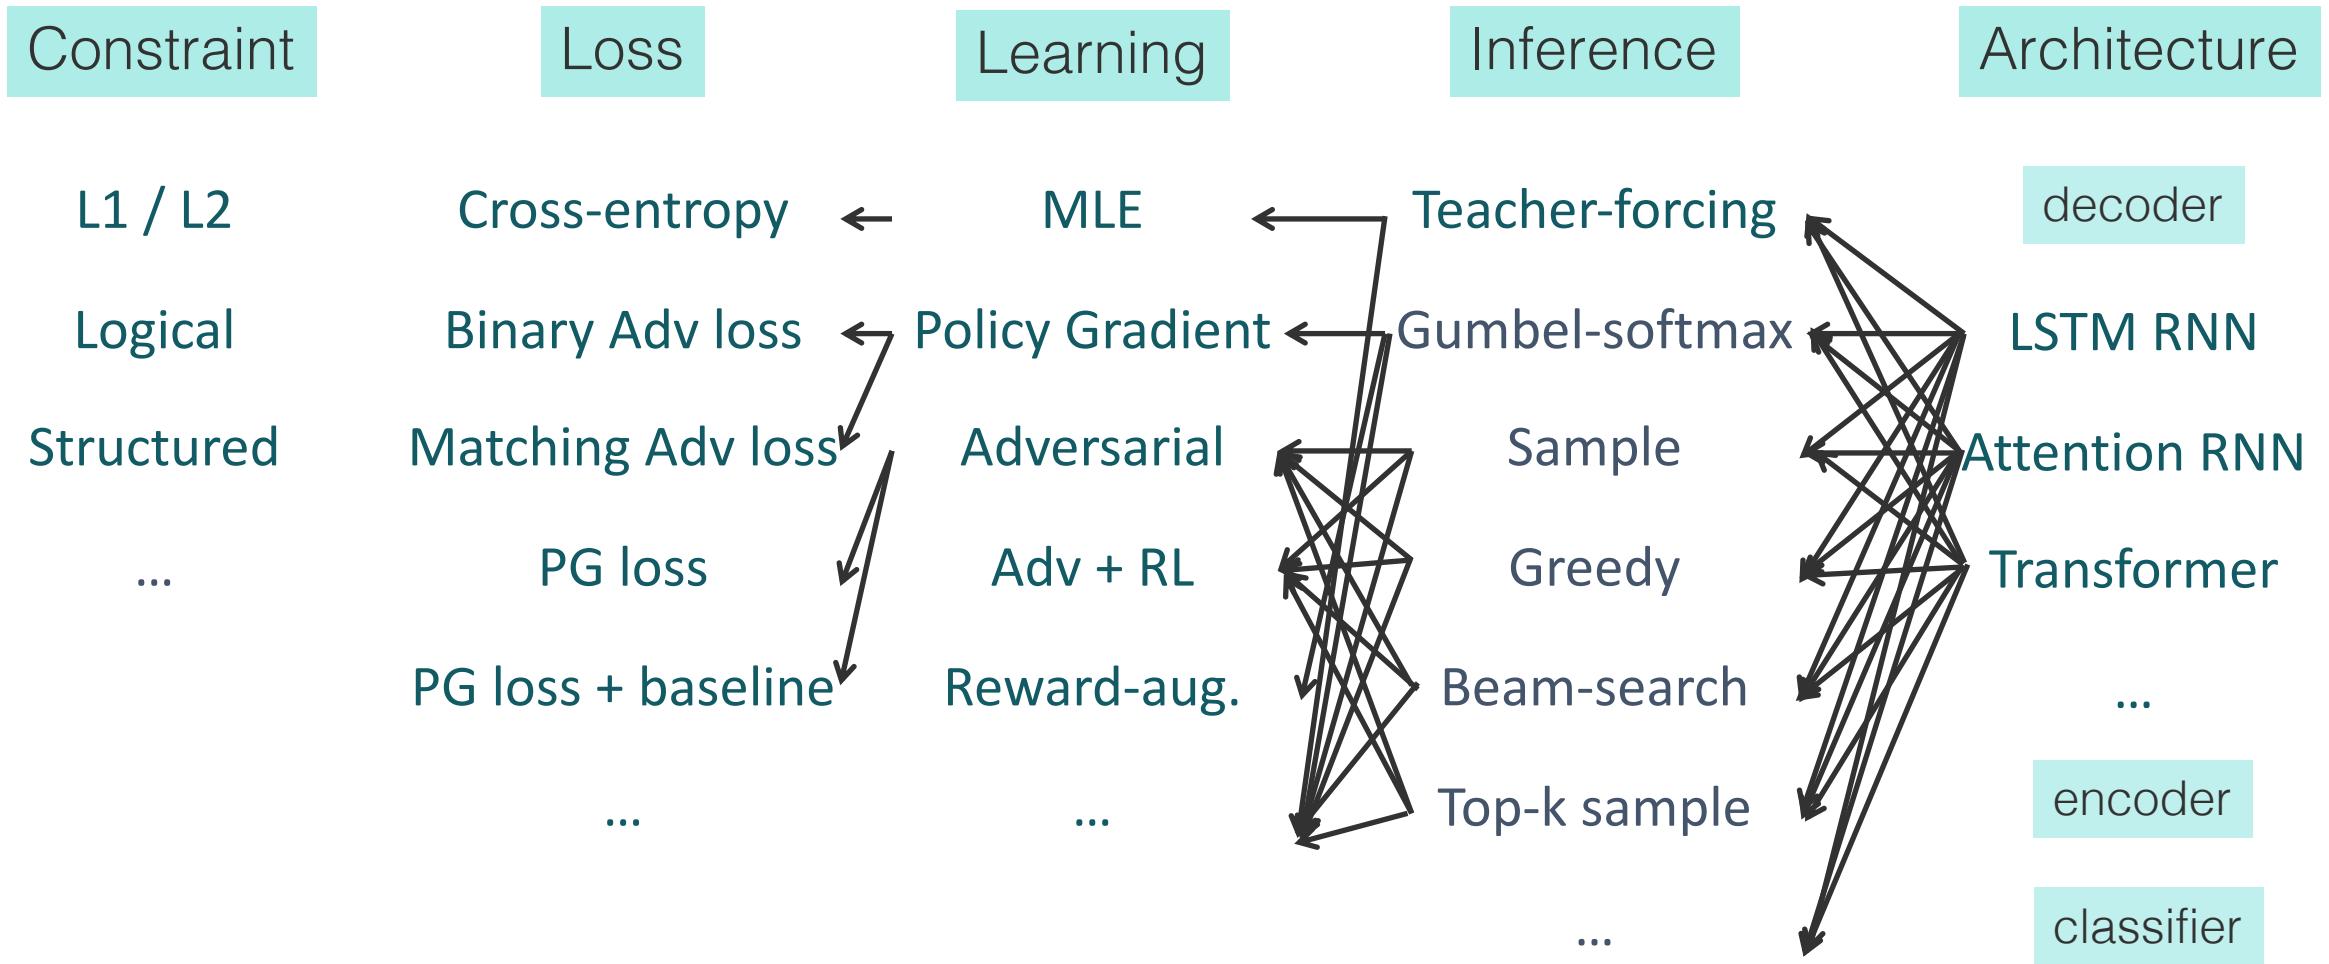
...

ML Components



Constraint	Loss	Learning	Inference	Architecture
L1 / L2	Cross-entropy	MLE	Teacher-forcing	decoder
Logical	Binary Adv loss	Policy Gradient	Gumbel-softmax	LSTM RNN
Structured	Matching Adv loss	Adversarial	Sample	Attention RNN
...	PG loss	Adv + RL	Greedy	Transformer
	PG loss + baseline	Reward-aug.	Beam-search	...
	Top-k sample	encoder
			...	classifier

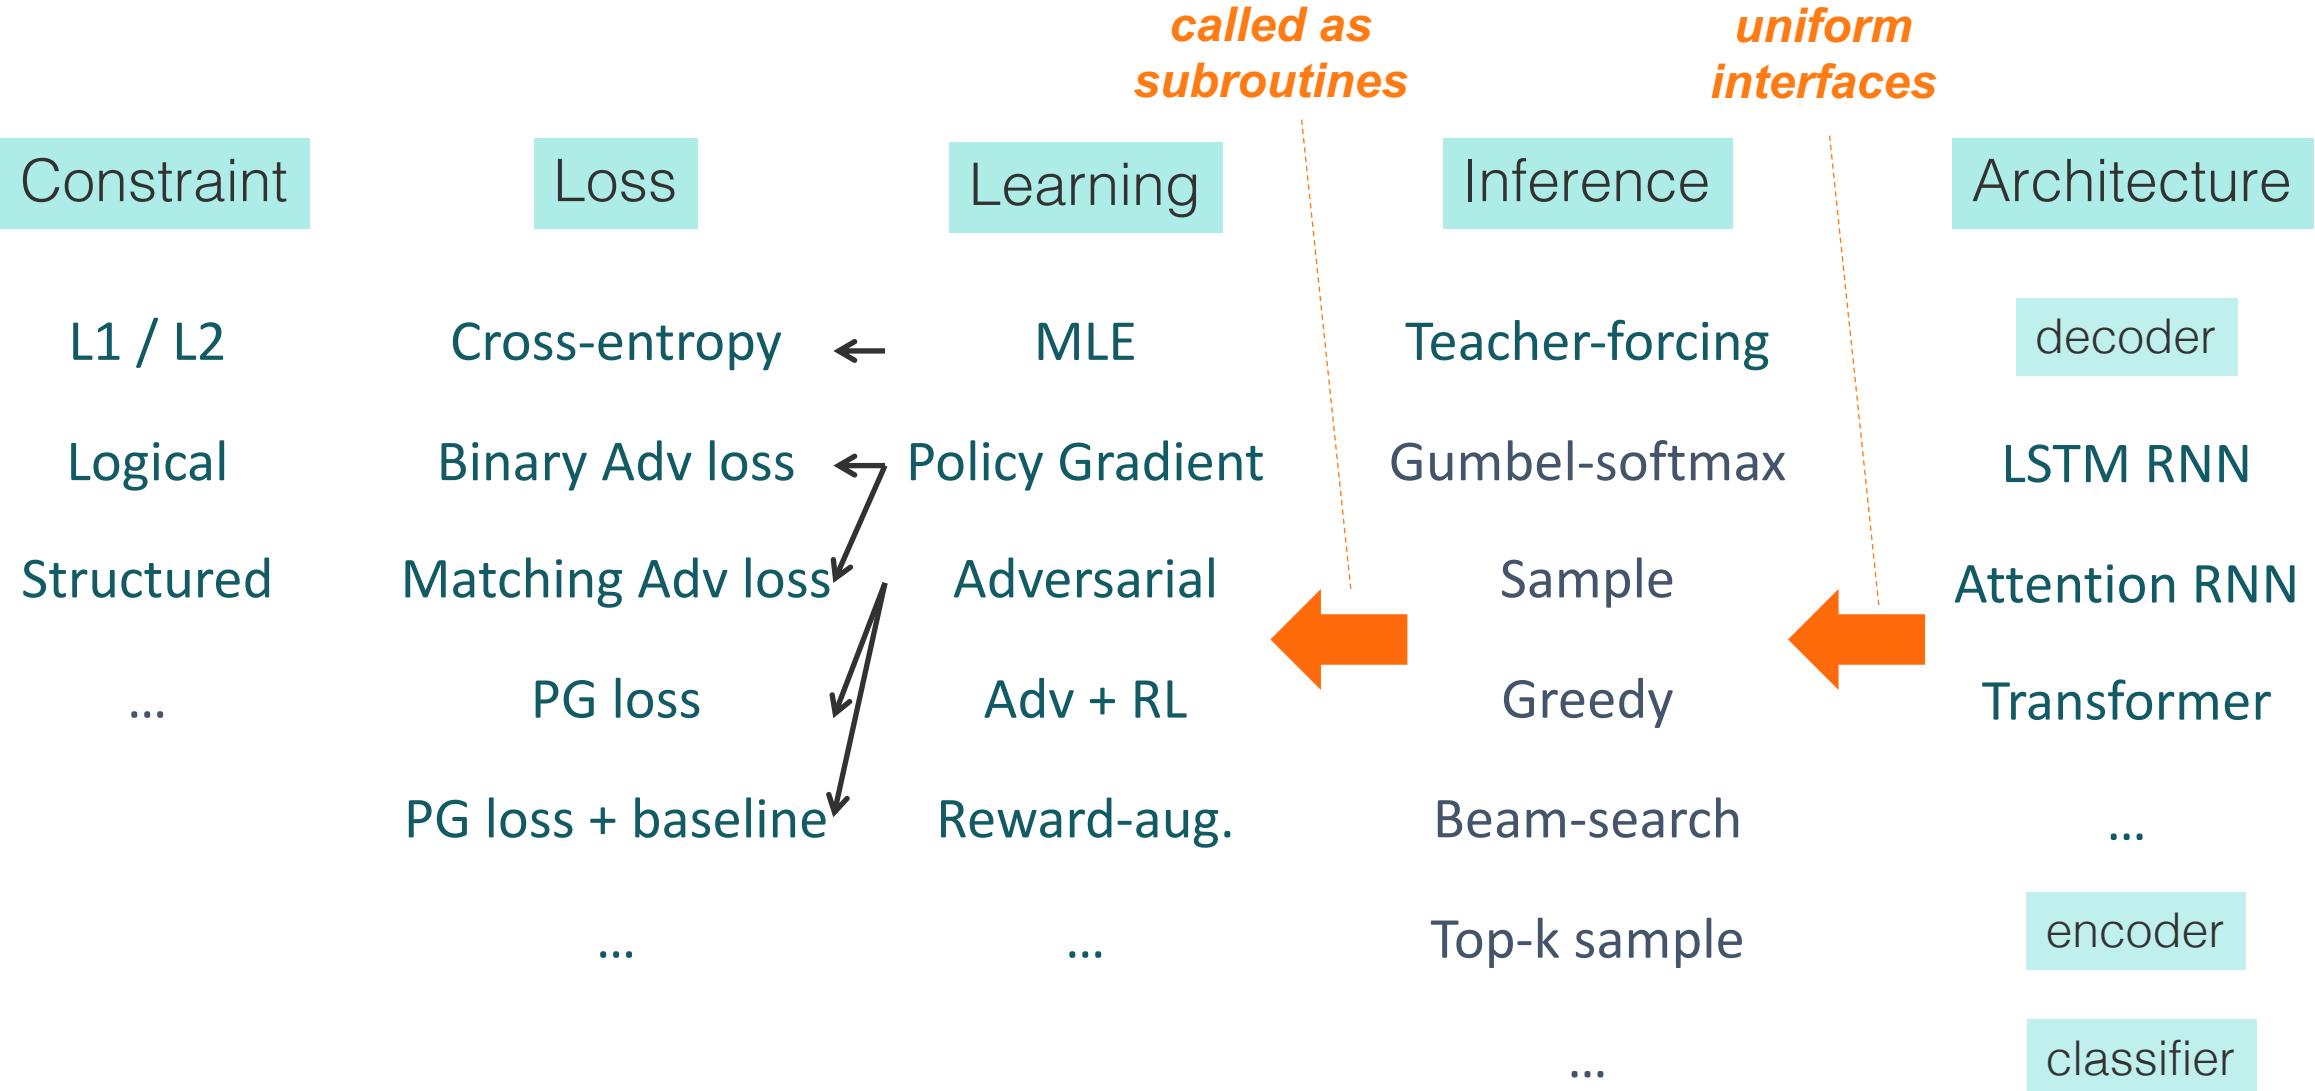
ML Components



Operationalize Composable ML with



Texar

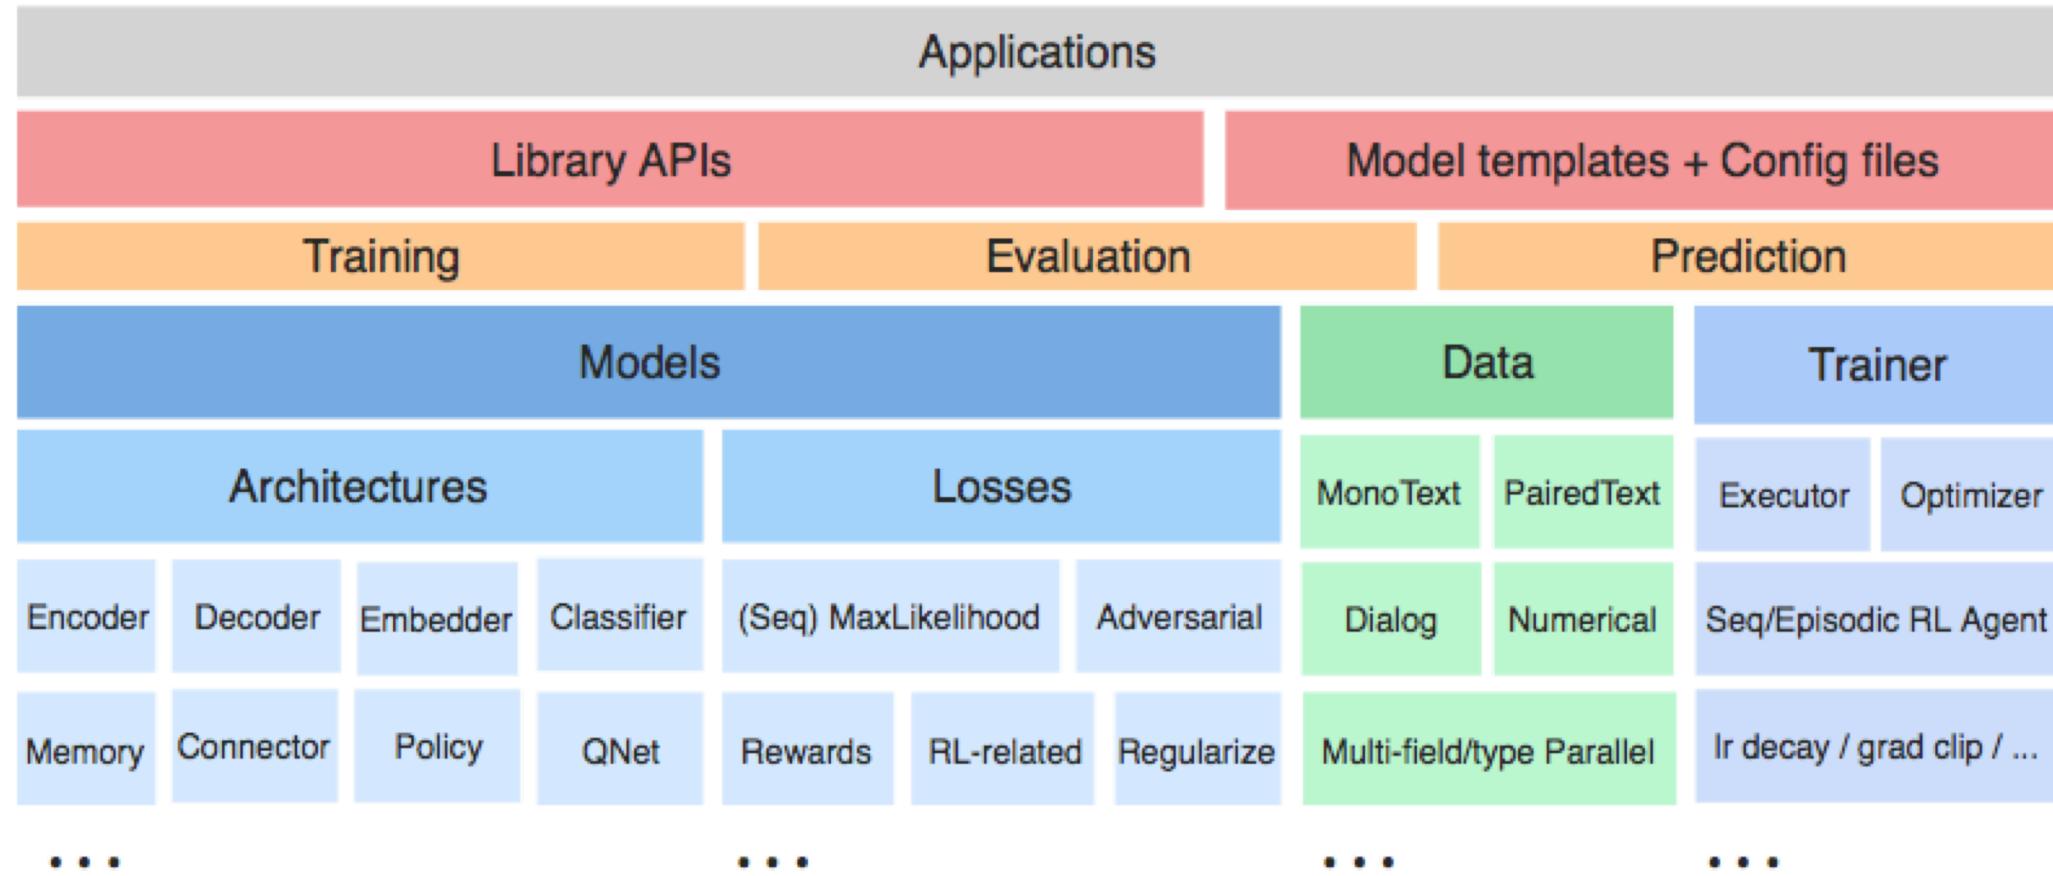


Operationalize Composable ML with

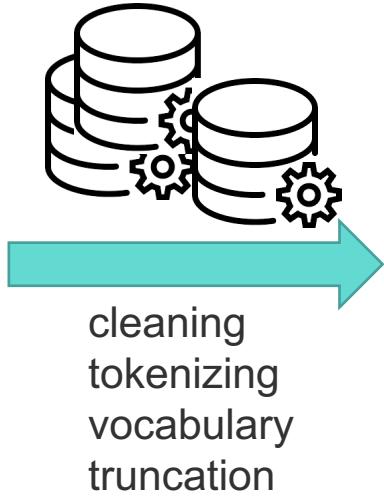


Texar

 **Texar stack**



Running Example: Machine Translation



source.dat

I like this movie.
Lovely and poignant
Insanely hilarious!
...

target.dat

Ich mag diesen film.
Schön und ergreifend
Wahnsinnig witzig!
...

evaluation
post-processing

...

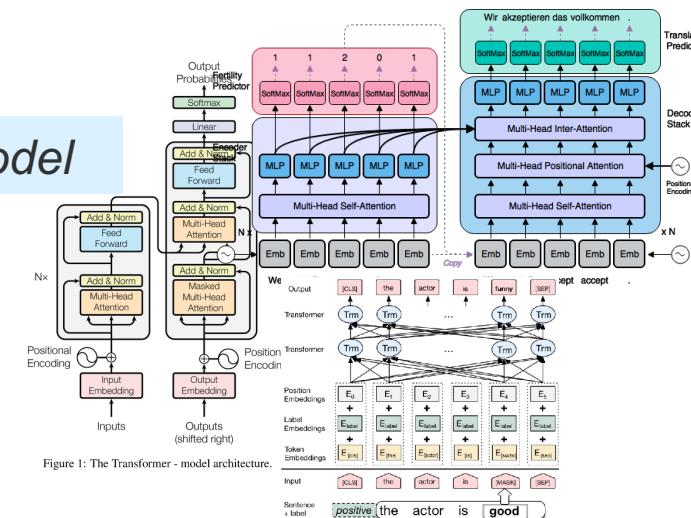
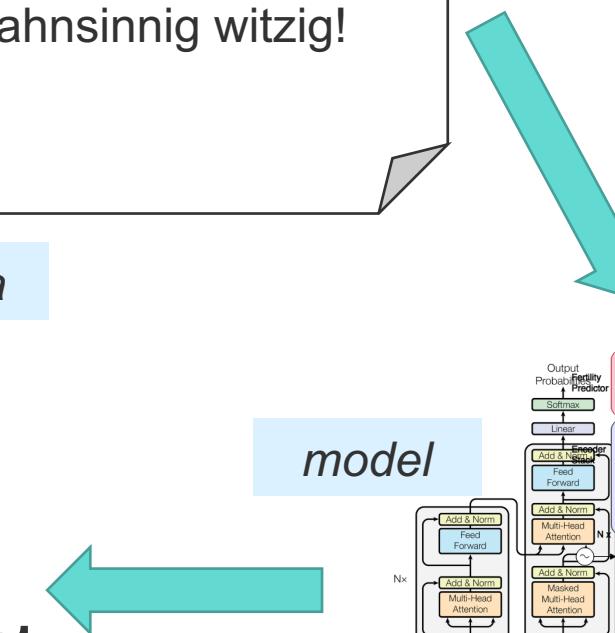
training

Maximum likelihood training

Reinforcement learning

Adversarial learning

Finetuning



Implementation with



Running Example: Machine Translation



Running Example: Machine Translation

Data

```
1 # Read data
2 dataset = PairedTextData(data_hparams)
3 batch = DataIterator(dataset).get_next()
```

YAML config files

```
1. data_hparams:
2.   batch_size: 64
3.   num_epochs: 10
4.   shuffle: True
5.   source_dataset:
6.     files: 'source.txt'
7.     vocab_file: 'vocab.txt'
8.     max_seq_length: 100
9.     bos_token: '<BOS>'
10.    eos_token: '<EOS>'
11.    target_dataset:
12.      ...
13.      ...
```

Running Example: Machine Translation

Data

```
1 # Read data  
2 dataset = PairedTextData(data_hparams)  
3 batch = DataIterator(dataset).get_next()
```

Architecture
& Inference

YAML config files

Running Example: Machine Translation

Data {

```
1 # Read data
2 dataset = PairedTextData(data_hparams)
3 batch = DataIterator(dataset).get_next()
4 # Encode
5 embedder = WordEmbedder(dataset.vocab.size, hparams=embedder_hparams)
```

YAML config files

- Architecture & Inference {
1. embedder_hparams:
 2. embedding_dim: 256
 3. dropout_rate: 0.9
 4. regularization: 'L1L2'
 5. ...

Running Example: Machine Translation

```
1 # Read data
2 dataset = PairedTextData(data_hparams)
3 batch = DataIterator(dataset).get_next()
4 # Encode
5 embedder = WordEmbedder(dataset.vocab.size, hparams=embedder_hparams)
6 encoder = TransformerEncoder(hparams=encoder_hparams)
```

Data
Architecture & Inference

YAML config files

```
1. encoder_hparams:
2.   num_blocks: 16
3.   num_heads: 8
4.   hidden_dim: 256
5.   output_dim: 128
6.   dropout_rate: 0.8
7.   ...
```

Running Example: Machine Translation

```
1 # Read data
2 dataset = PairedTextData(data_hparams)
3 batch = DataIterator(dataset).get_next()
4 # Encode
5 embedder = WordEmbedder(dataset.vocab.size, hparams=embedder_hparams)
6 encoder = TransformerEncoder(hparams=encoder_hparams)
7 enc_outputs = encoder(embedder(batch['source_text_ids']),
8                      batch['source_length'])
```

Data
Architecture & Inference



Running Example: Machine Translation

```
1 # Read data
2 dataset = PairedTextData(data_hparams)
3 batch = DataIterator(dataset).get_next()
4 # Encode
5 embedder = WordEmbedder(dataset.vocab.size, hparams=embedder_hparams)
6 encoder = TransformerEncoder(hparams=encoder_hparams)
7 enc_outputs = encoder(embedder(batch['source_text_ids']),
8                      batch['source_length'])
9 # Build decoder
10 decoder = AttentionRNNDDecoder(memory=enc_outputs,
11                                 hparams=decoder_hparams)
12 # Maximum Likelihood Estimation
13 ## Teacher-forcing decoding
14 outputs, length, _ = decoder(decoding_strategy='teacher-forcing',
15                             inputs=embedder(batch['target_text_ids']),
16                             seq_length=batch['target_length']-1)
```

Data
Architecture & Inference

Running Example: Machine Translation

```
1 # Read data
2 dataset = PairedTextData(data_hparams)
3 batch = DataIterator(dataset).get_next()
4 # Encode
5 embedder = WordEmbedder(dataset.vocab.size, hparams=embedder_hparams)
6 encoder = TransformerEncoder(hparams=encoder_hparams)
7 enc_outputs = encoder(embedder(batch['source_text_ids']),
8                      batch['source_length'])
9 # Build decoder
10 decoder = AttentionRNNDDecoder(memory=enc_outputs,
11                                 hparams=decoder_hparams)
12 # Maximum Likelihood Estimation
13 ## Teacher-forcing decoding
14 outputs, length, _ = decoder(decoding_strategy='teacher-forcing',
15                             inputs=embedder(batch['target_text_ids']),
16                             seq_length=batch['target_length']-1)
17 ## Cross-entropy loss
18 loss = sequence_sparse_softmax_cross_entropy(
19     labels=batch['target_text_ids'][:, 1:], logits=outputs.logits, seq_length=length)
20
```

Architecture
& InferenceLearning
Loss

Running Example: Machine Translation

```

1 # Read data
2 dataset = PairedTextData(data_hparams)
3 batch = DataIterator(dataset).get_next()
4 # Encode
5 embedder = WordEmbedder(dataset.vocab.size, hparams=embedder_hparams)
6 encoder = TransformerEncoder(hparams=encoder_hparams)
7 enc_outputs = encoder(embedder(batch['source_text_ids']),
8                      batch['source_length'])
9 # Build decoder
10 decoder = AttentionRNNDDecoder(memory=enc_outputs,
11                                 hparams=decoder_hparams)
12 # Maximum Likelihood Estimation
13 ## Teacher-forcing decoding
14 outputs, length, _ = decoder(decoding_strategy='teacher-forcing',
15                             inputs=embedder(batch['target_text_ids']),
16                             seq_length=batch['target_length']-1)
17 ## Cross-entropy loss
18 loss = sequence_sparse_softmax_cross_entropy(
19     labels=batch['target_text_ids'][:,1:], logits=outputs.logits, seq_length=length)
20

```

**Maximum likelihood
Estimation**

Switching between Learning Algorithms

Data

```

1 # Read data
2 dataset = PairedTextData(data_hparams)
3 batch = DataIterator(dataset).get_next()
4 # Encode
5 embedder = WordEmbedder(dataset.vocab.size, hparams=embedder_hparams)
6 encoder = TransformerEncoder(hparams=encoder_hparams)
7 enc_outputs = encoder(embedder(batch['source_text_ids']),
8                         batch['source_length'])

```

Architecture & Inference

```

9 # Build decoder
10 decoder = AttentionRNNDDecoder(memory=enc_outputs,
11                                 hparams=decoder_hparams)
12 # Maximum Likelihood Estimation
13 ## Teacher-forcing decoding
14 outputs, length, _ = decoder(decoding_strategy='teacher-forcing',
15                             inputs=embedder(batch['target_text_ids']),
16                             seq_length=batch['target_length']-1)

```

Learning Loss

```

17 ## Cross-entropy loss
18 loss = sequence_sparse_softmax_cross_entropy(
19     labels=batch['target_text_ids'][:, 1:], logits=outputs.logits, seq_length=length)
20

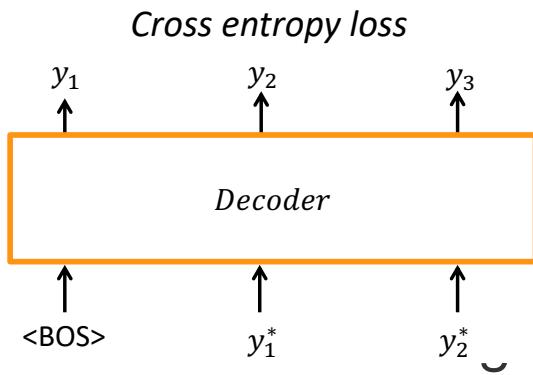
```

Keep unchanged

Maximum likelihood Estimation

Switching from MLE to Reinforcement Learning

- Maximum likelihood

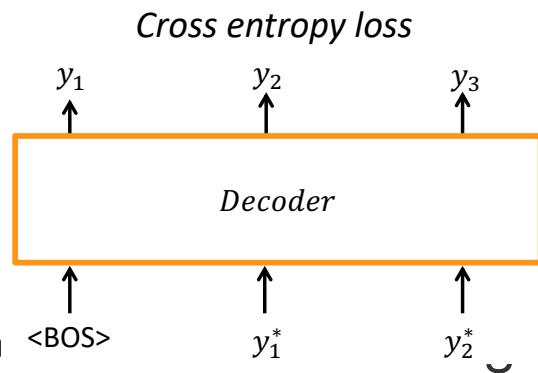


```
# Teacher-forcing decoding
outputs, length, _ = decoder(decoding_strategy='teacher-forcing',
                               inputs=embedder(batch['target_text_ids']),
                               seq_length=batch['target_length']-1)

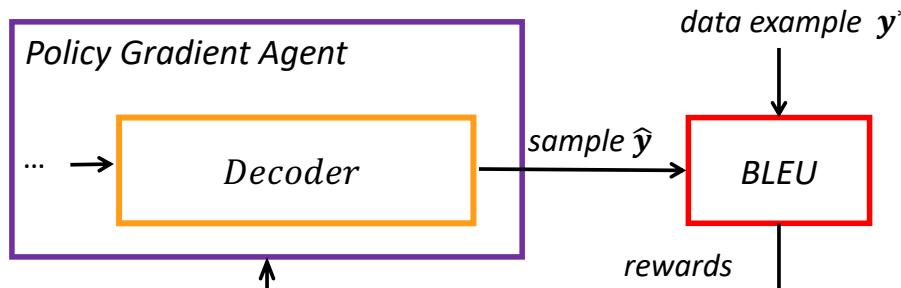
# Cross-entropy loss
loss = sequence_sparse_softmax_cross_entropy(
    labels=batch['target_text_ids'][:,1:], logits=outputs.logits, seq_length=length)
```

Switching from MLE to Reinforcement Learning

- Maximum likelihood



- R_t



```
# Teacher-forcing decoding
outputs, length, _ = decoder(decoding_strategy='teacher-forcing',
                             inputs=embedder(batch['target_text_ids']),
                             seq_length=batch['target_length']-1)

# Cross-entropy loss
loss = sequence_sparse_softmax_cross_entropy(
    labels=batch['target_text_ids'][:,1:], logits=outputs.logits, seq_length=length)

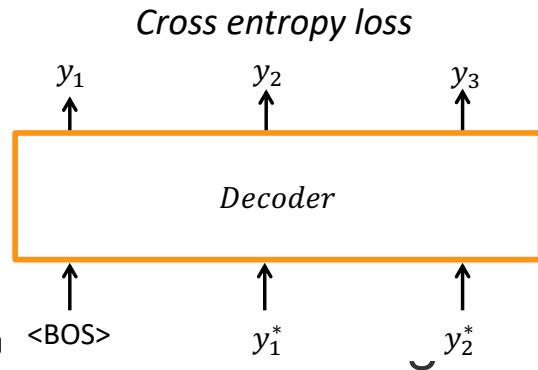
# Random sample decoding
outputs, length, _ = decoder(decoding_strategy='random_sample',
                             start_tokens=[BOS]*batch_size, end_token=EOS,
                             embedding=embedder)

# Policy gradient agent for learning
agent = SeqPGAgent(
    samples=outputs.sample_id, logits=outputs.logits, seq_length=length)

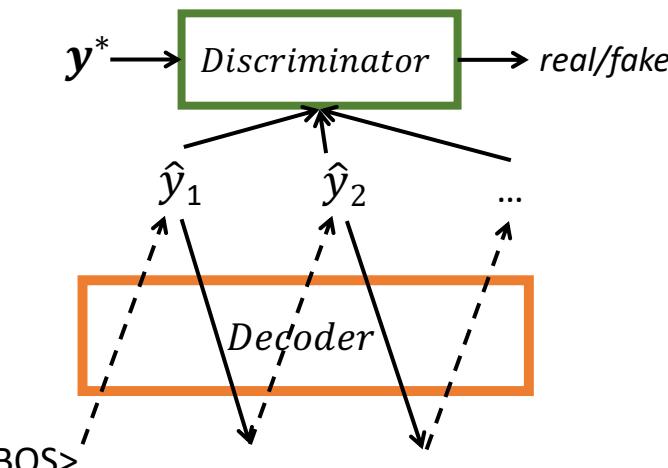
for _ in range(STEPS):
    samples = agent.get_samples()
    rewards = BLEU(batch['target_text_ids'], samples) # Reward
    agent.observe(rewards)
```

Switching from MLE to Adversarial Learning

- Maximum likelihood



- A



```
# Teacher-forcing decoding  
outputs, length, _ = decoder(decoding_strategy='teacher-forcing',  
                             inputs=embedder(batch['target_text_ids']),  
                             seq_length=batch['target_length']-1)  
  
# Cross-entropy loss  
loss = sequence_sparse_softmax_cross_entropy(  
    labels=batch['target_text_ids'][:, 1:], logits=outputs.logits, seq_length=length)
```

```
# Gumbel-softmax decoding  
outputs, _, _ = decoder(decoding_strategy='gumbel-softmax',  
                        start_tokens=[BOS]*batch_size, end_token=EOS,  
                        embedding=embedder)
```

```
discriminator = Conv1DClassifier(hparams=conv_hparams)
```

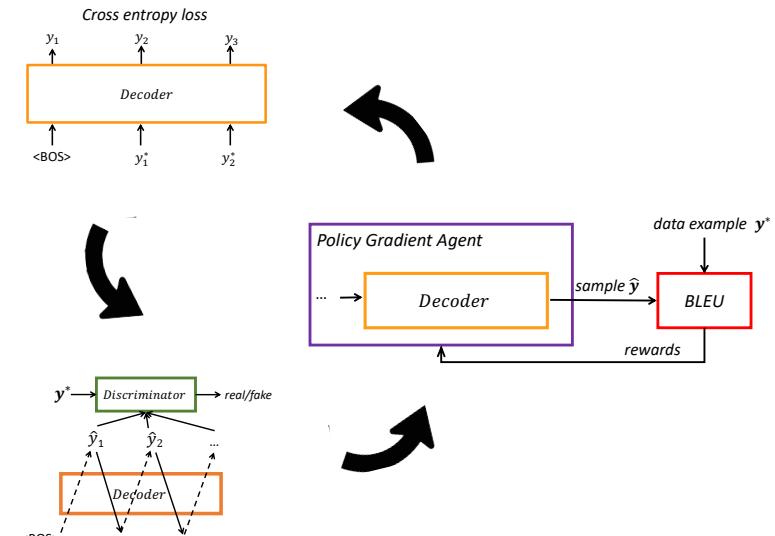
```
# Binary adversarial loss
```

```
G_loss, D_loss = binary_adversarial_losses(  
    embedder(batch['target_text_ids'][:, 1:]),  
    embedder(soft_ids=softmax(outputs.logits)),  
    discriminator)
```

Summary of MT in Texar

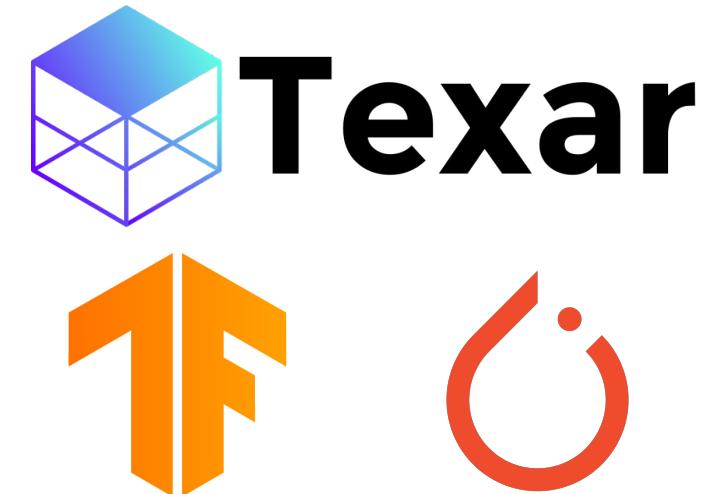
- Highly modularized programming
 - Data, architecture, loss, inference, learning, ...
 - Intuitive conceptual-level APIs
- Easy switch between learning algorithms
 - Plug in & out modules
 - No changes to irrelevant parts

```
1 # Read data
2 dataset = PairedTextData(data_hparams)
3 batch = DataIterator(dataset).get_next()
4 # Encode
5 embedder = WordEmbedder(dataset.vocab.size, hparams=embedder_hparams)
6 encoder = TransformerEncoder(hparams=encoder_hparams)
7 enc_outputs = encoder(embedder(batch['source_text_ids']),
8                      batch['source_length'])
9 # Build decoder
10 decoder = AttentionRNNDecoder(memory=enc_outputs,
11                                 hparams=decoder_hparams)
12 # Maximum Likelihood Estimation
13 ## Teacher-forcing decoding
14 outputs, length, _ = decoder(decoding_strategy='teacher-forcing',
15                             inputs=embedder(batch['target_text_ids']),
16                             seq_length=batch['target_length']-1)
17 ## Cross-entropy loss
18 loss = sequence_sparse_softmax_cross_entropy(
19     labels=batch['target_text_ids'][:,1:], logits=outputs.logits, seq_length=length)
20
```

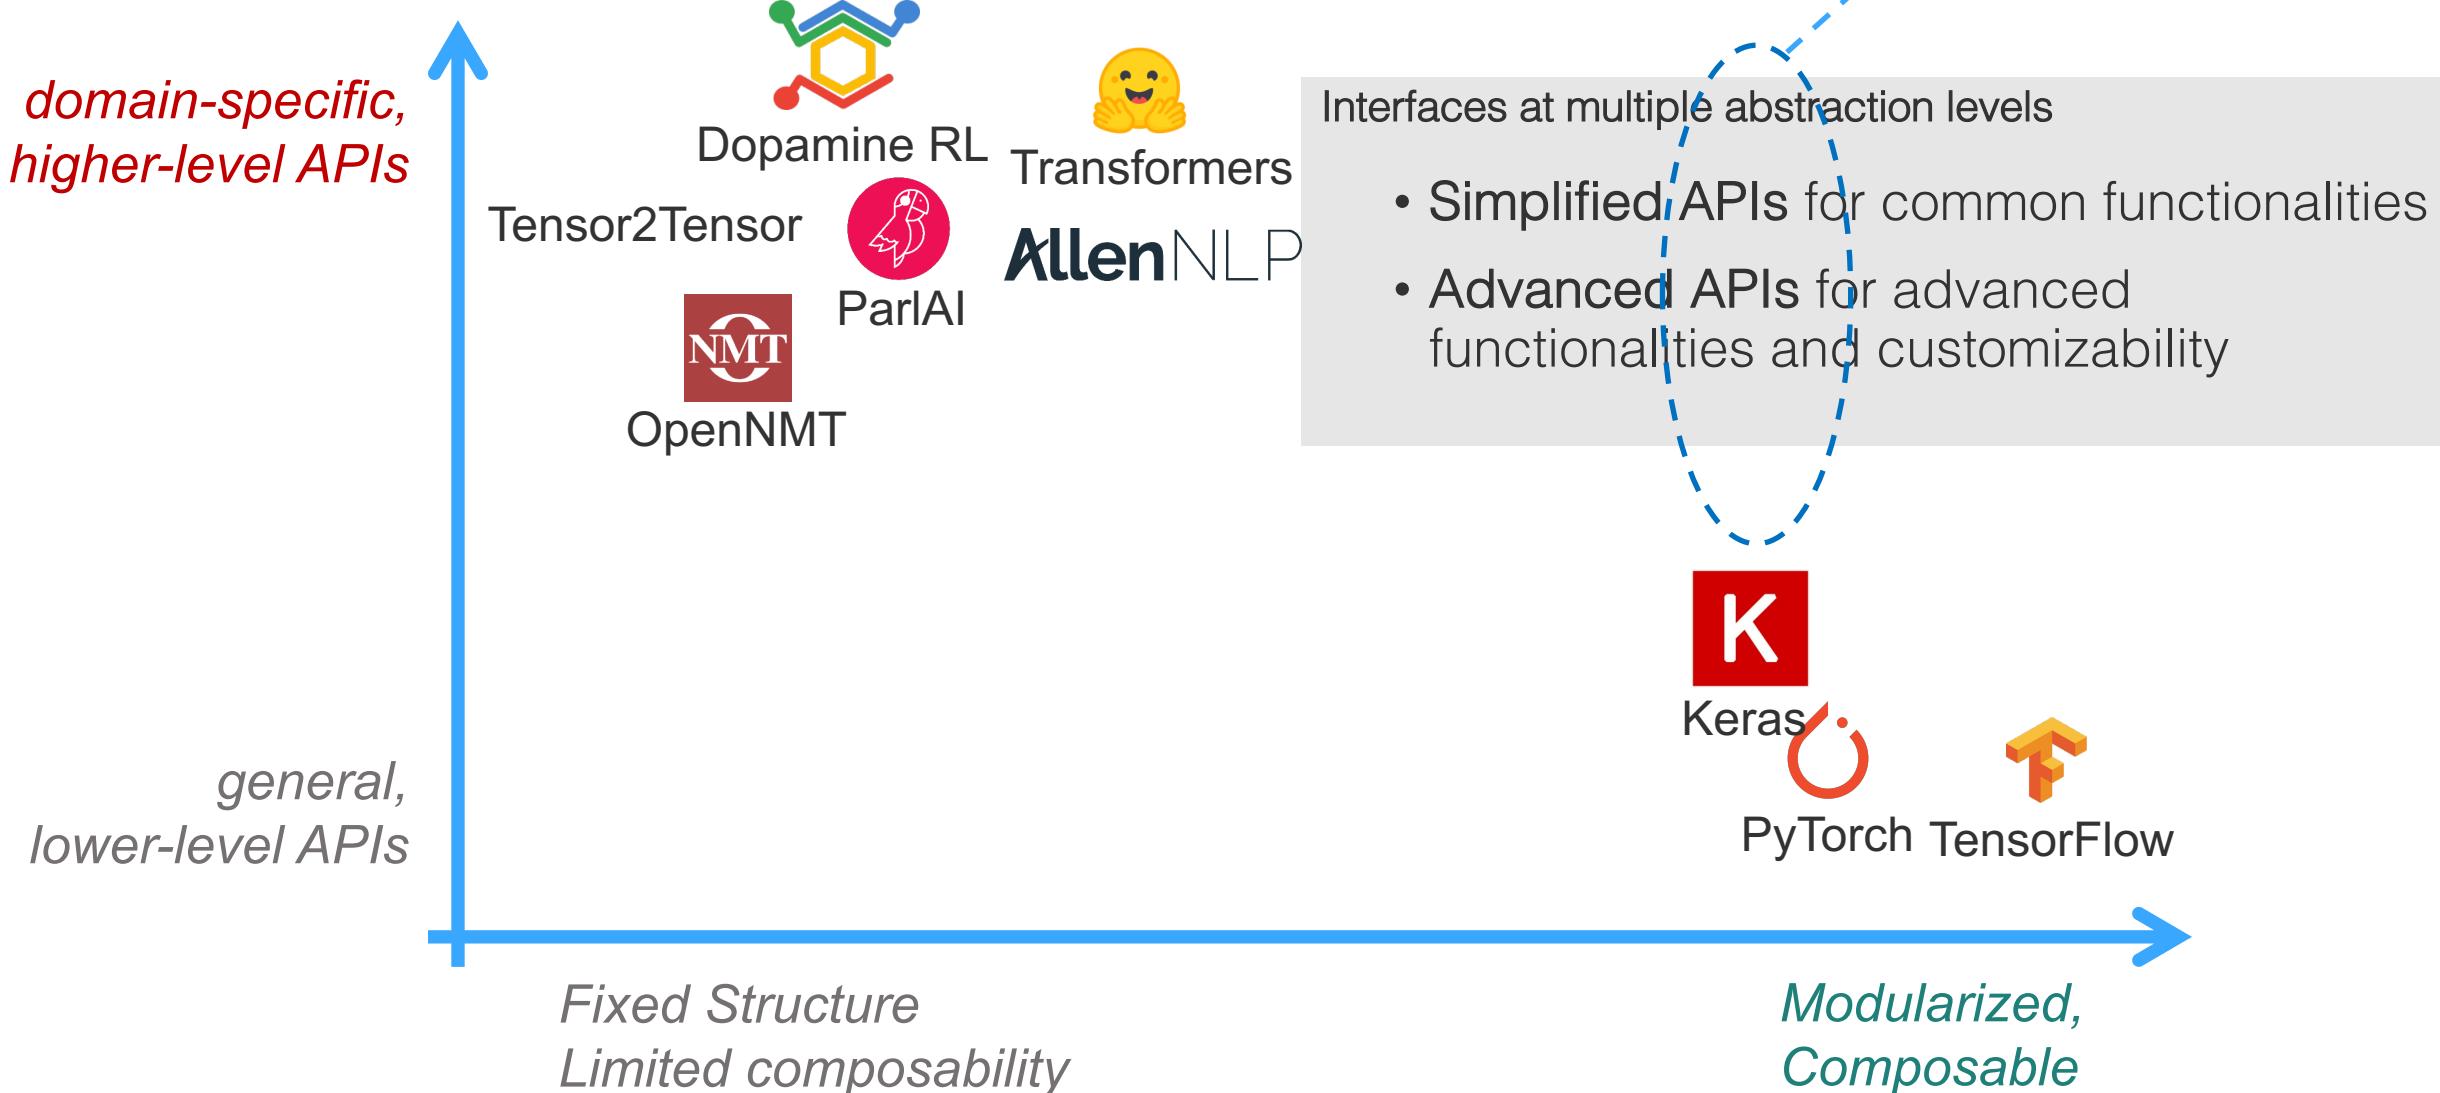


Support of TensorFlow and PyTorch

- Texar is built upon TF and PyTorch
 - Texar-TF & Texar-PyTorch: mostly the same interfaces!
 - Higher-level intuitive APIs without loss of flexibility
 - Lots of ML components ready to use
- Combine the best design of TF and PyTorch
 - TF:
 - Easy and efficient data processing APIs
 - Excellent factorization of ML modules
 - Turnkey model training processor
 - PyTorch:
 - Intuitive programming interfaces
 - Transparent variable scope and sharing to users



Spectrum of Existing Tools

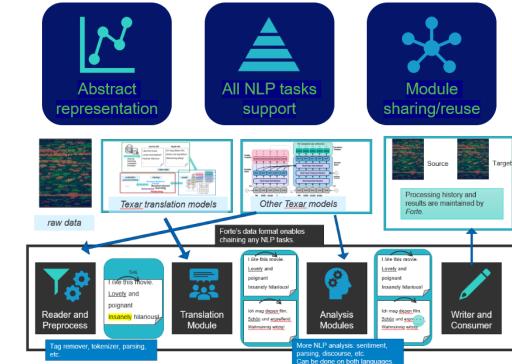


Applications of Texar

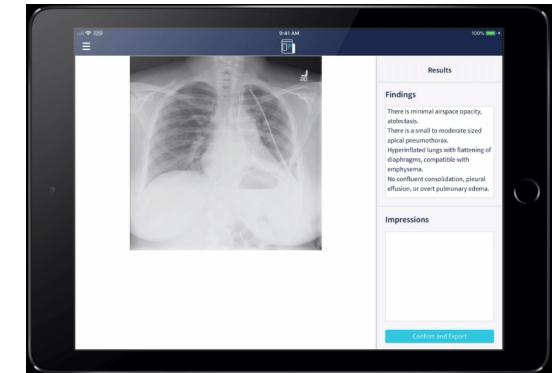
Many products built on Texar

- FORTE – templates for larger complex NLP applications
- Chest X-Ray report writer
- Medical Registry report writer
- ICD coding system
- Financial knowledge base builder
- Financial summary/report writer
- Multi-Lingual Cognitive Chat Bots
 - For Call Center Support
 - For Retail In-Store Assistance

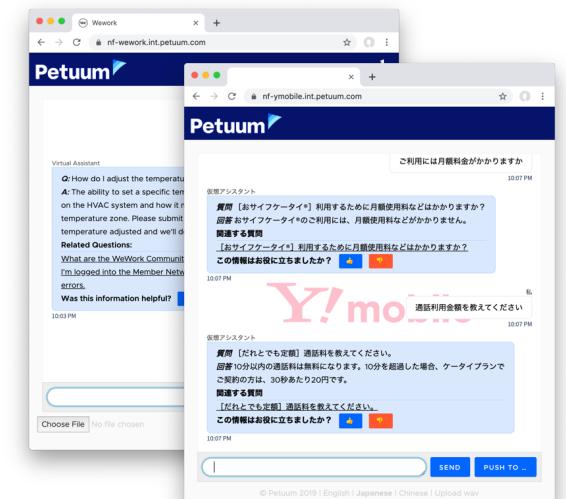
FORTE



Chest X-Ray Report Writer



Multi-Lingual Cognitive Chat Bots

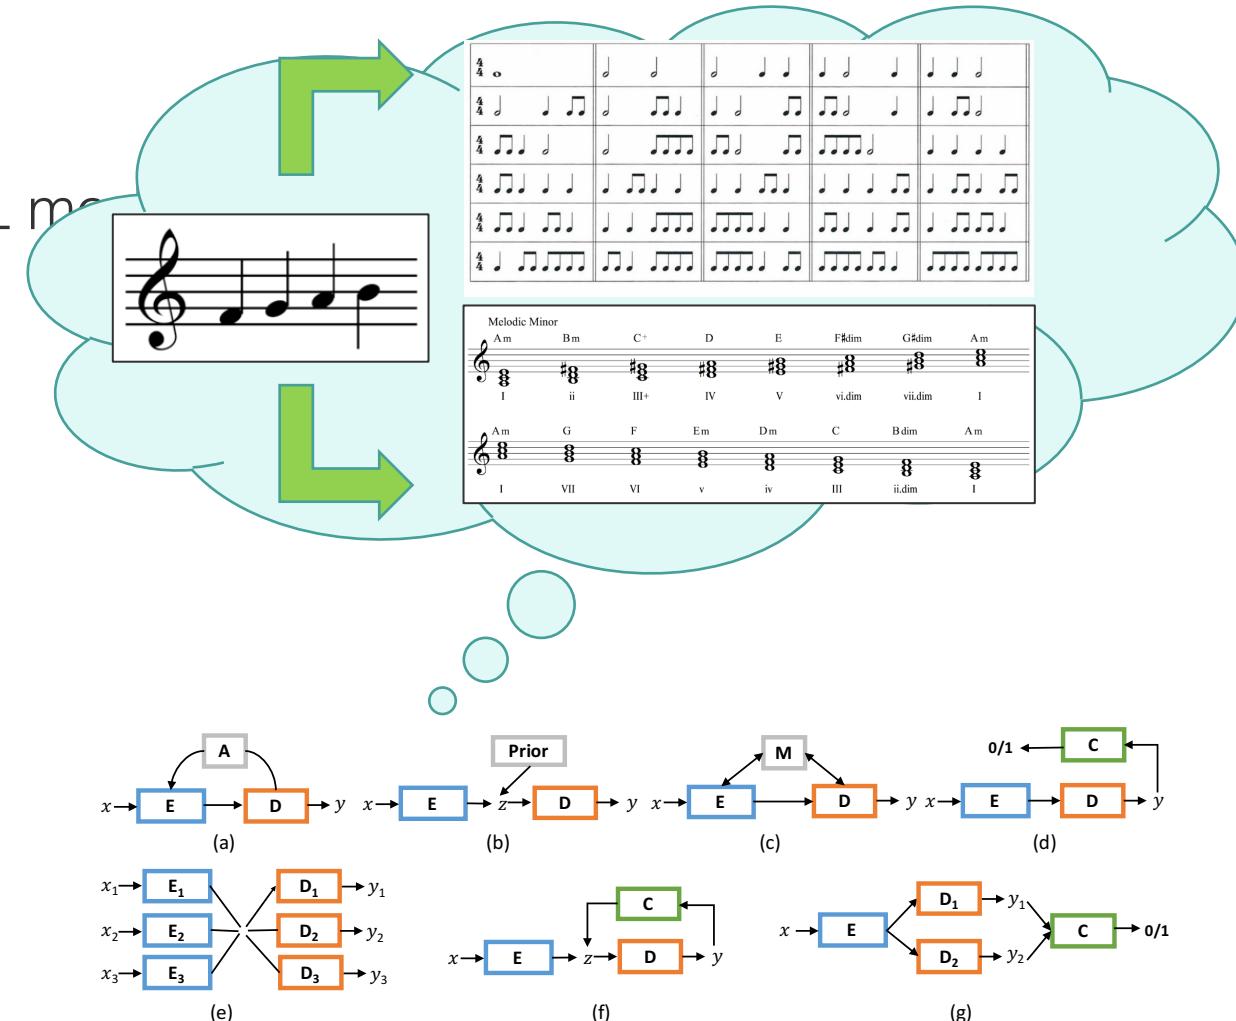


Composable ML – Take-Home Message

- Composable ML

- Basic “musical notes” for complex ML models
- (or just think of it as Lego for ML)

-  **Texar**



Agenda

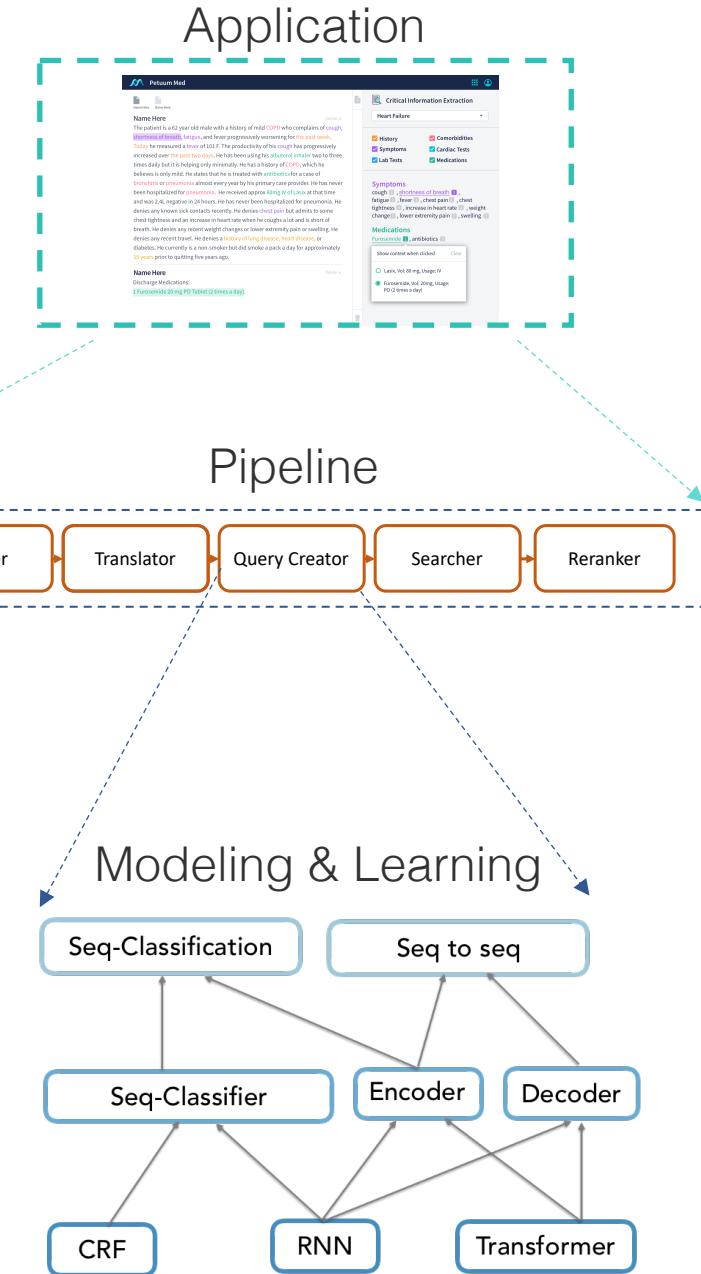


Forte



Texar

- Natural Language Processing Overview (10mins)
- Modularizing NLP Pipeline (35mins)
 - Complexity of NLP pipeline
 - A standardized view of NLP pipeline
 - A standardized implementation of NLP pipeline
- Short break & QA (5mins)
- Modularizing NLP Model & Learning (30mins)
 - Composable ML
- QA (10mins)



What's Next

- Data Manipulation
 - Data augmentation
 - Data visualization
 - Multi-modal data manipulation
- Task Inter-operation
 - Joint learning
 - Joint inference
- Automate ML workflow
 - Automate model composing, learning

Petuum Med

Name Here Name Here

Name Here

The patient is a 62 year old male with a history of mild COPD who complains of cough, shortness of breath, fatigue, and fever progressively worsening for the past week. Today he measured a fever of 101 F. The productivity of his cough has progressively increased over the past two days. He has been using his albuterol inhaler two to three times daily but it is helping only minimally. He has a history of COPD, which he believes is only mild. He states that he is treated with antibiotics for a case of bronchitis or pneumonia almost every year by his primary care provider. He has never been hospitalized for pneumonia. He received approx 80mg IV of Lasix at that time and was 2.4L negative in 24 hours. He has never been hospitalized for pneumonia. He denies any known sick contacts recently. He denies chest pain but admits to some chest tightness and an increase in heart rate when he coughs a lot and is short of breath. He denies any recent weight changes or lower extremity pain or swelling. He denies any recent travel. He denies a history of lung disease, heart disease, or diabetes. He currently is a non-smoker but did smoke a pack a day for approximately 15 years prior to quitting five years ago.

Symptoms

cough , shortness of breath , fatigue , fever , chest pain , chest tightness , increase in heart rate , weight change , lower extremity pain , swelling

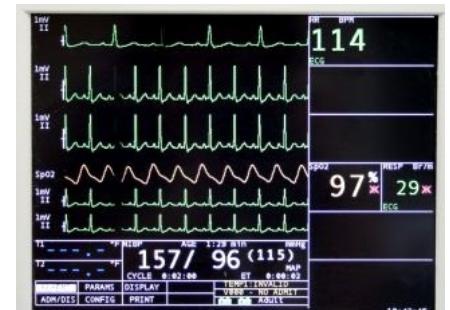
Medications

Furosemide , antibiotics

Show context when clicked Clear

Lasix, Vol: 80 mg, Usage: IV

Furosemide, Vol: 20mg, Usage: PO (2 times a day)





PETUUM

Thank you



Carnegie Mellon University
School of Computer Science